



Database Compability for Oracle® Developers Built-in Package Guide

Release 12

EDB Postgres™ Advanced Server 12

Mar 17, 2020

1	Introduction	1
1.1	What's New	2
2	Packages	3
2.1	Package Components	4
2.1.1	Package Specification Syntax	4
2.1.2	Package Body Syntax	7
2.2	Creating Packages	12
2.2.1	Creating the Package Specification	12
2.2.2	Creating the Package Body	13
2.3	Referencing a Package	15
2.4	Using Packages With User Defined Types	16
2.5	Dropping a Package	20
3	Built-In Packages	21
3.1	DBMS_ALERT	22
3.1.1	REGISTER	22
3.1.2	REMOVE	23
3.1.3	REMOVEALL	23
3.1.4	SIGNAL	23
3.1.5	WAITANY	24
3.1.6	WAITONE	25
3.1.7	Comprehensive Example	27
3.2	DBMS_AQ	29
3.2.1	ENQUEUE	32
3.2.2	DEQUEUE	35
3.2.3	REGISTER	39
3.2.4	UNREGISTER	41
3.3	DBMS_AQADM	43
3.3.1	ALTER_QUEUE	45
3.3.2	ALTER_QUEUE_TABLE	46
3.3.3	CREATE_QUEUE	47
3.3.4	CREATE_QUEUE_TABLE	49
3.3.5	DROP_QUEUE	51
3.3.6	DROP_QUEUE_TABLE	52
3.3.7	PURGE_QUEUE_TABLE	53
3.3.8	START_QUEUE	54

3.3.9	STOP_QUEUE	55
3.4	DBMS_CRYPTO	56
3.4.1	DECRYPT	57
3.4.2	ENCRYPT	59
3.4.3	HASH	61
3.4.4	MAC	62
3.4.5	RANDOMBYTES	63
3.4.6	RANDOMINTEGER	64
3.4.7	RANDOMNUMBER	65
3.5	DBMS_JOB	66
3.5.1	BROKEN	68
3.5.2	CHANGE	69
3.5.3	INTERVAL	70
3.5.4	NEXT_DATE	71
3.5.5	REMOVE	72
3.5.6	RUN	73
3.5.7	SUBMIT	74
3.5.8	WHAT	75
3.6	DBMS_LOB	76
3.6.1	APPEND	78
3.6.2	COMPARE	79
3.6.3	CONVERTTOBLOB	80
3.6.4	CONVERTTOCLOB	81
3.6.5	COPY	82
3.6.6	ERASE	83
3.6.7	GET_STORAGE_LIMIT	84
3.6.8	GETLENGTH	85
3.6.9	INSTR	86
3.6.10	READ	87
3.6.11	SUBSTR	88
3.6.12	TRIM	89
3.6.13	WRITE	90
3.6.14	WRITEAPPEND	91
3.7	DBMS_LOCK	92
3.7.1	SLEEP	92
3.8	DBMS_MVIEW	93
3.8.1	GET_MV_DEPENDENCIES	94
3.8.2	REFRESH	95
3.8.3	REFRESH_ALL_MVIEWS	97
3.8.4	REFRESH_DEPENDENT	98
3.9	DBMS_OUTPUT	100
3.9.1	CHARARR	100
3.9.2	DISABLE	101
3.9.3	ENABLE	102
3.9.4	GET_LINE	103
3.9.5	GET_LINES	105
3.9.6	NEW_LINE	107
3.9.7	PUT	108
3.9.8	PUT_LINE	109
3.9.9	SERVEROUTPUT	110
3.10	DBMS_PIPE	111
3.10.1	CREATE_PIPE	112
3.10.2	NEXT_ITEM_TYPE	113
3.10.3	PACK_MESSAGE	115

3.10.4	PURGE	116
3.10.5	RECEIVE_MESSAGE	118
3.10.6	REMOVE_PIPE	119
3.10.7	RESET_BUFFER	121
3.10.8	SEND_MESSAGE	122
3.10.9	UNIQUE_SESSION_NAME	123
3.10.10	UNPACK_MESSAGE	124
3.10.11	Comprehensive Example	125
3.11	DBMS_PROFILER	128
3.11.1	FLUSH_DATA	129
3.11.2	GET_VERSION	129
3.11.3	INTERNAL_VERSION_CHECK	129
3.11.4	PAUSE_PROFILER	129
3.11.5	RESUME_PROFILER	130
3.11.6	START_PROFILER	130
3.11.7	STOP_PROFILER	130
3.11.8	Using DBMS_PROFILER	131
3.11.8.1	Querying the DBMS_PROFILER Tables and View	133
3.11.8.2	DBMS_PROFILER - Reference	140
3.11.8.2.1	PLSQL_PROFILER_RUNS	140
3.11.8.2.2	PLSQL_PROFILER_UNITS	140
3.11.8.2.3	PLSQL_PROFILER_DATA	141
3.11.8.2.4	PLSQL_PROFILER_RAWDATA	141
3.12	DBMS_RANDOM	151
3.12.1	INITIALIZE	151
3.12.2	NORMAL	151
3.12.3	RANDOM	152
3.12.4	SEED	152
3.12.5	SEED	152
3.12.6	STRING	153
3.12.7	TERMINATE	153
3.12.8	VALUE	153
3.12.9	VALUE	154
3.13	DBMS_REDACT	155
3.13.1	Using DBMS_REDACT Constants and Function Parameters	156
3.13.2	ADD_POLICY	168
3.13.3	ALTER_POLICY	171
3.13.4	DISABLE_POLICY	174
3.13.5	ENABLE_POLICY	175
3.13.6	DROP_POLICY	176
3.13.7	UPDATE_FULL_REDACTION_VALUES	177
3.14	DBMS_RLS	180
3.14.1	ADD_POLICY	183
3.14.2	DROP_POLICY	190
3.14.3	ENABLE_POLICY	191
3.15	DBMS_SCHEDULER	192
3.15.1	Using Calendar Syntax to Specify a Repeating Interval	194
3.15.2	CREATE_JOB	196
3.15.3	CREATE_PROGRAM	198
3.15.4	CREATE_SCHEDULE	200
3.15.5	DEFINE_PROGRAM_ARGUMENT	201
3.15.6	DISABLE	203
3.15.7	DROP_JOB	204
3.15.8	DROP_PROGRAM	205

3.15.9	DROP_PROGRAM_ARGUMENT	206
3.15.10	DROP_SCHEDULE	207
3.15.11	ENABLE	208
3.15.12	EVALUATE_CALENDAR_STRING	209
3.15.13	RUN_JOB	210
3.15.14	SET_JOB_ARGUMENT_VALUE	211
3.16	DBMS_SESSION	212
3.16.1	SET_ROLE	212
3.17	DBMS_SQL	213
3.17.1	BIND_VARIABLE	215
3.17.2	BIND_VARIABLE_CHAR	217
3.17.3	BIND_VARIABLE_RAW	218
3.17.4	CLOSE_CURSOR	219
3.17.5	COLUMN_VALUE	220
3.17.6	COLUMN_VALUE_CHAR	222
3.17.7	COLUMN_VALUE_RAW	223
3.17.8	DEFINE_COLUMN	224
3.17.9	DEFINE_COLUMN_CHAR	226
3.17.10	DEFINE_COLUMN_RAW	227
3.17.11	DESCRIBE_COLUMNS	228
3.17.12	EXECUTE	229
3.17.13	EXECUTE_AND_FETCH	230
3.17.14	FETCH_ROWS	232
3.17.15	IS_OPEN	234
3.17.16	LAST_ROW_COUNT	235
3.17.17	OPEN_CURSOR	237
3.17.18	PARSE	238
3.18	DBMS_UTILITY	240
3.18.1	LNAME_ARRAY	241
3.18.2	UNCL_ARRAY	241
3.18.3	ANALYZE_DATABASE, ANALYZE_SCHEMA and ANALYZE_PART_OBJECT	242
3.18.4	CANONICALIZE	244
3.18.5	COMMA_TO_TABLE	246
3.18.6	DB_VERSION	247
3.18.7	EXEC_DDL_STATEMENT	248
3.18.8	FORMAT_CALL_STACK	249
3.18.9	GET_CPU_TIME	250
3.18.10	GET_DEPENDENCY	251
3.18.11	GET_HASH_VALUE	252
3.18.12	GET_PARAMETER_VALUE	253
3.18.13	GET_TIME	254
3.18.14	NAME_TOKENIZE	255
3.18.15	TABLE_TO_COMMA	258
3.19	UTL_ENCODE	260
3.19.1	BASE64_DECODE	261
3.19.2	BASE64_ENCODE	262
3.19.3	MIMEHEADER_DECODE	263
3.19.4	MIMEHEADER_ENCODE	264
3.19.5	QUOTED_PRINTABLE_DECODE	265
3.19.6	QUOTED_PRINTABLE_ENCODE	266
3.19.7	TEXT_DECODE	267
3.19.8	TEXT_ENCODE	268
3.19.9	UUDECODE	269
3.19.10	UUENCODE	270

3.20	UTL_FILE	272
3.20.1	Setting File Permissions with utl_file.umask	273
3.20.2	FCLOSE	274
3.20.3	FCLOSE_ALL	274
3.20.4	FCOPY	274
3.20.5	FFLUSH	276
3.20.6	FOPEN	277
3.20.7	FREMOVE	277
3.20.8	FRENAME	278
3.20.9	GET_LINE	280
3.20.10	IS_OPEN	281
3.20.11	NEW_LINE	281
3.20.12	PUT	282
3.20.13	PUT_LINE	284
3.20.14	PUTF	285
3.21	UTL_HTTP	287
3.21.1	HTML_PIECES	291
3.21.2	REQ	291
3.21.3	RESP	291
3.21.4	BEGIN_REQUEST	291
3.21.5	END_REQUEST	292
3.21.6	END_RESPONSE	292
3.21.7	GET_BODY_CHARSET	292
3.21.8	GET_FOLLOW_REDIRECT	293
3.21.9	GET_HEADER	293
3.21.10	GET_HEADER_BY_NAME	295
3.21.11	GET_HEADER_COUNT	295
3.21.12	GET_RESPONSE	296
3.21.13	GET_RESPONSE_ERROR_CHECK	296
3.21.14	GET_TRANSFER_TIMEOUT	296
3.21.15	READ_LINE	297
3.21.16	READ_RAW	299
3.21.17	READ_TEXT	299
3.21.18	REQUEST	300
3.21.19	REQUEST_PIECES	301
3.21.20	SET_BODY_CHARSET	301
3.21.21	SET_FOLLOW_REDIRECT	301
3.21.22	SET_HEADER	302
3.21.23	SET_RESPONSE_ERROR_CHECK	302
3.21.24	SET_TRANSFER_TIMEOUT	302
3.21.25	WRITE_LINE	303
3.21.26	WRITE_RAW	304
3.21.27	WRITE_TEXT	305
3.22	UTL_MAIL	306
3.22.1	SEND	307
3.22.2	SEND_ATTACH_RAW	308
3.22.3	SEND_ATTACH_VARCHAR2	310
3.23	UTL_RAW	311
3.23.1	CAST_TO_RAW	311
3.23.2	CAST_TO_VARCHAR2	312
3.23.3	CONCAT	313
3.23.4	CONVERT	314
3.23.5	LENGTH	315
3.23.6	SUBSTR	316

3.24	UTL_SMTP	317
3.24.1	CONNECTION	318
3.24.2	REPLY/REPLIES	318
3.24.3	CLOSE_DATA	318
3.24.4	COMMAND	318
3.24.5	COMMAND_REPLIES	319
3.24.6	DATA	319
3.24.7	EHLO	320
3.24.8	HELO	320
3.24.9	HELP	320
3.24.10	MAIL	321
3.24.11	NOOP	321
3.24.12	OPEN_CONNECTION	321
3.24.13	OPEN_DATA	322
3.24.14	QUIT	322
3.24.15	RCPT	322
3.24.16	RSET	322
3.24.17	VERFY	323
3.24.18	WRITE_DATA	323
3.24.19	Comprehensive Example	323
3.25	UTL_URL	325
3.25.1	ESCAPE	325
3.25.2	UNESCAPE	327
4	Acknowledgements	328
5	Conclusion	329
	Index	330

CHAPTER 1

Introduction

Database Compatibility for Oracle means that an application runs in an Oracle environment as well as in the EDB Postgres Advanced Server (Advanced Server) environment with minimal or no changes to the application code. This guide focuses solely on the features that are related to the package support provided by Advanced Server.

For more information about using other compatibility features offered by Advanced Server, please see the complete set of Advanced Server guides, available at:

<https://www.enterprisedb.com/edb-docs/>

1.1 What's New

The following database compatibility for Oracle features have been added to Advanced Server 11 to create Advanced Server 12:

- Advanced Server introduces `COMPOUND TRIGGERS`, which are stored as a PL block that executes in response to a specified triggering event. For information, see the *Database Compatibility for Oracle Developer's Guide*.
- Advanced Server now supports new `DATA DICTIONARY VIEWS` that provide information compatible with the Oracle data dictionary views. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added the `LISTAGG` function to support string aggregation that concatenates data from multiple rows into a single row in an ordered manner. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server now supports `CAST (MULTISET)` function, allowing subquery output to be `CAST` to a nested table type. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added the `MEDIAN` function to calculate a median value from the set of provided values. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added the `SYS_GUID` function to generate and return a globally unique identifier in the form of 16-bytes of RAW data. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server now supports an Oracle-compatible `SELECT UNIQUE` clause in addition to an existing `SELECT DISTINCT` clause. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has re-implemented `default_with_rowids` to create a table that includes a `ROWID` column in the newly created table. For information, see the *EDB Postgres Advanced Server Guide*.
- Advanced Server now supports logical decoding on the standby server, which allows creating a logical replication slot on a standby, independently of a primary server. For information, see the *EDB Postgres Advanced Server Guide*.
- Advanced Server introduces `INTERVAL PARTITIONING`, which allows a database to automatically create partitions of a specified interval as new data is inserted into a table. For information, see the *Database Compatibility for Oracle Developer's Guide*.

Note: Database Compatibility for Oracle Developer's Guide, Database Compatibility for Oracle Developer's Reference Guide, and EDB Postgres Advanced Server Guides are available at:

<https://www.enterprisedb.com/edb-docs/>

Packages

This chapter discusses the concept of packages in Advanced Server. A *package* is a named collection of functions, procedures, variables, cursors, user-defined record types, and records that are referenced using a common qualifier – the package identifier. Packages have the following characteristics:

- Packages provide a convenient means of organizing the functions and procedures that perform a related purpose. Permission to use the package functions and procedures is dependent upon one privilege granted to the entire package. All of the package programs must be referenced with a common name.
- Certain functions, procedures, variables, types, etc. in the package can be declared as *public*. Public entities are visible and can be referenced by other programs that are given `EXECUTE` privilege on the package. For public functions and procedures, only their signatures are visible - the program names, parameters if any, and return types of functions. The SPL code of these functions and procedures is not accessible to others, therefore applications that utilize a package are dependent only upon the information available in the signature – not in the procedural logic itself.
- Other functions, procedures, variables, types, etc. in the package can be declared as *private*. Private entities can be referenced and used by function and procedures within the package, but not by other external applications. Private entities are for use only by programs within the package.
- Function and procedure names can be overloaded within a package. One or more functions/procedures can be defined with the same name, but with different signatures. This provides the capability to create identically named programs that perform the same job, but on different types of input.

2.1 Package Components

Packages consist of two main components:

- The *package specification*: This is the public interface, (these are the elements which can be referenced outside the package). We declare all database objects that are to be a part of our package within the specification.
- The *package body*: This contains the actual implementation of all the database objects declared within the package specification.

The package body implements the specifications in the package specification. It contains implementation details and private declarations which are invisible to the application. You can debug, enhance or replace a package body without changing the specifications. Similarly, you can change the body without recompiling the calling programs because the implementation details are invisible to the application.

2.1.1 Package Specification Syntax

The package specification defines the user interface for a package (the API). The specification lists the functions, procedures, types, exceptions and cursors that are visible to a user of the package.

The syntax used to define the interface for a package is:

```
CREATE [ OR REPLACE ] PACKAGE <package_name>
[ <authorization_clause> ]
{ IS | AS }
[ <declaration>; ] ...
[ <procedure_or_function_declaration> ] ...
END [ <package_name> ] ;
```

Where <authorization_clause> :=

```
{ AUTHID DEFINER } | { AUTHID CURRENT_USER }
```

Where <procedure_or_function_declaration> :=

```
<procedure_declaration> | <function_declaration>
```

Where <procedure_declaration> :=

```
PROCEDURE <proc_name> [ <argument_list> ]; [ <restriction_pragma>; ]
```

Where <function_declaration> :=

```
FUNCTION <func_name> [ <argument_list> ]
RETURN <rettype> [ DETERMINISTIC ]; [ <restriction_pragma>; ]
```

Where <argument_list> :=

```
( <argument_declaration> [, ...] )
```

Where <argument_declaration> :=

```
<argname> [ IN | IN OUT | OUT ] <argtype> [ DEFAULT <value> ]
```

Where <restriction_pragma> :=

```
PRAGMA RESTRICT_REFERENCES ( <name>, <restrictions> )
```

Where <restrictions> :=

```
<restriction> [, ... ]
```

Parameters

<package_name>

<package_name> is an identifier assigned to the package - each package must have a name unique within the schema.

AUTHID DEFINER

If you omit the AUTHID clause or specify AUTHID DEFINER, the privileges of the package owner are used to determine access privileges to database objects.

AUTHID CURRENT_USER

If you specify AUTHID CURRENT_USER, the privileges of the current user executing a program in the package are used to determine access privileges.

<declaration>

<declaration> is an identifier of a public variable. A public variable can be accessed from outside of the package using the syntax <package_name.variable>. There can be zero, one, or more public variables. Public variable definitions must come before procedure or function declarations.

<declaration> can be any of the following:

- Variable Declaration
- Record Declaration
- Collection Declaration
- REF CURSOR and Cursor Variable Declaration
- TYPE Definitions for Records, Collections, and REF CURSORS
- Exception
- Object Variable Declaration

<proc_name>

The name of a public procedure.

<argname>

The name of an argument. The argument is referenced by this name within the function or procedure body.

IN | IN OUT | OUT

The argument mode. IN declares the argument for input only. This is the default. IN OUT allows the argument to receive a value as well as return a value. OUT specifies the argument is for output only.

<argtype>

The data type(s) of an argument. An argument type may be a base data type, a copy of the type of an existing column using %TYPE, or a user-defined type such as a nested table or an object type. A length must not be specified for any base type - for example, specify VARCHAR2, not VARCHAR2 (10).

The type of a column is referenced by writing <tablename.columnname> %TYPE; using this can sometimes help make a procedure independent from changes to the definition of a table.

DEFAULT <value>

The DEFAULT clause supplies a default value for an input argument if one is not supplied in the invocation. DEFAULT may not be specified for arguments with modes IN OUT or OUT.

<func_name>

The name of a public function.

<rettype>

The return data type.

DETERMINISTIC

DETERMINISTIC is a synonym for IMMUTABLE. A DETERMINISTIC function cannot modify the database and always reaches the same result when given the same argument values; it does not do database lookups or otherwise use information not directly present in its argument list. If you include this clause, any call of the function with all-constant arguments can be immediately replaced with the function value.

<restriction>

The following keywords are accepted for compatibility and ignored:

RNDS

RNPS

TRUST

WNDS

WNPS

2.1.2 Package Body Syntax

Package implementation details reside in the package body; the package body may contain objects that are not visible to the package user. Advanced Server supports the following syntax for the package body:

```
CREATE [ OR REPLACE ] PACKAGE BODY <package_name>
{ IS | AS }
[ <private_declaration>; ] ...
[ <procedure_or_function_definition> ] ...
[ <package_initializer> ]
END [ <package_name> ] ;
```

Where <procedure_or_function_definition> :=

<procedure_definition> | <function_definition>

Where <procedure_definition> :=

```
PROCEDURE <proc_name> [ <argument_list> ]
[ <options_list> ]
{ IS | AS }
<procedure_body>
END [ <proc_name> ] ;
```

Where <procedure_body> :=

```
[ PRAGMA AUTONOMOUS_TRANSACTION; ]
[ <declaration>; ][...]
BEGIN
    <statement>;[...]
    [ EXCEPTION
        {WHEN <exception> [OR <exception>] [...]] THEN
        <statement>; }
    [...]
]
```

Where <function_definition> :=

```
FUNCTION <func_name> [ <argument_list> ]
    RETURN <rettype> [ DETERMINISTIC ]
[ <options_list> ]
{ IS | AS }
<function_body>
END [ <func_name> ] ;
```

Where <function_body> :=

```
[ PRAGMA AUTONOMOUS_TRANSACTION; ]
[ <declaration>; ][...]
BEGIN
```

```

        <statement>; [...]
    [ EXCEPTION
        { WHEN <exception> [ OR <exception> ] [...] THEN <statement>; }
        [...]
    ]
Where <argument_list> :=
    ( <argument_declaration> [, ...] )
Where <argument_declaration> :=
    <argname> [ IN | IN OUT | OUT ] <argtype> [ DEFAULT <value> ]
Where <options_list> :=
    <option> [ ... ]
Where <option> :=
    STRICT
    LEAKPROOF
    PARALLEL { UNSAFE | RESTRICTED | SAFE }
    COST <execution_cost>
    ROWS <result_rows>
    SET <config_param> { TO <value> | = <value> | FROM CURRENT }
Where <package_initializer> :=
    BEGIN
        <statement>; [...]
    END;

```

Parameters

<package_name>

<package_name> is the name of the package for which this is the package body. There must be an existing package specification with this name.

<private_declaration>

<private_declaration> is an identifier of a private variable that can be accessed by any procedure or function within the package. There can be zero, one, or more private variables. <private_declaration> can be any of the following:

- Variable Declaration
- Record Declaration
- Collection Declaration
- REF CURSOR and Cursor Variable Declaration
- TYPE Definitions for Records, Collections, and REF CURSORS
- Exception
- Object Variable Declaration

<proc_name>

The name of the procedure being created.

PRAGMA AUTONOMOUS_TRANSACTION

PRAGMA AUTONOMOUS_TRANSACTION is the directive that sets the procedure as an autonomous transaction.

<declaration>

A variable, type, REF CURSOR, or subprogram declaration. If subprogram declarations are included, they must be declared after all other variable, type, and REF CURSOR declarations.

<statement>

An SPL program statement. Note that a DECLARE – BEGIN – END block is considered an SPL statement unto itself. Thus, the function body may contain nested blocks.

<exception>

An exception condition name such as NO_DATA_FOUND, OTHERS, etc.

<func_name>

The name of the function being created.

<rettype>

The return data type, which may be any of the types listed for <argtype>. As for <argtype>, a length must not be specified for <rettype>.

DETERMINISTIC

Include DETERMINISTIC to specify that the function will always return the same result when given the same argument values. A DETERMINISTIC function must not modify the database.

Note: The DETERMINISTIC keyword is equivalent to the PostgreSQL IMMUTABLE option.

Note: If DETERMINISTIC is specified for a public function in the package body, it must also be specified for the function declaration in the package specification. (For private functions, there is no function declaration in the package specification.)

PRAGMA AUTONOMOUS_TRANSACTION

PRAGMA AUTONOMOUS_TRANSACTION is the directive that sets the function as an autonomous transaction.

<declaration>

A variable, type, REF CURSOR, or subprogram declaration. If subprogram declarations are included, they must be declared after all other variable, type, and REF CURSOR declarations.

<argname>

The name of a formal argument. The argument is referenced by this name within the procedure body.

IN | IN OUT | OUT

The argument mode. IN declares the argument for input only. This is the default. IN OUT allows the argument to receive a value as well as return a value. OUT specifies the argument is for output only.

<argtype>

The data type(s) of an argument. An argument type may be a base data type, a copy of the type of an existing column using %TYPE, or a user-defined type such as a nested table or an object type. A length must not be specified for any base type - for example, specify VARCHAR2, not VARCHAR2 (10).

The type of a column is referenced by writing <tablename> . <columnname>`%TYPE; using this can sometimes help make a procedure independent from changes to the definition of a table.

DEFAULT <value>

The DEFAULT clause supplies a default value for an input argument if one is not supplied in the procedure call. DEFAULT may not be specified for arguments with modes IN OUT or OUT.

Please note: the following options are not compatible with Oracle databases; they are extensions to Oracle package syntax provided by Advanced Server only.

STRICT

The STRICT keyword specifies that the function will not be executed if called with a NULL argument; instead the function will return NULL.

LEAKPROOF

The LEAKPROOF keyword specifies that the function will not reveal any information about arguments, other than through a return value.

PARALLEL { UNSAFE | RESTRICTED | SAFE }

The PARALLEL clause enables the use of parallel sequential scans (parallel mode). A parallel sequential scan uses multiple workers to scan a relation in parallel during a query in contrast to a serial sequential scan.

When set to UNSAFE, the procedure or function cannot be executed in parallel mode. The presence of such a procedure or function forces a serial execution plan. This is the default setting if the PARALLEL clause is omitted.

When set to RESTRICTED, the procedure or function can be executed in parallel mode, but the execution is restricted to the parallel group leader. If the qualification for any particular relation has anything that is parallel restricted, that relation won't be chosen for parallelism.

When set to SAFE, the procedure or function can be executed in parallel mode with no restriction.

<execution_cost>

<execution_cost> specifies a positive number giving the estimated execution cost for the function, in units of cpu_operator_cost. If the function returns a set, this is the cost per returned row. The default is 0.0025.

`<result_rows>

<result_rows> is the estimated number of rows that the query planner should expect the function to return. The default is 1000.

SET

Use the SET clause to specify a parameter value for the duration of the function:

<config_param> specifies the parameter name.

<value> specifies the parameter value.

FROM CURRENT guarantees that the parameter value is restored when the function ends.

<package_initializer>

The statements in the <package_initializer> are executed once per user's session when the package is first referenced.

Note: The STRICT, LEAKPROOF, PARALLEL, COST, ROWS and SET keywords provide extended functionality for Advanced Server and are not supported by Oracle.

2.2 Creating Packages

A package is not an executable piece of code; rather it is a repository of code. When you use a package, you actually execute or make reference to an element within a package.

2.2.1 Creating the Package Specification

The package specification contains the definition of all the elements in the package that can be referenced from outside of the package. These are called the public elements of the package, and they act as the package interface. The following code sample is a package specification:

```
--
-- Package specification for the 'emp_admin' package.
--
CREATE OR REPLACE PACKAGE emp_admin
IS
    FUNCTION get_dept_name (
        p_deptno NUMBER DEFAULT 10
    )
    RETURN VARCHAR2;
    FUNCTION update_emp_sal (
        p_empno NUMBER,
        p_raise NUMBER
    )
    RETURN NUMBER;
    PROCEDURE hire_emp (
        p_empno          NUMBER,
        p_ename           VARCHAR2,
        p_job             VARCHAR2,
        p_sal             NUMBER,
        p_hiredate        DATE          DEFAULT sysdate,
        p_comm            NUMBER        DEFAULT 0,
        p_mgr             NUMBER,
        p_deptno          NUMBER        DEFAULT 10
    );
    PROCEDURE fire_emp (
        p_empno NUMBER
    );
END emp_admin;
```

This code sample creates the `emp_admin` package specification. This package specification consists of two functions and two stored procedures. We can also add the `OR REPLACE` clause to the `CREATE PACKAGE` statement for convenience.

2.2.2 Creating the Package Body

The body of the package contains the actual implementation behind the package specification. For the above emp_admin package specification, we shall now create a package body which will implement the specifications. The body will contain the implementation of the functions and stored procedures in the specification.

```
--
-- Package body for the 'emp_admin' package.
--
CREATE OR REPLACE PACKAGE BODY emp_admin
IS
--
    -- Function that queries the 'dept' table based on the department
    -- number and returns the corresponding department name.
    --
    FUNCTION get_dept_name (
        p_deptno IN NUMBER DEFAULT 10
    )
    RETURN VARCHAR2
    IS
        v_dname VARCHAR2(14);
    BEGIN
        SELECT dname INTO v_dname FROM dept WHERE deptno = p_deptno;
        RETURN v_dname;
    EXCEPTION
        WHEN NO_DATA_FOUND THEN
            DBMS_OUTPUT.PUT_LINE('Invalid department number ' || p_deptno);
            RETURN '';
    END;
--
    -- Function that updates an employee's salary based on the
    -- employee number and salary increment/decrement passed
    -- as IN parameters. Upon successful completion the function
    -- returns the new updated salary.
    --
    FUNCTION update_emp_sal (
        p_empno      IN NUMBER,
        p_raise       IN NUMBER
    )
    RETURN NUMBER
    IS
        v_sal        NUMBER := 0;
    BEGIN
        SELECT sal INTO v_sal FROM emp WHERE empno = p_empno;
        v_sal := v_sal + p_raise;
        UPDATE emp SET sal = v_sal WHERE empno = p_empno;
        RETURN v_sal;
    EXCEPTION
        WHEN NO_DATA_FOUND THEN
            DBMS_OUTPUT.PUT_LINE('Employee ' || p_empno || ' not found');
            RETURN -1;
        WHEN OTHERS THEN
            DBMS_OUTPUT.PUT_LINE('The following is SQLERRM:');
            DBMS_OUTPUT.PUT_LINE(SQLERRM);
            DBMS_OUTPUT.PUT_LINE('The following is SQLCODE:');
            DBMS_OUTPUT.PUT_LINE(SQLCODE);
            RETURN -1;
    END;
```

```
--
-- Procedure that inserts a new employee record into the 'emp' table.
--
PROCEDURE hire_emp (
  p_empno      NUMBER,
  p_ename      VARCHAR2,
  p_job        VARCHAR2,
  p_sal        NUMBER,
  p_hiredate    DATE   DEFAULT sysdate,
  p_comm        NUMBER DEFAULT 0,
  p_mgr        NUMBER,
  p_deptno     NUMBER DEFAULT 10
)
AS
BEGIN
  INSERT INTO emp(empno, ename, job, sal, hiredate, comm, mgr, deptno)
    VALUES(p_empno, p_ename, p_job, p_sal,
           p_hiredate, p_comm, p_mgr, p_deptno);
END;
--
-- Procedure that deletes an employee record from the 'emp' table based
-- on the employee number.
--
PROCEDURE fire_emp (
  p_empno NUMBER
)
AS
BEGIN
  DELETE FROM emp WHERE empno = p_empno;
END;
END;
```

2.3 Referencing a Package

To reference the types, items and subprograms that are declared within a package specification, we use the dot notation. For example:

`<package_name>.<type_name>`

`<package_name>.<item_name>`

`<package_name>.<subprogram_name>`

To invoke a function from the `emp_admin` package specification, we will execute the following SQL command.

```
SELECT emp_admin.get_dept_name(10) FROM DUAL;
```

Here we are invoking the `get_dept_name` function declared within the package `emp_admin`. We are passing the department number as an argument to the function, which will return the name of the department. Here the value returned should be `ACCOUNTING`, which corresponds to department number 10.

2.4 Using Packages With User Defined Types

The following example incorporates the various user-defined types discussed in earlier chapters within the context of a package.

The package specification of emp_rpt shows the declaration of a record type, emprec_typ, and a weakly-typed REF CURSOR, emp_refcur, as publicly accessible along with two functions and two procedures. Function, open_emp_by_dept, returns the REF CURSOR type, EMP_REFCUR. Procedures, fetch_emp and close_refcur, both declare a weakly-typed REF CURSOR as a formal parameter.

```
CREATE OR REPLACE PACKAGE emp_rpt
IS
    TYPE emprec_typ IS RECORD (
        empno      NUMBER(4),
        ename      VARCHAR(10)
    );
    TYPE emp_refcur IS REF CURSOR;

    FUNCTION get_dept_name (
        p_deptno    IN NUMBER
    ) RETURN VARCHAR2;
    FUNCTION open_emp_by_dept (
        p_deptno    IN emp.deptno%TYPE
    ) RETURN EMP_REFCUR;
    PROCEDURE fetch_emp (
        p_refcur    IN OUT SYS_REFCURSOR
    );
    PROCEDURE close_refcur (
        p_refcur    IN OUT SYS_REFCURSOR
    );
END emp_rpt;
```

The package body shows the declaration of several private variables - a static cursor, dept_cur, a table type, depttab_typ, a table variable, t_dept, an integer variable, t_dept_max, and a record variable, r_emp.

```
CREATE OR REPLACE PACKAGE BODY emp_rpt
IS
    CURSOR dept_cur IS SELECT * FROM dept;
    TYPE depttab_typ IS TABLE of dept%ROWTYPE
        INDEX BY BINARY_INTEGER;
    t_dept          DEPTTAB_TYP;
    t_dept_max      INTEGER := 1;
    r_emp           EMPREC_TYP;

    FUNCTION get_dept_name (
        p_deptno    IN NUMBER
    ) RETURN VARCHAR2
    IS
    BEGIN
        FOR i IN 1..t_dept_max LOOP
            IF p_deptno = t_dept(i).deptno THEN
                RETURN t_dept(i).dname;
            END IF;
        END LOOP;
        RETURN 'Unknown';
    END;
    FUNCTION open_emp_by_dept (
        p_deptno    IN emp.deptno%TYPE
```

```

) RETURN EMP_REFCUR
IS
    emp_by_dept EMP_REFCUR;
BEGIN
    OPEN emp_by_dept FOR SELECT empno, ename FROM emp
        WHERE deptno = p_deptno;
    RETURN emp_by_dept;
END;

PROCEDURE fetch_emp (
    p_refcur      IN OUT SYS_REFCURSOR
)
IS
BEGIN
    DBMS_OUTPUT.PUT_LINE('EMPNO ENAME');
    DBMS_OUTPUT.PUT_LINE('-----');
    LOOP
        FETCH p_refcur INTO r_emp;
        EXIT WHEN p_refcur%NOTFOUND;
        DBMS_OUTPUT.PUT_LINE(r_emp.empno || ' ' || r_emp.ename);
    END LOOP;
END;

PROCEDURE close_refcur (
    p_refcur      IN OUT SYS_REFCURSOR
)
IS
BEGIN
    CLOSE p_refcur;
END;
BEGIN
    OPEN dept_cur;
    LOOP
        FETCH dept_cur INTO t_dept(t_dept_max);
        EXIT WHEN dept_cur%NOTFOUND;
        t_dept_max := t_dept_max + 1;
    END LOOP;
    CLOSE dept_cur;
    t_dept_max := t_dept_max - 1;
END emp_rpt;

```

This package contains an initialization section that loads the private table variable, `t_dept`, using the private static cursor, `dept_cur`. `t_dept` serves as a department name lookup table in function, `get_dept_name`.

Function, `open_emp_by_dept` returns a `REF CURSOR` variable for a result set of employee numbers and names for a given department. This `REF CURSOR` variable can then be passed to procedure, `fetch_emp`, to retrieve and list the individual rows of the result set. Finally, procedure, `close_refcur`, can be used to close the `REF CURSOR` variable associated with this result set.

The following anonymous block runs the package function and procedures. In the anonymous block's declaration section, note the declaration of cursor variable, `v_emp_cur`, using the package's public `REF CURSOR` type, `EMP_REFCUR`. `v_emp_cur` contains the pointer to the result set that is passed between the package function and procedures.


```

DECLARE
    v_deptno dept.deptno%TYPE DEFAULT 30;
    v_emp_cur emp_rpt.EMP_REFCUR;
BEGIN
    v_emp_cur := emp_rpt.open_emp_by_dept(v_deptno);
    DBMS_OUTPUT.PUT_LINE('EMPLOYEES IN DEPT #' || v_deptno ||
        ': ' || emp_rpt.get_dept_name(v_deptno));
    emp_rpt.fetch_emp(v_emp_cur);
    DBMS_OUTPUT.PUT_LINE('*****');
    DBMS_OUTPUT.PUT_LINE(v_emp_cur%ROWCOUNT || ' rows were
    retrieved');
    emp_rpt.close_refcur(v_emp_cur);
END;

```

The following is the result of this anonymous block.

```

EMPLOYEES IN DEPT #30: SALES
EMPNO ENAME
-----
7499 ALLEN
7521 WARD
7654 MARTIN
7698 BLAKE
7844 TURNER
7900 JAMES
*****
6 rows were retrieved

```

The following anonymous block illustrates another means of achieving the same result. Instead of using the package procedures, `fetch_emp` and `close_refcur`, the logic of these programs is coded directly into the anonymous block. In the anonymous block's declaration section, note the addition of record variable, `r_emp`, declared using the package's public record type, `EMPREC_TYP`.

```

DECLARE
    v_deptno      dept.deptno%TYPE DEFAULT 30;
    v_emp_cur      emp_rpt.EMP_REFCUR;
    r_emp          emp_rpt.EMPREC_TYP;
BEGIN
    v_emp_cur := emp_rpt.open_emp_by_dept(v_deptno);
    DBMS_OUTPUT.PUT_LINE('EMPLOYEES IN DEPT #' || v_deptno ||
        ': ' || emp_rpt.get_dept_name(v_deptno));
    DBMS_OUTPUT.PUT_LINE('EMPNO ENAME');
    DBMS_OUTPUT.PUT_LINE('-----');
    LOOP
        FETCH v_emp_cur INTO r_emp;
        EXIT WHEN v_emp_cur%NOTFOUND;
        DBMS_OUTPUT.PUT_LINE(r_emp.empno || ' ' ||
            r_emp.ename);
    END LOOP;
    DBMS_OUTPUT.PUT_LINE('*****');
    DBMS_OUTPUT.PUT_LINE(v_emp_cur%ROWCOUNT || ' rows were retrieved');
    CLOSE v_emp_cur;
END;

```

The following is the result of this anonymous block.

```
EMPLOYEES IN DEPT #30: SALES
EMPNO  ENAME
-----
7499   ALLEN
7521   WARD
7654   MARTIN
7698   BLAKE
7844   TURNER
7900   JAMES
*****
6 rows were retrieved
```

2.5 Dropping a Package

The syntax for deleting an entire package or just the package body is as follows:

```
DROP PACKAGE [ BODY ] <package_name>;
```

If the keyword, `BODY`, is omitted, both the package specification and the package body are deleted - i.e., the entire package is dropped. If the keyword, `BODY`, is specified, then only the package body is dropped. The package specification remains intact. `<package_name>` is the identifier of the package to be dropped.

Following statement will destroy only the package body of `<emp_admin>`:

```
DROP PACKAGE BODY emp_admin;
```

The following statement will drop the entire `<emp_admin>` package:

```
DROP PACKAGE emp_admin;
```

CHAPTER 3

Built-In Packages

This chapter describes the built-in packages that are provided with Advanced Server. For certain packages, non-superusers must be explicitly granted the `EXECUTE` privilege on the package before using any of the package's functions or procedures. For most of the built-in packages, `EXECUTE` privilege has been granted to `PUBLIC` by default.

For information about using the `GRANT` command to provide access to a package, please see the *Database Compatibility for Oracle Developers Reference Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

All built-in packages are owned by the special `sys` user which must be specified when granting or revoking privileges on built-in packages:

```
GRANT EXECUTE ON PACKAGE SYS.UTL_FILE TO john;
```

3.1 DBMS_ALERT

The DBMS_ALERT package provides the capability to register for, send, and receive alerts. The following table lists the supported procedures:

Function/Procedure	Return Type	Description
REGISTER (<name>)	n/a	Register to be able to receive alerts named, <name>
REMOVE (<name>)	n/a	Remove registration for the alert named, <name>
REMOVEALL	n/a	Remove registration for all alerts.
SIGNAL (<name>, <message>)	n/a	Signals the alert named, <name>, with <message>
WAITANY (<name> OUT, <message> OUT, <status> OUT, <timeout>)	n/a	Wait for any registered alert to occur.
WAITONE (<name>, <message> OUT, <status> OUT, <timeout>)	n/a	Wait for the specified alert, <name>, to occur.

Advanced Server's implementation of DBMS_ALERT is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Advanced Server allows a maximum of 500 concurrent alerts. You can use the `dbms_alert.max_alerts` GUC variable (located in the `postgresql.conf` file) to specify the maximum number of concurrent alerts allowed on a system.

To set a value for the `dbms_alert.max_alerts` variable, open the `postgresql.conf` file (located by default in `/opt/PostgresPlus/10AS/data`) with your choice of editor, and edit the `dbms_alert.max_alerts` parameter as shown:

```
dbms_alert.max_alerts = <alert_count>
<alert_count>
```

`alert_count` specifies the maximum number of concurrent alerts. By default, the value of `dbms_alert.max_alerts` is 100. To disable this feature, set `dbms_alert.max_alerts` to 0.

For the `dbms_alert.max_alerts` GUC to function correctly, the `custom_variable_classes` parameter must contain `dbms_alerts`:

```
custom_variable_classes = 'dbms_alert, ...'
```

After editing the `postgresql.conf` file parameters, you must restart the server for the changes to take effect.

3.1.1 REGISTER

The REGISTER procedure enables the current session to be notified of the specified alert.

```
REGISTER (<name> VARCHAR2)
```

Parameters

<name>

Name of the alert to be registered.

Examples

The following anonymous block registers for an alert named, `alert_test`, then waits for the signal.

```
DECLARE
  v_name          VARCHAR2(30) := 'alert_test';
  v_msg           VARCHAR2(80);
  v_status        INTEGER;
  v_timeout       NUMBER(3) := 120;
BEGIN
  DBMS_ALERT.REGISTER(v_name);
  DBMS_OUTPUT.PUT_LINE('Registered for alert ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
  DBMS_ALERT.WAITONE(v_name,v_msg,v_status,v_timeout);
  DBMS_OUTPUT.PUT_LINE('Alert name      : ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Alert msg       : ' || v_msg);
  DBMS_OUTPUT.PUT_LINE('Alert status    : ' || v_status);
  DBMS_OUTPUT.PUT_LINE('Alert timeout:  ' || v_timeout || ' seconds');
  DBMS_ALERT.REMOVE(v_name);
END;
```

Registered for alert alert_test
Waiting for signal...

3.1.2 REMOVE

The `REMOVE` procedure unregisters the session for the named alert.

```
REMOVE (<name> VARCHAR2)
```

Parameters

<name>

Name of the alert to be unregistered.

3.1.3 REMOVEALL

The `REMOVEALL` procedure unregisters the session for all alerts.

```
REMOVEALL
```

3.1.4 SIGNAL

The `SIGNAL` procedure signals the occurrence of the named alert.

```
SIGNAL (<name> VARCHAR2, <message> VARCHAR2)
```

Parameters

<name>

Name of the alert.

<message>

Information to pass with this alert.

Examples

The following anonymous block signals an alert for alert_test.

```
DECLARE
    v_name    VARCHAR2(30) := 'alert_test';
BEGIN
    DBMS_ALERT.SIGNAL(v_name,'This is the message from ' || v_name);
    DBMS_OUTPUT.PUT_LINE('Issued alert for ' || v_name);
END;
Issued alert for alert_test
```

3.1.5 WAITANY

The WAITANY procedure waits for any of the registered alerts to occur.

```
WAITANY(<name> OUT VARCHAR2, <message> OUT VARCHAR2,
        <status> OUT INTEGER, <timeout> NUMBER)
```

Parameters

<name>

Variable receiving the name of the alert.

<message>

Variable receiving the message sent by the SIGNAL procedure.

<status>

Status code returned by the operation. Possible values are: 0 – alert occurred; 1 – timeout occurred.

<timeout>

Time to wait for an alert in seconds.

Examples

The following anonymous block uses the WAITANY procedure to receive an alert named, alert_test or any_alert:

```
DECLARE
    v_name          VARCHAR2(30);
    v_msg           VARCHAR2(80);
    v_status        INTEGER;
    v_timeout       NUMBER(3) := 120;
BEGIN
    DBMS_ALERT.REGISTER('alert_test');
    DBMS_ALERT.REGISTER('any_alert');
    DBMS_OUTPUT.PUT_LINE('Registered for alert alert_test and any_alert');
    DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
    DBMS_ALERT.WAITANY(v_name,v_msg,v_status,v_timeout);
    DBMS_OUTPUT.PUT_LINE('Alert name   : ' || v_name);
    DBMS_OUTPUT.PUT_LINE('Alert msg    : ' || v_msg);
    DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
    DBMS_OUTPUT.PUT_LINE('Alert timeout: ' || v_timeout || ' seconds');
    DBMS_ALERT.REMOVEALL;
END;
```

```
Registered for alert alert_test and any_alert
Waiting for signal...
```

An anonymous block in a second session issues a signal for any_alert:

```
DECLARE
    v_name  VARCHAR2(30) := 'any_alert';
BEGIN
    DBMS_ALERT.SIGNAL(v_name, 'This is the message from ' || v_name);
    DBMS_OUTPUT.PUT_LINE('Issued alert for ' || v_name);
END;

Issued alert for any_alert
```

Control returns to the first anonymous block and the remainder of the code is executed:

```
Registered for alert alert_test and any_alert
Waiting for signal...
Alert name      : any_alert
Alert msg       : This is the message from any_alert
Alert status    : 0
Alert timeout   : 120 seconds
```

3.1.6 WAITONE

The WAITONE procedure waits for the specified registered alert to occur.

```
WAITONE (<name> VARCHAR2, <message> OUT VARCHAR2,
        <status> OUT INTEGER, <timeout> NUMBER)
```

Parameters

<name>

Name of the alert.

<message>

Variable receiving the message sent by the SIGNAL procedure.

<status>

Status code returned by the operation. Possible values are: 0 – alert occurred; 1 – timeout occurred.

<timeout>

Time to wait for an alert in seconds.

Examples

The following anonymous block is similar to the one used in the WAITANY example except the WAITONE procedure is used to receive the alert named, alert_test.

```
DECLARE
    v_name          VARCHAR2(30) := 'alert_test';
    v_msg           VARCHAR2(80);
    v_status        INTEGER;
    v_timeout       NUMBER(3) := 120;
BEGIN
    DBMS_ALERT.REGISTER(v_name);
    DBMS_OUTPUT.PUT_  DBMS_ALERT.REGISTER(v_name);
    DBMS_OUTPUT.PUT_LINE('Registered for alert ' || v_name);
    DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
    DBMS_ALERT.WAITONE(v_name,v_msg,v_status,v_timeout);
    DBMS_OUTPUT.PUT_LINE('Alert name   : ' || v_name);
    DBMS_OUTPUT.PUT_LINE('Alert msg    : ' || v_msg);
    DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
    DBMS_OUTPUT.PUT_LINE('Alert timeout: ' || v_timeout || ' seconds');
    DBMS_ALERT.REMOVE(v_name);LINE('Registered for alert ' || v_name);
    DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
    DBMS_ALERT.WAITONE(v_name,v_msg,v_status,v_timeout);
    DBMS_OUTPUT.PUT_LINE('Alert name   : ' || v_name);
    DBMS_OUTPUT.PUT_LINE('Alert msg    : ' || v_msg);
    DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
    DBMS_OUTPUT.PUT_LINE('Alert timeout: ' || v_timeout || ' seconds');
    DBMS_ALERT.REMOVE(v_name);
END;

Registered for alert alert_test
Waiting for signal...
```

Signal sent for alert_test sent by an anonymous block in a second session:

```
DECLARE
    v_name          VARCHAR2(30) := 'alert_test';
BEGIN
    DBMS_ALERT.SIGNAL(v_name,'This is the message from ' || v_name);
    DBMS_OUTPUT.PUT_LINE('Issued alert for ' || v_name);
END;

Issued alert for alert_test
```

First session is alerted, control returns to the anonymous block, and the remainder of the code is executed:

```
Registered for alert alert_test
Waiting for signal...
Alert name   : alert_test
Alert msg    : This is the message from alert_test
Alert status : 0
Alert timeout: 120 seconds
```

3.1.7 Comprehensive Example

The following example uses two triggers to send alerts when the dept table or the emp table is changed. An anonymous block listens for these alerts and displays messages when an alert is received.

The following are the triggers on the dept and emp tables:

```
CREATE OR REPLACE TRIGGER dept_alert_trig
  AFTER INSERT OR UPDATE OR DELETE ON dept
DECLARE
  v_action          VARCHAR2(25);
BEGIN
  IF INSERTING THEN
    v_action := ' added department(s) ';
  ELSIF UPDATING THEN
    v_action := ' updated department(s) ';
  ELSIF DELETING THEN
    v_action := ' deleted department(s) ';
  END IF;
  DBMS_ALERT.SIGNAL('dept_alert',USER || v_action || 'on ' ||
    SYSDATE);
END;

CREATE OR REPLACE TRIGGER emp_alert_trig
  AFTER INSERT OR UPDATE OR DELETE ON emp
DECLARE
  v_action          VARCHAR2(25);
BEGIN
  IF INSERTING THEN
    v_action := ' added employee(s) ';
  ELSIF UPDATING THEN
    v_action := ' updated employee(s) ';
  ELSIF DELETING THEN
    v_action := ' deleted employee(s) ';
  END IF;
  DBMS_ALERT.SIGNAL('emp_alert',USER || v_action || 'on ' ||
    SYSDATE);
END;
```

The following anonymous block is executed in a session while updates to the dept and emp tables occur in other sessions:

```
DECLARE
  v_dept_alert      VARCHAR2(30) := 'dept_alert';
  v_emp_alert       VARCHAR2(30) := 'emp_alert';
  v_name            VARCHAR2(30);
  v_msg             VARCHAR2(80);
  v_status          INTEGER;
  v_timeout         NUMBER(3) := 60;
BEGIN
  DBMS_ALERT.REGISTER(v_dept_alert);
  DBMS_ALERT.REGISTER(v_emp_alert);
  DBMS_OUTPUT.PUT_LINE('Registered for alerts dept_alert and emp_alert');
  DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
  LOOP
    DBMS_ALERT.WAITANY(v_name,v_msg,v_status,v_timeout);
    EXIT WHEN v_status != 0;
    DBMS_OUTPUT.PUT_LINE('Alert name      : ' || v_name);
    DBMS_OUTPUT.PUT_LINE('Alert msg       : ' || v_msg);
```

```

        DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
        DBMS_OUTPUT.PUT_LINE('-----' ||
        '-----');
    END LOOP;
    DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
    DBMS_ALERT.REMOVEALL;
END;
```

Registered for alerts dept_alert and emp_alert Waiting for signal...

The following changes are made by user, mary:

```

INSERT INTO dept VALUES (50,'FINANCE','CHICAGO');
INSERT INTO emp (empno,ename,deptno) VALUES (9001,'JONES',50);
INSERT INTO emp (empno,ename,deptno) VALUES (9002,'ALICE',50);
```

The following change is made by user, john:

```

INSERT INTO dept VALUES (60,'HR','LOS ANGELES');
```

The following is the output displayed by the anonymous block receiving the signals from the triggers:

```

Registered for alerts dept_alert and emp_alert
Waiting for signal...
Alert name   : dept_alert
Alert msg    : mary added department(s) on 25-OCT-07 16:41:01
Alert status : 0
-----
Alert name   : emp_alert
Alert msg    : mary added employee(s) on 25-OCT-07 16:41:02
Alert status : 0
-----
Alert name   : dept_alert
Alert msg    : john added department(s) on 25-OCT-07 16:41:22
Alert status : 0
-----
Alert status : 1
```

3.2 DBMS_AQ

EDB Postgres Advanced Server Advanced Queueing provides message queueing and message processing for the Advanced Server database. User-defined messages are stored in a queue; a collection of queues is stored in a queue table. Procedures in the DBMS_AQADM package create and manage message queues and queue tables. Use the DBMS_AQ package to add messages to a queue or remove messages from a queue, or register or unregister a PL/SQL callback procedure.

Advanced Server also provides extended (non-compatible) functionality for the DBMS_AQ package with SQL commands. Please see the *Database Compability for Oracle Developers Reference Guide* for detailed information about the following SQL commands:

- ALTER QUEUE
- ALTER QUEUE TABLE
- CREATE QUEUE
- CREATE QUEUE TABLE
- DROP QUEUE
- DROP QUEUE TABLE

The DBMS_AQ package provides procedures that allow you to enqueue a message, dequeue a message, and manage callback procedures. The supported procedures are:

Function/Procedure	Return Type	Description
ENQUEUE	n/a	Post a message to a queue.
DEQUEUE	n/a	Retrieve a message from a queue if or when a message is available.
REGISTER	n/a	Register a callback procedure.
UNREGISTER	n/a	Unregister a callback procedure.

Advanced Server's implementation of DBMS_AQ is a partial implementation when compared to Oracle's version. Only those procedures listed in the table above are supported.

Advanced Server supports use of the constants listed below:

Constant	Description	For Parameters
DBMS_AQ.BROWSE (0)	Read the message without locking.	dequeue_options_t.dequeue_mode
DBMS_AQ.LOCKED (1)	This constant is defined, but will return an error if used.	dequeue_options_t.dequeue_mode
DBMS_AQ.REMOVE (2)	Delete the message after reading; the default.	dequeue_options_t.dequeue_mode
DBMS_AQ.REMOVE_NODATA (3)	This constant is defined, but will return an error if used.	dequeue_options_t.dequeue_mode
DBMS_AQ.FIRST_MESSAGE (0)	Return the first available message that matches the search terms.	dequeue_options_t.navigation
DBMS_AQ.NEXT_MESSAGE (1)	Return the next available message that matches the search terms.	dequeue_options_t.navigation

Continued on next page

Table 3.1 – continued from previous page

DBMS_AQ.NEXT_TRANSACTION (2)	This constant is defined, but will return an error if used.	dequeue_options_t.navigation
DBMS_AQ.FOREVER (-1)	Wait forever if a message that matches the search term is not found, the default.	dequeue_options_t.wait
DBMS_AQ.NO_WAIT (0)	Do not wait if a message that matches the search term is not found.	dequeue_options_t.wait
DBMS_AQ.ON_COMMIT (0)	The dequeue is part of the current transaction.	enqueue_options_t.visibility, dequeue_options_t.visibility
DBMS_AQ.IMMEDIATE (1)	This constant is defined, but will return an error if used.	enqueue_options_t.visibility, dequeue_options_t.visibility
DBMS_AQ.PERSISTENT (0)	The message should be stored in a table.	enqueue_options_t.delivery_mode
DBMS_AQ.BUFFERED (1)	This constant is defined, but will return an error if used.	enqueue_options_t.delivery_mode
DBMS_AQ.READY (0)	Specifies that the message is ready to process.	message_properties_t.state
DBMS_AQ.WAITING (1)	Specifies that the message is waiting to be processed.	message_properties_t.state
DBMS_AQ.PROCESSED (2)	Specifies that the message has been processed.	message_properties_t.state
DBMS_AQ.EXPIRED (3)	Specifies that the message is in the exception queue.	message_properties_t.state
DBMS_AQ.NO_DELAY (0)	This constant is defined, but will return an error if used	message_properties_t.delay
DBMS_AQ.NEVER (NULL)	This constant is defined, but will return an error if used	message_properties_t.expiration
DBMS_AQ.NAMESPACE_AQ (0)	Accept notifications from DBMS_AQ queues.	sys.aq\$_reg_info.namespace
DBMS_AQ.NAMESPACE_ANONYMOUS (1)	This constant is defined, but will return an error if used	sys.aq\$_reg_info.namespace

The DBMS_AQ configuration parameters listed in the following table can be defined in the `postgresql.conf` file. After the configuration parameters are defined, you can invoke the DBMS_AQ package to use and manage messages held in queues and queue tables.

Parameter	Description
<code>dbms_aq.max_workers</code>	The maximum number of workers to run.
<code>dbms_aq.max_idle_time</code>	The idle time a worker must wait before exiting.
<code>dbms_aq.min_work_time</code>	The minimum time a worker can run before exiting.
<code>dbms_aq.launch_delay</code>	The minimum time between creating workers.
<code>dbms_aq.batch_size</code>	The maximum number of messages to process in a single transaction. The default batch size is 10.
<code>dbms_aq.max_databases</code>	The size of DBMS_AQ's hash table of databases. The default value is 1024.
<code>dbms_aq.max_pending_retries</code>	The size of DBMS_AQ's hash table of pending retries. The default value is 1024.

3.2.1 ENQUEUE

The ENQUEUE procedure adds an entry to a queue. The signature is:

```
ENQUEUE (
    <queue_name> IN VARCHAR2,
    <enqueue_options> IN DBMS_AQ.ENQUEUE_OPTIONS_T,
    <message_properties> IN DBMS_AQ.MESSAGE_PROPERTIES_T,
    <payload IN <type_name>,
    <msgid> OUT RAW)
```

Parameters

<queue_name>

The name (optionally schema-qualified) of an existing queue. If you omit the schema name, the server will use the schema specified in the SEARCH_PATH. Please note that unlike Oracle, unquoted identifiers are converted to lower case before storing. To include special characters or use a case-sensitive name, enclose the name in double quotes.

For detailed information about creating a queue, please see DBMS_AQADM.CREATE_QUEUE.

<enqueue_options>

<enqueue_options> is a value of the type, enqueue_options_t:

```
DBMS_AQ.ENQUEUE_OPTIONS_T IS RECORD (
    visibility BINARY_INTEGER DEFAULT ON_COMMIT,
    relative_msgid RAW(16) DEFAULT NULL,
    sequence_deviation BINARY_INTEGER DEFAULT NULL,
    transformation VARCHAR2(61) DEFAULT NULL,
    delivery_mode PLS_INTEGER NOT NULL DEFAULT PERSISTENT);
```

Currently, the only supported parameter values for enqueue_options_t are:

visibility	ON_COMMIT.
delivery_mode	PERSISTENT
sequence_deviation	NULL
transformation	NULL
relative_msgid	NULL

<message_properties>

<message_properties> is a value of the type, message_properties_t:

```
message_properties_t IS RECORD (
    priority INTEGER,
    delay INTEGER,
    expiration INTEGER,
    correlation CHARACTER VARYING(128) COLLATE pg_catalog."C",
    attempts INTEGER,
    recipient_list "AQ$RECIPIENT_LIST_T",
    exception_queue CHARACTER VARYING(61) COLLATE pg_catalog."C",
    enqueue_time TIMESTAMP WITHOUT TIME ZONE,
    state INTEGER,
    original_msgid BYTEA,
```

```
transaction_group CHARACTER VARYING(30) COLLATE pg_catalog."C",
delivery_mode INTEGER
DBMS_AQ.PERSISTENT);
```

The supported values for `message_properties_t` are:

<code>priority</code>	If the queue table definition includes a <code>sort_list</code> that references <code>priority</code> , this parameter affects the order that messages are dequeued. A lower value indicates a higher dequeue priority.
<code>delay</code>	Specify the number of seconds that will pass before a message is available for dequeuing or <code>NO_DELAY</code> .
<code>expiration</code>	Use the expiration parameter to specify the number of seconds until a message expires.
<code>correlation</code>	Use correlation to specify a message that will be associated with the entry; the default is <code>NULL</code> .
<code>attempts</code>	This is a system-maintained value that specifies the number of attempts to dequeue the message.
<code>recipient_list</code>	This parameter is not supported.
<code>exception_queue</code>	Use the <code>exception_queue</code> parameter to specify the name of an exception queue to which a message will be moved if it expires or is dequeued by a transaction that rolls back too many times.
<code>enqueue_time</code>	<code>enqueue_time</code> is the time the record was added to the queue; this value is provided by the system.
<code>state</code>	This parameter is maintained by <code>DBMS_AQ</code> ; state can be: <code>DBMS_AQ.WAITING</code> – the delay has not been reached. <code>DBMS_AQ.READY</code> – the queue entry is ready for processing. <code>DBMS_AQ.PROCESSED</code> – the queue entry has been processed. <code>DBMS_AQ.EXPIRED</code> – the queue entry has been moved to the exception queue.
<code>original_msgid</code>	This parameter is accepted for compatibility and ignored.
<code>transaction_group</code>	This parameter is accepted for compatibility and ignored.
<code>delivery_mode</code>	This parameter is not supported; specify a value of <code>DBMS_AQ.PERSISTENT</code> .

<payload>

Use the <payload> parameter to provide the data that will be associated with the queue entry. The payload type must match the type specified when creating the corresponding queue table (see `DBMS_AQADM.CREATE_QUEUE_TABLE`).

<msgid>

Use the <msgid> parameter to retrieve a unique (system-generated) message identifier.

Example

The following anonymous block calls `DBMS_AQ.ENQUEUE`, adding a message to a queue named `work_order`:

```
DECLARE
  enqueue_options    DBMS_AQ.ENQUEUE_OPTIONS_T;
  message_properties DBMS_AQ.MESSAGE_PROPERTIES_T;
  message_handle      raw(16);
  payload             work_order;

BEGIN

  payload := work_order('Smith', 'system upgrade');
```



```
DBMS_AQ.ENQUEUE (  
  queue_name      => 'work_order',  
  enqueue_options => enqueue_options,  
  message_properties => message_properties,  
  payload         => payload,  
  msgid          => message_handle  
);  
END;
```

3.2.2 DEQUEUE

The DEQUEUE procedure dequeues a message. The signature is:

```
DEQUEUE (
    <queue_name> IN VARCHAR2,
    <dequeue_options> IN DBMS_AQ.DEQUEUE_OPTIONS_T,
    <message_properties> OUT DBMS_AQ.MESSAGE_PROPERTIES_T,
    <payload> OUT type_name,
    <msgid> OUT RAW)
```

Parameters

<queue_name>

The name (optionally schema-qualified) of an existing queue. If you omit the schema name, the server will use the schema specified in the `SEARCH_PATH`. Please note that unlike Oracle, unquoted identifiers are converted to lower case before storing. To include special characters or use a case-sensitive name, enclose the name in double quotes.

For detailed information about creating a queue, please see `DBMS_AQADM.CREATE_QUEUE`.

<dequeue_options>

<dequeue_options> is a value of the type, `dequeue_options_t`:

```
DEQUEUE_OPTIONS_T IS RECORD (
    consumer_name CHARACTER VARYING(30),
    dequeue_mode INTEGER,
    navigation INTEGER,
    visibility INTEGER,
    wait INTEGER,
    msgid BYTEA,
    correlation CHARACTER VARYING(128),
    deq_condition CHARACTER VARYING(4000),
    transformation CHARACTER VARYING(61),
    delivery_mode INTEGER);
```

Currently, the supported parameter values for `dequeue_options_t` are:

<code>consumer_name</code>	Must be NULL.
<code>dequeue_mode</code>	The locking behavior of the dequeue operation. Must be either: <code>DBMS_AQ.BROWSE</code> – Read the message without obtaining a lock. <code>DBMS_AQ.LOCKED</code> – Read the message after acquiring a lock. <code>DBMS_AQ.REMOVE</code> – Read the message before deleting the message. <code>DBMS_AQ.REMOVE_NODATA</code> – Read the message, but do not delete the message.
<code>navigation</code>	Identifies the message that will be retrieved. Must be either: <code>FIRST_MESSAGE</code> – The first message within the queue that matches the search term. <code>NEXT_MESSAGE</code> – The next message that is available that matches the first term.
<code>visibility</code>	Must be <code>ON_COMMIT</code> – if you roll back the current transaction the dequeued item will remain in the queue.
<code>wait</code>	Must be a number larger than 0, or: <code>DBMS_AQ.FOREVER</code> – Wait indefinitely. <code>DBMS_AQ.NO_WAIT</code> – Do not wait.
<code>msgid</code>	The message ID of the message that will be dequeued.
<code>correlation</code>	Accepted for compatibility, and ignored.
<code>deq_condition</code>	A <code>VARCHAR2</code> expression that evaluates to a <code>BOOLEAN</code> value indicating if the message should be dequeued
<code>transformation</code>	Accepted for compatibility, and ignored.
<code>delivery_mode</code>	Must be <code>PERSISTENT</code> ; buffered messages are not supported at this time.

<message_properties>

<message_properties> is a value of the type, `message_properties_t`:

```
message_properties_t IS RECORD(
    priority INTEGER,
    delay INTEGER,
    expiration INTEGER,
    correlation CHARACTER VARYING(128) COLLATE pg_catalog."C",
    attempts INTEGER,
    recipient_list "AQ$RECIPIENT_LIST_T",
    exception_queue CHARACTER VARYING(61) COLLATE pg_catalog."C",
    enqueue_time TIMESTAMP WITHOUT TIME ZONE,
    state INTEGER,
    original_msgid BYTEA,
    transaction_group CHARACTER VARYING(30) COLLATE pg_catalog."C",
    delivery_mode INTEGER
    DBMS_AQ.PERSISTENT);
```

The supported values for `message_properties_t` are:

priority	If the queue table definition includes a <code>sort_list</code> that references <code>priority</code> , this parameter affects the order that messages are dequeued. A lower value indicates a higher dequeue priority.
delay	Specify the number of seconds that will pass before a message is available for dequeuing or <code>NO_DELAY</code> .
expiration	Use the <code>expiration</code> parameter to specify the number of seconds until a message expires.
correlation	Use <code>correlation</code> to specify a message that will be associated with the entry; the default is <code>NULL</code> .
attempts	This is a system-maintained value that specifies the number of attempts to dequeue the message.
recipient_list	This parameter is not supported.
exception_queue	Use the <code>exception_queue</code> parameter to specify the name of an exception queue to which a message will be moved if it expires or is dequeued by a transaction that rolls back too many times.
enqueue_time	<code>enqueue_time</code> is the time the record was added to the queue; this value is provided by the system.
state	This parameter is maintained by <code>DBMS_AQ</code> ; state can be: <code>DBMS_AQ.WAITING</code> – the delay has not been reached. <code>DBMS_AQ.READY</code> – the queue entry is ready for processing. <code>DBMS_AQ.PROCESSED</code> – the queue entry has been processed. <code>DBMS_AQ.EXPIRED</code> – the queue entry has been moved to the exception queue.
original_msgid	This parameter is accepted for compatibility and ignored.
transaction_group	This parameter is accepted for compatibility and ignored.
delivery_mode	This parameter is not supported; specify a value of <code>DBMS_AQ.PERSISTENT</code> .

<payload>

Use the <payload> parameter to retrieve the payload of a message with a dequeue operation. The payload type must match the type specified when creating the queue table.

<msgid>

Use the <msgid> parameter to retrieve a unique message identifier.

Example

The following anonymous block calls `DBMS_AQ.DEQUEUE`, retrieving a message from the queue and a payload:

```
DECLARE
    dequeue_options    DBMS_AQ.DEQUEUE_OPTIONS_T;
    message_properties DBMS_AQ.MESSAGE_PROPERTIES_T;
    message_handle      raw(16);
    payload             work_order;

BEGIN
    dequeue_options.dequeue_mode := DBMS_AQ.BROWSE;

    DBMS_AQ.DEQUEUE (
        queue_name      => 'work_queue',
        dequeue_options => dequeue_options,
        message_properties => message_properties,
        payload          => payload,
        msgid            => message_handle
    );
```

```
DBMS_OUTPUT.PUT_LINE(  
  'The next work order is [' || payload.subject || '].'  
);  
END;
```

The payload is displayed by `DBMS_OUTPUT.PUT_LINE`.

3.2.3 REGISTER

Use the REGISTER procedure to register an email address, procedure or URL that will be notified when an item is enqueued or dequeued. The signature is:

```
REGISTER(
    <reg_list> IN SYS.AQ$_REG_INFO_LIST,
    <count> IN NUMBER)
```

Parameters

<reg_list>

<reg_list> is a list of type AQ\$_REG_INFO_LIST; that provides information about each subscription that you would like to register. Each entry within the list is of the type AQ\$_REG_INFO, and may contain:

Attribute	Type	Description
name	VARCHAR2 (128)	The (optionally schema-qualified) name of the subscription.
namespace	NUMERIC	The only supported value is DBMS_AQ.NAMESPACE_AQ (0)
callback	VARCHAR2 (4000)	Describes the action that will be performed upon notification. Currently, only calls to PL/SQL procedures are supported. The call should take the form: plsql://schema.procedure Where: schema specifies the schema in which the procedure resides. procedure specifies the name of the procedure that will be notified.
context	RAW (16)	Any user-defined value required by the callback procedure.

<count>

<count> is the number of entries in <reg_list>.

Example

The following anonymous block calls DBMS_AQ.REGISTER, registering procedures that will be notified when an item is added to or removed from a queue. A set of attributes (of sys.aq\$_reg_info type) is provided for each subscription identified in the DECLARE section:

```
DECLARE
    subscription1 sys.aq$_reg_info;
    subscription2 sys.aq$_reg_info;
    subscription3 sys.aq$_reg_info;
    subscriptionlist sys.aq$_reg_info_list;
BEGIN
    subscription1 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
    'plsql://assign_worker?PR=0',HEXTORAW('FFFF'));
    subscription2 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
    'plsql://add_to_history?PR=1',HEXTORAW('FFFF'));
    subscription3 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
    'plsql://reserve_parts?PR=2',HEXTORAW('FFFF'));
```

```
subscriptionlist := sys.aq$_reg_info_list(subscription1,  
subscription2, subscription3);  
dbms_aq.register(subscriptionlist, 3);  
commit;  
  
END;  
/
```

The `subscriptionlist` is of type `sys.aq$_reg_info_list`, and contains the previously described `sys.aq$_reg_info` objects. The list name and an object count are passed to `dbms_aq.register`.

3.2.4 UNREGISTER

Use the UNREGISTER procedure to turn off notifications related to enqueueing and dequeueing. The signature is:

```
UNREGISTER (
    <reg_list> IN SYS.AQ$_REG_INFO_LIST,
    <count> IN NUMBER)
```

Parameter

<reg_list>

<reg_list> is a list of type AQ\$_REG_INFO_LIST; that provides information about each subscription that you would like to register. Each entry within the list is of the type AQ\$_REG_INFO, and may contain:

Attribute	Type	Description
name	VARCHAR2 (128)	The (optionally schema-qualified) name of the subscription.
namespace	NUMERIC	The only supported value is DBMS_AQ.NAMESPACE_AQ (0)
callback	VARCHAR2 (4000)	Describes the action that will be performed upon notification. Currently, only calls to PL/SQL procedures are supported. The call should take the form: plsql://schema.procedure Where: schema specifies the schema in which the procedure resides. procedure specifies the name of the procedure that will be notified.
context	RAW (16)	Any user-defined value required by the procedure.

<count>

<count> is the number of entries in <reg_list>.

Example

The following anonymous block calls DBMS_AQ.UNREGISTER, disabling the notifications specified in the example for DBMS_AQ.REGISTER:

```
DECLARE

    subscription1 sys.aq$_reg_info;
    subscription2 sys.aq$_reg_info;
    subscription3 sys.aq$_reg_info;
    subscriptionlist sys.aq$_reg_info_list;

BEGIN

    subscription1 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
    'plsql://assign_worker?PR=0',HEXTORAW('FFFF'));
    subscription2 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
    'plsql://add_to_history?PR=1',HEXTORAW('FFFF'));
    subscription3 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
    'plsql://reserve_parts?PR=2',HEXTORAW('FFFF'));

    subscriptionlist := sys.aq$_reg_info_list(subscription1,
```



```
subscription2, subscription3);

    dbms_aq.unregister(subscriptionlist, 3);
    commit;
END;
/
```

The `subscriptionlist` is of type `sys.aq$_reg_info_list`, and contains the previously described `sys.aq$_reg_info` objects. The list name and an object count are passed to `dbms_aq.unregister`.

3.3 DBMS_AQADM

EDB Postgres Advanced Server Advanced Queueing provides message queueing and message processing for the Advanced Server database. User-defined messages are stored in a queue; a collection of queues is stored in a queue table. Procedures in the DBMS_AQADM package create and manage message queues and queue tables. Use the DBMS_AQ package to add messages to a queue or remove messages from a queue, or register or unregister a PL/SQL callback procedure.

Advanced Server also provides extended (non-compatible) functionality for the DBMS_AQ package with SQL commands. Please see the *Database Compability for Oracle Developers Reference Guide* for detailed information about the following SQL commands:

- ALTER QUEUE
- ALTER QUEUE TABLE
- CREATE QUEUE
- CREATE QUEUE TABLE
- DROP QUEUE
- DROP QUEUE TABLE

The DBMS_AQADM package provides procedures that allow you to create and manage queues and queue tables.

Function/Procedure	Return Type	Description
ALTER_QUEUE	n/a	Modify an existing queue.
ALTER_QUEUE_TABLE	n/a	Modify an existing queue table.
CREATE_QUEUE	n/a	Create a queue.
CREATE_QUEUE_TABLE	n/a	Create a queue table.
DROP_QUEUE	n/a	Drop an existing queue.
DROP_QUEUE_TABLE	n/a	Drop an existing queue table.
PURGE_QUEUE_TABLE	n/a	Remove one or more messages from a queue table.
START_QUEUE	n/a	Make a queue available for enqueueing and dequeueing procedures.
STOP_QUEUE	n/a	Make a queue unavailable for enqueueing and dequeueing procedures

Advanced Server's implementation of DBMS_AQADM is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Advanced Server supports use of the arguments listed below:

Constant	Description	For Parameters
DBMS_AQADM.TRANSACTIONAL (1)	This constant is defined, but will return an error if used.	message_grouping
DBMS_AQADM.NONE (0)	Use to specify message grouping for a queue table.	message_grouping
DBMS_AQADM.NORMAL_QUEUE (0)	Use with create_queue to specify queue_type.	queue_type
DBMS_AQADM.EXCEPTION_QUEUE (1)	Use with create_queue to specify queue_type.	queue_type
DBMS_AQADM.INFINITE (-1)	Use with create_queue to specify retention_time.	retention_time
DBMS_AQADM.PERSISTENT (0)	The message should be stored in a table.	enqueue_options_t. delivery_mode
DBMS_AQADM.BUFFERED (1)	This constant is defined, but will return an error if used.	enqueue_options_t. delivery_mode
DBMS_AQADM.PERSISTENT_OR_BUFFERED (2)	This constant is defined, but will return an error if used.	enqueue_options_t. delivery_mode

3.3.1 ALTER_QUEUE

Use the ALTER_QUEUE procedure to modify an existing queue. The signature is:

```
ALTER_QUEUE (
    <max_retries> IN NUMBER DEFAULT NULL,
    <retry_delay> IN NUMBER DEFAULT 0
    <retention_time> IN NUMBER DEFAULT 0,
    <auto_commit> IN BOOLEAN DEFAULT TRUE)
    <comment> IN VARCHAR2 DEFAULT NULL,
```

Parameters

<queue_name>

The name of the new queue.

<max_retries>

<max_retries> specifies the maximum number of attempts to remove a message with a dequeue statement. The value of <max_retries> is incremented with each ROLLBACK statement. When the number of failed attempts reaches the value specified by <max_retries>, the message is moved to the exception queue. Specify 0 to indicate that no retries are allowed.

<retry_delay>

<retry_delay> specifies the number of seconds until a message is scheduled for re-processing after a ROLLBACK. Specify 0 to indicate that the message should be retried immediately (the default).

<retention_time>

<retention_time> specifies the length of time (in seconds) that a message will be stored after being dequeued. You can also specify 0 (the default) to indicate the message should not be retained after dequeuing, or INFINITE to retain the message forever.

<auto_commit>

This parameter is accepted for compatibility and ignored.

<comment>

<comment> specifies a comment associated with the queue.

Example

The following command alters a queue named work_order, setting the retry_delay parameter to 5 seconds:

```
EXEC DBMS_AQADM.ALTER_QUEUE(queue_name => 'work_order', retry_delay
=> 5);
```

3.3.2 ALTER_QUEUE_TABLE

Use the ALTER_QUEUE_TABLE procedure to modify an existing queue table.

The signature is:

```
ALTER_QUEUE_TABLE (  
    <queue_table> IN VARCHAR2,  
    <comment> IN VARCHAR2 DEFAULT NULL,  
    <primary_instance> IN BINARY_INTEGER DEFAULT 0,  
    <secondary_instance> IN BINARY_INTEGER DEFAULT 0,
```

Parameters

<queue_table>

The (optionally schema-qualified) name of the queue table.

<comment>

Use the <comment> parameter to provide a comment about the queue table.

<primary_instance>

<primary_instance> is accepted for compatibility and stored, but is ignored.

<secondary_instance>

<secondary_instance> is accepted for compatibility, but is ignored.

Example

The following command modifies a queue table named work_order_table:

```
EXEC DBMS_AQADM.ALTER_QUEUE_TABLE  
    (queue_table => 'work_order_table', comment => 'This queue table  
contains work orders for the shipping department.');
```

The queue table is named work_order_table; the command adds a comment to the definition of the queue table.

3.3.3 CREATE_QUEUE

Use the CREATE_QUEUE procedure to create a queue in an existing queue table. The signature is:

```
CREATE_QUEUE (
    <queue_name> IN VARCHAR2
    <queue_table> IN VARCHAR2,
    <queue_type> IN BINARY_INTEGER DEFAULT NORMAL_QUEUE,
    <max_retries> IN NUMBER DEFAULT 5,
    <retry_delay> IN NUMBER DEFAULT 0
    <retention_time> IN NUMBER DEFAULT 0,
    <dependency_tracking> IN BOOLEAN DEFAULT FALSE,
    <comment> IN VARCHAR2 DEFAULT NULL,
    <auto_commit> IN BOOLEAN DEFAULT TRUE)
```

Parameters

<queue_name>

The name of the new queue.

<queue_table>

The name of the table in which the new queue will reside.

<queue_type>

The type of the new queue. The valid values for <queue_type> are:

DBMS_AQADM.NORMAL_QUEUE – This value specifies a normal queue (the default).

DBMS_AQADM.EXCEPTION_QUEUE – This value specifies that the new queue is an exception queue. An exception queue will support only dequeue operations.

<max_retries>

<max_retries> specifies the maximum number of attempts to remove a message with a dequeue statement. The value of <max_retries> is incremented with each ROLLBACK statement. When the number of failed attempts reaches the value specified by <max_retries>, the message is moved to the exception queue. The default value for a system table is 0; the default value for a user created table is 5.

<retry_delay>

<retry_delay> specifies the number of seconds until a message is scheduled for re-processing after a ROLLBACK. Specify 0 to indicate that the message should be retried immediately (the default).

<retention_time>

<retention_time> specifies the length of time (in seconds) that a message will be stored after being dequeued. You can also specify 0 (the default) to indicate the message should not be retained after dequeuing, or INFINITE to retain the message forever.

<dependency_tracking>

This parameter is accepted for compatibility and ignored.

<comment>

<comment> specifies a comment associated with the queue.

<auto_commit>

This parameter is accepted for compatibility and ignored.

Example

The following anonymous block creates a queue named `work_order` in the `work_order_table` table:

```
BEGIN
DBMS_AQADM.CREATE_QUEUE ( queue_name => 'work_order', queue_table
=> 'work_order_table', comment => 'This queue contains pending work
orders. ');
END;
```

3.3.4 CREATE_QUEUE_TABLE

Use the CREATE_QUEUE_TABLE procedure to create a queue table. The signature is:

```
CREATE_QUEUE_TABLE (
    <queue_table> IN VARCHAR2,
    <queue_payload_type> IN VARCHAR2,
    <storage_clause> IN VARCHAR2 DEFAULT NULL,
    <sort_list> IN VARCHAR2 DEFAULT NULL,
    <multiple_consumers> IN BOOLEAN DEFAULT FALSE,
    <message_grouping> IN BINARY_INTEGER DEFAULT NONE,
    <comment> IN VARCHAR2 DEFAULT NULL,
    <auto_commit> IN BOOLEAN DEFAULT TRUE,
    <primary_instance> IN BINARY_INTEGER DEFAULT 0,
    <secondary_instance> IN BINARY_INTEGER DEFAULT 0,
    <compatible> IN VARCHAR2 DEFAULT NULL,
    <secure> IN BOOLEAN DEFAULT FALSE)
```

Parameters

<queue_table>

The (optionally schema-qualified) name of the queue table.

<queue_payload_type>

The user-defined type of the data that will be stored in the queue table. Please note that to specify a RAW data type, you must create a user-defined type that identifies a RAW type.

<storage_clause>

Use the <storage_clause> parameter to specify attributes for the queue table. Please note that only the TABLESPACE option is enforced; all others are accepted for compatibility and ignored. Use the TABLESPACE clause to specify the name of a tablespace in which the table will be created.

<storage_clause> may be one or more of the following:

TABLESPACE <tablespace_name>, PCTFREE integer, PCTUSED integer,
INITTRANS integer, MAXTRANS integer or STORAGE <storage_option>.

<storage_option> may be one or more of the following:

MINEXTENTS integer, MAXEXTENTS integer, PCTINCREASE integer,
INITIAL <size_clause>, NEXT, FREELISTS integer, OPTIMAL
<size_clause>, BUFFER_POOL {KEEP|RECYCLE|DEFAULT}.

<sort_list>

<sort_list> controls the dequeuing order of the queue; specify the names of the column(s) that will be used to sort the queue (in ascending order). The currently accepted values are the following combinations of enq_time and priority:

enq_time, priority


```

    priority, enq_time
    priority
    enq_time
<multiple_consumers>
    <multiple_consumers> queue tables is not supported.
<message_grouping>
    If specified, <message_grouping> must be NONE.
<comment>
    Use the <comment> parameter to provide a comment about the queue table.
<auto_commit>
    <auto_commit> is accepted for compatibility, but is ignored.
<primary_instance>
    <primary_instance> is accepted for compatibility and stored, but is ignored.
<secondary_instance>
    <secondary_instance> is accepted for compatibility, but is ignored.
<compatible>
    <compatible> is accepted for compatibility, but is ignored.
<secure>
    <secure> is accepted for compatibility, but is ignored.

```

Example

The following anonymous block first creates a type (`work_order`) with attributes that hold a name (a `VARCHAR2`), and a project description (a `TEXT`). The block then uses that type to create a queue table:

```

BEGIN

CREATE TYPE work_order AS (name VARCHAR2, project TEXT, completed BOOLEAN);

EXEC DBMS_AQADM.CREATE_QUEUE_TABLE
    (queue_table => 'work_order_table',
     queue_payload_type => 'work_order',
     comment => 'Work order message queue table');

END;

```

The queue table is named `work_order_table`, and contains a payload of a type `work_order`. A comment notes that this is the Work order message queue table.

3.3.5 DROP_QUEUE

Use the DROP_QUEUE procedure to delete a queue. The signature is:

```
DROP_QUEUE (  
    <queue_name> IN VARCHAR2,  
    <auto_commit> IN BOOLEAN DEFAULT TRUE)
```

Parameters

<queue_name>

The name of the queue that you wish to drop.

<auto_commit>

<auto_commit> is accepted for compatibility, but is ignored.

Example

The following anonymous block drops the queue named work_order:

```
BEGIN  
DBMS_AQADM.DROP_QUEUE(queue_name => 'work_order');  
END;
```

3.3.6 DROP_QUEUE_TABLE

Use the `DROP_QUEUE_TABLE` procedure to delete a queue table. The signature is:

```
DROP_QUEUE_TABLE (  
    <queue_table> IN VARCHAR2,  
    <force> IN BOOLEAN default FALSE,  
    <auto_commit> IN BOOLEAN default TRUE)
```

Parameters

<queue_table>

The (optionally schema-qualified) name of the queue table.

<force>

The <force> keyword determines the behavior of the `DROP_QUEUE_TABLE` command when dropping a table that contain entries:

If the target table contains entries and force is `FALSE`, the command will fail, and the server will issue an error.

If the target table contains entries and force is `TRUE`, the command will drop the table and any dependent objects.

<auto_commit>

<auto_commit> is accepted for compatibility, but is ignored.

Example

The following anonymous block drops a table named `work_order_table`:

```
BEGIN  
    DBMS_AQADM.DROP_QUEUE_TABLE ('work_order_table', force => TRUE);  
END;
```

3.3.7 PURGE_QUEUE_TABLE

Use the PURGE_QUEUE_TABLE procedure to delete messages from a queue table. The signature is:

```
PURGE_QUEUE_TABLE (
    <queue_table> IN VARCHAR2,
    <purge_condition> IN VARCHAR2,
    <purge_options> IN aq$_purge_options_t)
```

Parameters

<queue_table>

<queue_table> specifies the name of the queue table from which you are deleting a message.

<purge_condition>

Use <purge_condition> to specify a condition (a SQL WHERE clause) that the server will evaluate when deciding which messages to purge.

<purge_options>

<purge_options> is an object of the type aq\$_purge_options_t. An aq\$_purge_options_t object contains:

Attribute	Type	Description
Block	Boolean	Specify TRUE if an exclusive lock should be held on all queues within the table; the default is FALSE.
delivery_mode	INTEGER	<delivery_mode> specifies the type of message that will be purged. The only accepted value is DBMS_AQ.PERSISTENT.

Example

The following anonymous block removes any messages from the work_order_table with a value in the completed column of YES:

```
DECLARE
    purge_options dbms_aqadm.aq$_purge_options_t;
BEGIN
    dbms_aqadm.purge_queue_table('work_order_table', 'completed = YES',
    purge_options);
END;
```

3.3.8 START_QUEUE

Use the `START_QUEUE` procedure to make a queue available for enqueueing and dequeueing.

The signature is:

```
START_QUEUE (  
    <queue_name> IN VARCHAR2,  
    <enqueue> IN BOOLEAN DEFAULT TRUE,  
    <dequeue> IN BOOLEAN DEFAULT TRUE)
```

Parameters

<queue_name>

<queue_name> specifies the name of the queue that you are starting.

<enqueue>

Specify `TRUE` to enable enqueueing (the default), or `FALSE` to leave the current setting unchanged.

<dequeue>

Specify `TRUE` to enable dequeueing (the default), or `FALSE` to leave the current setting unchanged.

Example

The following anonymous block makes a queue named `work_order` available for enqueueing:

```
BEGIN  
DBMS_AQADM.START_QUEUE  
(queue_name => 'work_order');  
END;
```

3.3.9 STOP_QUEUE

Use the STOP_QUEUE procedure to disable enqueueing or dequeuing on a specified queue.

The signature is:

```
STOP_QUEUE (  
    <queue_name> IN VARCHAR2,  
    <enqueue> IN BOOLEAN DEFAULT TRUE,  
    <dequeue> IN BOOLEAN DEFAULT TRUE,  
    <wait> IN BOOLEAN DEFAULT TRUE)
```

Parameters

<queue_name>

<queue_name> specifies the name of the queue that you are stopping.

<enqueue>

Specify TRUE to disable enqueueing (the default), or FALSE to leave the current setting unchanged.

<dequeue>

Specify TRUE to disable dequeuing (the default), or FALSE to leave the current setting unchanged.

<wait>

Specify TRUE to instruct the server to wait for any uncompleted transactions to complete before applying the specified changes; while waiting to stop the queue, no transactions are allowed to enqueue or dequeue from the specified queue. Specify FALSE to stop the queue immediately.

Example

The following anonymous block disables enqueueing and dequeuing from the queue named work_order:

```
BEGIN  
DBMS_AQADM.STOP_QUEUE(queue_name =>'work_order', enqueue=>TRUE,  
dequeue=>TRUE, wait=>TRUE);  
END;
```

Enqueueing and dequeuing will stop after any outstanding transactions complete.

3.4 DBMS_CRYPTO

The DBMS_CRYPTO package provides functions and procedures that allow you to encrypt or decrypt RAW, BLOB or CLOB data. You can also use DBMS_CRYPTO functions to generate cryptographically strong random values.

The following table lists the DBMS_CRYPTO Functions and Procedures.

Function/Procedure	Return Type	Description
DECRYPT(src, typ, key, iv)	RAW	Decrypts RAW data.
DECRYPT(dst INOUT, src, typ, key, iv)	N/A	Decrypts BLOB data.
DECRYPT(dst INOUT, src, typ, key, iv)	N/A	Decrypts CLOB data.
ENCRYPT(src, typ, key, iv)	RAW	Encrypts RAW data.
ENCRYPT(dst INOUT, src, typ, key, iv)	N/A	Encrypts BLOB data.
ENCRYPT(dst INOUT, src, typ, key, iv)	N/A	Encrypts CLOB data.
HASH(src, typ)	RAW	Applies a hash algorithm to RAW data.
HASH(src)	RAW	Applies a hash algorithm to CLOB data.
MAC(src, typ, key)	RAW	Returns the hashed MAC value of the given RAW data using the specified hash algorithm and key.
MAC(src, typ, key)	RAW	Returns the hashed MAC value of the given CLOB data using the specified hash algorithm and key.
RANDOMBYTES(number_bytes)	RAW	Returns a specified number of cryptographically strong random bytes.
RANDOMINTEGER()	INTEGER	Returns a random INTEGER.
RANDOMNUMBER()	NUMBER	Returns a random NUMBER.

DBMS_CRYPTO functions and procedures support the following error messages:

```
ORA-28239 - DBMS_CRYPTO.KeyNull
ORA-28829 - DBMS_CRYPTO.CipherSuiteNull
ORA-28827 - DBMS_CRYPTO.CipherSuiteInvalid
```

Unlike Oracle, Advanced Server will not return error ORA-28233 if you re-encrypt previously encrypted information.

Please note that RAW and BLOB are synonyms for the PostgreSQL BYTEA data type, and CLOB is a synonym for TEXT.

3.4.1 DECRYPT

The DECRYPT function or procedure decrypts data using a user-specified cipher algorithm, key and optional initialization vector. The signature of the DECRYPT function is:

```
DECRYPT
  (<src> IN RAW, <typ> IN INTEGER, <key> IN RAW, <iv> IN RAW
  DEFAULT NULL) RETURN RAW
```

The signature of the DECRYPT procedure is:

```
DECRYPT
  (<dst> INOUT BLOB, <src> IN BLOB, <typ> IN INTEGER, <key> IN RAW,
  <iv> IN RAW DEFAULT NULL)
```

or

```
DECRYPT
  (<dst> INOUT CLOB, <src> IN CLOB, <typ> IN INTEGER, <key> IN RAW,
  <iv> IN RAW DEFAULT NULL)
```

When invoked as a procedure, DECRYPT returns BLOB or CLOB data to a user-specified BLOB.

Parameters

<dst>

<dst> specifies the name of a BLOB to which the output of the DECRYPT procedure will be written. The DECRYPT procedure will overwrite any existing data currently in <dst>.

<src>

<src> specifies the source data that will be decrypted. If you are invoking DECRYPT as a function, specify RAW data; if invoking DECRYPT as a procedure, specify BLOB or CLOB data.

<typ>

<typ> specifies the block cipher type and any modifiers. This should match the type specified when the <src> was encrypted. Advanced Server supports the following block cipher algorithms, modifiers and cipher suites:

Block Cipher Algorithms	
ENCRYPT_DES	CONSTANT INTEGER := 1;
ENCRYPT_3DES	CONSTANT INTEGER := 3;
ENCRYPT_AES	CONSTANT INTEGER := 4;
ENCRYPT_AES128	CONSTANT INTEGER := 6;
Block Cipher Modifiers	
CHAIN_CBC	CONSTANT INTEGER := 256;
CHAIN_ECB	CONSTANT INTEGER := 768;
Block Cipher Padding Modifiers	
PAD_PKCS5	CONSTANT INTEGER := 4096;
PAD_NONE	CONSTANT INTEGER := 8192;
Block Cipher Suites	
DES_CBC_PKCS5	CONSTANT INTEGER := ENCRYPT_DES + CHAIN_CBC + PAD_PKCS5;
DES3_CBC_PKCS5	CONSTANT INTEGER := ENCRYPT_3DES + CHAIN_CBC + PAD_PKCS5;
AES_CBC_PKCS5	CONSTANT INTEGER := ENCRYPT_AES + CHAIN_CBC + PAD_PKCS5;

<key>

<key> specifies the user-defined decryption key. This should match the key specified when the <src> was encrypted.

<iv>

<iv> (optional) specifies an initialization vector. If an initialization vector was specified when the <src> was encrypted, you must specify an initialization vector when decrypting the <src>. The default is NULL.

Examples

The following example uses the `DBMS_CRYPTO.DECRYPT` function to decrypt an encrypted password retrieved from the passwords table:

```
CREATE TABLE passwords
(
  principal VARCHAR2(90) PRIMARY KEY, -- username
  ciphertext RAW(9) -- encrypted password
);

CREATE FUNCTION get_password(username VARCHAR2) RETURN RAW AS
  typ      INTEGER := DBMS_CRYPTO.DES_CBC_PKCS5;
  key      RAW(128) := 'my secret key';
  iv       RAW(100) := 'my initialization vector';
  password RAW(2048);
BEGIN
  SELECT ciphertext INTO password FROM passwords WHERE principal =
    username;
  RETURN dbms_crypto.decrypt(password, typ, key, iv);
END;
```

Note that when calling `DECRYPT`, you must pass the same cipher type, key value and initialization vector that was used when `ENCRYPTING` the target.

3.4.2 ENCRYPT

The ENCRYPT function or procedure uses a user-specified algorithm, key, and optional initialization vector to encrypt RAW, BLOB or CLOB data. The signature of the ENCRYPT function is:

```
ENCRYPT
  (<src> IN RAW, <typ> IN INTEGER, <key> IN RAW,
   <iv> IN RAW DEFAULT NULL) RETURN RAW
```

The signature of the ENCRYPT procedure is:

```
ENCRYPT
  (<dst> INOUT BLOB, <src> IN BLOB, <typ> IN INTEGER, <key> IN RAW,
   <iv> IN RAW DEFAULT NULL)
```

or

```
ENCRYPT
  (<dst> INOUT BLOB, <src> IN CLOB, <typ> IN INTEGER, <key> IN RAW,
   <iv> IN RAW DEFAULT NULL)
```

When invoked as a procedure, ENCRYPT returns BLOB or CLOB data to a user-specified BLOB.

Parameters

<dst>

<dst> specifies the name of a BLOB to which the output of the ENCRYPT procedure will be written. The ENCRYPT procedure will overwrite any existing data currently in <dst>.

<src>

<src> specifies the source data that will be encrypted. If you are invoking ENCRYPT as a function, specify RAW data; if invoking ENCRYPT as a procedure, specify BLOB or CLOB data.

<typ>

<typ> specifies the block cipher type that will be used by ENCRYPT, and any modifiers. Advanced Server supports the block cipher algorithms, modifiers and cipher suites listed below:

Block Cipher Algorithms	
ENCRYPT_DES	CONSTANT INTEGER := 1;
ENCRYPT_3DES	CONSTANT INTEGER := 3;
ENCRYPT_AES	CONSTANT INTEGER := 4;
ENCRYPT_AES128	CONSTANT INTEGER := 6;
Block Cipher Modifiers	
CHAIN_CBC	CONSTANT INTEGER := 256;
CHAIN_ECB	CONSTANT INTEGER := 768;
Block Cipher Padding Modifiers	
“PAD_PKCS5”	CONSTANT INTEGER := 4096;
PAD_NONE	CONSTANT INTEGER := 8192;
Block Cipher Suites	
DES_CBC_PKCS5	CONSTANT INTEGER := ENCRYPT_DES + CHAIN_CBC + PAD_PKCS5;
DES3_CBC_PKCS5	CONSTANT INTEGER := ENCRYPT_3DES + CHAIN_CBC + PAD_PKCS5;
AES_CBC_PKCS5	CONSTANT INTEGER := ENCRYPT_AES + CHAIN_CBC + PAD_PKCS5;

<key>

<key> specifies the encryption key.

<iv>

<iv> (optional) specifies an initialization vector. By default, <iv> is NULL.

Examples

The following example uses the DBMS_CRYPTO.DES_CBC_PKCS5 Block Cipher Suite (a pre-defined set of algorithms and modifiers) to encrypt a value retrieved from the passwords table:

```
CREATE TABLE passwords
(
  principal  VARCHAR2(90) PRIMARY KEY, -- username
  ciphertext RAW(9) -- encrypted password
);
CREATE PROCEDURE set_password(username VARCHAR2, cleartext RAW) AS
  typ          INTEGER := DBMS_CRYPTO.DES_CBC_PKCS5;
  key          RAW(128) := 'my secret key';
  iv          RAW(100) := 'my initialization vector';
  encrypted    RAW(2048);
BEGIN
  encrypted := dbms_crypto.encrypt(cleartext, typ, key, iv);
  UPDATE passwords SET ciphertext = encrypted WHERE principal =
  username;
END;
```

ENCRYPT uses a key value of my secret key and an initialization vector of my initialization vector when encrypting the password; specify the same key and initialization vector when decrypting the password.

3.4.3 HASH

The HASH function uses a user-specified algorithm to return the hash value of a RAW or CLOB value. The HASH function is available in three forms:

```
HASH
  (<src> IN RAW, <typ> IN INTEGER) RETURN RAW
```

```
HASH
  (<src> IN CLOB, <typ> IN INTEGER) RETURN RAW
```

Parameters

<src>

<src> specifies the value for which the hash value will be generated. You can specify a RAW, a BLOB, or a CLOB value.

<typ>

<typ> specifies the HASH function type. Advanced Server supports the HASH function types listed below:

HASH Functions	
HASH_MD4	CONSTANT INTEGER := 1;
HASH_MD5	CONSTANT INTEGER := 2;
HASH_SH1	CONSTANT INTEGER := 3;

Examples

The following example uses DBMS_CRYPTO.HASH to find the md5 hash value of the string, cleartext source:

```
DECLARE
  typ INTEGER := DBMS_CRYPTO.HASH_MD5;
  hash_value RAW(100);
BEGIN

  hash_value := DBMS_CRYPTO.HASH('cleartext source', typ);

END;
```

3.4.4 MAC

The MAC function uses a user-specified MAC function to return the hashed MAC value of a RAW or CLOB value. The MAC function is available in three forms:

```
MAC
    (<src> IN RAW, <typ> IN INTEGER, <key> IN RAW) RETURN RAW
```

```
MAC
    (<src> IN CLOB, <typ> IN INTEGER, <key> IN RAW) RETURN RAW
```

Parameters

<src>

<src> specifies the value for which the MAC value will be generated. Specify a RAW, BLOB, or CLOB value.

<typ>

<typ> specifies the MAC function used. Advanced Server supports the MAC functions listed below.

MAC Functions	
HMAC_MD5	CONSTANT INTEGER := 1;
HMAC_SH1	CONSTANT INTEGER := 2;

<key>

<key> specifies the key that will be used to calculate the hashed MAC value.

Examples

The following example finds the hashed MAC value of the string cleartext source:

```
DECLARE
    typ INTEGER := DBMS_CRYPTO.HMAC_MD5;
    key RAW(100) := 'my secret key';
    mac_value RAW(100);
BEGIN
    mac_value := DBMS_CRYPTO.MAC('cleartext source', typ, key);
END;
```

DBMS_CRYPTO.MAC uses a key value of my secret key when calculating the MAC value of cleartext source.

3.4.5 RANDOMBYTES

The RANDOMBYTES function returns a RAW value of the specified length, containing cryptographically random bytes. The signature is:

```
RANDOMBYTES  
  (<number_bytes> IN INTEGER) RETURNS RAW
```

Parameter

<number_bytes>

<number_bytes> specifies the number of random bytes to be returned

Examples

The following example uses RANDOMBYTES to return a value that is 1024 bytes long:

```
DECLARE  
  result RAW(1024);  
BEGIN  
  result := DBMS_CRYPTO.RANDOMBYTES(1024);  
END;
```

3.4.6 RANDOMINTEGER

The `RANDOMINTEGER()` function returns a random `INTEGER` between 0 and 268,435,455. The signature is:

`RANDOMINTEGER()` RETURNS `INTEGER`

Examples

The following example uses the `RANDOMINTEGER` function to return a cryptographically strong random `INTEGER` value:

```
DECLARE
    result INTEGER;
BEGIN
    result := DBMS_CRYPTO.RANDOMINTEGER();
    DBMS_OUTPUT.PUT_LINE(result);
END;
```

3.4.7 RANDOMNUMBER

The `RANDOMNUMBER()` function returns a random `NUMBER` between 0 and 268,435,455. The signature is:

`RANDOMNUMBER()` RETURNS `NUMBER`

Examples

The following example uses the `RANDOMNUMBER` function to return a cryptographically strong random number:

```
DECLARE
    result NUMBER;
BEGIN
    result := DBMS_CRYPTO.RANDOMNUMBER();
    DBMS_OUTPUT.PUT_LINE(result);
END;
```


3.5 DBMS_JOB

The DBMS_JOB package provides for the creation, scheduling, and managing of jobs. A job runs a stored procedure which has been previously stored in the database. The SUBMIT procedure is used to create and store a job definition. A job identifier is assigned to a job along with its associated stored procedure and the attributes describing when and how often the job is to be run.

This package relies on the pgAgent scheduler. By default, the Advanced Server installer installs pgAgent, but you must start the pgAgent service manually prior to using DBMS_JOB. If you attempt to use this package to schedule a job after un-installing pgAgent, DBMS_JOB will throw an error. DBMS_JOB verifies that pgAgent is installed, but does not verify that the service is running.

The following table lists the supported DBMS_JOB procedures:

Function/Procedure	Return Type	Description
BROKEN(<job>, <broken> [, <next_date>])	n/a	Specify that a given job is either broken or not broken.
CHANGE(<job>, <what>, <next_date>, <interval>, instance, force>)	n/a	Change the job's parameters.
INTERVAL(<job>, <interval>)	n/a	Set the execution frequency by means of a date function that is recalculated each time the job is run. This value becomes the next date/time for execution.
NEXT_DATE(<job>, <next_date>)	n/a	Set the next date/time the job is to be run.
REMOVE(<job>)	n/a	Delete the job definition from the database.
RUN(<job>)	n/a	Forces execution of a job even if it is marked broken.
SUBMIT(<job> OUT, <what> [, <next_date> [, <interval> [, <no_parse>]]])	n/a	Creates a job and stores its definition in the database.
WHAT(<job>, <what>)	n/a	Change the stored procedure run by a job.

Advanced Server's implementation of DBMS_JOB is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Before using DBMS_JOB, a database superuser must create the pgAgent extension. Use the psql client to connect to a database and invoke the command:

```
CREATE EXTENSION pgagent;
```

When and how often a job is run is dependent upon two interacting parameters – <next_date> and <interval>. The <next_date> parameter is a date/time value that specifies the next date/time when the job is to be executed. The <interval> parameter is a string that contains a date function that evaluates to a date/time value.

Just prior to any execution of the job, the expression in the <interval> parameter is evaluated. The resulting value replaces the <next_date> value stored with the job. The job is then executed. In this manner, the expression in <interval> is repeatedly re-evaluated prior to each job execution, supplying the <next_date> date/time for the next execution.

The following examples use the following stored procedure, job_proc, which simply inserts a timestamp into table, jobrun, containing a single VARCHAR2 column.

```
CREATE TABLE jobrun (  
    runtime VARCHAR2(40)  
);  
  
CREATE OR REPLACE PROCEDURE job_proc  
IS  
BEGIN  
    INSERT INTO jobrun VALUES ('job_proc run at ' || TO_CHAR(SYSDATE,  
        'yyyy-mm-dd hh24:mi:ss'));  
END;
```

3.5.1 BROKEN

The **BROKEN** procedure sets the state of a job to either broken or not broken. A broken job cannot be executed except by using the **RUN** procedure.

BROKEN(**<job>** BINARY_INTEGER, **<broken>** BOOLEAN [, **<next_date>** DATE])

Parameters

<job>

Identifier of the job to be set as broken or not broken.

<broken>

If set to **TRUE** the job's state is set to broken. If set to **FALSE** the job's state is set to not broken. Broken jobs cannot be run except by using the **RUN** procedure.

<next_date>

Date/time when the job is to be run. The default is **SYSDATE**.

Examples

Set the state of a job with job identifier 104 to broken:

```
BEGIN
    DBMS_JOB.BROKEN(104,true);
END;
```

Change the state back to not broken:

```
BEGIN
    DBMS_JOB.BROKEN(104,false);
END;
```

3.5.2 CHANGE

The **CHANGE** procedure modifies certain job attributes including the stored procedure to be run, the next date/time the job is to be run, and how often it is to be run.

```
CHANGE(<job> BINARY_INTEGER <what> VARCHAR2, <next_date> DATE,
      <interval> VARCHAR2, <instance> BINARY_INTEGER, <force> BOOLEAN)
```

Parameters

<job>

Identifier of the job to modify.

<what>

Stored procedure name. Set this parameter to null if the existing value is to remain unchanged.

<next_date>

Date/time when the job is to be run next. Set this parameter to null if the existing value is to remain unchanged.

<interval>

Date function that when evaluated, provides the next date/time the job is to run. Set this parameter to null if the existing value is to remain unchanged.

<instance>

This argument is ignored, but is included for compatibility.

<force>

This argument is ignored, but is included for compatibility.

Examples

Change the job to run next on December 13, 2007. Leave other parameters unchanged.

```
BEGIN
    DBMS_JOB.CHANGE(104, NULL, TO_DATE('13-DEC-07', 'DD-MON-YY'), NULL, NULL,
    NULL);
END;
```

3.5.3 INTERVAL

The `INTERVAL` procedure sets the frequency of how often a job is to be run.

`INTERVAL(<job> BINARY_INTEGER, <interval> VARCHAR2)`

Parameters

`<job>`

Identifier of the job to modify.

`<interval>`

Date function that when evaluated, provides the next date/time the job is to be run. If `<interval>` is `NULL` and the job is complete, the job is removed from the queue.

Examples

Change the job to run once a week:

```
BEGIN

    DBMS_JOB.INTERVAL(104, 'SYSDATE + 7');

END;
```

3.5.4 NEXT_DATE

The NEXT_DATE procedure sets the date/time of when the job is to be run next.

NEXT_DATE (<job> BINARY_INTEGER, <next_date> DATE)

Parameters

<job>

Identifier of the job whose next run date is to be set.

<next_date>

Date/time when the job is to be run next.

Examples

Change the job to run next on December 14, 2007:

```
BEGIN

    DBMS_JOB.NEXT_DATE(104, TO_DATE('14-DEC-07', 'DD-MON-YY'));

END;
```

3.5.5 REMOVE

The `REMOVE` procedure deletes the specified job from the database. The job must be resubmitted using the `SUBMIT` procedure in order to have it executed again. Note that the stored procedure that was associated with the job is not deleted.

```
REMOVE (<job> BINARY_INTEGER)
```

Parameter

<job>

Identifier of the job that is to be removed from the database.

Examples

Remove a job from the database:

```
BEGIN

    DBMS_JOB.REMOVE(104);

END;
```

3.5.6 RUN

The RUN procedure forces the job to be run, even if its state is broken.

RUN (<job> BINARY_INTEGER)

Parameter

<job>

Identifier of the job to be run.

Examples

Force a job to be run.

```
BEGIN  
  
    DBMS_JOB.RUN (104) ;  
  
END;
```


3.5.7 SUBMIT

The `SUBMIT` procedure creates a job definition and stores it in the database. A job consists of a job identifier, the stored procedure to be executed, when the job is to be first run, and a date function that calculates the next date/time the job is to be run.

```
SUBMIT(<job> OUT BINARY_INTEGER, <what> VARCHAR2
[, <next_date> DATE [, <interval> VARCHAR2 [, <no_parse> BOOLEAN ]]])
```

Parameters

<job>

Identifier assigned to the job.

<what>

Name of the stored procedure to be executed by the job.

<next_date>

Date/time when the job is to be run next. The default is `SYSDATE`.

<interval>

Date function that when evaluated, provides the next date/time the job is to run. If <interval> is set to null, then the job is run only once. Null is the default.

<no_parse>

If set to `TRUE`, do not syntax-check the stored procedure upon job creation – check only when the job first executes. If set to `FALSE`, check the procedure upon job creation. The default is `FALSE`.

Note: The <no_parse> option is not supported in this implementation of `SUBMIT()`. It is included for compatibility only.

Examples

The following example creates a job using stored procedure, `job_proc`. The job will execute immediately and run once a day thereafter as set by the <interval> parameter, `SYSDATE + 1`.

```
DECLARE
    jobid INTEGER;
BEGIN
    DBMS_JOB.SUBMIT(jobid, 'job_proc;', SYSDATE,
        'SYSDATE + 1');
    DBMS_OUTPUT.PUT_LINE('jobid: ' || jobid);
END;

jobid: 104
```

The job immediately executes procedure, `job_proc`, populating table, `jobrun`, with a row:

```
SELECT * FROM jobrun;

          runtime

-----
job_proc run at 2007-12-11 11:43:25
(1 row)
```

3.5.8 WHAT

The WHAT procedure changes the stored procedure that the job will execute.

WHAT (<job> BINARY_INTEGER, <what> VARCHAR2)

Parameters

<job>

Identifier of the job for which the stored procedure is to be changed.

<what>

Name of the stored procedure to be executed.

Examples

Change the job to run the list_emp procedure:

```
BEGIN
    DBMS_JOB.WHAT(104, 'list_emp;');
END;
```

3.6 DBMS_LOB

The DBMS_LOB package provides the capability to operate on large objects. The following table lists the supported functions and procedures:

Function/Procedure	Return Type	Description
APPEND(<dest_lob> IN OUT, <src_lob>)	n/a	Appends one large object to another.
COMPARE(<lob_1>, <lob_2> [, <amount> [, <offset_1> [, <offset_2>]]])	INTEGER	Compares two large objects.
CONVERTTOBLOB(<dest_lob> IN OUT, <src_clob>, <amount>, <dest_offset> IN OUT, <src_offset> IN OUT, <blob_csid>, <lang_context> IN OUT, <warning> OUT)	n/a	Converts character data to binary.
CONVERTTOCLOB(<dest_lob> IN OUT, <src_blob>, <amount>, <dest_offset> IN OUT, <src_offset> IN OUT, <blob_csid>, <lang_context> IN OUT, <warning> OUT)	n/a	Converts binary data to character.
COPY(<dest_lob> IN OUT, <src_lob>, <amount> [, <dest_offset> [, <src_offset>]])	n/a	Copies one large object to another.
ERASE(lob_loc IN OUT, <amount> IN OUT [, <offset>])	n/a	Erase a large object.
GET_STORAGE_LIMIT(<lob_loc>)	INTEGER	Get the storage limit for large objects.
GETLENGTH(<lob_loc>)	INTEGER	Get the length of the large object.
INSTR(<lob_loc>, <pattern> [, <offset> [, <nth>]])	INTEGER	Get the position of the nth occurrence of a pattern in the large object starting at <offset>
READ(<lob_loc>, <amount> IN OUT, <offset>, <buffer> OUT)	n/a	Read a large object.
SUBSTR(<lob_loc> [, <amount> [, <offset>]])	RAW, VARCHAR2	Get part of a large object.
TRIM(<lob_loc> IN OUT, <newlen>)	n/a	Trim a large object to the specified length.
WRITE(<lob_loc> IN OUT, <amount>, <offset>, <buffer>)	n/a	Write data to a large object.
WRITEAPPEND(<lob_loc> IN OUT, <amount>, <buffer>)	n/a	Write data from the buffer to the end of a large object.

Advanced Server's implementation of DBMS_LOB is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The following table lists the public variables available in the package.

Public Variables	Data Type	Value
compress off	INTEGER	0
compress_on	INTEGER	1
deduplicate_off	INTEGER	0
deduplicate_on	INTEGER	4
default_csid	INTEGER	0
default_lang_ctx	INTEGER	0
encrypt_off	INTEGER	0
encrypt_on	INTEGER	1
file_readonly	INTEGER	0
lobmaxsize	INTEGER	1073741823
lob_readonly	INTEGER	0
lob_readwrite	INTEGER	1
no_warning	INTEGER	0
opt_compress	INTEGER	1
opt_deduplicate	INTEGER	4
opt_encrypt	INTEGER	2
warn_inconvertible_char	INTEGER	1

In the following sections, lengths and offsets are measured in bytes if the large objects are BLOBs. Lengths and offsets are measured in characters if the large objects are CLOBs.

3.6.1 APPEND

The APPEND procedure provides the capability to append one large object to another. Both large objects must be of the same type.

```
APPEND(<dest_lob> IN OUT { BLOB | CLOB }, <src_lob> { BLOB | CLOB })
```

Parameters

<dest_lob>

Large object locator for the destination object. Must be the same data type as <src_lob>.

<src_lob>

Large object locator for the source object. Must be the same data type as <dest_lob>.

3.6.2 COMPARE

The COMPARE procedure performs an exact byte-by-byte comparison of two large objects for a given length at given offsets. The large objects being compared must be the same data type.

```
<status> INTEGER COMPARE(<lob_1> { BLOB | CLOB },  
    <lob_2> { BLOB | CLOB }  
    [, <amount> INTEGER [, <offset_1> INTEGER [, <offset_2> INTEGER ]]])
```

Parameters

<lob_1>

Large object locator of the first large object to be compared. Must be the same data type as <lob_2>.

<lob_2>

Large object locator of the second large object to be compared. Must be the same data type as <lob_1>.

<amount>

If the data type of the large objects is BLOB, then the comparison is made for <amount> bytes. If the data type of the large objects is CLOB, then the comparison is made for <amount> characters. The default is the maximum size of a large object.

<offset_1>

Position within the first large object to begin the comparison. The first byte/character is offset 1. The default is 1.

<offset_2>

Position within the second large object to begin the comparison. The first byte/character is offset 1. The default is 1.

<status>

Zero if both large objects are exactly the same for the specified length for the specified offsets. Non-zero, if the objects are not the same. <NULL> if <amount>, <offset_1>, or <offset_2> are less than zero.

3.6.3 CONVERTTLOB

The CONVERTTLOB procedure provides the capability to convert character data to binary.

```
CONVERTTLOB(<dest_lob> IN OUT BLOB, <src_clob> CLOB,
            <amount> INTEGER, <dest_offset> IN OUT INTEGER,
            <src_offset> IN OUT INTEGER, <blob_csid> NUMBER,
            <lang_context> IN OUT INTEGER, <warning> OUT INTEGER)
```

Parameters

<dest_lob>

BLOB large object locator to which the character data is to be converted.

<src_clob>

CLOB large object locator of the character data to be converted.

<amount>

Number of characters of <src_clob> to be converted.

<dest_offset> IN

Position in bytes in the destination BLOB where writing of the source CLOB should begin. The first byte is offset 1.

<dest_offset> OUT

Position in bytes in the destination BLOB after the write operation completes. The first byte is offset 1.

<src_offset> IN

Position in characters in the source CLOB where conversion to the destination BLOB should begin. The first character is offset 1.

<src_offset> OUT

Position in characters in the source CLOB after the conversion operation completes. The first character is offset 1.

<blob_csid>

Character set ID of the converted, destination BLOB.

<lang_context> IN

Language context for the conversion. The default value of 0 is typically used for this setting.

<lang_context> OUT

Language context after the conversion completes.

<warning>

0 if the conversion was successful, 1 if an inconvertible character was encountered.

3.6.4 CONVERTTOCLOB

The CONVERTTOCLOB procedure provides the capability to convert binary data to character.

```
CONVERTTOCLOB(<dest_lob> IN OUT CLOB, <src_blob> BLOB,
              <amount> INTEGER, <dest_offset> IN OUT INTEGER,
              <src_offset> IN OUT INTEGER, <blob_csid> NUMBER,
              <lang_context> IN OUT INTEGER, <warning> OUT INTEGER)
```

Parameters

<dest_lob>

CLOB large object locator to which the binary data is to be converted.

<src_blob>

BLOB large object locator of the binary data to be converted.

<amount>

Number of bytes of <src_blob> to be converted.

<dest_offset> IN

Position in characters in the destination CLOB where writing of the source BLOB should begin. The first character is offset 1.

<dest_offset> OUT

Position in characters in the destination CLOB after the write operation completes. The first character is offset 1.

<src_offset> IN

Position in bytes in the source BLOB where conversion to the destination CLOB should begin. The first byte is offset 1.

<src_offset> OUT

Position in bytes in the source BLOB after the conversion operation completes. The first byte is offset 1.

<blob_csid>

Character set ID of the converted, destination CLOB.

<lang_context> IN

Language context for the conversion. The default value of 0 is typically used for this setting.

<lang_context> OUT

Language context after the conversion completes.

<warning>

0 if the conversion was successful, 1 if an inconvertible character was encountered.

3.6.5 COPY

The `COPY` procedure provides the capability to copy one large object to another. The source and destination large objects must be the same data type.

```
COPY(<dest_lob> IN OUT { BLOB | CLOB }, <src_lob>
{ BLOB | CLOB },
    <amount> INTEGER
    [, <dest_offset> INTEGER [, <src_offset> INTEGER ]])
```

Parameters

`<dest_lob>`

Large object locator of the large object to which `<src_lob>` is to be copied. Must be the same data type as `<src_lob>`.

`<src_lob>`

Large object locator of the large object to be copied to `<dest_lob>`. Must be the same data type as `<dest_lob>`.

`<amount>`

Number of bytes/characters of `<src_lob>` to be copied.

`<dest_offset>`

Position in the destination large object where writing of the source large object should begin. The first position is offset 1. The default is 1.

`<src_offset>`

Position in the source large object where copying to the destination large object should begin. The first position is offset 1. The default is 1.

3.6.6 ERASE

The ERASE procedure provides the capability to erase a portion of a large object. To erase a large object means to replace the specified portion with zero-byte fillers for BLOBs or with spaces for CLOBs. The actual size of the large object is not altered.

```
ERASE( <lob_loc> IN OUT { BLOB | CLOB }, <amount> IN OUT INTEGER  
      [, <offset> INTEGER ])
```

Parameters

<lob_loc>

Large object locator of the large object to be erased.

<amount> IN

Number of bytes/characters to be erased.

<amount> OUT

Number of bytes/characters actually erased. This value can be smaller than the input value if the end of the large object is reached before <amount> bytes/characters have been erased.

<offset>

Position in the large object where erasing is to begin. The first byte/character is position 1. The default is 1.

3.6.7 GET_STORAGE_LIMIT

The GET_STORAGE_LIMIT function returns the limit on the largest allowable large object.

```
<size> INTEGER GET_STORAGE_LIMIT(<lob_loc> BLOB)
```

```
<size> INTEGER GET_STORAGE_LIMIT(<lob_loc> CLOB)
```

Parameters

<size>

Maximum allowable size of a large object in this database.

<lob_loc>

This parameter is ignored, but is included for compatibility.

3.6.8 GETLENGTH

The GETLENGTH function returns the length of a large object.

```
<amount> INTEGER GETLENGTH (<lob_loc> BLOB)
```

```
<amount> INTEGER GETLENGTH (<lob_loc> CLOB)
```

Parameters

<lob_loc>

Large object locator of the large object whose length is to be obtained.

<amount>

Length of the large object in bytes for BLOBs or characters for CLOBs.

3.6.9 INSTR

The INSTR function returns the location of the nth occurrence of a given pattern within a large object.

```
<position> INTEGER INSTR(<lob_loc> { BLOB | CLOB },  
    <pattern> { RAW | VARCHAR2 } [, <offset> INTEGER [, <nth> INTEGER])
```

Parameters

<lob_loc>

Large object locator of the large object in which to search for pattern.

<pattern>

Pattern of bytes or characters to match against the large object, lob. <pattern> must be RAW if <lob_loc> is a BLOB. pattern must be VARCHAR2 if <lob_loc> is a CLOB.

<offset>

Position within <lob_loc> to start search for <pattern>. The first byte/character is position 1. The default is 1.

<nth>

Search for <pattern>, <nth> number of times starting at the position given by <offset>. The default is 1.

<position>

Position within the large object where <pattern> appears the nth time specified by <nth> starting from the position given by <offset>.

3.6.10 READ

The READ procedure provides the capability to read a portion of a large object into a buffer.

```
READ(<lob_loc> { BLOB | CLOB }, <amount> IN OUT BINARY_INTEGER,  
    <offset> INTEGER, <buffer> OUT { RAW | VARCHAR2 })
```

Parameters

<lob_loc>

Large object locator of the large object to be read.

<amount> IN

Number of bytes/characters to read.

<amount> OUT

Number of bytes/characters actually read. If there is no more data to be read, then <amount> returns 0 and a DATA_NOT_FOUND exception is thrown.

<offset>

Position to begin reading. The first byte/character is position 1.

<buffer>

Variable to receive the large object. If <lob_loc> is a BLOB, then <buffer> must be RAW. If <lob_loc> is a CLOB, then <buffer> must be VARCHAR2.

3.6.11 SUBSTR

The SUBSTR function provides the capability to return a portion of a large object.

```
<data> { RAW | VARCHAR2 } SUBSTR(<lob_loc> { BLOB | CLOB }  
    [, <amount> INTEGER [, <offset> INTEGER ]])
```

Parameters

<lob_loc>

Large object locator of the large object to be read.

<amount>

Number of bytes/characters to be returned. Default is 32,767.

<offset>

Position within the large object to begin returning data. The first byte/character is position 1. The default is 1.

<data>

Returned portion of the large object to be read. If <lob_loc> is a BLOB, the return data type is RAW. If <lob_loc> is a CLOB, the return data type is VARCHAR2.

3.6.12 TRIM

The TRIM procedure provides the capability to truncate a large object to the specified length.

```
TRIM(<lob_loc> IN OUT { BLOB | CLOB }, <newlen> INTEGER)
```

Parameters

<lob_loc>

Large object locator of the large object to be trimmed.

<newlen>

Number of bytes/characters to which the large object is to be trimmed.

3.6.13 WRITE

The `WRITE` procedure provides the capability to write data into a large object. Any existing data in the large object at the specified offset for the given length is overwritten by data given in the buffer.

```
WRITE(<lob_loc> IN OUT { BLOB | CLOB },  
      <amount> BINARY_INTEGER,  
      <offset> INTEGER, <buffer> { RAW | VARCHAR2 })
```

Parameters

`<lob_loc>`

Large object locator of the large object to be written.

`<amount>`

The number of bytes/characters in `<buffer>` to be written to the large object.

`<offset>`

The offset in bytes/characters from the beginning of the large object (origin is 1) for the write operation to begin.

`<buffer>`

Contains data to be written to the large object. If `<lob_loc>` is a BLOB, then `<buffer>` must be RAW. If `<lob_loc>` is a CLOB, then `<buffer>` must be VARCHAR2.

3.6.14 WRITEAPPEND

The WRITEAPPEND procedure provides the capability to add data to the end of a large object.

```
WRITEAPPEND(<lob_loc> IN OUT { BLOB | CLOB },  
            <amount> BINARY_INTEGER, <buffer> { RAW | VARCHAR2 })
```

Parameters

<lob_loc>

Large object locator of the large object to which data is to be appended.

<amount>

Number of bytes/characters from <buffer> to be appended the large object.

<buffer>

Data to be appended to the large object. If <lob_loc> is a BLOB, then <buffer> must be RAW. If <lob_loc> is a CLOB, then <buffer> must be VARCHAR2.

3.7 DBMS_LOCK

Advanced Server provides support for the DBMS_LOCK . SLEEP procedure.

Function/Procedure	Return Type	Description
SLEEP (<seconds>)	n/a	Suspends a session for the specified number of <seconds>.

Advanced Server's implementation of DBMS_LOCK is a partial implementation when compared to Oracle's version. Only DBMS_LOCK . SLEEP is supported.

3.7.1 SLEEP

The SLEEP procedure suspends the current session for the specified number of seconds.

SLEEP (<seconds> NUMBER)

Parameters

<seconds>

<seconds> specifies the number of seconds for which you wish to suspend the session. <seconds> can be a fractional value; for example, enter 1 . 75 to specify one and three-fourths of a second.

3.8 DBMS_MVIEW

Use procedures in the DBMS_MVIEW package to manage and refresh materialized views and their dependencies. Advanced Server provides support for the following DBMS_MVIEW procedures:

Procedure	Re- turn Type	Description
GET_MV_DEPENDENCIES(<list> VARCHAR2, <deplist> VARCHAR2);	n/a	The GET_MV_DEPENDENCIES procedure returns a list of dependencies for a specified view.
REFRESH(<list> VARCHAR2, <method> VARCHAR2, <rollback_seg> VARCHAR2, <push_deferred_rpc> BOOLEAN, <refresh_after_errors> BOOLEAN, <purge_option> NUMBER, <parallelism> NUMBER, <heap_size> NUMBER, <atomic_refresh> BOOLEAN, <nested> BOOLEAN);	n/a	This variation of the REFRESH procedure refreshes all views named in a comma-separated list of view names.
REFRESH(<tab> dbms_utility.uncl_array, <method> VARCHAR2, <rollback_seg> VARCHAR2, <push_deferred_rpc> BOOLEAN, <refresh_after_errors> BOOLEAN, <purge_option> NUMBER, <parallelism> NUMBER, <heap_size> NUMBER, <atomic_refresh> BOOLEAN, <nested> BOOLEAN);	n/a	This variation of the REFRESH procedure refreshes all views named in a table of dbms_utility.uncl_array values.
REFRESH_ALL_MVIEWS(<number_of_failures> BINARY_INTEGER, <method> VARCHAR2, <rollback_seg> VARCHAR2, <refresh_after_errors> BOOLEAN, <atomic_refresh> BOOLEAN);	n/a	The REFRESH_ALL_MVIEWS procedure refreshes all materialized views.
REFRESH_DEPENDENT(<number_of_failures> BINARY_INTEGER, <list> VARCHAR2, <method> VARCHAR2, <rollback_seg> VARCHAR2, <refresh_after_errors> BOOLEAN, <atomic_refresh> BOOLEAN, <nested> BOOLEAN);	n/a	This variation of the REFRESH_DEPENDENT procedure refreshes all views that are dependent on the views listed in a comma-separated list.
REFRESH_DEPENDENT(<number_of_failures> BINARY_INTEGER, <tab> dbms_utility.uncl_array, <method> VARCHAR2, <rollback_seg> VARCHAR2, <refresh_after_errors> BOOLEAN, <atomic_refresh> BOOLEAN, <nested> BOOLEAN);	n/a	This variation of the REFRESH_DEPENDENT procedure refreshes all views that are dependent on the views listed in a table of dbms_utility.uncl_array values.

Advanced Server's implementation of DBMS_MVIEW is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

3.8.1 GET_MV_DEPENDENCIES

When given the name of a materialized view, GET_MV_DEPENDENCIES returns a list of items that depend on the specified view. The signature is:

```
GET_MV_DEPENDENCIES (  
    <list> IN VARCHAR2,  
    <deplist> OUT VARCHAR2);
```

Parameters

<list>

<list> specifies the name of a materialized view, or a comma-separated list of materialized view names.

<deplist>

<deplist> is a comma-separated list of schema-qualified dependencies. <deplist> is a VARCHAR2 value.

Examples

The following example:

```
DECLARE  
    deplist VARCHAR2(1000);  
BEGIN  
    DBMS_MVIEW.GET_MV_DEPENDENCIES('public.emp_view', deplist);  
    DBMS_OUTPUT.PUT_LINE('deplist: ' || deplist);  
END;
```

Displays a list of the dependencies on a materialized view named `public.emp_view`.

3.8.2 REFRESH

Use the REFRESH procedure to refresh all views specified in either a comma-separated list of view names, or a table of DBMS_UTILITY.UNCL_ARRAY values. The procedure has two signatures; use the first form when specifying a comma-separated list of view names:

```
REFRESH (
    <list> IN VARCHAR2,
    <method> IN VARCHAR2 DEFAULT NULL,
    <rollback_seg> IN VARCHAR2 DEFAULT NULL,
    <push_deferred_rpc> IN BOOLEAN DEFAULT TRUE,
    <refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
    <purge_option> IN NUMBER DEFAULT 1,
    <parallelism> IN NUMBER DEFAULT 0,
    <heap_size> IN NUMBER DEFAULT 0,
    <atomic_refresh> IN BOOLEAN DEFAULT TRUE,
    <nested> IN BOOLEAN DEFAULT FALSE);
```

Use the second form to specify view names in a table of DBMS_UTILITY.UNCL_ARRAY values:

```
REFRESH (
    <tab> IN OUT DBMS_UTILITY.UNCL_ARRAY,
    <method> IN VARCHAR2 DEFAULT NULL,
    <rollback_seg> IN VARCHAR2 DEFAULT NULL,
    <push_deferred_rpc> IN BOOLEAN DEFAULT TRUE,
    <refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
    <purge_option> IN NUMBER DEFAULT 1,
    <parallelism> IN NUMBER DEFAULT 0,
    <heap_size> IN NUMBER DEFAULT 0,
    <atomic_refresh> IN BOOLEAN DEFAULT TRUE,
    <nested> IN BOOLEAN DEFAULT FALSE);
```

Parameters

<list>

<list> is a VARCHAR2 value that specifies the name of a materialized view, or a comma-separated list of materialized view names. The names may be schema-qualified.

<tab>

<tab> is a table of DBMS_UTILITY.UNCL_ARRAY values that specify the name (or names) of a materialized view.

<method>

<method> is a VARCHAR2 value that specifies the refresh method that will be applied to the specified view (or views). The only supported method is C; this performs a complete refresh of the view.

<rollback_seg>

<rollback_seg> is accepted for compatibility and ignored. The default is NULL.

<push_deferred_rpc>

<push_deferred_rpc> is accepted for compatibility and ignored. The default is TRUE.

<refresh_after_errors>

<refresh_after_errors> is accepted for compatibility and ignored. The default is FALSE.

<purge_option>

<purge_option> is accepted for compatibility and ignored. The default is 1.

<parallelism>

<parallelism> is accepted for compatibility and ignored. The default is 0.

<heap_size> IN NUMBER DEFAULT 0,

<heap_size> is accepted for compatibility and ignored. The default is 0.

<atomic_refresh>

<atomic_refresh> is accepted for compatibility and ignored. The default is TRUE.

<nested>

<nested> is accepted for compatibility and ignored. The default is FALSE.

Examples

The following example uses `DBMS_MVIEW.REFRESH` to perform a COMPLETE refresh on the `public.emp_view` materialized view:

```
EXEC DBMS_MVIEW.REFRESH(list => 'public.emp_view', method => 'C');
```

3.8.3 REFRESH_ALL_MVIEWS

Use the REFRESH_ALL_MVIEWS procedure to refresh any materialized views that have not been refreshed since the table or view on which the view depends has been modified. The signature is:

```
REFRESH_ALL_MVIEWS (
    <number_of_failures> OUT BINARY_INTEGER,
    <method> IN VARCHAR2 DEFAULT NULL,
    <rollback_seg> IN VARCHAR2 DEFAULT NULL,
    <refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
    <atomic_refresh> IN BOOLEAN DEFAULT TRUE);
```

Parameters

<number_of_failures>

<number_of_failures> is a BINARY_INTEGER that specifies the number of failures that occurred during the refresh operation.

<method>

<method> is a VARCHAR2 value that specifies the refresh method that will be applied to the specified view (or views). The only supported method is C; this performs a complete refresh of the view.

<rollback_seg>

<rollback_seg> is accepted for compatibility and ignored. The default is NULL.

<refresh_after_errors>

<refresh_after_errors> is accepted for compatibility and ignored. The default is FALSE.

<atomic_refresh>

<atomic_refresh> is accepted for compatibility and ignored. The default is TRUE.

Examples

The following example performs a COMPLETE refresh on all materialized views:

```
DECLARE
    errors INTEGER;
BEGIN
    DBMS_MVIEW.REFRESH_ALL_MVIEWS(errors, method => 'C');
END;
```

Upon completion, errors contains the number of failures.

3.8.4 REFRESH_DEPENDENT

Use the `REFRESH_DEPENDENT` procedure to refresh all material views that are dependent on the views specified in the call to the procedure. You can specify a comma-separated list or provide the view names in a table of `DBMS_UTILITY.UNCL_ARRAY` values.

Use the first form of the procedure to refresh all material views that are dependent on the views specified in a comma-separated list:

```
REFRESH_DEPENDENT (
    <number_of_failures> OUT BINARY_INTEGER,
    <list> IN VARCHAR2,
    <method> IN VARCHAR2 DEFAULT NULL,
    <rollback_seg> IN VARCHAR2 DEFAULT NULL
    <refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
    <atomic_refresh> IN BOOLEAN DEFAULT TRUE,
    <nested> IN BOOLEAN DEFAULT FALSE);
```

Use the second form of the procedure to refresh all material views that are dependent on the views specified in a table of `DBMS_UTILITY.UNCL_ARRAY` values:

```
REFRESH_DEPENDENT (
    <number_of_failures> OUT BINARY_INTEGER,
    <tab> IN DBMS_UTILITY.UNCL_ARRAY,
    <method> IN VARCHAR2 DEFAULT NULL,
    <rollback_seg> IN VARCHAR2 DEFAULT NULL,
    <refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
    <atomic_refresh> IN BOOLEAN DEFAULT TRUE,
    <nested> IN BOOLEAN DEFAULT FALSE);
```

Parameters

`<number_of_failures>`

`<number_of_failures>` is a `BINARY_INTEGER` that contains the number of failures that occurred during the refresh operation.

`<list>`

`<list>` is a `VARCHAR2` value that specifies the name of a materialized view, or a comma-separated list of materialized view names. The names may be schema-qualified.

`<tab>`

`<tab>` is a table of `DBMS_UTILITY.UNCL_ARRAY` values that specify the name (or names) of a materialized view.

`<method>`

`<method>` is a `VARCHAR2` value that specifies the refresh method that will be applied to the specified view (or views). The only supported method is `C`; this performs a complete refresh of the view.

`<rollback_seg>`

<rollback_seg> is accepted for compatibility and ignored. The default is NULL.

<refresh_after_errors>

<refresh_after_errors> is accepted for compatibility and ignored. The default is FALSE.

<atomic_refresh>

<atomic_refresh> is accepted for compatibility and ignored. The default is TRUE.

<nested>

<nested> is accepted for compatibility and ignored. The default is FALSE.

Examples

The following example performs a COMPLETE refresh on all materialized views dependent on a materialized view named emp_view that resides in the public schema:

```
DECLARE
    errors INTEGER;
BEGIN
    DBMS_MVIEW.REFRESH_DEPENDENT(errors, list => 'public.emp_view',
    method =>
    'C');
END;
```

Upon completion, errors contains the number of failures.

3.9 DBMS_OUTPUT

The DBMS_OUTPUT package provides the capability to send messages (lines of text) to a message buffer, or get messages from the message buffer. A message buffer is local to a single session. Use the DBMS_PIPE package to send messages between sessions.

The procedures and functions available in the DBMS_OUTPUT package are listed in the following table.

Function/Procedure	Return Type	Description
DISABLE	n/a	Disable the capability to send and receive messages.
ENABLE(<buffer_size>)	n/a	Enable the capability to send and receive messages.
GET_LINE(<line> OUT, <status> OUT)	n/a	Get a line from the message buffer.
GET_LINES(<lines> OUT, <numlines> IN OUT)	n/a	Get multiple lines from the message buffer.
NEW_LINE	n/a	Puts an end-of-line character sequence.
PUT(<item>)	n/a	Puts a partial line without an end-of-line character sequence.
PUT_LINE(<item>)	n/a	Puts a complete line with an end-of-line character sequence.
SERVEROUTPUT(<stdout>)	n/a	Direct messages from PUT, PUT_LINE, or NEW_LINE to either standard output or the message buffer.

The following table lists the public variables available in the DBMS_OUTPUT package.

Public Variables	Data Type	Value	Description
chararr	TABLE		For message lines.

3.9.1 CHARARR

The CHARARR is for storing multiple message lines.

```
TYPE chararr IS TABLE OF VARCHAR2(32767) INDEX BY BINARY_INTEGER;
```

3.9.2 DISABLE

The `DISABLE` procedure clears out the message buffer. Any messages in the buffer at the time the `DISABLE` procedure is executed will no longer be accessible. Any messages subsequently sent with the `PUT`, `PUT_LINE`, or `NEW_LINE` procedures are discarded. No error is returned to the sender when the `PUT`, `PUT_LINE`, or `NEW_LINE` procedures are executed and messages have been disabled.

Use the `ENABLE` procedure or `SERVEROUTPUT (TRUE)` procedure to re-enable the sending and receiving of messages.

`DISABLE`

Examples

This anonymous block disables the sending and receiving messages in the current session.

```
BEGIN
    DBMS_OUTPUT.DISABLE;
END;
```

3.9.3 ENABLE

The `ENABLE` procedure enables the capability to send messages to the message buffer or retrieve messages from the message buffer. Running `SERVEROUTPUT (TRUE)` also implicitly performs the `ENABLE` procedure.

The destination of a message sent with `PUT`, `PUT_LINE`, or `NEW_LINE` depends upon the state of `SERVEROUTPUT`.

- If the last state of `SERVEROUTPUT` is `TRUE`, the message goes to standard output of the command line.
- If the last state of `SERVEROUTPUT` is `FALSE`, the message goes to the message buffer.

```
ENABLE [ (<buffer_size> INTEGER) ]
```

Parameter

<buffer_size>

Maximum length of the message buffer in bytes. If a <buffer_size> of less than 2000 is specified, the buffer size is set to 2000.

Examples

The following anonymous block enables messages. Setting `SERVEROUTPUT (TRUE)` forces them to standard output.

```
BEGIN
  DBMS_OUTPUT.ENABLE;
  DBMS_OUTPUT.SERVEROUTPUT (TRUE);
  DBMS_OUTPUT.PUT_LINE ('Messages enabled');
END;

Messages enabled
```

The same effect could have been achieved by simply using `SERVEROUTPUT (TRUE)`.

```
BEGIN
  DBMS_OUTPUT.SERVEROUTPUT (TRUE);
  DBMS_OUTPUT.PUT_LINE ('Messages enabled');
END;

Messages enabled
```

The following anonymous block enables messages, but setting `SERVEROUTPUT (FALSE)` directs messages to the message buffer.

```
BEGIN
  DBMS_OUTPUT.ENABLE;
  DBMS_OUTPUT.SERVEROUTPUT (FALSE);
  DBMS_OUTPUT.PUT_LINE ('Message sent to buffer');
END;
```

3.9.4 GET_LINE

The `GET_LINE` procedure provides the capability to retrieve a line of text from the message buffer. Only text that has been terminated by an end-of-line character sequence is retrieved – that is complete lines generated using `PUT_LINE`, or by a series of `PUT` calls followed by a `NEW_LINE` call.

```
GET_LINE(<line> OUT VARCHAR2, <status> OUT INTEGER)
```

Parameters

<line>

Variable receiving the line of text from the message buffer.

<status>

0 if a line was returned from the message buffer, 1 if there was no line to return.

Examples

The following anonymous block writes the `emp` table out to the message buffer as a comma-delimited string for each row.

```
EXEC DBMS_OUTPUT.SERVEROUTPUT (FALSE);

DECLARE
    v_emprec          VARCHAR2(120);
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    DBMS_OUTPUT.ENABLE;
    FOR i IN emp_cur LOOP
        v_emprec := i.empno || ',' || i.ename || ',' || i.job || ',' ||
            NVL(LTRIM(TO_CHAR(i.mgr, '9999')), '') || ',' || i.hiredate ||
            ',' || i.sal || ',' ||
            NVL(LTRIM(TO_CHAR(i.comm, '9990.99')), '') || ',' || i.deptno;
        DBMS_OUTPUT.PUT_LINE(v_emprec);
    END LOOP;
END;
```

The following anonymous block reads the message buffer and inserts the messages written by the prior example into a table named `messages`. The rows in `messages` are then displayed.

```
CREATE TABLE messages (
    status          INTEGER,
    msg             VARCHAR2(100)
);

DECLARE
    v_line          VARCHAR2(100);
    v_status        INTEGER := 0;
BEGIN
    DBMS_OUTPUT.GET_LINE(v_line, v_status);
    WHILE v_status = 0 LOOP
        INSERT INTO messages VALUES(v_status, v_line);
        DBMS_OUTPUT.GET_LINE(v_line, v_status);
    END LOOP;
END;

SELECT msg FROM messages;
```

msg

```

-----
7369, SMITH, CLERK, 7902, 17-DEC-80 00:00:00, 800.00, , 20
7499, ALLEN, SALESMAN, 7698, 20-FEB-81 00:00:00, 1600.00, 300.00, 30
7521, WARD, SALESMAN, 7698, 22-FEB-81 00:00:00, 1250.00, 500.00, 30
7566, JONES, MANAGER, 7839, 02-APR-81 00:00:00, 2975.00, , 20
7654, MARTIN, SALESMAN, 7698, 28-SEP-81 00:00:00, 1250.00, 1400.00, 30
7698, BLAKE, MANAGER, 7839, 01-MAY-81 00:00:00, 2850.00, , 30
7782, CLARK, MANAGER, 7839, 09-JUN-81 00:00:00, 2450.00, , 10
7788, SCOTT, ANALYST, 7566, 19-APR-87 00:00:00, 3000.00, , 20
7839, KING, PRESIDENT, , 17-NOV-81 00:00:00, 5000.00, , 10
7844, TURNER, SALESMAN, 7698, 08-SEP-81 00:00:00, 1500.00, 0.00, 30
7876, ADAMS, CLERK, 7788, 23-MAY-87 00:00:00, 1100.00, , 20
7900, JAMES, CLERK, 7698, 03-DEC-81 00:00:00, 950.00, , 30
7902, FORD, ANALYST, 7566, 03-DEC-81 00:00:00, 3000.00, , 20
7934, MILLER, CLERK, 7782, 23-JAN-82 00:00:00, 1300.00, , 10
(14 rows)

```

3.9.5 GET_LINES

The `GET_LINES` procedure provides the capability to retrieve one or more lines of text from the message buffer into a collection. Only text that has been terminated by an end-of-line character sequence is retrieved – that is complete lines generated using `PUT_LINE`, or by a series of `PUT` calls followed by a `NEW_LINE` call.

```
GET_LINES(<lines> OUT CHARARR, <numlines> IN OUT INTEGER)
```

Parameters

<lines>

Table receiving the lines of text from the message buffer. See `CHARARR` for a description of <lines>.

<numlines> IN

Number of lines to be retrieved from the message buffer.

<numlines> OUT

Actual number of lines retrieved from the message buffer. If the output value of <numlines> is less than the input value, then there are no more lines left in the message buffer.

Examples

The following example uses the `GET_LINES` procedure to store all rows from the `emp` table that were placed on the message buffer, into an array.

```
EXEC DBMS_OUTPUT.SERVEROUTPUT(FALSE);

DECLARE
    v_emprec          VARCHAR2(120);
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    DBMS_OUTPUT.ENABLE;
    FOR i IN emp_cur LOOP
        v_emprec := i.empno || ',' || i.ename || ',' || i.job || ',' ||
            NVL(LTRIM(TO_CHAR(i.mgr,'9999')),') || ',' || i.hiredate ||
            ',' || i.sal || ',' ||
            NVL(LTRIM(TO_CHAR(i.comm,'9990.99')),') || ',' || i.deptno;
        DBMS_OUTPUT.PUT_LINE(v_emprec);
    END LOOP;
END;

DECLARE
    v_lines           DBMS_OUTPUT.CHARARR;
    v_numlines        INTEGER := 14;
    v_status           INTEGER := 0;
BEGIN
    DBMS_OUTPUT.GET_LINES(v_lines,v_numlines);
    FOR i IN 1..v_numlines LOOP
        INSERT INTO messages VALUES(v_numlines, v_lines(i));
    END LOOP;
END;

SELECT msg FROM messages;

-----
msg
-----
7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
```



```
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
(14 rows)
```

3.9.6 NEW_LINE

The NEW_LINE procedure writes an end-of-line character sequence in the message buffer.

NEW_LINE

Parameter

The NEW_LINE procedure expects no parameters.

3.9.7 PUT

The PUT procedure writes a string to the message buffer. No end-of-line character sequence is written at the end of the string. Use the NEW_LINE procedure to add an end-of-line character sequence.

```
PUT(<item> VARCHAR2)
```

Parameter

<item>

Text written to the message buffer.

Examples

The following example uses the PUT procedure to display a comma-delimited list of employees from the emp table.

```
DECLARE
  CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
  FOR i IN emp_cur LOOP
    DBMS_OUTPUT.PUT(i.empno);
    DBMS_OUTPUT.PUT(', ');
    DBMS_OUTPUT.PUT(i.ename);
    DBMS_OUTPUT.PUT(', ');
    DBMS_OUTPUT.PUT(i.job);
    DBMS_OUTPUT.PUT(', ');
    DBMS_OUTPUT.PUT(i.mgr);
    DBMS_OUTPUT.PUT(', ');
    DBMS_OUTPUT.PUT(i.hiredate);
    DBMS_OUTPUT.PUT(', ');
    DBMS_OUTPUT.PUT(i.sal);
    DBMS_OUTPUT.PUT(', ');
    DBMS_OUTPUT.PUT(i.comm);
    DBMS_OUTPUT.PUT(', ');
    DBMS_OUTPUT.PUT(i.deptno);
    DBMS_OUTPUT.NEW_LINE;
  END LOOP;
END;
```

```
7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
```

3.9.8 PUT_LINE

The PUT_LINE procedure writes a single line to the message buffer including an end-of-line character sequence.

```
PUT_LINE(<item> VARCHAR2)
```

Parameter

<item>

Text to be written to the message buffer.

Examples

The following example uses the PUT_LINE procedure to display a comma-delimited list of employees from the emp table.

```
DECLARE
    v_emprec          VARCHAR2(120);
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    FOR i IN emp_cur LOOP
        v_emprec := i.empno || ',' || i.ename || ',' || i.job || ',' ||
            NVL(LTRIM(TO_CHAR(i.mgr,'9999')),') || ',' || i.hiredate ||
            ',' || i.sal || ',' ||
            NVL(LTRIM(TO_CHAR(i.comm,'9990.99')),') || ',' || i.deptno;
        DBMS_OUTPUT.PUT_LINE(v_emprec);
    END LOOP;
END;
```

```
7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
```

3.9.9 SERVEROUTPUT

The `SERVEROUTPUT` procedure provides the capability to direct messages to standard output of the command line or to the message buffer. Setting `SERVEROUTPUT (TRUE)` also performs an implicit execution of `ENABLE`.

The default setting of `SERVEROUTPUT` is implementation dependent. For example, in Oracle SQL*Plus, `SERVEROUTPUT (FALSE)` is the default. In PSQL, `SERVEROUTPUT (TRUE)` is the default. Also note that in Oracle SQL*Plus, this setting is controlled using the SQL*Plus `SET` command, not by a stored procedure as implemented in Advanced Server.

```
SERVEROUTPUT (<stdout> BOOLEAN)
```

Parameter

<stdout>

Set to `TRUE` if subsequent `PUT`, `PUT_LINE`, or `NEW_LINE` commands are to send text directly to standard output of the command line. Set to `FALSE` if text is to be sent to the message buffer.

Examples

The following anonymous block sends the first message to the command line and the second message to the message buffer.

```
BEGIN
  DBMS_OUTPUT.SERVEROUTPUT (TRUE);
  DBMS_OUTPUT.PUT_LINE ('This message goes to the command line');
  DBMS_OUTPUT.SERVEROUTPUT (FALSE);
  DBMS_OUTPUT.PUT_LINE ('This message goes to the message buffer');
END;
```

This message goes to the command line

If within the same session, the following anonymous block is executed, the message stored in the message buffer from the prior example is flushed and displayed on the command line as well as the new message.

```
BEGIN
  DBMS_OUTPUT.SERVEROUTPUT (TRUE);
  DBMS_OUTPUT.PUT_LINE ('Flush messages from the buffer');
END;
```

This message goes to the message buffer
Flush messages from the buffer

3.10 DBMS_PIPE

The DBMS_PIPE package provides the capability to send messages through a pipe within or between sessions connected to the same database cluster.

The procedures and functions available in the DBMS_PIPE package are listed in the following table:

Function/Procedure	Return Type	Description
CREATE_PIPE(<pipename> [, <maxpipesize>] [, <private>])	INTEGER	Explicitly create a private pipe if <private> is “true” (the default) or a public pipe if <private> is “false”.
NEXT_ITEM_TYPE	INTEGER	Determine the data type of the next item in a received message.
PACK_MESSAGE(<item>)	n/a	Place <item> in the session’s local message buffer.
PURGE(<pipename>)	n/a	Remove unreceived messages from the specified pipe.
RECEIVE_MESSAGE(<pipename> [, <timeout>])	INTEGER	Get a message from a specified pipe.
REMOVE_PIPE(<pipename>)	INTEGER	Delete an explicitly created pipe.
RESET_BUFFER	n/a	Reset the local message buffer.
SEND_MESSAGE(<pipename> [, <timeout>] [, <maxpipesize>])	INTEGER	Send a message on a pipe.
UNIQUE_SESSION_NAME	VARCHAR2	Obtain a unique session name.
UNPACK_MESSAGE(<item> OUT)	n/a	Retrieve the next data item from a message into a type-compatible variable, <item>.

Pipes are categorized as implicit or explicit. An *implicit pipe* is created if a reference is made to a pipe name that was not previously created by the CREATE_PIPE function. For example, if the SEND_MESSAGE function is executed using a non-existent pipe name, a new implicit pipe is created with that name. An *explicit pipe* is created using the CREATE_PIPE function whereby the first parameter specifies the pipe name for the new pipe.

Pipes are also categorized as private or public. A *private pipe* can only be accessed by the user who created the pipe. Even a superuser cannot access a private pipe that was created by another user. A *public pipe* can be accessed by any user who has access to the DBMS_PIPE package.

A public pipe can only be created by using the CREATE_PIPE function with the third parameter set to FALSE. The CREATE_PIPE function can be used to create a private pipe by setting the third parameter to TRUE or by omitting the third parameter. All implicit pipes are private.

The individual data items or “lines” of a message are first built-in a *local message buffer*, unique to the current session. The PACK_MESSAGE procedure builds the message in the session’s local message buffer. The SEND_MESSAGE function is then used to send the message through the pipe.

Receipt of a message involves the reverse operation. The RECEIVE_MESSAGE function is used to get a message from the specified pipe. The message is written to the session’s local message buffer. The UNPACK_MESSAGE procedure is then used to transfer the message data items from the message buffer to program variables. If a pipe contains multiple messages, RECEIVE_MESSAGE gets the messages in *FIFO* (first-in-first-out) order.

Each session maintains separate message buffers for messages created with the PACK_MESSAGE procedure and messages retrieved by the RECEIVE_MESSAGE function. Thus messages can be both built and received in the same session. However, if consecutive RECEIVE_MESSAGE calls are made, only the message from the last RECEIVE_MESSAGE call will be preserved in the local message buffer.

3.10.1 CREATE_PIPE

The CREATE_PIPE function creates an explicit public pipe or an explicit private pipe with a specified name.

```
<status> INTEGER CREATE_PIPE(<pipename> VARCHAR2
                             [, <maxpipesize> INTEGER ] [, <private> BOOLEAN ])
```

Parameters

<pipename>

Name of the pipe.

<maxpipesize>

Maximum capacity of the pipe in bytes. Default is 8192 bytes.

<private>

Create a public pipe if set to FALSE. Create a private pipe if set to TRUE . This is the default.

<status>

Status code returned by the operation. 0 indicates successful creation.

Examples

The following example creates a private pipe named messages :

```
DECLARE
    v_status      INTEGER;
BEGIN
    v_status := DBMS_PIPE.CREATE_PIPE('messages');
    DBMS_OUTPUT.PUT_LINE('CREATE_PIPE status: ' || v_status);
END;
CREATE_PIPE status: 0
```

The following example creates a public pipe named mailbox :

```
DECLARE
    v_status      INTEGER;
BEGIN
    v_status := DBMS_PIPE.CREATE_PIPE('mailbox',8192,FALSE);
    DBMS_OUTPUT.PUT_LINE('CREATE_PIPE status: ' || v_status);
END;
CREATE_PIPE status: 0
```

3.10.2 NEXT_ITEM_TYPE

The `NEXT_ITEM_TYPE` function returns an integer code identifying the data type of the next data item in a message that has been retrieved into the session's local message buffer. As each item is moved off of the local message buffer with the `UNPACK_MESSAGE` procedure, the `NEXT_ITEM_TYPE` function will return the data type code for the next available item. A code of 0 is returned when there are no more items left in the message.

<typecode> INTEGER `NEXT_ITEM_TYPE`

Parameters

<typecode>

Code identifying the data type of the next data item as shown in the following table.

Type Code	Data Type
0	No more data items
9	NUMBER
11	VARCHAR2
13	DATE
23	RAW

Note: The type codes list in the table are not compatible with Oracle databases. Oracle assigns a different numbering sequence to the data types.

Examples

The following example shows a pipe packed with a NUMBER item, a VARCHAR2 item, a DATE item, and a RAW item. A second anonymous block then uses the `NEXT_ITEM_TYPE` function to display the type code of each item.

```

DECLARE
    v_number          NUMBER := 123;
    v_varchar         VARCHAR2(20) := 'Character data';
    v_date            DATE := SYSDATE;
    v_raw             RAW(4) := '21222324';
    v_status          INTEGER;
BEGIN
    DBMS_PIPE.PACK_MESSAGE(v_number);
    DBMS_PIPE.PACK_MESSAGE(v_varchar);
    DBMS_PIPE.PACK_MESSAGE(v_date);
    DBMS_PIPE.PACK_MESSAGE(v_raw);
    v_status := DBMS_PIPE.SEND_MESSAGE('datatypes');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);
EXCEPTION
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
        DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;

SEND_MESSAGE status: 0

DECLARE
    v_number          NUMBER;
    v_varchar         VARCHAR2(20);
    v_date            DATE;
    v_timestamp       TIMESTAMP;

```



```

v_raw          RAW(4);
v_status       INTEGER;
BEGIN
  v_status := DBMS_PIPE.RECEIVE_MESSAGE('datatypes');
  DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
  DBMS_OUTPUT.PUT_LINE('-----');

  v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
  DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
  DBMS_PIPE.UNPACK_MESSAGE(v_number);
  DBMS_OUTPUT.PUT_LINE('NUMBER Item    : ' || v_number);
  DBMS_OUTPUT.PUT_LINE('-----');

  v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
  DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
  DBMS_PIPE.UNPACK_MESSAGE(v_varchar);
  DBMS_OUTPUT.PUT_LINE('VARCHAR2 Item : ' || v_varchar);
  DBMS_OUTPUT.PUT_LINE('-----');

  v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
  DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
  DBMS_PIPE.UNPACK_MESSAGE(v_date);
  DBMS_OUTPUT.PUT_LINE('DATE Item      : ' || v_date);
  DBMS_OUTPUT.PUT_LINE('-----');

  v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
  DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
  DBMS_PIPE.UNPACK_MESSAGE(v_raw);
  DBMS_OUTPUT.PUT_LINE('RAW Item       : ' || v_raw);
  DBMS_OUTPUT.PUT_LINE('-----');

  v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
  DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
  DBMS_OUTPUT.PUT_LINE('-----');
EXCEPTION
  WHEN OTHERS THEN
    DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
    DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;

RECEIVE_MESSAGE status: 0
-----
NEXT_ITEM_TYPE: 9
NUMBER Item    : 123
-----
NEXT_ITEM_TYPE: 11
VARCHAR2 Item : Character data
-----
NEXT_ITEM_TYPE: 13
DATE Item      : 02-OCT-07 11:11:43
-----
NEXT_ITEM_TYPE: 23
RAW Item       : 21222324
-----
NEXT_ITEM_TYPE: 0

```

3.10.3 PACK_MESSAGE

The `PACK_MESSAGE` procedure places an item of data in the session's local message buffer. `PACK_MESSAGE` must be executed at least once before issuing a `SEND_MESSAGE` call.

```
PACK_MESSAGE(<item> { DATE | NUMBER | VARCHAR2 | RAW })
```

Use the `UNPACK_MESSAGE` procedure to obtain data items once the message is retrieved using a `RECEIVE_MESSAGE` call.

Parameters

<item>

An expression evaluating to any of the acceptable parameter data types. The value is added to the session's local message buffer.

3.10.4 PURGE

The PURGE procedure removes the unreceived messages from a specified implicit pipe.

PURGE (<pipename> VARCHAR2)

Use the REMOVE_PIPE function to delete an explicit pipe.

Parameter

<pipename>

Name of the pipe.

Examples

Two messages are sent on a pipe:

```
DECLARE
    v_status      INTEGER;
BEGIN
    DBMS_PIPE.PACK_MESSAGE('Message #1');
    v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);

    DBMS_PIPE.PACK_MESSAGE('Message #2');
    v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);
END;

SEND_MESSAGE status: 0
SEND_MESSAGE status: 0
```

Receive the first message and unpack it:

```
DECLARE
    v_item        VARCHAR2(80);
    v_status      INTEGER;
BEGIN
    v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
    DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
    DBMS_PIPE.UNPACK_MESSAGE(v_item);
    DBMS_OUTPUT.PUT_LINE('Item: ' || v_item);
END;

RECEIVE_MESSAGE status: 0
Item: Message #1
```

Purge the pipe:

```
EXEC DBMS_PIPE.PURGE('pipe');
```

Try to retrieve the next message. The `RECEIVE_MESSAGE` call returns status code 1 indicating it timed out because no message was available.

```
DECLARE
    v_item          VARCHAR2(80);
    v_status        INTEGER;
BEGIN
    v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
    DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
END;

RECEIVE_MESSAGE status: 1
```

3.10.5 RECEIVE_MESSAGE

The RECEIVE_MESSAGE function obtains a message from a specified pipe.

```
<status> INTEGER RECEIVE_MESSAGE(<pipename> VARCHAR2  
    [, <timeout> INTEGER ])
```

Parameters

<pipename>

Name of the pipe.

<timeout>

Wait time (seconds). Default is 86400000 (1000 days).

<status>

Status code returned by the operation.

The possible status codes are:

Status Code	Description
0	Success
1	Time out
2	Message too large .for the buffer

3.10.6 REMOVE_PIPE

The REMOVE_PIPE function deletes an explicit private or explicit public pipe.

<status> INTEGER REMOVE_PIPE(<pipename> VARCHAR2)

Use the REMOVE_PIPE function to delete explicitly created pipes – i.e., pipes created with the CREATE_PIPE function.

Parameters

<pipename>

Name of the pipe.

<status>

Status code returned by the operation. A status code of 0 is returned even if the named pipe is non-existent.

Examples

Two messages are sent on a pipe:

```
DECLARE
    v_status          INTEGER;
BEGIN
    v_status := DBMS_PIPE.CREATE_PIPE('pipe');
    DBMS_OUTPUT.PUT_LINE('CREATE_PIPE status : ' || v_status);

    DBMS_PIPE.PACK_MESSAGE('Message #1');
    v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);

    DBMS_PIPE.PACK_MESSAGE('Message #2');
    v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);
END;

CREATE_PIPE status : 0
SEND_MESSAGE status: 0
SEND_MESSAGE status: 0
```

Receive the first message and unpack it:

```
DECLARE
    v_item            VARCHAR2(80);
    v_status          INTEGER;
BEGIN
    v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
    DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
    DBMS_PIPE.UNPACK_MESSAGE(v_item);
    DBMS_OUTPUT.PUT_LINE('Item: ' || v_item);
END;

RECEIVE_MESSAGE status: 0
Item: Message #1
```

Remove the pipe:

```
SELECT DBMS_PIPE.REMOVE_PIPE('pipe') FROM DUAL;
```

```
remove_pipe
```

```
-----
```

```
0
```

```
(1 row)
```

Try to retrieve the next message. The `RECEIVE_MESSAGE` call returns status code 1 indicating it timed out because the pipe had been deleted.

```
DECLARE
    v_item          VARCHAR2(80);
    v_status        INTEGER;
BEGIN
    v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
    DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
END;
```

RECEIVE_MESSAGE status: 1

3.10.7 RESET_BUFFER

The RESET_BUFFER procedure resets a “pointer” to the session’s local message buffer back to the beginning of the buffer. This has the effect of causing subsequent PACK_MESSAGE calls to overwrite any data items that existed in the message buffer prior to the RESET_BUFFER call.

RESET_BUFFER

Examples

A message to John is written to the local message buffer. It is replaced by a message to Bob by calling RESET_BUFFER. The message is sent on the pipe.

```
DECLARE
    v_status      INTEGER;
BEGIN
    DBMS_PIPE.PACK_MESSAGE('Hi, John');
    DBMS_PIPE.PACK_MESSAGE('Can you attend a meeting at 3:00, today?');
    DBMS_PIPE.PACK_MESSAGE('If not, is tomorrow at 8:30 ok with you?');
    DBMS_PIPE.RESET_BUFFER;
    DBMS_PIPE.PACK_MESSAGE('Hi, Bob');
    DBMS_PIPE.PACK_MESSAGE('Can you attend a meeting at 9:30, tomorrow?');
    v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);
END;

SEND_MESSAGE status: 0
```

The message to Bob is in the received message.

```
DECLARE
    v_item        VARCHAR2(80);
    v_status      INTEGER;
BEGIN
    v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
    DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
    DBMS_PIPE.UNPACK_MESSAGE(v_item);
    DBMS_OUTPUT.PUT_LINE('Item: ' || v_item);
    DBMS_PIPE.UNPACK_MESSAGE(v_item);
    DBMS_OUTPUT.PUT_LINE('Item: ' || v_item);
END;

RECEIVE_MESSAGE status: 0
Item: Hi, Bob
Item: Can you attend a meeting at 9:30, tomorrow?
```


3.10.8 SEND_MESSAGE

The SEND_MESSAGE function sends a message from the session's local message buffer to the specified pipe.

```
<status> SEND_MESSAGE(<pipename> VARCHAR2 [, <timeout> INTEGER ]  
                        [, <maxpipesize> INTEGER ])
```

Parameters

<pipename>

Name of the pipe.

<timeout>

Wait time (seconds). Default is 86400000 (1000 days).

<maxpipesize>

Maximum capacity of the pipe in bytes. Default is 8192 bytes.

<status>

Status code returned by the operation.

The possible status codes are:

Status Code	Description
0	Success
1	Time out
3	Function interrupted

3.10.9 UNIQUE_SESSION_NAME

The UNIQUE_SESSION_NAME function returns a name, unique to the current session.

<name> VARCHAR2 UNIQUE_SESSION_NAME

Parameters

<name>

Unique session name.

Examples

The following anonymous block retrieves and displays a unique session name.

```
DECLARE
    v_session      VARCHAR2(30);
BEGIN
    v_session := DBMS_PIPE.UNIQUE_SESSION_NAME;
    DBMS_OUTPUT.PUT_LINE('Session Name: ' || v_session);
END;
```

Session Name: PG\$PIPE\$5\$2752

3.10.10 UNPACK_MESSAGE

The UNPACK_MESSAGE procedure copies the data items of a message from the local message buffer to a specified program variable. The message must be placed in the local message buffer with the RECEIVE_MESSAGE function before using UNPACK_MESSAGE.

```
UNPACK_MESSAGE(<item> OUT { DATE | NUMBER | VARCHAR2 | RAW })
```

Parameter

<item>

Type-compatible variable that receives a data item from the local message buffer.

3.10.11 Comprehensive Example

The following example uses a pipe as a “mailbox”. The procedures to create the mailbox, add a multi-item message to the mailbox (up to three items), and display the full contents of the mailbox are enclosed in a package named, mailbox.

```
CREATE OR REPLACE PACKAGE mailbox
IS
    PROCEDURE create_mailbox;
    PROCEDURE add_message (
        p_mailbox    VARCHAR2,
        p_item_1     VARCHAR2,
        p_item_2     VARCHAR2 DEFAULT 'END',
        p_item_3     VARCHAR2 DEFAULT 'END'
    );
    PROCEDURE empty_mailbox (
        p_mailbox    VARCHAR2,
        p_waittime   INTEGER DEFAULT 10
    );
END mailbox;

CREATE OR REPLACE PACKAGE BODY mailbox
IS
    PROCEDURE create_mailbox
    IS
        v_mailbox    VARCHAR2(30);
        v_status     INTEGER;
    BEGIN
        v_mailbox := DBMS_PIPE.UNIQUE_SESSION_NAME;
        v_status := DBMS_PIPE.CREATE_PIPE(v_mailbox,1000,FALSE);
        IF v_status = 0 THEN
            DBMS_OUTPUT.PUT_LINE('Created mailbox: ' || v_mailbox);
        ELSE
            DBMS_OUTPUT.PUT_LINE('CREATE_PIPE failed - status: ' ||
                                v_status);
        END IF;
    END create_mailbox;

    PROCEDURE add_message (
        p_mailbox    VARCHAR2,
        p_item_1     VARCHAR2,
        p_item_2     VARCHAR2 DEFAULT 'END',
        p_item_3     VARCHAR2 DEFAULT 'END'
    )
    IS
        v_item_cnt   INTEGER := 0;
        v_status     INTEGER;
    BEGIN
        DBMS_PIPE.PACK_MESSAGE(p_item_1);
        v_item_cnt := 1;
        IF p_item_2 != 'END' THEN
            DBMS_PIPE.PACK_MESSAGE(p_item_2);
            v_item_cnt := v_item_cnt + 1;
        END IF;
        IF p_item_3 != 'END' THEN
            DBMS_PIPE.PACK_MESSAGE(p_item_3);
            v_item_cnt := v_item_cnt + 1;
        END IF;
    END add_message;
END mailbox;
```

```

v_status := DBMS_PIPE.SEND_MESSAGE(p_mailbox);
IF v_status = 0 THEN
    DBMS_OUTPUT.PUT_LINE('Added message with ' || v_item_cnt ||
        ' item(s) to mailbox ' || p_mailbox);
ELSE
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE in add_message failed - ' ||
        'status: ' || v_status);
END IF;
END add_message;

PROCEDURE empty_mailbox (
    p_mailbox    VARCHAR2,
    p_waittime   INTEGER DEFAULT 10
)
IS
    v_msgno      INTEGER DEFAULT 0;
    v_itemno     INTEGER DEFAULT 0;
    v_item       VARCHAR2(100);
    v_status     INTEGER;
BEGIN
    v_status := DBMS_PIPE.RECEIVE_MESSAGE(p_mailbox,p_waittime);
    WHILE v_status = 0 LOOP
        v_msgno := v_msgno + 1;
        DBMS_OUTPUT.PUT_LINE('***** Start message #' || v_msgno ||
            ' *****');
        BEGIN
            LOOP
                v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
                EXIT WHEN v_status = 0;
                DBMS_PIPE.UNPACK_MESSAGE(v_item);
                v_itemno := v_itemno + 1;
                DBMS_OUTPUT.PUT_LINE('Item #' || v_itemno || ': ' ||
                    v_item);
            END LOOP;
            DBMS_OUTPUT.PUT_LINE('***** End message #' || v_msgno ||
                ' *****');
            DBMS_OUTPUT.PUT_LINE('*');
            v_itemno := 0;
            v_status := DBMS_PIPE.RECEIVE_MESSAGE(p_mailbox,1);
        END;
    END LOOP;
    DBMS_OUTPUT.PUT_LINE('Number of messages received: ' || v_msgno);
    v_status := DBMS_PIPE.REMOVE_PIPE(p_mailbox);
    IF v_status = 0 THEN
        DBMS_OUTPUT.PUT_LINE('Deleted mailbox ' || p_mailbox);
    ELSE
        DBMS_OUTPUT.PUT_LINE('Could not delete mailbox - status: '
            || v_status);
    END IF;
END empty_mailbox;
END mailbox;

```

The following demonstrates the execution of the procedures in `mailbox`. The first procedure creates a public pipe using a name generated by the `UNIQUE_SESSION_NAME` function.

```
EXEC mailbox.create_mailbox;  
Created mailbox: PG$PIPE$13$3940
```

Using the mailbox name, any user in the same database with access to the `mailbox` package and `DBMS_PIPE` package can add messages:

```
EXEC mailbox.add_message('PG$PIPE$13$3940','Hi, John','Can you attend a meeting at_  
↪3:00, today?','-- Mary');  
  
Added message with 3 item(s) to mailbox PG$PIPE$13$3940  
  
EXEC mailbox.add_message('PG$PIPE$13$3940','Don''t forget to submit your report',  
↪'Thanks,','-- Joe');  
  
Added message with 3 item(s) to mailbox PG$PIPE$13$3940
```

Finally, the contents of the mailbox can be emptied:

```
EXEC mailbox.empty_mailbox('PG$PIPE$13$3940');  
  
***** Start message #1 *****  
Item #1: Hi, John  
Item #2: Can you attend a meeting at 3:00, today?  
Item #3: -- Mary  
***** End message #1 *****  
*  
***** Start message #2 *****  
Item #1: Don't forget to submit your report  
Item #2: Thanks,  
Item #3: Joe  
***** End message #2 *****  
*  
Number of messages received: 2  
Deleted mailbox PG$PIPE$13$3940
```

3.11 DBMS_PROFILER

The DBMS_PROFILER package collects and stores performance information about the PL/pgSQL and SPL statements that are executed during a performance profiling session; use the functions and procedures listed below to control the profiling tool.

Function/Procedure	Return Type	Description
FLUSH_DATA	Status Code or Exception	Flushes performance data collected in the current session without terminating the session (profiling continues).
GET_VERSION(<major> OUT, <minor> OUT)	n/a	Returns the version number of this package.
INTERNAL_VERSION_CHECK	Status Code	Confirms that the current version of the profiler will work with the current database.
PAUSE_PROFILER	Status Code or Exception	Pause data collection.
RESUME_PROFILER	Status Code or Exception	Resume data collection.
START_PROFILER(<run_comment>, <run_comment1> [, <run_number> OUT])	Status Code or Exception	Start data collection.
STOP_PROFILER	Status Code or Exception	Stop data collection and flush performance data to the PLSQL_PROFILER_RAWDATA table.

The functions within the DBMS_PROFILER package return a status code to indicate success or failure; the DBMS_PROFILER procedures raise an exception only if they encounter a failure. The status codes and messages returned by the functions, and the exceptions raised by the procedures are listed in the table below.

Status Code	Message	Exception	Description
-1	error version	version_mismatch	The profiler version and the database are incompatible.
0	success	n/a	The operation completed successfully.
1	error_param	profiler_error	The operation received an incorrect parameter.
2	error_io	profiler_error	The data flush operation has failed.

3.11.1 FLUSH_DATA

The `FLUSH_DATA` function/procedure flushes the data collected in the current session without terminating the profiler session. The data is flushed to the tables described in the Advanced Server Performance Features Guide. The function and procedure signatures are:

```
<status> INTEGER FLUSH_DATA  
FLUSH_DATA
```

Parameter

<status>

Status code returned by the operation.

3.11.2 GET_VERSION

The `GET_VERSION` procedure returns the version of `DBMS_PROFILER`. The procedure signature is:

```
GET_VERSION(<major> OUT INTEGER, <minor> OUT INTEGER)
```

Parameters

<major>

The major version number of `DBMS_PROFILER`.

<minor>

The minor version number of `DBMS_PROFILER`.

3.11.3 INTERNAL_VERSION_CHECK

The `INTERNAL_VERSION_CHECK` function confirms that the current version of `DBMS_PROFILER` will work with the current database. The function signature is:

```
<status> INTEGER INTERNAL_VERSION_CHECK
```

Parameter

<status>

Status code returned by the operation.

3.11.4 PAUSE_PROFILER

The `PAUSE_PROFILER` function/procedure pauses a profiling session. The function and procedure signatures are:

```
<status> INTEGER PAUSE_PROFILER  
PAUSE_PROFILER
```

Parameter

<status>

Status code returned by the operation.

3.11.5 RESUME_PROFILER

The RESUME_PROFILER function/procedure pauses a profiling session. The function and procedure signatures are:

```
<status> INTEGER RESUME_PROFILER
RESUME_PROFILER
```

Parameter

<status>
Status code returned by the operation.

3.11.6 START_PROFILER

The START_PROFILER function/procedure starts a data collection session. The function and procedure signatures are:

```
<status> INTEGER START_PROFILER(<run_comment> TEXT := SYSDATE,
    <run_comment1> TEXT := '' [, <run_number> OUT INTEGER ])
START_PROFILER(<run_comment> TEXT := SYSDATE,
    <run_comment1> TEXT := '' [, <run_number> OUT INTEGER ])
```

Parameters

<run_comment>
A user-defined comment for the profiler session. The default value is SYSDATE.

<run_comment1>
An additional user-defined comment for the profiler session. The default value is ''.

<run_number>
The session number of the profiler session.

<status>
Status code returned by the operation.

3.11.7 STOP_PROFILER

The STOP_PROFILER function/procedure stops a profiling session and flushes the performance information to the DBMS_PROFILER tables and view. The function and procedure signatures are:

```
<status> INTEGER STOP_PROFILER
STOP_PROFILER
```

Parameter

<status>
Status code returned by the operation.

3.11.8 Using DBMS_PROFILER

The `DBMS_PROFILER` package collects and stores performance information about the PL/pgSQL and SPL statements that are executed during a profiling session; you can review the performance information in the tables and views provided by the profiler.

`DBMS_PROFILER` works by recording a set of performance-related counters and timers for each line of PL/pgSQL or SPL statement that executes within a profiling session. The counters and timers are stored in a table named `SYS.PLSQL_PROFILER_DATA`. When you complete a profiling session, `DBMS_PROFILER` will write a row to the performance statistics table for each line of PL/pgSQL or SPL code that executed within the session. For example, if you execute the following function:

```
1 - CREATE OR REPLACE FUNCTION getBalance(acctNumber INTEGER)
2 - RETURNS NUMERIC AS $$
3 - DECLARE
4 -   result NUMERIC;
5 - BEGIN
6 -   SELECT INTO result balance FROM acct WHERE id = acctNumber;
7 -
8 -   IF (result IS NULL) THEN
9 -     RAISE INFO 'Balance is null';
10-   END IF;
11-
12-   RETURN result;
13- END;
14- $$ LANGUAGE 'plpgsql';
```

`DBMS_PROFILER` adds one `PLSQL_PROFILER_DATA` entry for each line of code within the `getBalance()` function (including blank lines and comments). The entry corresponding to the `SELECT` statement executed exactly one time; and required a very small amount of time to execute. On the other hand, the entry corresponding to the `RAISE INFO` statement executed once or not at all (depending on the value for the `balance` column).

Some of the lines in this function contain no executable code so the performance statistics for those lines will always contain zero values.

To start a profiling session, invoke the `DBMS_PROFILER.START_PROFILER` function (or procedure). Once you've invoked `START_PROFILER`, Advanced Server will profile every PL/pgSQL or SPL function, procedure, trigger, or anonymous block that your session executes until you either stop or pause the profiler (by calling `STOP_PROFILER` or `PAUSE_PROFILER`).

It is important to note that when you start (or resume) the profiler, the profiler will only gather performance statistics for functions/procedures/triggers that start after the call to `START_PROFILER` (or `RESUME_PROFILER`).

While the profiler is active, Advanced Server records a large set of timers and counters in memory; when you invoke the `STOP_PROFILER` (or `FLUSH_DATA`) function/procedure, `DBMS_PROFILER` writes those timers and counters to a set of three tables:

- `SYS.PLSQL_PROFILER_RAWDATA`

Contains the performance counters and timers for each statement executed within the session.

- `SYS.PLSQL_PROFILER_RUNS`

Contains a summary of each run (aggregating the information found in `PLSQL_PROFILER_RAWDATA`).

- `SYS.PLSQL_PROFILER_UNITS`

Contains a summary of each code unit (function, procedure, trigger, or anonymous block) executed within a session.

In addition, `DBMS_PROFILER` defines a view, `SYS.PLSQL_PROFILER_DATA`, which contains a subset of the `PLSQL_PROFILER_RAWDATA` table.

Please note that a non-superuser may gather profiling information, but may not view that profiling information unless a superuser grants specific privileges on the profiling tables (stored in the `SYS` schema). This permits a non-privileged user to gather performance statistics without exposing information that the administrator may want to keep secret.

3.11.8.1 Querying the DBMS_PROFILER Tables and View

The following step-by-step example uses DBMS_PROFILER to retrieve performance information for procedures, functions, and triggers included in the sample data distributed with Advanced Server.

1. Open the EDB-PSQL command line, and establish a connection to the Advanced Server database. Use an EXEC statement to start the profiling session:

```
acctg=# EXEC dbms_profiler.start_profiler('profile list_emp');

EDB-SPL Procedure successfully completed
```

Note: (The call to `start_profiler()` includes a comment that DBMS_PROFILER associates with the profiler session).

2. Then call the `list_emp` function:

```
acctg=# SELECT list_emp();
INFO:  EMPNO      ENAME
INFO:  -----
INFO:  7369      SMITH
INFO:  7499      ALLEN
INFO:  7521      WARD
INFO:  7566      JONES
INFO:  7654      MARTIN
INFO:  7698      BLAKE
INFO:  7782      CLARK
INFO:  7788      SCOTT
INFO:  7839      KING
INFO:  7844      TURNER
INFO:  7876      ADAMS
INFO:  7900      JAMES
INFO:  7902      FORD
INFO:  7934      MILLER
list_emp
-----

(1 row)
```

3. Stop the profiling session with a call to `dbms_profiler.stop_profiler`:

```
acctg=# EXEC dbms_profiler.stop_profiler;

EDB-SPL Procedure successfully completed
```

4. Start a new session with the `dbms_profiler.start_profiler` function (followed by a new comment):

```
acctg=# EXEC dbms_profiler.start_profiler('profile get_dept_name and
emp_sal_trig');

EDB-SPL Procedure successfully completed
```

5. Invoke the `get_dept_name` function:

```
acctg=# SELECT get_dept_name(10);
       get_dept_name
-----
ACCOUNTING
(1 row)
```

6. Execute an UPDATE statement that causes a trigger to execute:

```

acctg=# UPDATE memp SET sal = 500 WHERE empno = 7902;
INFO: Updating employee 7902
INFO: ..Old salary: 3000.00
INFO: ..New salary: 500.00
INFO: ..Raise: -2500.00
INFO: User enterprisedb updated employee(s) on 04-FEB-14
UPDATE 1

```

7. Terminate the profiling session and flush the performance information to the profiling tables:

```
acctg=# EXEC dbms_profiler.stop_profiler;
```

EDB-SPL Procedure successfully completed

8. Now, query the `plsql_profiler_runs` table to view a list of the profiling sessions, arranged by `runid`:

```

acctg=# SELECT * FROM plsql_profiler_runs;
 runid | related_run | run_owner | run_date | run_
      ↳ comment
|run_total_time | run_system_info | run_comment1 | spare1
-----+-----+-----+-----+-----+-----+-----+-----
      ↳ -----+-----+-----+-----+-----+-----+-----
      ↳ 1 | | enterprisedb | 04-FEB-14 09:32:48.874315 | profile list_emp
      ↳ 4154 | |
      ↳ 2 | | enterprisedb | 04-FEB-14 09:41:30.546503 | profile get_dept_
      ↳ name and emp_sal_trig | 2088 | |
(2 rows)

```

9. Query the `plsql_profiler_units` table to view the amount of time consumed by each unit (each function, procedure, or trigger):

```

acctg=# SELECT * FROM plsql_profiler_units;
   runid | unit_number | unit_type | unit_owner |          unit_name          |
   unit_timestamp | total_time | spare1 | spare2
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
   1 |      16999 | FUNCTION | enterprisedb | list_emp() |
   2 |      17002 | FUNCTION | enterprisedb | user_audit_trig() |
   2 |      17000 | FUNCTION | enterprisedb | get_dept_name(p_deptno numeric) |
   2 |      17004 | FUNCTION | enterprisedb | emp_sal_trig() |
(4 rows)

```

10. Query the `plsql_profiler_rawdata` table to view a list of the wait event counters and wait event times:

```
acctg=# SELECT runid, sourcecode, func_oid, line_number, exec_count, tuples_returned,
↳time_total FROM plsql_profiler_rawdata;
```

runid	sourcecode				func_oid
↳line_number	exec_count	tuples_returned	time_total		
-----+-----+-----+-----+-----+-----					
↳ 1	DECLARE				16999
↳ 1	1	0	0	0	
↳ 1	v_empno	NUMERIC(4);			16999
↳ 2	0	0	0	0	
↳ 1	v_ename	VARCHAR(10);			16999
↳ 3	0	0	0	0	
↳ 1	emp_cur	CURSOR FOR			16999
↳ 4	0	0	0	0	
↳ 1	SELECT empno, ename	FROM memp ORDER BY empno;			16999
↳ 5	0	0	0	0	
↳ 1	BEGIN				16999
↳ 6	0	0	0	0	
↳ 1	OPEN emp_cur;				16999
↳ 7	0	0	0	0	
↳ 1	RAISE INFO 'EMPNO	ENAME';			16999
↳ 8	1	0	0.001621		
↳ 1	RAISE INFO '-----	-----';			16999
↳ 9	1	0	0.000301		
↳ 1	LOOP				16999
↳ 10	1	0	4.6e-05		
↳ 1	FETCH emp_cur	INTO v_empno, v_ename;			16999
↳ 11	1	0	0.001114		
↳ 1	EXIT WHEN NOT FOUND;				16999
↳ 12	15	0	0.000206		
↳ 1	RAISE INFO '%	%, v_empno, v_ename;			16999
↳ 13	15	0	8.3e-05		
↳ 1	END LOOP;				16999
↳ 14	14	0	0.000773		
↳ 1	CLOSE emp_cur;				16999
↳ 15	0	0	0		
↳ 1	RETURN;				16999
↳ 16	1	0	1e-05		
↳ 1	END;				16999
↳ 17	1	0	0		
↳ 1					16999
↳ 18	0	0	0		
↳ 2	DECLARE				17002
↳ 1	0	0	0		
↳ 2	v_action	VARCHAR(24);			17002
↳ 2	0	0	0		
↳ 2	v_text	TEXT;			17002
↳ 3	0	0	0		
↳ 2	BEGIN				17002
↳ 4	0	0	0		
↳ 2	IF TG_OP = 'INSERT' THEN				17002
↳ 5	0	0	0		
↳ 2	v_action := ' added employee(s) on ';				17002
↳ 6	1	0	0.000143		
↳ 2	ELSIF TG_OP = 'UPDATE' THEN				17002
↳ 7	0	0	0		
↳ 2	v_action := ' updated employee(s) on ';				17002
↳ 8	0	0	0		

2		ELSIF TG_OP = 'DELETE' THEN				17002		u
↪		9		1		0		3.2e-05
2		v_action := ' deleted employee(s) on ';				17002		u
↪		10		0		0		0
2		END IF;				17002		u
↪		11		0		0		0
2		v_text := 'User ' USER v_action CURRENT_DATE;				17002		u
↪		12		0		0		0
2		RAISE INFO ' %', v_text;				17002		u
↪		13		1		0		0.000383
2		RETURN NULL;				17002		u
↪		14		1		0		6.3e-05
2		END;				17002		u
↪		15		1		0		3.6e-05
2						17002		u
↪		16		0		0		0
2		DECLARE				17000		u
↪		1		0		0		0
2		v_dname VARCHAR(14);				17000		u
↪		2		0		0		0
2		BEGIN				17000		u
↪		3		0		0		0
2		SELECT INTO v_dname dname FROM dept WHERE deptno = p_deptno;				17000		u
↪		4		0		0		0
2		RETURN v_dname;				17000		u
↪		5		1		0		0.000647
2		IF NOT FOUND THEN				17000		u
↪		6		1		0		2.6e-05
2		RAISE INFO 'Invalid department number %', p_deptno;				17000		u
↪		7		0		0		0
2		RETURN '';				17000		u
↪		8		0		0		0
2		END IF;				17000		u
↪		9		0		0		0
2		END;				17000		u
↪		10		0		0		0
2						17000		u
↪		11		0		0		0
2		DECLARE				17004		u
↪		1		0		0		0
2		sal_diff NUMERIC(7,2);				17004		u
↪		2		0		0		0
2		BEGIN				17004		u
↪		3		0		0		0
2		IF TG_OP = 'INSERT' THEN				17004		u
↪		4		0		0		0
2		RAISE INFO 'Inserting employee %', NEW.empno;				17004		u
↪		5		1		0		8.4e-05
2		RAISE INFO '..New salary: %', NEW.sal;				17004		u
↪		6		0		0		0
2		RETURN NEW;				17004		u
↪		7		0		0		0
2		END IF;				17004		u
↪		8		0		0		0
2		IF TG_OP = 'UPDATE' THEN				17004		u
↪		9		0		0		0
2		sal_diff := NEW.sal - OLD.sal;				17004		u
↪		10		1		0		0.000355

2		RAISE INFO 'Updating employee %', OLD.empno;		17004		↵
↵	11	1 0 0.000177		17004		↵
2		RAISE INFO '..Old salary: %', OLD.sal;		17004		↵
↵	12	1 0 5.5e-05		17004		↵
2		RAISE INFO '..New salary: %', NEW.sal;		17004		↵
↵	13	1 0 3.1e-05		17004		↵
2		RAISE INFO '..Raise : %', sal_diff;		17004		↵
↵	14	1 0 2.8e-05		17004		↵
2		RETURN NEW;		17004		↵
↵	15	1 0 2.7e-05		17004		↵
2		END IF;		17004		↵
↵	16	1 0 1e-06		17004		↵
2		IF TG_OP = 'DELETE' THEN		17004		↵
↵	17	0 0 0		17004		↵
2		RAISE INFO 'Deleting employee %', OLD.empno;		17004		↵
↵	18	0 0 0		17004		↵
2		RAISE INFO '..Old salary: %', OLD.sal;		17004		↵
↵	19	0 0 0		17004		↵
2		RETURN OLD;		17004		↵
↵	20	0 0 0		17004		↵
2		END IF;		17004		↵
↵	21	0 0 0		17004		↵
2		END;		17004		↵
↵	22	0 0 0		17004		↵
2				17004		↵
↵	23	0 0 0		17004		↵
(68 rows)						

11. Query the `plssql_profiler_data` view to review a subset of the information found in `plssql_profiler_rawdata` table:

```
acctg=# SELECT * FROM plsql_profiler_data;
```

runid	unit_number	line#	total_occur	total_time	min_time	max_time	spare1	spare2	spare3	spare4									
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----																			
↵	1		16999		1		0		0		0		0		0		0		↵
↵	1		16999		2		0		0		0		0		0		0		↵
↵	1		16999		3		0		0		0		0		0		0		↵
↵	1		16999		4		0		0		0		0		0		0		↵
↵	1		16999		5		0		0		0		0		0		0		↵
↵	1		16999		6		0		0		0		0		0		0		↵
↵	1		16999		7		0		0		0		0		0		0		↵
↵	1		16999		8		1		0.001621		0.001621		0.001621		0.001621		0.001621		↵
↵	1		16999		9		1		0.000301		0.000301		0.000301		0.000301		0.000301		↵
↵	1		16999		10		1		4.6e-05		4.6e-05		4.6e-05		4.6e-05		4.6e-05		↵
↵	1		16999		11		1		0.001114		0.001114		0.001114		0.001114		0.001114		↵

↳	1	16999	12	15	0.000206	5e-06	7.8e-05	↳
↳	1	16999	13	15	8.3e-05	2e-06	4.7e-05	↳
↳	1	16999	14	14	0.000773	4.7e-05	0.000116	↳
↳	1	16999	15	0	0	0	0	↳
↳	1	16999	16	1	1e-05	1e-05	1e-05	↳
↳	1	16999	17	1	0	0	0	↳
↳	1	16999	18	0	0	0	0	↳
↳	2	17002	1	0	0	0	0	↳
↳	2	17002	2	0	0	0	0	↳
↳	2	17002	3	0	0	0	0	↳
↳	2	17002	4	0	0	0	0	↳
↳	2	17002	5	0	0	0	0	↳
↳	2	17002	6	1	0.000143	0.000143	0.000143	↳
↳	2	17002	7	0	0	0	0	↳
↳	2	17002	8	0	0	0	0	↳
↳	2	17002	9	1	3.2e-05	3.2e-05	3.2e-05	↳
↳	2	17002	10	0	0	0	0	↳
↳	2	17002	11	0	0	0	0	↳
↳	2	17002	12	0	0	0	0	↳
↳	2	17002	13	1	0.000383	0.000383	0.000383	↳
↳	2	17002	14	1	6.3e-05	6.3e-05	6.3e-05	↳
↳	2	17002	15	1	3.6e-05	3.6e-05	3.6e-05	↳
↳	2	17002	16	0	0	0	0	↳
↳	2	17000	1	0	0	0	0	↳
↳	2	17000	2	0	0	0	0	↳
↳	2	17000	3	0	0	0	0	↳
↳	2	17000	4	0	0	0	0	↳
↳	2	17000	5	1	0.000647	0.000647	0.000647	↳
↳	2	17000	6	1	2.6e-05	2.6e-05	2.6e-05	↳

2	17000	7	0	0	0	0
2	17000	8	0	0	0	0
2	17000	9	0	0	0	0
2	17000	10	0	0	0	0
2	17000	11	0	0	0	0
2	17004	1	0	0	0	0
2	17004	2	0	0	0	0
2	17004	3	0	0	0	0
2	17004	4	0	0	0	0
2	17004	5	1	8.4e-05	8.4e-05	8.4e-05
2	17004	6	0	0	0	0
2	17004	7	0	0	0	0
2	17004	8	0	0	0	0
2	17004	9	0	0	0	0
2	17004	10	1	0.000355	0.000355	0.000355
2	17004	11	1	0.000177	0.000177	0.000177
2	17004	12	1	5.5e-05	5.5e-05	5.5e-05
2	17004	13	1	3.1e-05	3.1e-05	3.1e-05
2	17004	14	1	2.8e-05	2.8e-05	2.8e-05
2	17004	15	1	2.7e-05	2.7e-05	2.7e-05
2	17004	16	1	1e-06	1e-06	1e-06
2	17004	17	0	0	0	0
2	17004	18	0	0	0	0
2	17004	19	0	0	0	0
2	17004	20	0	0	0	0
2	17004	21	0	0	0	0
2	17004	22	0	0	0	0
2	17004	23	0	0	0	0

(68 rows)

3.11.8.2 DBMS_PROFILER - Reference

The Advanced Server installer creates the following tables and views that you can query to review PL/SQL performance profile information:

Table Name	Description
PLSQL_PROFILER_RUNS	Table containing information about all profiler runs, organized by runid.
PLSQL_PROFILER_UNITS	Table containing information about all profiler runs, organized by unit.
PLSQL_PROFILER_DATA	View containing performance statistics.
PLSQL_PROFILER_RAWDATA	Table containing the performance statistics and the extended performance statistics for DRITA counters and timers.

3.11.8.2.1 PLSQL_PROFILER_RUNS

The PLSQL_PROFILER_RUNS table contains the following columns:

Column	Data Type	Description
runid	INTEGER (NOT NULL)	Unique identifier (plsql_profiler_runnumber)
related_run	INTEGER	The runid of a related run.
run_owner	TEXT	The role that recorded the profiling session.
run_date	TIMESTAMP WITHOUT TIME ZONE	The profiling session start time.
run_comment	TEXT	User comments relevant to this run
run_total_time	BIGINT	Run time (in microseconds)
run_system_info	TEXT	Currently Unused
run_comment1	TEXT	Additional user comments
spare1	TEXT	Currently Unused

3.11.8.2.2 PLSQL_PROFILER_UNITS

The PLSQL_PROFILER_UNITS table contains the following columns:

Column	Data Type	Description
runid	INTEGER	Unique identifier (plsql_profiler_runnumber)
unit_number	OID	Corresponds to the OID of the row in the pg_proc table that identifies the unit.
unit_type	TEXT	PL/SQL function, procedure, trigger or anonymous block
unit_owner	TEXT	The identity of the role that owns the unit.
unit_name	TEXT	The complete signature of the unit.
unit_timestamp	TIMESTAMP WITHOUT TIME ZONE	Creation date of the unit (currently NULL).
total_time	BIGINT	Time spent within the unit (in milliseconds)
spare1	BIGINT	Currently Unused
spare2	BIGINT	Currently Unused

3.11.8.2.3 PLSQL_PROFILER_DATA

The PLSQL_PROFILER_DATA view contains the following columns:

Column	Data Type	Description
runid	INTEGER	Unique identifier (plsql_profiler_runnumber)
unit_number	OID	Object ID of the unit that contains the current line.
line#	INTEGER	Current line number of the profiled workload.
total_occur	BIGINT	The number of times that the line was executed.
total_time	DOUBLE PRECISION	The amount of time spent executing the line (in seconds)
min_time	DOUBLE PRECISION	The minimum execution time for the line.
max_time	DOUBLE PRECISION	The maximum execution time for the line.
spare1	NUMBER	Currently Unused
spare2	NUMBER	Currently Unused
spare3	NUMBER	Currently Unused
spare4	NUMBER	Currently Unused

3.11.8.2.4 PLSQL_PROFILER_RAWDATA

The PLSQL_PROFILER_RAWDATA table contains the statistical and wait events information that is found in the PLSQL_PROFILER_DATA view, as well as the performance statistics returned by the DRITA counters and timers.

Column	Data Type	Description
runid	INTEGER	The run identifier (plsql_profiler_runnumber).
sourcecode	TEXT	The individual line of profiled code.
func_oid	OID	Object ID of the unit that contains the current line.
line_number	INTEGER	Current line number of the profiled workload.
exec_count	BIGINT	The number of times that the line was executed.
tuples_returned	BIGINT	Currently Unused
time_total	DOUBLE PRECISION	The amount of time spent executing the line (in seconds)
time_shortest	DOUBLE PRECISION	The minimum execution time for the line.
time_longest	DOUBLE PRECISION	The maximum execution time for the line.
num_scans	BIGINT	Currently Unused
tuples_fetched	BIGINT	Currently Unused
tuples_inserted	BIGINT	Currently Unused
tuples_updated	BIGINT	Currently Unused
tuples_deleted	BIGINT	Currently Unused
blocks_fetched	BIGINT	Currently Unused
blocks_hit	BIGINT	Currently Unused

Continued on next page

Table 3.2 – continued from previous page

Column	Data Type	Description
wal_write	BIGINT	A server has waited for a write to the write-ahead log buffer (expect this value to be high).
wal_flush	BIGINT	A server has waited for the write-ahead log to flush to disk.
wal_file_sync	BIGINT	A server has waited for the write-ahead log to sync to disk (related to the wal_sync_method parameter which, by default, is 'fsync' - better performance can be gained by changing this parameter to open_sync).
db_file_read	BIGINT	A server has waited for the completion of a read (from disk).
db_file_write	BIGINT	A server has waited for the completion of a write (to disk).
db_file_sync	BIGINT	A server has waited for the operating system to flush all changes to disk.
db_file_extend	BIGINT	A server has waited for the operating system while adding a new page to the end of a file.
sql_parse	BIGINT	Currently Unused.
query_plan	BIGINT	A server has generated a query plan.
other_lwlock_acquire	BIGINT	A server has waited for other light-weight lock to protect data.
shared_plan_cache_collision	BIGINT	A server has waited for the completion of the shared_plan_cache_collision event.
shared_plan_cache_insert	BIGINT	A server has waited for the completion of the shared_plan_cache_insert event.
shared_plan_cache_hit	BIGINT	A server has waited for the completion of the shared_plan_cache_hit event.
shared_plan_cache_miss	BIGINT	A server has waited for the completion of the shared_plan_cache_miss event.
shared_plan_cache_lock	BIGINT	A server has waited for the completion of the shared_plan_cache_lock event.
shared_plan_cache_busy	BIGINT	A server has waited for the completion of the shared_plan_cache_busy event.
shmemindexlock	BIGINT	A server has waited to find or allocate space in the shared memory.
oidgenlock	BIGINT	A server has waited to allocate or assign an OID.
xidgenlock	BIGINT	A server has waited to allocate or assign a transaction ID.
procarraylock	BIGINT	A server has waited to get a snapshot or clearing a transaction ID at transaction end.
sinvalreadlock	BIGINT	A server has waited to retrieve or remove messages from shared invalidation queue.
sinvalwritelock	BIGINT	A server has waited to add a message to the shared invalidation queue.

Continued on next page

Table 3.2 – continued from previous page

Column	Data Type	Description
walbufmappinglock	BIGINT	A server has waited to replace a page in WAL buffers.
walwritelock	BIGINT	A server has waited for WAL buffers to be written to disk.
controlfilelock	BIGINT	A server has waited to read or update the control file or creation of a new WAL file.
checkpointlock	BIGINT	A server has waited to perform a checkpoint.
clogcontrollock	BIGINT	A server has waited to read or update the transaction status.
subtranscontrollock	BIGINT	A server has waited to read or update the sub-transaction information.
multixactgenlock	BIGINT	A server has waited to read or update the shared multixact state.
multixactoffsetcontrollock	BIGINT	A server has waited to read or update multixact offset mappings.
multixactmembercontrollock	BIGINT	A server has waited to read or update multixact member mappings.
relcacheinitlock	BIGINT	A server has waited to read or write the relation cache initialization file.
checkpointercommlock	BIGINT	A server has waited to manage the fsync requests.
twophasestatelock	BIGINT	A server has waited to read or update the state of prepared transactions.
tablespacecreatelock	BIGINT	A server has waited to create or drop the tablespace.
btreevacuumlock	BIGINT	A server has waited to read or update the vacuum related information for a B-tree index.
addinshmeminitlock	BIGINT	A server has waited to manage space allocation in shared memory.
autovacuumlock	BIGINT	The autovacuum launcher waiting to read or update the current state of autovacuum workers.
autovacuumschedulelock	BIGINT	A server has waited to ensure that the table selected for a vacuum still needs vacuuming.
syncscanlock	BIGINT	A server has waited to get the start location of a scan on a table for synchronized scans.
relationmappinglock	BIGINT	A server has waited to update the relation map file used to store catalog to file node mapping.
asyncctllock	BIGINT	A server has waited to read or update shared notification state.
asyncqueuelock	BIGINT	A server has waited to read or update the notification messages.
serializableexacthashlock	BIGINT	A server has waited to retrieve or store information about serializable transactions.
serializablefinishedlistlock	BIGINT	A server has waited to access the list of finished serializable transactions.

Continued on next page

Table 3.2 – continued from previous page

Column	Data Type	Description
serializablepredicatelock	BIGINT	A server has waited to perform an operation on a list of locks held by serializable transactions.
oldserxidlock	BIGINT	A server has waited to read or record the conflicting serializable transactions.
syncreplock	BIGINT	A server has waited to read or update information about synchronous replicas.
backgroundworkerlock	BIGINT	A server has waited to read or update the background worker state.
dynamicsharedmemorycontrollock	BIGINT	A server has waited to read or update the dynamic shared memory state.
autofilelock	BIGINT	A server has waited to update the postgresql.auto.conf file.
replicationslotallocationlock	BIGINT	A server has waited to allocate or free a replication slot.
replicationslotcontrollock	BIGINT	A server has waited to read or update replication slot state.
committscontrollock	BIGINT	A server has waited to read or update transaction commit timestamps.
committslock	BIGINT	A server has waited to read or update the last value set for the transaction timestamp.
replicationoriginlock	BIGINT	A server has waited to set up, drop, or use replication origin.
multixacttruncationlock	BIGINT	A server has waited to read or truncate multixact information.
oldsnapshottimemaplock	BIGINT	A server has waited to read or update old snapshot control information.
backendrandomlock	BIGINT	A server has waited to generate a random number.
logicalrepworkerlock	BIGINT	A server has waited for the action on logical replication worker to finish.
clogtruncationlock	BIGINT	A server has waited to truncate the write-ahead log or waiting for write-ahead log truncation to finish.
bulkloadlock	BIGINT	A server has waited for the bulkloadlock to bulk upload the data.
edbresource managerlock	BIGINT	The edbresource managerlock provides detail about edb resource manager lock module.
wal_write_time	BIGINT	The amount of time that the server has waited for a wal_write wait event to write to the write-ahead log buffer (expect this value to be high).
wal_flush_time	BIGINT	The amount of time that the server has waited for a wal_flush wait event to write-ahead log to flush to disk.

Continued on next page

Table 3.2 – continued from previous page

Column	Data Type	Description
wal_file_sync_time	BIGINT	The amount of time that the server has waited for a wal_file_sync wait event to write-ahead log to sync to disk (related to the wal_sync_method parameter which, by default, is 'fsync' - better performance can be gained by changing this parameter to open_sync).
db_file_read_time	BIGINT	The amount of time that the server has waited for the db_file_read wait event for completion of a read (from disk).
db_file_write_time	BIGINT	The amount of time that the server has waited for the db_file_write wait event for completion of a write (to disk).
db_file_sync_time	BIGINT	The amount of time that the server has waited for the db_file_sync wait event to sync all changes to disk.
db_file_extend_time	BIGINT	The amount of time that the server has waited for the db_file_extend wait event while adding a new page to the end of a file.
sql_parse_time	BIGINT	The amount of time that the server has waited for the sql_parse wait event to parse a SQL statement.
query_plan_time	BIGINT	The amount of time that the server has waited for the query_plan wait event to compute the execution plan for a SQL statement.
other_lwlock_acquire_time	BIGINT	The amount of time that the server has waited for the other_lwlock_acquire wait event to protect data.
shared_plan_cache_collision_time	BIGINT	The amount of time that the server has waited for the shared_plan_cache_collision wait event.
shared_plan_cache_insert_time	BIGINT	The amount of time that the server has waited for the shared_plan_cache_insert wait event.
shared_plan_cache_hit_time	BIGINT	The amount of time that the server has waited for the shared_plan_cache_hit wait event.
shared_plan_cache_miss_time	BIGINT	The amount of time that the server has waited for the shared_plan_cache_miss wait event.

Continued on next page

Table 3.2 – continued from previous page

Column	Data Type	Description
shared_plan_cache_lock_time	BIGINT	The amount of time that the server has waited for the shared_plan_cache_lock wait event.
shared_plan_cache_busy_time	BIGINT	The amount of time that the server has waited for the shared_plan_cache_busy wait event.
shmemindexlock_time	BIGINT	The amount of time that the server has waited for the shmemindexlock wait event to find or allocate space in the shared memory.
oidgenlock_time	BIGINT	The amount of time that the server has waited for the oidgenlock wait event to allocate or assign an OID.
xidgenlock_time	BIGINT	The amount of time that the server has waited for xidgenlock wait event to allocate or assign a transaction ID.
procarraylock_time	BIGINT	The amount of time that the server has waited for a procarraylock wait event to clear a transaction ID at transaction end.
sinvalreadlock_time	BIGINT	The amount of time that the server has waited for a sinvalreadlock wait event to retrieve or remove messages from shared invalidation queue.
sinvalwritelock_time	BIGINT	The amount of time that the server has waited for a sinvalwritelock wait event to add a message to the shared invalidation queue.
walbufmappinglock_time	BIGINT	The amount of time that the server has waited for a walbufmappinglock wait event to replace a page in WAL buffers.
walwritelock_time	BIGINT	The amount of time that the server has waited for a walwritelock wait event to write the WAL buffers to disk.
controlfilelock_time	BIGINT	The amount of time that the server has waited for a controlfilelock wait event to read or update the control file or to create a new WAL file.
checkpointlock_time	BIGINT	The amount of time that the server has waited for a checkpointlock wait event to perform a checkpoint.
clogcontrollock_time	BIGINT	The amount of time that the server has waited for a clogcontrollock wait event to read or update the transaction status.
subtranscontrollock_time	BIGINT	The amount of time that the server has waited for the subtranscontrollock wait event to read or update the sub-transaction information.

Continued on next page

Table 3.2 – continued from previous page

Column	Data Type	Description
multixactgenlock_time	BIGINT	The amount of time that the server has waited for the multixactgenlock wait event to read or update the shared multixact state.
multixactoffsetcontrollock_time	BIGINT	The amount of time that the server has waited for the multixactoffsetcontrollock wait event to read or update multixact offset mappings.
multixactmembercontrollock_time	BIGINT	The amount of time that the server has waited for the multixactmembercontrollock wait event to read or update multixact member mappings.
relcacheinitlock_time	BIGINT	The amount of time that the server has waited for the relcacheinitlock wait event to read or write the relation cache initialization file.
checkpointercommlock_time	BIGINT	The amount of time that the server has waited for the checkpointercommlock wait event to manage the fsync requests.
twophasestatelock_time	BIGINT	The amount of time that the server has waited for the twophasestatelock wait event to read or update the state of prepared transactions.
tablespacecreatelock_time	BIGINT	The amount of time that the server has waited for the tablespacecreatelock wait event to create or drop the tablespace.
btreevacuumlock_time	BIGINT	The amount of time that the server has waited for the btreevacuumlock wait event to read or update the vacuum related information for a B-tree index.
addinshmeminitlock_time	BIGINT	The amount of time that the server has waited for the addinshmeminitlock wait event to manage space allocation in shared memory.
autovacuumlock_time	BIGINT	The amount of time that the server has waited for the autovacuumlock wait event to read or update the current state of autovacuum workers.
autovacuumschedulelock_time	BIGINT	The amount of time that the server has waited for the autovacuumschedulelock wait event to ensure that the table selected for a vacuum still needs vacuuming.
syncscanlock_time	BIGINT	The amount of time that the server has waited for the syncscanlock wait event to get the start location of a scan on a table for synchronized scans.

Continued on next page

Table 3.2 – continued from previous page

Column	Data Type	Description
relationmappinglock_time	BIGINT	The amount of time that the server has waited for the relationmappinglock wait event to update the relation map file used to store catalog to file node mapping.
asyncctllock_time	BIGINT	The amount of time that the server has waited for the asyncctllock wait event to read or update shared notification state.
asyncqueuelock_time	BIGINT	The amount of time that the server has waited for the asyncqueuelock wait event to read or update the notification messages.
serializableexacthashlock_time	BIGINT	The amount of time that the server has waited for the serializableexacthashlock wait event to retrieve or store information about serializable transactions.
serializablefinishedlistlock_time	BIGINT	The amount of time that the server has waited for the serializablefinishedlistlock wait event to access the list of finished serializable transactions.
serializablepredicatelocklistlock_time	BIGINT	The amount of time that the server has waited for the serializablepredicatelocklistlock wait event to perform an operation on a list of locks held by serializable transactions.
oldserxidlock_time	BIGINT	The amount of time that the server has waited for the oldserxidlock wait event to read or record the conflicting serializable transactions.
syncreplock_time	BIGINT	The amount of time that the server has waited for the syncreplock wait event to read or update information about synchronous replicas.
backgroundworkerlock_time	BIGINT	The amount of time that the server has waited for the backgroundworkerlock wait event to read or update the background worker state.
dynamicsharedmemorycontrollock_time	BIGINT	The amount of time that the server has waited for the dynamicsharedmemorycontrollock wait event to read or update the dynamic shared memory state.
autofilelock_time	BIGINT	The amount of time that the server has waited for the autofilelock wait event to update the postgresql.auto.conf file.

Continued on next page

Table 3.2 – continued from previous page

Column	Data Type	Description
replicationslotallocationlock_time	BIGINT	The amount of time that the server has waited for the replicationslotallocationlock wait event to allocate or free a replication slot.
replicationslotcontrollock_time	BIGINT	The amount of time that the server has waited for the replicationslotcontrollock wait event to read or update replication slot state.
committscontrollock_time	BIGINT	The amount of time that the server has waited for the committscontrollock wait event to read or update transaction commit timestamps.
committslock_time	BIGINT	The amount of time that the server has waited for the committslock wait event to read or update the last value set for the transaction timestamp.
replicationoriginlock_time	BIGINT	The amount of time that the server has waited for the replicationoriginlock wait event to set up, drop, or use replication origin.
multixacttruncationlock_time	BIGINT	The amount of time that the server has waited for the multixacttruncationlock wait event to read or truncate multixact information.
oldsnapshottimemaplock_time	BIGINT	The amount of time that the server has waited for the oldsnapshottimemaplock wait event to read or update old snapshot control information.
backendrandomlock_time	BIGINT	The amount of time that the server has waited for the backendrandomlock wait event to generate a random number.
logicalrepworkerlock_time	BIGINT	The amount of time that the server has waited for the logicalrepworkerlock wait event for an action on logical replication worker to finish.
clogtruncationlock_time	BIGINT	The amount of time that the server has waited for the clogtruncationlock wait event to truncate the write-ahead log or waiting for write-ahead log truncation to finish.
bulkloadlock_time	BIGINT	The amount of time that the server has waited for the bulkloadlock wait event to bulk upload the data.

Continued on next page

Table 3.2 – continued from previous page

Column	Data Type	Description
edbresource_manager_lock_time	BIGINT	The amount of time that the server has waited for the edbresource_manager_lock wait event.
total_waits	BIGINT	The total number of event waits.
total_wait_time	BIGINT	The total time spent waiting for an event.

3.12 DBMS_RANDOM

The DBMS_RANDOM package provides a number of methods to generate random values. The procedures and functions available in the DBMS_RANDOM package are listed in the following table.

Function/Procedure	Return Type	Description
INITIALIZE (<val>)	n/a	Initializes the DBMS_RANDOM package with the specified seed <value>. Deprecated, but supported for backward compatibility.
NORMAL ()	NUMBER	Returns a random NUMBER.
RANDOM	INTEGER	Returns a random INTEGER with a value greater than or equal to -2^{31} and less than 2^{31} . Deprecated, but supported for backward compatibility.
SEED (<val>)	n/a	Resets the seed with the specified <value>.
SEED (<val>)	n/a	Resets the seed with the specified <value>.
STRING (<opt>, <len>)	VARCHAR2	Returns a random string.
TERMINATE	n/a	TERMINATE has no effect. Deprecated, but supported for backward compatibility.
VALUE	NUMBER	Returns a random number with a value greater than or equal to 0 and less than 1, with 38 digit precision.
VALUE (<low>, <high>)	NUMBER	Returns a random number with a value greater than or equal to <low> and less than <high>.

3.12.1 INITIALIZE

The INITIALIZE procedure initializes the DBMS_RANDOM package with a seed value. The signature is:

```
INITIALIZE (<val> IN INTEGER)
```

This procedure should be considered deprecated; it is included for backward compatibility only.

Parameter

<val>

<val> is the seed value used by the DBMS_RANDOM package algorithm.

Example

The following code snippet demonstrates a call to the INITIALIZE procedure that initializes the DBMS_RANDOM package with the seed value, 6475.

```
DBMS_RANDOM.INITIALIZE (6475);
```

3.12.2 NORMAL

The NORMAL function returns a random number of type NUMBER. The signature is:

```
<result> NUMBER NORMAL ()
```

Parameter

<result>

<result> is a random value of type NUMBER.

Example

The following code snippet demonstrates a call to the NORMAL function:

```
x := DBMS_RANDOM.NORMAL();
```

3.12.3 RANDOM

The RANDOM function returns a random INTEGER value that is greater than or equal to -2^{31} and less than 2^{31} . The signature is:

```
<result> INTEGER RANDOM()
```

This function should be considered deprecated; it is included for backward compatibility only.

Parameter

<result>

<result> is a random value of type INTEGER.

Example

The following code snippet demonstrates a call to the RANDOM function. The call returns a random number:

```
x := DBMS_RANDOM.RANDOM();
```

3.12.4 SEED

The first form of the SEED procedure resets the seed value for the DBMS_RANDOM package with an INTEGER value. The SEED procedure is available in two forms; the signature of the first form is:

```
SEED(<val> IN INTEGER)
```

Parameter

<val>

<val> is the seed value used by the DBMS_RANDOM package algorithm.

Example

The following code snippet demonstrates a call to the SEED procedure; the call sets the seed value at 8495.

```
DBMS_RANDOM.SEED(8495);
```

3.12.5 SEED

The second form of the SEED procedure resets the seed value for the DBMS_RANDOM package with a string value. The SEED procedure is available in two forms; the signature of the second form is:

```
SEED(<val> IN VARCHAR2)
```

Parameter

<val>

<val> is the seed value used by the DBMS_RANDOM package algorithm.

Example

The following code snippet demonstrates a call to the SEED procedure; the call sets the seed value to abc123.

```
DBMS_RANDOM.SEED('abc123');
```

3.12.6 STRING

The STRING function returns a random VARCHAR2 string in a user-specified format. The signature of the STRING function is:

```
<result> VARCHAR2 STRING(<opt> IN CHAR, <len> IN NUMBER)
```

Parameters

<opt>

Formatting option for the returned string. <option> may be:

Option	Specifies Formatting Option
u or U	Uppercase alpha string
l or L	Lowercase alpha string
a or A	Mixed case string
x or X	Uppercase alpha-numeric string
p or P	Any printable characters

<len>

The length of the returned string.

<result>

<result> is a random value of type VARCHAR2.

Example

The following code snippet demonstrates a call to the STRING function; the call returns a random alpha-numeric character string that is 10 characters long.

```
x := DBMS_RANDOM.STRING('X', 10);
```

3.12.7 TERMINATE

The TERMINATE procedure has no effect. The signature is:

```
TERMINATE
```

The TERMINATE procedure should be considered deprecated; the procedure is supported for compatibility only.

3.12.8 VALUE

The VALUE function returns a random NUMBER that is greater than or equal to 0, and less than 1, with 38 digit precision. The VALUE function has two forms; the signature of the first form is:

```
<result> NUMBER VALUE()
```


Parameter

<result>

<result> is a random value of type NUMBER.

Example

The following code snippet demonstrates a call to the VALUE function. The call returns a random NUMBER:

```
x := DBMS_RANDOM.VALUE();
```

3.12.9 VALUE

The VALUE function returns a random NUMBER with a value that is between user-specified boundaries. The VALUE function has two forms; the signature of the second form is:

<result> NUMBER VALUE(<low> IN NUMBER, <high> IN NUMBER)

Parameters

<low>

<low> specifies the lower boundary for the random value. The random value may be equal to <low>.

<high>

<high> specifies the upper boundary for the random value; the random value will be less than <high>.

<result>

<result> is a random value of type NUMBER.

Example

The following code snippet demonstrates a call to the VALUE function. The call returns a random NUMBER with a value that is greater than or equal to 1 and less than 100:

```
x := DBMS_RANDOM.VALUE(1, 100);
```

3.13 DBMS_REDACT

The DBMS_REDACT package enables the redacting or masking of data returned by a query. The DBMS_REDACT package provides a procedure to create policies, alter policies, enable policies, disable policies, and drop policies. The procedures available in the DBMS_REDACT package are listed in the following table.

Function/Procedure	Function or Procedure	Return Type	Description
ADD_POLICY(<object_schema>, <object_name>, <policy_name>, <policy_description>, <column_name>, <column_description>, <function_type>, <function_parameters>, <expression>, <enable>, <regexp_pattern>, <regexp_replace_string>, <regexp_position>, <regexp_occurence>, <regexp_match_parameter>, <custom_function_expression>)	Procedure	n/a	Adds a data redaction policy.
ALTER_POLICY(<object_schema>, <object_name>, <policy_name>, <action>, <column_name>, <function_type>, <function_parameters>, <expression>, <regexp_pattern>, <regexp_replace_string>, <regexp_position>, <regexp_occurence>, <regexp_match_parameter>, <policy_description>, <column_description>, <custom_function_expression>)	Procedure	n/a	Alters the existing data redaction policy.
DISABLE_POLICY(<object_schema>, <object_name>, <policy_name>)	Procedure	n/a	Disables the existing data redaction policy.
ENABLE_POLICY(<object_schema>, <object_name>, <policy_name>)	Procedure	n/a	Enables a previously disabled data redaction policy.
DROP_POLICY(<object_schema>, <object_name>, <policy_name>)	Procedure	n/a	Drops a data redaction policy.
UPDATE_FULL_REDACTION_VALUES(<number_val>, <binfloat_val>, <bindouble_val>, <char_val>, <varchar_val>, <nchar_val>, <nvarchar_val>, <datecol_val>, <ts_val>, <tswtz_val>, <blob_val>, <clob_val>, <nclob_val>)	Procedure	n/a	Updates the full redaction default values for the specified datatype.

The data redaction feature uses the DBMS_REDACT package to define policies or conditions to redact data in a column based on the table column type and redaction type.

Note that you must be the owner of the table to create or change the data redaction policies. The users are exempted from all the column redaction policies, which the table owner or super-user is by default.

3.13.1 Using DBMS_REDACT Constants and Function Parameters

The DBMS_REDACT package uses the constants and redacts the column data by using any one of the data redaction types. The redaction type can be decided based on the `function_type` parameter of `dbms_redact.add_policy` and `dbms_redact.alter_policy` procedure. The below table highlights the values for `function_type` parameters of `dbms_redact.add_policy` and `dbms_redact.alter_policy`.

Constant	Type	Value	Description
NONE	INTEGER	0	No redaction, zero effect on the result of a query against table.
FULL	INTEGER	1	Full redaction, redacts full values of the column data.
PARTIAL	INTEGER	2	Partial redaction, redacts a portion of the column data.
RANDOM	INTEGER	4	Random redaction, each query results in a different random value depending on the datatype of the column.
REGEXP	INTEGER	5	Regular Expression based redaction, searches for the pattern of data to redact.
CUSTOM	INTEGER	99	Custom redaction type.

The following table shows the values for the `action` parameter of `dbms_redact.alter_policy`.

Constant	Type	Value	Description
ADD_COLUMN	INTEGER	1	Adds a column to the redaction policy.
DROP_COLUMN	INTEGER	2	Drops a column from the redaction policy.
MODIFY_EXPRESSION	INTEGER	3	Modifies the expression of a redaction policy. The redaction is applied when the expression evaluates to the BOOLEAN value to TRUE.
MODIFY_COLUMN	INTEGER	4	Modifies a column in the redaction policy to change the redaction function type or function parameter.
SET_POLICY_DESCRIPTION	INTEGER	5	Sets the redaction policy description.
SET_COLUMN_DESCRIPTION	INTEGER	6	Sets a description for the redaction performed on the column.

The partial data redaction enables you to redact only a portion of the column data. To use partial redaction, you must set the `dbms_redact.add_policy` procedure `function_type` parameter to `dbms_redact.partial` and use the `function_parameters` parameter to specify the partial redaction behavior.

The data redaction feature provides a predefined format to configure policies that use the following datatype:

- Character
- Number
- Datetime

The following table highlights the format descriptor for partial redaction with respect to datatype. The example described below shows how to perform a redaction for a string datatype (in this scenario, a Social Security Number (SSN)), a Number datatype, and a DATE datatype.

Datatype	Format Descriptor	Description	Examples
Character	REDACT_PARTIAL_INPUT_FORMAT	Specifies the input format. Enter V for each character from the input string to be possibly redacted. Enter F for each character from the input string that can be considered as a separator such as blank spaces or hyphens.	Consider 'VVVFVVVFVVVV, VVV-VV-VVVV, X, 1, 5' for masking first 5 digits of SSN strings such as 123-45-6789, adding hyphen to format it and thereby resulting in strings such as XXX-6789. The field value VVVVFVVVFVVVV for matching SSN strings such as 123-45-6789.
	REDACT_PARTIAL_OUTPUT_FORMAT	Specifies the output format. Enter V for each character from the input string to be possibly redacted. Replace each F character from the input format with a character such as a hyphen or any other separator.	The field value VVV-VV-VVVV can be used to redact SSN strings into XXX-XX-6789 where X comes from REDACT_PARTIAL_MASKCHAR field.
	REDACT_PARTIAL_MASKCHAR	Specifies the character to be used for redaction.	The value X for redacting SSN strings into XXX-XX-6789.
	REDACT_PARTIAL_MASKFROM	Specifies which V within the input format from which to start the redaction.	The value 1 for redacting SSN strings starting at the first V of the input format of VVVVFVVVFVVVV into strings such as XXX-XX-6789.
	REDACT_PARTIAL_MASKTO	Specifies which V within the input format at which to end the redaction.	The value 5 for redacting SSN strings up to and including the fifth V within the input format of VVVVFVVVFVVVV into strings such as XXX-XX-6789.
Number	REDACT_PARTIAL_MASKCHAR	Specifies the character to be displayed in the range between 0 and 9.	'9, 1, 5' for redacting the first five digits of the Social Security Number 123456789 into 999996789.
	REDACT_PARTIAL_MASKFROM	Specifies the start digit position for redaction.	
	REDACT_PARTIAL_MASKTO	Specifies the end digit position for redaction.	

Continued on next page

Table 3.3 – continued from previous page

Datetime	REDACT_PARTIAL_DATE_MONTH	'm' redacts the month. To mask a specific month, specify 'm#' where # indicates the month specified by its number between 1 and 12.	m3 displays as March.
	REDACT_PARTIAL_DATE_DAY	'd' redacts the day of the month. To mask with a day of the month, append 1–31 to a lowercase d.	d3 displays as 03.
	REDACT_PARTIAL_DATE_YEAR	'y' redacts the year. To mask with a year, append 1–9999 to a lowercase y.	y1960 displays as 60.
	REDACT_PARTIAL_DATE_HOUR	'h' redacts the hour. To mask with an hour, append 0–23 to a lowercase h.	h18 displays as 18.
	REDACT_PARTIAL_DATE_MINUTE	'm' redacts the minute. To mask with a minute, append 0–59 to a lowercase m.	m20 displays as 20.
	REDACT_PARTIAL_DATE_SECOND	's' redacts the second. To mask with a second, append 0–59 to a lowercase s.	s40 displays as 40.

The following table represents `function_parameters` values that can be used in partial redaction.

Function Parameter	Data Type	Value	Description
REDACT_US_SSN_F5	VARCHAR2	'VVVFVVVFVVV, VVV-VV-VVVV,X,1,5'	Redacts the first 5 numbers of SSN. Example: The number 123-45-6789 becomes XXX-XX-6789.
REDACT_US_SSN_L4	VARCHAR2	'VVVFVVVFVVV, VVV-VV-VVVV,X,6,9'	Redacts the last 4 numbers of SSN. Example: The number 123-45-6789 becomes 123-45-XXXX.

Continued on next page

Table 3.4 – continued from previous page

REDACT_US_SSN_ENTIRE	VARCHAR2	'VVVFVVVFVVV, VVV-VV-VVV,X,1,9'	Redacts the entire SSN. Example: The number 123-45-6789 becomes XXX-XX-XXXX.
REDACT_NUM_US_SSN_F5	VARCHAR2	'9,1,5'	Redacts the first 5 numbers of SSN when the column is a number datatype. Example: The number 123456789 becomes 999996789.
REDACT_NUM_US_SSN_L4	VARCHAR2	'9,6,9'	Redacts the last four numbers of SSN when the column is a number datatype. Example: The number 123456789 becomes 123459999.
REDACT_NUM_US_SSN_ENTIRE	VARCHAR2	'9,1,9'	Redacts the entire SSN when the column is a number datatype. Example: The number 123456789 becomes 999999999.
REDACT_ZIP_CODE	VARCHAR2	'VVVVV,VVVV,X,1, 5'	Redacts a 5 digit zip code. Example: 12345 becomes XXXXX.
REDACT_NUM_ZIP_CODE	VARCHAR2	'9,1,5'	Redacts a 5 digit zip code when the column is a number datatype. Example: 12345 becomes 99999.
REDACT_CCN16_F12	VARCHAR2	'VVVFVVVFVVVFVVVV, VVVV-VVVV-VVVV-VVV, *,1,12'	Redacts a 16 digit credit card number and displays only 4 digits. Example: 1234 5678 9000 2358 becomes ****-****-****-2358.
REDACT_DATE_MILLENNIUM	VARCHAR2	'm1d1y2000'	Redacts a date that is in the DD-MM-YY format. Example: Redacts all date to 01-JAN-2000.
REDACT_DATE_EPOCH	VARCHAR2	'm1d1y1970'	Redacts all dates to 01-JAN-70.
REDACT_AMEX_CCN_FORMATTED	VARCHAR2	'VVVFVVVFVVVFVVVV, VVVV-VVVVVV-VVVVV, *,1,10'	Redacts the Americian Express credit card number and replaces the digit with * except for the last 5 digits. Example: The credit card number 1234 567890 34500 becomes **** * 34500.

Continued on next page

Table 3.4 – continued from previous page

REDACT_AMEX_CCN_NUMBER	VARCHAR2	'0,1,10'	Redacts the Amercian Express credit card number and replaces the digit with 0 except the last 5 digits. Example: The credit card number 1234 567890 34500 becomes 0000 000000 34500.
REDACT_SIN_FORMATTED	VARCHAR2	'VVVFVVVFVVV, VVV-VVV-VVV,*,1,6'	Redacts the Social Insurance Number by replacing the first 6 digits by * Example: 123-456-789 becomes ***-***-789.
REDACT_SIN_NUMBER	VARCHAR2	'9,1,6'	Redacts the Social Insurance Number by replacing the first 6 digits by 9. Example: 123456789 becomes 999999789.
REDACT_SIN_UNFORMATTED	VARCHAR2	'VVVVVVVVV, VVVVVVVVV,*,1,6'	Redacts the Social Insurance Number by replacing the first 6 digits by *. Example: 123456789 becomes *****789.
REDACT_CCN_FORMATTED	VARCHAR2	'VVVFVVVFVVVFVVVFVVV, VVVV-VVVV-VVVV-VVVV, *,1,12'	Redacts a credit card number by * and displays only 4 digits. Example: The credit card number 1234-5678-9000-4671 becomes ****-****-****-4671.
REDACT_CCN_NUMBER	VARCHAR2	'9,1,12'	Redacts a credit card number by 0 except the last 4 digits. Example: The credit card number 1234567890004671 becomes 0000000000004671.
REDACT_NA_PHONE_FORMATTED	VARCHAR2	'VVVFVVVFVVV, VVV-VVV-VVVV,X,4, 10'	Redacts the North American phone number by X leaving the area code. Example: 123-456-7890 becomes 123-XXX-XXXX.
REDACT_NA_PHONE_NUMBER	VARCHAR2	'0,4,10'	Redacts the North American phone number by 0 leaving the area code. Example: 1234567890 becomes 1230000000.

Continued on next page

Table 3.4 – continued from previous page

REDACT_NA_PHONE_UNFORMATTED	VARCHAR2	'VVVVVVVVVV, VVVVVVVVVV, X, 4, 10 '	Redacts the North American phone number by X leaving the area code. Example: 1234567890 becomes 123XXXXXXX.
REDACT_UK_NIN_FORMATTED	VARCHAR2	'VVFVVFVVFVVFV, VV VV VV V, X, 3, 8 '	Redacts the UK National Insurance Number by X but leaving the alphabetic characters. Example: NY 22 01 34 D becomes NY XX XX XX D.
REDACT_UK_NIN_UNFORMATTED	VARCHAR2	'VVVVVVVVV, VVVVVVVVV, X, 3, 8 '	Redacts the UK National Insurance Number by X but leaving the alphabetic characters. Example: NY220134D becomes NYXXXXXXD.

A regular expression-based redaction searches for patterns of data to redact. The `regexp_pattern` search the values in order for the `regexp_replace_string` to change the value. The following table illustrates the `regexp_pattern` values that you can use during REGEXP based redaction.

Function Parameter and Description	Data Type	Value
RE_PATTERN_CC_L6_T4: Searches for the middle digits of a credit card number that includes 6 leading digits and 4 trailing digits. The <code>regexp_replace_string</code> setting to use with the format is <code>RE_REDACT_CC_MIDDLE_DIGITS</code> that replaces the identified pattern with the characters specified by the <code>RE_REDACT_CC_MIDDLE_DIGITS</code> parameter.	VARCHAR2	'(dddddd)(ddd*)(dddd)'

Continued on next page

Table 3.5 – continued from previous page

Function Parameter and Description	Data Type	Value
RE_PATTERN_ANY_DIGITS: Searches for any digit and replaces the identified pattern with the characters specified by the following values of the regexp_replace_string parameter. regexp_replace_string=> RE_REDACT_WITH_SINGLE_X (replaces any matched digit with the X character). regexp_replace_string=> RE_REDACT_WITH_SINGLE_1 (replaces any matched digit with the 1 character).	VARCHAR2	'd'
RE_PATTERN_US_PHONE: Searches for the U.S phone number and replaces the identified pattern with the characters specified by the regexp_replace_string parameter. regexp_replace_string=> RE_REDACT_US_PHONE_L7 (searches the phone number and then replaces the last 7 digits).	VARCHAR2	'((ddd)ddd)-(ddd)-(dddd)'
RE_PATTERN_EMAIL_ADDRESS: Searches for the email address and replaces the identified pattern with the characters specified by the following values of the regexp_replace_string parameter. regexp_replace_string=> RE_REDACT_EMAIL_NAME (finds the email address and redacts the email username). regexp_replace_string=> RE_REDACT_EMAIL_DOMAIN (finds the email address and redacts the email domain). regexp_replace_string=> RE_REDACT_EMAIL_ENTIRE (finds the email address and redacts the entire email address).	VARCHAR2	'([A-Za-z0-9._%+-]+)@([A-Za-z0-9.-]+.[A-Za-z]{2,4})'

Continued on next page

Table 3.5 – continued from previous page

Function Parameter and Description	Data Type	Value
RE_PATTERN_IP_ADDRESS: Searches for an IP address and replaces the identified pattern with the characters specified by the regexp_replace_string parameter. The regexp_replace_string parameter to be used is RE_REDACT_IP_L3 that replaces the last section of an IP address with 999 and indicates it is redacted.	VARCHAR2	'(d{1,3}.d{1,3}.d{1,3}).d{1,3}'
RE_PATTERN_AMEX_CCN: Searches for the American Express credit card number. The regexp_replace_string parameter to be used is RE_REDACT_AMEX_CCN that redacts all of the digits except the last 5.	VARCHAR2	'.*(dddddd)\$'
RE_PATTERN_CCN: Searches for the credit card number other than American Express credit cards. The regexp_replace_string parameter to be used is RE_REDACT_CCN that redacts all of the digits except the last 4.	VARCHAR2	'.*(dddd)\$'
RE_PATTERN_US_SSN: Searches the SSN number and replaces the identified pattern with the characters specified by the regexp_replace_string parameter. '1-XXX-XXXX' or 'XXX-XXX-3' will return 123-XXX-XXXX or XXX-XXX-6789 for the value '123-45-6789' respectively.	VARCHAR2	'(ddd)-(dd)-(ddd)'

Continued on next page

Table 3.5 – continued from previous page

Function Parameter and Description	Data Type	Value
RE_REDACT_CC_MIDDLE_DIGITS: Redacts the middle digits of a credit card number according to the regexp_pattern parameter with the RE_PATTERN_CC_L6_T4 format and replaces each redacted character with an X. Example: The credit card number 1234 5678 9000 2490 becomes 1234 56XX XXXX 2490.	VARCHAR2	'1XXXXXX3'
RE_REDACT_WITH_SINGLE_X: Replaces the data with a single X character for each matching pattern as specified by setting	VARCHAR2	'X'

The below table illustrates the regexp_replace_string values that you can use during REGEXP based redaction.

Function Parameter	Data Type	Value	Description
RE_REDACT_CC_MIDDLE_DIGITS	VARCHAR2	'1XXXXXX3'	Redacts the middle digits of a credit card number according to the regexp_pattern parameter with the RE_PATTERN_CC_L6_T4 format and replaces each redacted character with an X. Example: The credit card number 1234 5678 9000 2490 becomes 1234 56XX XXXX 2490.
RE_REDACT_WITH_SINGLE_X	VARCHAR2	'X'	Replaces the data with a single X character for each matching pattern as specified by setting the regexp_pattern parameter with the RE_PATTERN_ANY_DIGIT format. Example: The credit card number 1234 5678 9000 2490 becomes XXXX XXXX XXXX XXXX.

Continued on next page

Table 3.6 – continued from previous page

RE_REDACT_WITH_SINGLE_1	VARCHAR2	'1'	Replaces the data with a single 1 digit for each of the data digits as specified by setting the regexp_pattern parameter with the RE_PATTERN_ANY_DIGIT format. Example: The credit card number 1234 5678 9000 2490 becomes 1111 1111 1111 1111.
RE_REDACT_US_PHONE_L7	VARCHAR2	'1-XXX-XXXX'	Redacts the last 7 digits of U.S phone number according to the regexp_pattern parameter with the RE_PATTERN_US_PHONE format and replaces each redacted character with an X. Example: The phone number 123-444-5900 becomes 123-XXX-XXXX.
RE_REDACT_EMAIL_NAME	VARCHAR2	'xxxx@2'	Redacts the email name according to the regexp_pattern parameter with the RE_PATTERN_EMAIL_ADDRESS format and replaces the email username with the four x characters. Example: The email address sjohn@example.com becomes xxxx@example.com.
RE_REDACT_EMAIL_DOMAIN	VARCHAR2	'1@xxxxx.com'	Redacts the email domain name according to the regexp_pattern parameter with the RE_PATTERN_EMAIL_ADDRESS format and replaces the domain with the five x characters. Example: The email address sjohn@example.com becomes sjohn@xxxxx.com.
RE_REDACT_EMAIL_ENTIRE	VARCHAR2	'xxxx@xxxxx.com'	Redacts the entire email address according to the regexp_pattern parameter with the RE_PATTERN_EMAIL_ADDRESS format and replaces the email address with the x characters. Example: The email address sjohn@example.com becomes xxxx@xxxxx.com.

Continued on next page

Table 3.6 – continued from previous page

RE_REDACT_IP_L3	VARCHAR2	'1.999'	Redacts the last 3 digits of an IP address according to the reg-exp_pattern parameter with the RE_PATTERN_IP_ADDRESS format. Example: The IP address 172.0.1.258 becomes 172.0.1.999, which is an invalid IP address.
RE_REDACT_AMEX_CCN	VARCHAR2	'*****1'	Redacts the first 10 digits of an American Express credit card number according to the reg-exp_pattern parameter with the RE_PATTERN_AMEX_CCN format. Example: 123456789062816 becomes *****62816.
RE_REDACT_CCN	VARCHAR2	'*****1'	Redacts the first 12 digits of a credit card number as specified by the regexp_pattern parameter with the RE_PATTERN_CCN format. Example: 8749012678345671 becomes *****5671.

The following tables show the `regexp_position` value and `regexp_occurrence` values that you can use during REGEXP based redaction.

Function Parameter	Data Type	Value	Description
RE_BEGINNING	INTEGER	1	Specifies the position of a character where search must begin. By default, the value is 1 that indicates the search begins at the first character of source_char.

Function Parameter	Data Type	Value	Description
RE_ALL	INTEGER	0	Specifies the replacement occurrence of a substring. If the value is 0, then the replacement of each matching substring occurs.
RE_FIRST	INTEGER	1	Specifies the replacement occurrence of a substring. If the value is 1, then the replacement of the first matching substring occurs.

The following table shows the `regexp_match_parameter` values that you can use during REGEXP based redaction which lets you change the default matching behavior of a function.

Function Parameter	Data Type	Value	Description
RE_CASE_SENSITIVE	VARCHAR2	'c'	Specifies the case-sensitive matching.
RE_CASE_INSENSITIVE	VARCHAR2	'i'	Specifies the case-insensitive matching.
RE_MULTIPLE_LINES	VARCHAR2	'm'	Treats the source string as multiple lines but if you omit this parameter, then it indicates as a single line.
RE_NEWLINE_WILDCARD	VARCHAR2	'n'	Specifies the period (.), but if you omit this parameter, then the period does not match the newline character.
RE_IGNORE_WHITESPACE	VARCHAR2	'x'	Ignores the whitespace characters.

Note: If you create a redaction policy based on a numeric type column, then make sure that the result after redaction is a number and accordingly set the replacement string to avoid runtime errors.

Note: If you create a redaction policy based on a character type column, then make sure that a length of the result after redaction is compatible with the column type and accordingly set the replacement string to avoid runtime errors.

3.13.2 ADD_POLICY

The `add_policy` procedure creates a new data redaction policy for a table.

```
PROCEDURE add_policy (
object_schema          IN VARCHAR2 DEFAULT NULL,
object_name            IN VARCHAR2,
policy_name            IN VARCHAR2,
policy_description     IN VARCHAR2 DEFAULT NULL,
column_name           IN VARCHAR2 DEFAULT NULL,
column_description     IN VARCHAR2 DEFAULT NULL,
function_type          IN INTEGER DEFAULT DBMS_REDACT.FULL,
function_parameters    IN VARCHAR2 DEFAULT NULL,
expression             IN VARCHAR2,
enable                IN BOOLEAN DEFAULT TRUE,
regexp_pattern         IN VARCHAR2 DEFAULT NULL,
regexp_replace_string  IN VARCHAR2 DEFAULT NULL,
regexp_position       IN INTEGER DEFAULT DBMS_REDACT.RE_BEGINNING,
regexp_occurrence     IN INTEGER DEFAULT      DBMS_REDACT.RE_ALL,
regexp_match_parameter IN VARCHAR2 DEFAULT NULL,
custom_function_expression IN VARCHAR2 DEFAULT NULL
)
```

Parameters

<object_schema>

Specifies the name of the schema in which the object resides and on which the data redaction policy will be applied. If you specify `NULL` then the given object is searched by the order specified by `search_path` setting.

<object_name>

Name of the table on which the data redaction policy is created.

<policy_name>

Name of the policy to be added. Ensure that the `policy_name` is unique for the table on which the policy is created.

<policy_description>

Specify the description of a redaction policy.

<column_name>

Name of the column to which the redaction policy applies. To redact more than one column, use the `alter_policy` procedure to add additional columns.

<column_description>

Description of the column to be redacted. The `column_description` is not supported, but if you specify the description for a column then, you will get a warning message.

<function_type>

The type of redaction function to be used. The possible values are `NONE`, `FULL`, `PARTIAL`, `RANDOM`, `REGEXP`, and `CUSTOM`.

<function_parameters>

Specifies the function parameters for the partition redaction and is applicable only for partial redaction.

<expression>

Specifies the Boolean expression for the table and determines how the policy is to be applied. The redaction occurs if this policy expression is evaluated to TRUE.

<enable>

When set to TRUE, the policy is enabled upon creation. The default is set as TRUE. When set to FALSE, the policy is disabled but the policy can be enabled by calling the `enable_policy` procedure.

<regexp_pattern>

Specifies the regular expression pattern to redact data. If the `regexp_pattern` does not match, then the NULL value is returned.

<regexp_replace_string>

Specifies the replacement string value.

<regexp_position>

Specifies the position of a character where search must begin. By default, the function parameter is RE_BEGINNING.

<regexp_occurrence>

Specifies the replacement occurrence of a substring. If the constant is RE_ALL, then the replacement of each matching substring occurs. If the constant is RE_FIRST, then the replacement of the first matching substring occurs.

<regexp_match_parameter>

Changes the default matching behavior of a function. The possible `regexp_match_parameter` constants can be 'RE_CASE_SENSITIVE', 'RE_CASE_INSENSITIVE', 'RE_MULTIPLE_LINES', 'RE_NEWLINE_WILDCARD', 'RE_IGNORE_WHITESPACE'.

Note: For more information on constants, function_parameters, or regexp (regular expressions) see, Using DBMS_REDACT Constants and Function Parameters.

<custom_function_expression>

The `custom_function_expression` is applicable only for the CUSTOM redaction type. The `custom_function_expression` is a function expression that is, schema-qualified function with a parameter such as `schema_name.function_name (argument1, ...)` that allows a user to use their redaction logic to redact the column data.

Example

The following example illustrates how to create a policy and use full redaction for values in the `payment_details_tab` table `customer_id` column.

```
edb=# CREATE TABLE payment_details_tab (
customer_id NUMBER          NOT NULL,
card_string VARCHAR2(19) NOT NULL);
CREATE TABLE

edb=# BEGIN
  INSERT INTO payment_details_tab VALUES (4000, '1234-1234-1234-1234');
  INSERT INTO payment_details_tab VALUES (4001, '2345-2345-2345-2345');
END;

EDB-SPL Procedure successfully completed
```



```

edb=# CREATE USER redact_user;
CREATE ROLE
edb=# GRANT SELECT ON payment_details_tab TO redact_user;
GRANT

\c edb base_user

BEGIN
  DBMS_REDACT.add_policy(
    object_schema      => 'public',
    object_name        => 'payment_details_tab',
    policy_name        => 'redactPolicy_001',
    policy_description  => 'redactPolicy_001 for payment_details_tab table',
    column_name        => 'customer_id',
    function_type       => DBMS_REDACT.full,
    expression         => '1=1',
    enable             => TRUE);
END;

```

Redacted Result:

```

edb=# \c edb redact_user
You are now connected to database "edb" as user "redact_user".

edb=> select customer_id from payment_details_tab order by 1;
 customer_id
-----
           0
           0
(2 rows)

```

3.13.3 ALTER_POLICY

The `alter_policy` procedure alters or modifies an existing data redaction policy for a table.

```
PROCEDURE alter_policy (
object_schema                IN VARCHAR2 DEFAULT NULL,
object_name                  IN VARCHAR2,
policy_name                  IN VARCHAR2,
action                       IN INTEGER DEFAULT DBMS_REDACT.ADD_COLUMN,
column_name                  IN VARCHAR2 DEFAULT NULL,
function_type                IN INTEGER DEFAULT DBMS_REDACT.FULL,
function_parameters          IN VARCHAR2 DEFAULT NULL,
expression                   IN VARCHAR2 DEFAULT NULL,
regexp_pattern               IN VARCHAR2 DEFAULT NULL,
regexp_replace_string        IN VARCHAR2 DEFAULT NULL,
regexp_position              IN INTEGER DEFAULT DBMS_REDACT.RE_BEGINNING,
regexp_occurrence            IN INTEGER DEFAULT DBMS_REDACT.RE_ALL,
regexp_match_parameter       IN VARCHAR2 DEFAULT NULL,
policy_description           IN VARCHAR2 DEFAULT NULL,
column_description           IN VARCHAR2 DEFAULT NULL,
custom_function_expression   IN VARCHAR2 DEFAULT NULL
)
```

Parameters

<object_schema>

Specifies the name of the schema in which the object resides and on which the data redaction policy will be altered. If you specify `NULL` then the given object is searched by the order specified by `search_path` setting.

<object_name>

Name of the table to which to alter a data redaction policy.

<policy_name>

Name of the policy to be altered.

<action>

The action to perform. For more information about action parameters see, “DBMS_REDACT Constants and Function Parameters”.

<column_name>

Name of the column to which the redaction policy applies.

<function_type>

The type of redaction function to be used. The possible values are `NONE`, `FULL`, `PARTIAL`, `RANDOM`, `REGEXP`, and `CUSTOM`.

<function_parameters>

Specifies the function parameters for the redaction function.

<expression>

Specifies the Boolean expression for the table and determines how the policy is to be applied. The redaction occurs if this policy expression is evaluated to `TRUE`.

<regexp_pattern>

Enables the use of regular expressions to redact data. If the `regexp_pattern` does not match the data, then the NULL value is returned.

<regexp_replace_string>

Specifies the replacement string value.

<regexp_position>

Specifies the position of a character where search must begin. By default, the function parameter is `RE_BEGINNING`.

<regexp_occurence>

Specifies the replacement occurrence of a substring. If the constant is `RE_ALL`, then the replacement of each matching substring occurs. If the constant is `RE_FIRST`, then the replacement of the first matching substring occurs.

<regexp_match_parameter>

Changes the default matching behavior of a function. The possible `regexp_match_parameter` constants can be `'RE_CASE_SENSITIVE'`, `'RE_CASE_INSENSITIVE'`, `'RE_MULTIPLE_LINES'`, `'RE_NEWLINE_WILDCARD'`, `'RE_IGNORE_WHITESPACE'`.

Note: For more information on constants, function_parameters, or regexp (regular expressions) see, Using DBMS_REDACT Constants and Function Parameters.

<policy_description>

Specify the description of a redaction policy.

<column_description>

Description of the column to be redacted. The `column_description` is not supported, but if you specify the description for a column then, you will get a warning message.

<custom_function_expression>

The `custom_function_expression` is applicable only for the `CUSTOM` redaction type. The `custom_function_expression` is a function expression that is, schema-qualified function with a parameter such as `schema_name.function_name (argument1,...)` that allows a user to use their redaction logic to redact the column data.

Example

The following example illustrates to alter a policy partial redaction for values in the `payment_details_tab` table `card_string` (usually a credit card number) column.

```
\c edb base _user

BEGIN
  DBMS_REDACT.alter_policy (
    object_schema      => 'public',
    object_name        => 'payment_details_tab',
    policy_name        => 'redactPolicy_001',
    action              => DBMS_REDACT.ADD_COLUMN,
    column_name        => 'card_string',
    function_type       => DBMS_REDACT.partial,
```

```
function_parameters => DBMS_REDACT.REDACT_CCN16_F12);  
END;
```

Redacted Result:

```
edb=# \c - redact_user  
You are now connected to database "edb" as user "redact_user".  
edb=> SELECT * FROM payment_details_tab;  
 customer_id |      card_string  
-----+-----  
           0 | ****-****-****-1234  
           0 | ****-****-****-2345  
(2 rows)
```

3.13.4 DISABLE_POLICY

The `disable_policy` procedure disables an existing data redaction policy.

```
PROCEDURE disable_policy(  
    <object_schema> IN VARCHAR2 DEFAULT NULL,  
    <object_name> IN VARCHAR2,  
    <policy_name> IN VARCHAR2  
)
```

Parameters

<object_schema>

Specifies the name of the schema in which the object resides and on which the data redaction policy will be applied. If you specify `NULL` then the given object is searched by the order specified by `search_path` setting.

<object_name>

Name of the table for which to disable a data redaction policy.

<policy_name>

Name of the policy to be disabled.

Example

The following example illustrates how to disable a policy.

```
\c edb base_user  
  
BEGIN  
    DBMS_REDACT.disable_policy(  
        object_schema => 'public',  
        object_name => 'payment_details_tab',  
        policy_name => 'redactPolicy_001');  
END;
```

Redacted Result: Data is no longer redacted after disabling a policy.

3.13.5 ENABLE_POLICY

The `enable_policy` procedure enables the previously disabled data redaction policy.

```
PROCEDURE enable_policy (  
    <object_schema> IN VARCHAR2 DEFAULT NULL,  
    <object_name> IN VARCHAR2,  
    <policy_name> IN VARCHAR2  
)
```

Parameters

<object_schema>

Specifies the name of the schema in which the object resides and on which the data redaction policy will be applied. If you specify `NULL` then the given object is searched by the order specified by `search_path` setting.

<object_name>

Name of the table to which to enable a data redaction policy.

<policy_name>

Name of the policy to be enabled.

Example

The following example illustrates how to enable a policy.

```
\c edb base_user  
  
BEGIN  
    DBMS_REDACT.enable_policy(  
        object_schema => 'public',  
        object_name => 'payment_details_tab',  
        policy_name => 'redactPolicy_001');  
END;
```

Redacted Result: Data is redacted after enabling a policy.

3.13.6 DROP_POLICY

The `drop_policy` procedure drops a data redaction policy by removing the masking policy from a table.

```
PROCEDURE drop_policy(  
    <object_schema> IN VARCHAR2 DEFAULT NULL,  
    <object_name> IN VARCHAR2,  
    <policy_name> IN VARCHAR2  
)
```

Parameters

<object_schema>

Specifies the name of the schema in which the object resides and on which the data redaction policy will be applied. If you specify `NULL` then the given object is searched by the order specified by `search_path` setting.

<object_name>

Name of the table from which to drop a data redaction policy.

<policy_name>

Name of the policy to be dropped.

Example

The following example illustrates how to drop a policy.

```
\c edb base_user  
  
BEGIN  
    DBMS_REDACT.drop_policy(  
        object_schema => 'public',  
        object_name => 'payment_details_tab',  
        policy_name => 'redactPolicy_001');  
END
```

Redacted Result: The server drops the specified policy.

3.13.7 UPDATE_FULL_REDACTION_VALUES

The `update_full_redaction_values` procedure updates the default displayed values for a data redaction policy and these default values can be viewed using the `redaction_values_for_type_full` view that use the full redaction type.

```
PROCEDURE update_full_redaction_values (
number_val          IN NUMBER          DEFAULT NULL,
binfloat_val        IN FLOAT4          DEFAULT NULL,
bindouble_val       IN FLOAT8          DEFAULT NULL,
char_val            IN CHAR            DEFAULT NULL,
varchar_val         IN VARCHAR2        DEFAULT NULL,
nchar_val           IN NCHAR           DEFAULT NULL,
nvarchar_val        IN NVARCHAR2       DEFAULT NULL,
datecol_val         IN DATE            DEFAULT NULL,
ts_val              IN TIMESTAMP       DEFAULT NULL,
tswtz_val           IN TIMESTAMPTZ     DEFAULT NULL,
blob_val            IN BLOB            DEFAULT NULL,
clob_val            IN CLOB            DEFAULT NULL,
nclob_val           IN CLOB            DEFAULT NULL
)
```

Parameters

<number_val>

Updates the default value for columns of the NUMBER datatype.

<binfloat_val>

The FLOAT4 datatype is a random value. The binary float datatype is not supported.

<bindouble_val>

The FLOAT8 datatype is a random value. The binary double datatype is not supported.

<char_val>

Updates the default value for columns of the CHAR datatype.

<varchar_val>

Updates the default value for columns of the VARCHAR2 datatype.

<nchar_val>

The `nchar_val` is mapped to CHAR datatype and returns the CHAR value.

<nvarchar_val>

The `nvarchar_val` is mapped to VARCHAR2 datatype and returns the VARCHAR value.

<datecol_val>

Updates the default value for columns of the DATE datatype.

<ts_val>

Updates the default value for columns of the TIMESTAMP datatype.

<tswtz_val>

Updates the default value for columns of the TIMESTAMPTZ datatype.

<blob_val>

Updates the default value for columns of the BLOB datatype.

<clob_val>

Updates the default value for columns of the CLOB datatype.

<nclob_val>

The nclob_val is mapped to CLOB datatype and returns the CLOB value.

Example

The following example illustrates how to update the full redaction values but before updating the values, you can:

1. View the default values using redaction_values_for_type_full view as shown below:

```
edb=# \x
Expanded display is on.
edb=# SELECT number_value, char_value, varchar_value, date_value,
        timestamp_value, timestamp_with_time_zone_value, blob_value, clob_value
FROM redaction_values_for_type_full;
-[ RECORD 1 ]-----+-----
number_value          | 0
char_value            | 
varchar_value         | 
date_value            | 01-JAN-01 00:00:00
timestamp_value       | 01-JAN-01 01:00:00
timestamp_with_time_zone_value | 31-DEC-00 20:00:00 -05:00
blob_value            | \x5b72656461637465645d
clob_value            | [redacted]
(1 row)
```

2. Now, update the default values for full redaction type. The NULL values will be ignored. c edb base_user

```
edb=# BEGIN
DBMS_REDACT.update_full_redaction_values (
  number_val => 9999999,
  char_val => 'Z',
  varchar_val => 'V',
  datecol_val => to_date('17/10/2018', 'DD/MM/YYYY'),
  ts_val => to_timestamp('17/10/2018 11:12:13', 'DD/MM/YYYY HH24:MI:SS'),
  tswtz_val => NULL,
  blob_val => 'NEW REDACTED VALUE',
  clob_val => 'NEW REDACTED VALUE');
END;
```

3. You can now see the updated values using redaction_values_for_type_full view.

```
EDB-SPL Procedure successfully completed
edb=# SELECT number_value, char_value, varchar_value, date_value,
        timestamp_value, timestamp_with_time_zone_value, blob_value, clob_value
FROM redaction_values_for_type_full;
-[ RECORD 1 ]-----+-----
number_value          | 9999999
char_value            | Z
varchar_value         | V
date_value            | 17-OCT-18 00:00:00
timestamp_value       | 17-OCT-18 11:12:13
timestamp_with_time_zone_value | 31-DEC-00 20:00:00 -05:00
```

blob_value	\x4e45572052454441435445442056414c5545
clob_value	NEW REDACTED VALUE
(1 row)	

Redacted Result:

edb=# \c edb redact_user	
You are now connected to database "edb" as user "redact_user".	
edb=> select * from payment_details_tab order by 1;	
customer_id	card_string
-----+-----	
99999999	V
99999999	V
(2 rows)	

3.14 DBMS_RLS

The DBMS_RLS package enables the implementation of Virtual Private Database on certain Advanced Server database objects.

Function/Procedure	Function or Procedure	Return Type	Description
ADD_POLICY(<object_schema>, <object_name>, <policy_name>, <function_schema>, <policy_function> [, <statement_types> [, <update_check> [, <enable> [, <static_policy> [, <policy_type> [, <long_predicate> [, <sec_relevant_cols> [, <sec_relevant_cols_opt>]]]]]])	Procedure	n/a	Add a security policy to a database object.
DROP_POLICY(<object_schema>, <object_name>, <policy_name>)	Procedure	n/a	Remove a security policy from a database object.
ENABLE_POLICY(<object_schema>, <object_name>, <policy_name>, <enable>)	Procedure	n/a	Enable or disable a security policy.

Advanced Server's implementation of DBMS_RLS is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Virtual Private Database is a type of fine-grained access control using security policies. *Fine-grained access control* in Virtual Private Database means that access to data can be controlled down to specific rows as defined by the security policy.

The rules that encode a security policy are defined in a *policy function*, which is an SPL function with certain input parameters and return value. The *security policy* is the named association of the policy function to a particular database object, typically a table.

Note: In Advanced Server, the policy function can be written in any language supported by Advanced Server such as SQL, PL/pgSQL and SPL.

Note: The database objects currently supported by Advanced Server Virtual Private Database are tables. Policies cannot be applied to views or synonyms.

The advantages of using Virtual Private Database are the following:

- Provides a fine-grained level of security. Database object level privileges given by the GRANT command determine access privileges to the entire instance of a database object, while Virtual Private Database provides access control for the individual rows of a database object instance.
- A different security policy can be applied depending upon the type of SQL command (INSERT, UPDATE, DELETE, or SELECT).
- The security policy can vary dynamically for each applicable SQL command affecting the database object depending upon factors such as the session user of the application accessing the database object.
- Invocation of the security policy is transparent to all applications that access the database object and thus, individual applications do not have to be modified to apply the security policy.
- Once a security policy is enabled, it is not possible for any application (including new applications) to circumvent the security policy except by the system privilege noted by the following.

- Even superusers cannot circumvent the security policy except by the system privilege noted by the following.

Note: The only way security policies can be circumvented is if the `EXEMPT ACCESS POLICY` system privilege has been granted to a user. The `EXEMPT ACCESS POLICY` privilege should be granted with extreme care as a user with this privilege is exempted from all policies in the database.

The `DBMS_RLS` package provides procedures to create policies, remove policies, enable policies, and disable policies.

The process for implementing Virtual Private Database is as follows:

- Create a policy function. The function must have two input parameters of type `VARCHAR2`. The first input parameter is for the schema containing the database object to which the policy is to apply and the second input parameter is for the name of that database object. The function must have a `VARCHAR2` return type. The function must return a string in the form of a `WHERE` clause predicate. This predicate is dynamically appended as an `AND` condition to the SQL command that acts upon the database object. Thus, rows that do not satisfy the policy function predicate are filtered out from the SQL command result set.
- Use the `ADD_POLICY` procedure to define a new policy, which is the association of a policy function with a database object. With the `ADD_POLICY` procedure, you can also specify the types of SQL commands (`INSERT`, `UPDATE`, `DELETE`, or `SELECT`) to which the policy is to apply, whether or not to enable the policy at the time of its creation, and if the policy should apply to newly inserted rows or the modified image of updated rows.
- Use the `ENABLE_POLICY` procedure to disable or enable an existing policy.
- Use the `DROP_POLICY` procedure to remove an existing policy. The `DROP_POLICY` procedure does not drop the policy function or the associated database object.

Once policies are created, they can be viewed in the catalog views, compatible with Oracle databases: `ALL_POLICIES`, `DBA_POLICIES`, or `USER_POLICIES`. The supported compatible views are listed in the *Database Compability for Oracle Developers Reference Guide*, available at the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs/>

The `SYS_CONTEXT` function is often used with `DBMS_RLS`. The signature is:

```
SYS_CONTEXT (<namespace>, <attribute>)
```

Where:

<namespace> is a `VARCHAR2`; the only accepted value is `USERENV`. Any other value will return `NULL`.

<attribute> is a `VARCHAR2`. <attribute> may be:

attribute Value	Equivalent Value
<code>SESSION_USER</code>	<code>pg_catalog.session_user</code>
<code>CURRENT_USER</code>	<code>pg_catalog.current_user</code>
<code>CURRENT_SCHEMA</code>	<code>pg_catalog.current_schema</code>
<code>HOST</code>	<code>pg_catalog.inet_host</code>
<code>IP_ADDRESS</code>	<code>pg_catalog.inet_client_addr</code>
<code>SERVER_HOST</code>	<code>pg_catalog.inet_server_addr</code>

Note: The examples used to illustrate the `DBMS_RLS` package are based on a modified copy of the sample `emp` table provided with Advanced Server along with a role named `salesmgr` that is granted all privileges on the table. You can create the modified copy of the `emp` table named `vpemp` and the `salesmgr` role as shown by the following:

```

CREATE TABLE public.vpemp AS SELECT empno, ename, job, sal, comm, deptno FROM emp;
ALTER TABLE vpemp ADD authid VARCHAR2(12);
UPDATE vpemp SET authid = 'researchmgr' WHERE deptno = 20;
UPDATE vpemp SET authid = 'salesmgr' WHERE deptno = 30;
SELECT * FROM vpemp;

```

empno	ename	job	sal	comm	deptno	authid
7782	CLARK	MANAGER	2450.00		10	
7839	KING	PRESIDENT	5000.00		10	
7934	MILLER	CLERK	1300.00		10	
7369	SMITH	CLERK	800.00		20	researchmgr
7566	JONES	MANAGER	2975.00		20	researchmgr
7788	SCOTT	ANALYST	3000.00		20	researchmgr
7876	ADAMS	CLERK	1100.00		20	researchmgr
7902	FORD	ANALYST	3000.00		20	researchmgr
7499	ALLEN	SALESMAN	1600.00	300.00	30	salesmgr
7521	WARD	SALESMAN	1250.00	500.00	30	salesmgr
7654	MARTIN	SALESMAN	1250.00	1400.00	30	salesmgr
7698	BLAKE	MANAGER	2850.00		30	salesmgr
7844	TURNER	SALESMAN	1500.00	0.00	30	salesmgr
7900	JAMES	CLERK	950.00		30	salesmgr

(14 rows)

```

CREATE ROLE salesmgr WITH LOGIN PASSWORD 'password';
GRANT ALL ON vpemp TO salesmgr;

```

3.14.1 ADD_POLICY

The ADD_POLICY procedure creates a new policy by associating a policy function with a database object.

You must be a superuser to execute this procedure.

```
ADD_POLICY(<object_schema> VARCHAR2, <object_name> VARCHAR2,
          <policy_name> VARCHAR2, <function_schema> VARCHAR2,
          <policy_function> VARCHAR2
          [, <statement_types> VARCHAR2
          [, <update_check> BOOLEAN
          [, <enable> BOOLEAN
          [, <static_policy> BOOLEAN
          [, <policy_type> INTEGER
          [, <long_predicate> BOOLEAN
          [, <sec_relevant_cols> VARCHAR2
          [, <sec_relevant_cols_opt> INTEGER ]]]]]]))
```

Parameters

<object_schema>

Name of the schema containing the database object to which the policy is to be applied.

<object_name>

Name of the database object to which the policy is to be applied. A given database object may have more than one policy applied to it.

<policy_name>

Name assigned to the policy. The combination of database object (identified by <object_schema> and <object_name>) and policy name must be unique within the database.

<function_schema>

Name of the schema containing the policy function.

Note: The policy function may belong to a package in which case <function_schema> must contain the name of the schema in which the package is defined.

<policy_function>

Name of the SPL function that defines the rules of the security policy. The same function may be specified in more than one policy.

Note: The policy function may belong to a package in which case <policy_function> must also contain the package name in dot notation (that is, <package_name>.<function_name>).

<statement_types>

Comma-separated list of SQL commands to which the policy applies. Valid SQL commands are INSERT, UPDATE, DELETE, and SELECT. The default is INSERT, UPDATE, DELETE, SELECT.

Note: Advanced Server accepts INDEX as a statement type, but it is ignored. Policies are not applied to index operations in Advanced Server.

<update_check>

Applies to INSERT and UPDATE SQL commands only.

When set to TRUE, the policy is applied to newly inserted rows and to the modified image of updated rows. If any of the new or modified rows do not qualify according to the policy function predicate, then the INSERT or UPDATE command throws an exception and no rows are inserted or modified by the INSERT or UPDATE command.

When set to FALSE, the policy is not applied to newly inserted rows or the modified image of updated rows. Thus, a newly inserted row may not appear in the result set of a subsequent SQL command that invokes the same policy. Similarly, rows which qualified according to the policy prior to an UPDATE command may not appear in the result set of a subsequent SQL command that invokes the same policy.

The default is FALSE.

<enable>

When set to TRUE, the policy is enabled and applied to the SQL commands given by the <statement_types> parameter. When set to FALSE the policy is disabled and not applied to any SQL commands. The policy can be enabled using the ENABLE_POLICY procedure. The default is TRUE.

<static_policy>

In Oracle, when set to TRUE, the policy is *static*, which means the policy function is evaluated once per database object the first time it is invoked by a policy on that database object. The resulting policy function predicate string is saved in memory and reused for all invocations of that policy on that database object while the database server instance is running.

When set to FALSE, the policy is *dynamic*, which means the policy function is re-evaluated and the policy function predicate string regenerated for all invocations of the policy.

The default is FALSE.

Note: In Oracle 10g, the <policy_type> parameter was introduced, which is intended to replace the <static_policy> parameter. In Oracle, if the <policy_type> parameter is not set to its default value of NULL, the <policy_type> parameter setting overrides the <static_policy> setting.

Note: The setting of <static_policy> is ignored by Advanced Server. Advanced Server implements only the dynamic policy, regardless of the setting of the <static_policy> parameter.

<policy_type>

In Oracle, determines when the policy function is re-evaluated, and hence, if and when the predicate string returned by the policy function changes. The default is NULL.

Note: The setting of this parameter is ignored by Advanced Server. Advanced Server always assumes a dynamic policy.

<long_predicate>

In Oracle, allows predicates up to 32K bytes if set to TRUE, otherwise predicates are limited to 4000 bytes. The default is FALSE.

Note: The setting of this parameter is ignored by Advanced Server. An Advanced Server policy function can return a predicate of unlimited length for all practical purposes.

<sec_relevant_cols>

Comma-separated list of columns of <object_name>. Provides *column-level Virtual Private Database* for the listed columns. The policy is enforced if any of the listed columns are referenced in a SQL command of a type listed in <statement_types>. The policy is not enforced if no such columns are referenced.

The default is NULL, which has the same effect as if all of the database object's columns were included in <sec_relevant_cols>.

<sec_relevant_cols_opt>

In Oracle, if <sec_relevant_cols_opt> is set to DBMS_RLS.ALL_ROWS (INTEGER constant of value 1), then the columns listed in <sec_relevant_cols> return NULL on all rows where the applied policy predicate is false. (If <sec_relevant_cols_opt> is not set to DBMS_RLS.ALL_ROWS, these rows would not be returned at all in the result set.) The default is NULL.

Note: Advanced Server does not support the DBMS_RLS.ALL_ROWS functionality. Advanced Server throws an error if sec_relevant_cols_opt is set to DBMS_RLS.ALL_ROWS (INTEGER value of 1).

Examples

This example uses the following policy function:

```
CREATE OR REPLACE FUNCTION verify_session_user (
  p_schema      VARCHAR2,
  p_object      VARCHAR2
)
RETURN VARCHAR2
IS
BEGIN
  RETURN 'authid = SYS_CONTEXT(''USERENV'', ''SESSION_USER'')';
END;
```

This function generates the predicate `authid = SYS_CONTEXT(''USERENV'', 'SESSION_USER')`, which is added to the WHERE clause of any SQL command of the type specified in the ADD_POLICY procedure.

This limits the effect of the SQL command to those rows where the content of the authid column is the same as the session user.

Note: This example uses the SYS_CONTEXT function to return the login user name. In Oracle the SYS_CONTEXT function is used to return attributes of an *application context*. The first parameter of the SYS_CONTEXT function is the name of an application context while the second parameter is the name of an attribute set within the application context. USERENV is a special built-in namespace that describes the current session. Advanced Server does not support application contexts, but only this specific usage of the SYS_CONTEXT function.

The following anonymous block calls the ADD_POLICY procedure to create a policy named `secure_update` to be applied to the `vpemp` table using function `verify_session_user` whenever an INSERT, UPDATE, or DELETE SQL command is given referencing the `vpemp` table.

```
DECLARE
  v_object_schema    VARCHAR2(30) := 'public';
  v_object_name      VARCHAR2(30) := 'vpemp';
  v_policy_name      VARCHAR2(30) := 'secure_update';
  v_function_schema  VARCHAR2(30) := 'enterprisedb';
  v_policy_function   VARCHAR2(30) := 'verify_session_user';
  v_statement_types  VARCHAR2(30) := 'INSERT,UPDATE,DELETE';
  v_update_check     BOOLEAN      := TRUE;
  v_enable           BOOLEAN      := TRUE;
BEGIN
  DBMS_RLS.ADD_POLICY(
    v_object_schema,
    v_object_name,
    v_policy_name,
    v_function_schema,
    v_policy_function,
    v_statement_types,
    v_update_check,
    v_enable
  );
END;
```

After successful creation of the policy, a terminal session is started by user `salesmgr`. The following query shows the content of the `vpemp` table:

```
edb=# \c edb salesmgr
Password for user salesmgr:
You are now connected to database "edb" as user "salesmgr".
edb=> SELECT * FROM vpemp;
 empno | ename  | job      | sal   | comm  | deptno | authid
-----+-----+-----+-----+-----+-----+-----
  7782 | CLARK  | MANAGER | 2450.00 |      |    10 |
  7839 | KING   | PRESIDENT | 5000.00 |      |    10 |
  7934 | MILLER | CLERK    | 1300.00 |      |    10 |
  7369 | SMITH  | CLERK    |  800.00 |      |    20 | researchmgr
  7566 | JONES  | MANAGER | 2975.00 |      |    20 | researchmgr
  7788 | SCOTT  | ANALYST  | 3000.00 |      |    20 | researchmgr
  7876 | ADAMS  | CLERK    | 1100.00 |      |    20 | researchmgr
  7902 | FORD   | ANALYST  | 3000.00 |      |    20 | researchmgr
  7499 | ALLEN  | SALESMAN | 1600.00 | 300.00 |    30 | salesmgr
  7521 | WARD   | SALESMAN | 1250.00 | 500.00 |    30 | salesmgr
  7654 | MARTIN | SALESMAN | 1250.00 | 1400.00 |    30 | salesmgr
  7698 | BLAKE  | MANAGER | 2850.00 |      |    30 | salesmgr
  7844 | TURNER | SALESMAN | 1500.00 |  0.00 |    30 | salesmgr
  7900 | JAMES  | CLERK    |  950.00 |      |    30 | salesmgr
(14 rows)
```

An unqualified UPDATE command (no WHERE clause) is issued by the salesmgr user:

```
edb=> UPDATE vpemp SET comm = sal * .75;
UPDATE 6
```

Instead of updating all rows in the table, the policy restricts the effect of the update to only those rows where the authid column contains the value salesmgr as specified by the policy function predicate authid = SYS_CONTEXT('USERENV', 'SESSION_USER').

The following query shows that the comm column has been changed only for those rows where authid contains salesmgr. All other rows are unchanged.

```
edb=> SELECT * FROM vpemp;
 empno | ename | job | sal | comm | deptno | authid
-----+-----+-----+-----+-----+-----+-----
 7782 | CLARK | MANAGER | 2450.00 | | 10 |
 7839 | KING | PRESIDENT | 5000.00 | | 10 |
 7934 | MILLER | CLERK | 1300.00 | | 10 |
 7369 | SMITH | CLERK | 800.00 | | 20 | researchmgr
 7566 | JONES | MANAGER | 2975.00 | | 20 | researchmgr
 7788 | SCOTT | ANALYST | 3000.00 | | 20 | researchmgr
 7876 | ADAMS | CLERK | 1100.00 | | 20 | researchmgr
 7902 | FORD | ANALYST | 3000.00 | | 20 | researchmgr
 7499 | ALLEN | SALESMAN | 1600.00 | 1200.00 | 30 | salesmgr
 7521 | WARD | SALESMAN | 1250.00 | 937.50 | 30 | salesmgr
 7654 | MARTIN | SALESMAN | 1250.00 | 937.50 | 30 | salesmgr
 7698 | BLAKE | MANAGER | 2850.00 | 2137.50 | 30 | salesmgr
 7844 | TURNER | SALESMAN | 1500.00 | 1125.00 | 30 | salesmgr
 7900 | JAMES | CLERK | 950.00 | 712.50 | 30 | salesmgr
(14 rows)
```

Furthermore, since the <update_check> parameter was set to TRUE in the ADD_POLICY procedure, the following INSERT command throws an exception since the value given for the authid column, researchmgr, does not match the session user, which is salesmgr, and hence, fails the policy.

```
edb=> INSERT INTO vpemp VALUES (9001, 'SMITH', 'ANALYST', 3200.00, NULL, 20, 'researchmgr
↪');
ERROR: policy with check option violation
DETAIL: Policy predicate was evaluated to FALSE with the updated values
```

If <update_check> was set to FALSE, the preceding INSERT command would have succeeded.

The following example illustrates the use of the <sec_relevant_cols> parameter to apply a policy only when certain columns are referenced in the SQL command. The following policy function is used for this example, which selects rows where the employee salary is less than 2000.

```
CREATE OR REPLACE FUNCTION sal_lt_2000 (
  p_schema VARCHAR2,
  p_object VARCHAR2
)
RETURN VARCHAR2
IS
BEGIN
  RETURN 'sal < 2000';
END
```

The policy is created so that it is enforced only if a SELECT command includes columns `sal` or `comm`:

```
DECLARE
    v_object_schema    VARCHAR2(30) := 'public';
    v_object_name       VARCHAR2(30) := 'vpemp';
    v_policy_name       VARCHAR2(30) := 'secure_salary';
    v_function_schema   VARCHAR2(30) := 'enterprisedb';
    v_policy_function   VARCHAR2(30) := 'sal_lt_2000';
    v_statement_types   VARCHAR2(30) := 'SELECT';
    v_sec_relevant_cols VARCHAR2(30) := 'sal,comm';
BEGIN
    DBMS_RLS.ADD_POLICY(
        v_object_schema,
        v_object_name,
        v_policy_name,
        v_function_schema,
        v_policy_function,
        v_statement_types,
        sec_relevant_cols => v_sec_relevant_cols
    );
END;
```

If a query does not reference columns `sal` or `comm`, then the policy is not applied. The following query returns all 14 rows of table `vpemp`:

```
edb=# SELECT empno, ename, job, deptno, authid FROM vpemp;
 empno | ename  | job      | deptno | authid
-----+-----+-----+-----+-----
  7782 | CLARK  | MANAGER  |    10  |
  7839 | KING   | PRESIDENT |    10  |
  7934 | MILLER | CLERK    |    10  |
  7369 | SMITH  | CLERK    |    20  | researchmgr
  7566 | JONES  | MANAGER  |    20  | researchmgr
  7788 | SCOTT  | ANALYST  |    20  | researchmgr
  7876 | ADAMS  | CLERK    |    20  | researchmgr
  7902 | FORD   | ANALYST  |    20  | researchmgr
  7499 | ALLEN  | SALESMAN |    30  | salesmgr
  7521 | WARD   | SALESMAN |    30  | salesmgr
  7654 | MARTIN | SALESMAN |    30  | salesmgr
  7698 | BLAKE  | MANAGER  |    30  | salesmgr
  7844 | TURNER | SALESMAN |    30  | salesmgr
  7900 | JAMES  | CLERK    |    30  | salesmgr
(14 rows)
```

If the query references the `sal` or `comm` columns, then the policy is applied to the query eliminating any rows where `sal` is greater than or equal to 2000 as shown by the following:

```
edb=# SELECT empno, ename, job, sal, comm, deptno, authid FROM vpemp;
 empno | ename  | job   | sal   | comm  | deptno | authid
-----+-----+-----+-----+-----+-----+-----
 7934 | MILLER | CLERK | 1300.00 |      | 10 |
 7369 | SMITH  | CLERK | 800.00 |      | 20 | researchmgr
 7876 | ADAMS  | CLERK | 1100.00 |      | 20 | researchmgr
 7499 | ALLEN  | SALESMAN | 1600.00 | 1200.00 | 30 | salesmgr
 7521 | WARD   | SALESMAN | 1250.00 | 937.50 | 30 | salesmgr
 7654 | MARTIN | SALESMAN | 1250.00 | 937.50 | 30 | salesmgr
 7844 | TURNER | SALESMAN | 1500.00 | 1125.00 | 30 | salesmgr
 7900 | JAMES  | CLERK | 950.00 | 712.50 | 30 | salesmgr
(8 rows)
```

3.14.2 DROP_POLICY

The `DROP_POLICY` procedure deletes an existing policy. The policy function and database object associated with the policy are not deleted by the `DROP_POLICY` procedure.

You must be a superuser to execute this procedure.

```
DROP_POLICY(<object_schema> VARCHAR2, <object_name> VARCHAR2,  
            <policy_name> VARCHAR2)
```

Parameters

<object_schema>

Name of the schema containing the database object to which the policy applies.

<object_name>

Name of the database object to which the policy applies.

<policy_name>

Name of the policy to be deleted.

Examples

The following example deletes policy `secure_update` on table `public.vpemp`:

```
DECLARE  
    v_object_schema    VARCHAR2(30) := 'public';  
    v_object_name      VARCHAR2(30) := 'vpemp';  
    v_policy_name      VARCHAR2(30) := 'secure_update';  
BEGIN  
    DBMS_RLS.DROP_POLICY(  
        v_object_schema,  
        v_object_name,  
        v_policy_name  
    );  
END;
```

3.14.3 ENABLE_POLICY

The `ENABLE_POLICY` procedure enables or disables an existing policy on the specified database object.

You must be a superuser to execute this procedure.

```
ENABLE_POLICY(<object_schema> VARCHAR2, <object_name> VARCHAR2,  
             <policy_name> VARCHAR2, <enable> BOOLEAN)
```

Parameters

`<object_schema>`

Name of the schema containing the database object to which the policy applies.

`<object_name>`

Name of the database object to which the policy applies.

`<policy_name>`

Name of the policy to be enabled or disabled.

`<enable>`

When set to `TRUE`, the policy is enabled. When set to `FALSE`, the policy is disabled.

Examples

The following example disables policy `secure_update` on table `public.vpemp`:

```
DECLARE  
    v_object_schema    VARCHAR2(30) := 'public';  
    v_object_name      VARCHAR2(30) := 'vpemp';  
    v_policy_name      VARCHAR2(30) := 'secure_update';  
    v_enable           BOOLEAN := FALSE;  
BEGIN  
    DBMS_RLS.ENABLE_POLICY(  
        v_object_schema,  
        v_object_name,  
        v_policy_name,  
        v_enable  
    );  
END;
```

3.15 DBMS_SCHEDULER

The DBMS_SCHEDULER package provides a way to create and manage Oracle-styled jobs, programs and job schedules. The DBMS_SCHEDULER package implements the following functions and procedures:

Function/Procedure	Return Type	Description
CREATE_JOB(<job_name>, <job_type>, <job_action>, <number_of_arguments>, <start_date>, <repeat_interval>, <end_date>, <job_class>, <enabled>, <auto_drop>, <comments>)	n/a	Use the first form of the CREATE_JOB procedure to create a job, specifying program and schedule details by means of parameters.
CREATE_JOB(<job_name>, <program_name>, <schedule_name>, <job_class>, <enabled>, <auto_drop>, <comments>)	n/a	Use the second form of CREATE_JOB to create a job that uses a named program and named schedule.
CREATE_PROGRAM(<program_name>, <program_type>, <program_action>, <number_of_arguments>, <enabled>, <comments>)	n/a	Use CREATE_PROGRAM to create a program.
CREATE_SCHEDULE(<schedule_name>, <start_date>, <repeat_interval>, <end_date>, <comments>)	n/a	Use the CREATE_SCHEDULE procedure to create a schedule.
DEFINE_PROGRAM_ARGUMENT(<program_name>, <argument_position>, <argument_name>, <argument_type>, <default_value>, <Mout_argument>)	n/a	Use the first form of the DEFINE_PROGRAM_ARGUMENT procedure to define a program argument that has a default value.
DEFINE_PROGRAM_ARGUMENT(<program_name>, <argument_position>, <argument_name>, <argument_type>, <out_argument>)	n/a	Use the first form of the DEFINE_PROGRAM_ARGUMENT procedure to define a program argument that does not have a default value.
DISABLE(<name>, <force>, <commit_semantics>)	n/a	Use the DISABLE procedure to disable a job or program.
DROP_JOB(<job_name>, <force>, <defer>, <commit_semantics>)	n/a	Use the DROP_JOB procedure to drop a job.
DROP_PROGRAM(<program_name>, <force>)	n/a	Use the DROP_PROGRAM procedure to drop a program.
DROP_PROGRAM_ARGUMENT(<program_name>, <argument_position>)	n/a	Use the first form of DROP_PROGRAM_ARGUMENT to drop a program argument by specifying the argument position.
DROP_PROGRAM_ARGUMENT(<program_name>, <argument_name>)	n/a	Use the second form of DROP_PROGRAM_ARGUMENT to drop a program argument by specifying the argument name.
DROP_SCHEDULE(<schedule_name>, <force>)	n/a	Use the DROP SCHEDULE procedure to drop a schedule.
ENABLE(<name>, <commit_semantics>)	n/a	Use the ENABLE command to enable a program or job.
EVALUATE_CALENDAR_STRING(<calendar_string>, <start_date>, <return_date_after>, <next_run_date>)	n/a	Use EVALUATE_CALENDAR_STRING to review the execution date described by a user-defined calendar schedule.
RUN_JOB(<job_name>, <use_current_session>, <manually>)	n/a	Use the RUN_JOB procedure to execute a job immediately.

Continued on next page

Table 3.7 – continued from previous page

SET_JOB_ARGUMENT_VALUE(<job_name>, <argument_position>, <argument_value>)	n/a	Use the first form of SET_JOB_ARGUMENT value to set the value of a job argument described by the argument's position.
SET_JOB_ARGUMENT_VALUE(<job_name>, <argument_name>, <argument_value>)	n/a	Use the second form of SET_JOB_ARGUMENT value to set the value of a job argument described by the argument's name.

Advanced Server's implementation of DBMS_SCHEDULER is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The DBMS_SCHEDULER package is dependent on the pgAgent service; you must have a pgAgent service installed and running on your server before using DBMS_SCHEDULER.

Before using DBMS_SCHEDULER, a database superuser must create the catalog tables in which the DBMS_SCHEDULER programs, schedules and jobs are stored. Use the `psql` client to connect to the database, and invoke the command:

```
CREATE EXTENSION dbms_scheduler;
```

By default, the `dbms_scheduler` extension resides in the `contrib/dbms_scheduler_ext` subdirectory (under the Advanced Server installation).

Note that after creating the DBMS_SCHEDULER tables, only a superuser will be able to perform a dump or reload of the database.

3.15.1 Using Calendar Syntax to Specify a Repeating Interval

The `CREATE_JOB` and `CREATE_SCHEDULE` procedures use Oracle-styled calendar syntax to define the interval with which a job or schedule is repeated. You should provide the scheduling information in the `<repeat_interval>` parameter of each procedure.

`<repeat_interval>` is a value (or series of values) that define the interval between the executions of the scheduled job. Each value is composed of a token, followed by an equal sign, followed by the unit (or units) on which the schedule will execute. Multiple token values must be separated by a semi-colon (;).

For example, the following value:

```
FREQ=DAILY;BYDAY=MON,TUE,WED,THU,FRI;BYHOUR=17;BYMINUTE=45
```

Defines a schedule that is executed each weeknight at 5:45.

The token types and syntax described in the table below are supported by Advanced Server:

Token type	Syntax	Valid Values
FREQ	FREQ=<predefined_interval>	Where <predefined_interval> is one of the following: YEARLY, MONTHLY, WEEKLY, DAILY, HOURLY, MINUTELY. The SECONDLY keyword is not supported.
BYMONTH	BYMONTH=<month>(<month>)...	Where <month> is the three-letter abbreviation of the month name: JAN FEB MAR APR MAY JUN JUL AUG SEP OCT NOV DEC
BYMONTH	BYMONTH=<month>(<month>)...	Where <month> is the numeric value representing the month: 1 2 3 4 5 6 7 8 9 10 11 12
BYMONTHDAY	BYMONTHDAY=<day_of_month>	Where <day_of_month> is a value from 1 through 31
BYDAY	BYDAY=<weekday>	Where <weekday> is a three-letter abbreviation or single-digit value representing the day of the week.
		Monday MON 1
		Tuesday TUE 2
		Wednesday WED 3
		Thursday THU 4
		Friday FRI 5
		Saturday SAT 6
		Sunday SUN 7
BYDATE	BYDATE=<date>(<date>)...	Where <date> is <YYYYMMDD>. YYYY is a four-digit year representation of the year, MM is a two-digit representation of the month, and DD is a two-digit day representation of the day.

Continued on next page

Table 3.8 – continued from previous page

Token type	Syntax	Valid Values
BYDATE	BYDATE=<date> (, <date>)...	Where <date> is <MMDD>. MM is a two-digit representation of the month, and DD is a two-digit day representation of the day
BYHOUR	BYHOUR=<hour>	Where <hour> is a value from 0 through 23.
BYMINUTE	BYMINUTE=<minute>	Where <minute> is a value from 0 through 59.

3.15.2 CREATE_JOB

Use the `CREATE_JOB` procedure to create a job. The procedure comes in two forms; the first form of the procedure specifies a schedule within the job definition, as well as a job action that will be invoked when the job executes:

```
create_job(
  <job_name> IN VARCHAR2,
  <job_type> IN VARCHAR2,
  <job_action> IN VARCHAR2,
  <number_of_arguments> IN PLS_INTEGER DEFAULT 0,
  <start_date> IN TIMESTAMP WITH TIME ZONE DEFAULT NULL,
  <repeat_interval> IN VARCHAR2 DEFAULT NULL,
  <end_date> IN TIMESTAMP WITH TIME ZONE DEFAULT NULL,
  <job_class> IN VARCHAR2 DEFAULT 'DEFAULT_JOB_CLASS',
  <enabled> IN BOOLEAN DEFAULT FALSE,
  <auto_drop> IN BOOLEAN DEFAULT TRUE,
  <comments> IN VARCHAR2 DEFAULT NULL)
```

The second form uses a job schedule to specify the schedule on which the job will execute, and specifies the name of a program that will execute when the job runs:

```
create_job(
  <job_name> IN VARCHAR2,
  <program_name> IN VARCHAR2,
  <schedule_name> IN VARCHAR2,
  <job_class> IN VARCHAR2 DEFAULT 'DEFAULT_JOB_CLASS',
  <enabled> IN BOOLEAN DEFAULT FALSE,
  <auto_drop> IN BOOLEAN DEFAULT TRUE,
  <comments> IN VARCHAR2 DEFAULT NULL)
```

Parameters

<job_ame>

<job_name> specifies the optionally schema-qualified name of the job being created.

<job_type>

<job_type> specifies the type of job. The current implementation of `CREATE_JOB` supports a job type of `PLSQL_BLOCK` or `STORED_PROCEDURE`.

<job_action>

If <job_type> is `PLSQL_BLOCK`, <job_action> specifies the content of the PL/SQL block that will be invoked when the job executes. The block must be terminated with a semi-colon (;).

If <job_type> is `STORED_PROCEDURE`, <job_action> specifies the optionally schema-qualified name of the procedure.

<number_of_arguments>

<number_of_arguments> is an `INTEGER` value that specifies the number of arguments expected by the job. The default is 0.

<start_date>

<start_date> is a `TIMESTAMP WITH TIME ZONE` value that specifies the first time that the job is scheduled to execute. The default value is `NULL`, indicating that the job should be scheduled to execute when the job is enabled.

<repeat_interval>

<repeat_interval> is a VARCHAR2 value that specifies how often the job will repeat. If a <repeat_interval> is not specified, the job will execute only once. The default value is NULL.

<end_date>

<end_date> is a TIMESTAMP WITH TIME ZONE value that specifies a time after which the job will no longer execute. If a date is specified, the <end_date> must be after <start_date>. The default value is NULL.

Please note that if an <end_date> is not specified and a <repeat_interval> is specified, the job will repeat indefinitely until it is disabled.

<program_name>

<program_name> is the name of a program that will be executed by the job.

<schedule_name>

<schedule_name> is the name of the schedule associated with the job.

<job_class>

<job_class> is accepted for compatibility and ignored.

<enabled>

<enabled> is a BOOLEAN value that specifies if the job is enabled when created. By default, a job is created in a disabled state, with <enabled> set to FALSE. To enable a job, specify a value of TRUE when creating the job, or enable the job with the DBMS_SCHEDULER.ENABLE procedure.

<auto_drop>

The <auto_drop> parameter is accepted for compatibility and is ignored. By default, a job's status will be changed to DISABLED after the time specified in <end_date>.

<comments>

Use the <comments> parameter to specify a comment about the job.

Example

The following example demonstrates a call to the CREATE_JOB procedure:

```
EXEC
  DBMS_SCHEDULER.CREATE_JOB (
    job_name      => 'update_log',
    job_type      => 'PLSQL_BLOCK',
    job_action    => 'BEGIN INSERT INTO my_log VALUES(current_timestamp);END;',
    start_date    => '01-JUN-15 09:00:00.000000',
    repeat_interval => 'FREQ=DAILY;BYDAY=MON,TUE,WED,THU,FRI;BYHOUR=17;',
    end_date      => NULL,
    enabled       => TRUE,
    comments      => 'This job adds a row to the my_log table.';
```

The code fragment creates a job named update_log that executes each weeknight at 5:00. The job executes a PL/SQL block that inserts the current timestamp into a logfile (my_log). Since no end_date is specified, the job will execute until it is disabled by the DBMS_SCHEDULER.DISABLE procedure.

3.15.3 CREATE_PROGRAM

Use the `CREATE_PROGRAM` procedure to create a `DBMS_SCHEDULER` program. The signature is:

```
CREATE_PROGRAM (
  <program_name> IN VARCHAR2,
  <program_type> IN VARCHAR2,
  <program_action> IN VARCHAR2,
  <number_of_arguments> IN PLS_INTEGER DEFAULT 0,
  <enabled> IN BOOLEAN DEFAULT FALSE,
  <comments> IN VARCHAR2 DEFAULT NULL)
```

Parameters

`<program_name>`

`<program_name>` specifies the name of the program that is being created.

`<program_type>`

`<program_type>` specifies the type of program. The current implementation of `CREATE_PROGRAM` supports a `<program_type>` of `PLSQL_BLOCK` or `PROCEDURE`.

`<program_action>`

If `<program_type>` is `PLSQL_BLOCK`, `<program_action>` contains the PL/SQL block that will execute when the program is invoked. The PL/SQL block must be terminated with a semi-colon (;).

If `<program_type>` is `PROCEDURE`, `<program_action>` contains the name of the stored procedure.

`<number_of_arguments>`

If `<program_type>` is `PLSQL_BLOCK`, this argument is ignored.

If `<program_type>` is `PROCEDURE`, `<number_of_arguments>` specifies the number of arguments required by the procedure. The default value is 0.

`<enabled>`

`<enabled>` specifies if the program is created enabled or disabled:

- If `<enabled>` is `TRUE`, the program is created enabled.
- If `<enabled>` is `FALSE`, the program is created disabled; use the `DBMS_SCHEDULER.ENABLE` program to enable a disabled program.

The default value is `FALSE`.

`<comments>`

Use the `<comments>` parameter to specify a comment about the program; by default, this parameter is `NULL`.

Example

The following call to the CREATE_PROGRAM procedure creates a program named update_log:

```
EXEC
  DBMS_SCHEDULER.CREATE_PROGRAM (
program_name      => 'update_log',
  program_type     => 'PLSQL_BLOCK',
  program_action   => 'BEGIN INSERT INTO my_log VALUES(current_timestamp);END;',
  enabled          => TRUE,
  comment          => 'This program adds a row to the my_log table.');
```

update_log is a PL/SQL block that adds a row containing the current date and time to the my_log table. The program will be enabled when the CREATE_PROGRAM procedure executes.

3.15.4 CREATE_SCHEDULE

Use the CREATE_SCHEDULE procedure to create a job schedule. The signature of the CREATE_SCHEDULE procedure is:

```
create_schedule(
  <schedule_name> IN VARCHAR2,
  <start_date> IN TIMESTAMP WITH TIME ZONE DEFAULT NULL,
  <repeat_interval> IN VARCHAR2,
  <end_date> IN TIMESTAMP WITH TIME ZONE DEFAULT NULL,
  <comments> IN VARCHAR2 DEFAULT NULL)
```

Parameters

<schedule_name>

<schedule_name> specifies the name of the schedule.

<start_date>

<start_date> is a TIMESTAMP WITH TIME ZONE value that specifies the date and time that the schedule is eligible to execute. If a <start_date> is not specified, the date that the job is enabled is used as the <start_date>. By default, <start_date> is NULL.

<repeat_interval>

<repeat_interval> is a VARCHAR2 value that specifies how often the job will repeat. If a <repeat_interval> is not specified, the job will execute only once, on the date specified by <start_date>.

Note: You must provide a value for either <start_date> or <repeat_interval>; if both <start_date> and <repeat_interval> are NULL, the server will return an error.

<end_date> IN TIMESTAMP WITH TIME ZONE DEFAULT NULL

<end_date> is a TIMESTAMP WITH TIME ZONE value that specifies a time after which the schedule will no longer execute. If a date is specified, the <end_date> must be after the <start_date>. The default value is NULL.

Please note that if a <repeat_interval> is specified and an <end_date> is not specified, the schedule will repeat indefinitely until it is disabled.

<comments> IN VARCHAR2 DEFAULT NULL)

Use the <comments> parameter to specify a comment about the schedule; by default, this parameter is NULL.

Example

The following code fragment calls CREATE_SCHEDULE to create a schedule named weeknights_at_5:

```
EXEC
DBMS_SCHEDULER.CREATE_SCHEDULE (
  schedule_name    => 'weeknights_at_5',
  start_date       => '01-JUN-13 09:00:00.000000',
  repeat_interval  => 'FREQ=DAILY;BYDAY=MON,TUE,WED,THU,FRI;BYHOUR=17;',
  comments         => 'This schedule executes each weeknight at 5:00');
```

The schedule executes each weeknight, at 5:00 pm, effective after June 1, 2013. Since no end_date is specified, the schedule will execute indefinitely until it is disabled with DBMS_SCHEDULER.DISABLE.

3.15.5 DEFINE_PROGRAM_ARGUMENT

Use the `DEFINE_PROGRAM_ARGUMENT` procedure to define a program argument. The `DEFINE_PROGRAM_ARGUMENT` procedure comes in two forms; the first form defines an argument with a default value:

```
DEFINE_PROGRAM_ARGUMENT (
  <program_name> IN VARCHAR2,
  <argument_position> IN PLS_INTEGER,
  <argument_name> IN VARCHAR2 DEFAULT NULL,
  <argument_type> IN VARCHAR2,
  <default_value> IN VARCHAR2,
  <out_argument> IN BOOLEAN DEFAULT FALSE)
```

The second form defines an argument without a default value:

```
DEFINE_PROGRAM_ARGUMENT (
  <program_name> IN VARCHAR2,
  <argument_position> IN PLS_INTEGER,
  <argument_name> IN VARCHAR2 DEFAULT NULL,
  <argument_type> IN VARCHAR2,
  <out_argument> IN BOOLEAN DEFAULT FALSE)
```

Parameters

<program_name>

<program_name> is the name of the program to which the arguments belong.

<argument_position>

<argument_position> specifies the position of the argument as it is passed to the program.

<argument_name>

<argument_name> specifies the optional name of the argument. By default, <argument_name> is NULL.

<argument_type> IN VARCHAR2

<argument_type> specifies the data type of the argument.

<default_value>

<default_value> specifies the default value assigned to the argument. <default_value> will be overridden by a value specified by the job when the job executes.

<out_argument> IN BOOLEAN DEFAULT FALSE

<out_argument> is not currently used; if specified, the value must be FALSE.

Example

The following code fragment uses the `DEFINE_PROGRAM_ARGUMENT` procedure to define the first and second arguments in a program named `add_emp`:

```
EXEC
  DBMS_SCHEDULER.DEFINE_PROGRAM_ARGUMENT (
    program_name      => 'add_emp',
    argument_position  => 1,
    argument_name      => 'dept_no',
    argument_type      => 'INTEGER',
    default_value      => '20');
EXEC
  DBMS_SCHEDULER.DEFINE_PROGRAM_ARGUMENT (
    program_name      => 'add_emp',
    argument_position  => 2,
    argument_name      => 'emp_name',
    argument_type      => 'VARCHAR2');
```

The first argument is an `INTEGER` value named “ dept_no “ that has a default value of 20. The second argument is a `VARCHAR2` value named `emp_name`; the second argument does not have a default value.

3.15.6 DISABLE

Use the `DISABLE` procedure to disable a program or a job. The signature of the `DISABLE` procedure is:

```
DISABLE (  
  <name> IN VARCHAR2,  
  <force> IN BOOLEAN DEFAULT FALSE,  
  <commit_semantics> IN VARCHAR2 DEFAULT 'STOP_ON_FIRST_ERROR')
```

Parameters

<name>

<name> specifies the name of the program or job that is being disabled.

<force>

<force> is accepted for compatibility, and ignored.

<commit_semantics>

<commit_semantics> instructs the server how to handle an error encountered while disabling a program or job. By default, <commit_semantics> is set to `STOP_ON_FIRST_ERROR`, instructing the server to stop when it encounters an error. Any programs or jobs that were successfully disabled prior to the error will be committed to disk.

The `TRANSACTIONAL` and `ABSORB_ERRORS` keywords are accepted for compatibility, and ignored.

Example

The following call to the `DISABLE` procedure disables a program named `update_emp`:

```
DBMS_SCHEDULER.DISABLE ('update_emp');
```

3.15.7 DROP_JOB

Use the `DROP_JOB` procedure to DROP a job, DROP any arguments that belong to the job, and eliminate any future job executions. The signature of the procedure is:

```
DROP_JOB (  
  <job_name> IN VARCHAR2,  
  <force> IN BOOLEAN DEFAULT FALSE,  
  <defer> IN BOOLEAN DEFAULT FALSE,  
  <commit_semantics> IN VARCHAR2 DEFAULT 'STOP_ON_FIRST_ERROR')
```

Parameters

<job_name>

<job_name> specifies the name of the job that is being dropped.

<force>

<force> is accepted for compatibility, and ignored.

<defer>

<defer> is accepted for compatibility, and ignored.

<commit_semantics>

<commit_semantics> instructs the server how to handle an error encountered while dropping a program or job. By default, <commit_semantics> is set to `STOP_ON_FIRST_ERROR`, instructing the server to stop when it encounters an error.

The `TRANSACTIONAL` and `ABSORB_ERRORS` keywords are accepted for compatibility, and ignored.

Example

The following call to `DROP_JOB` drops a job named `update_log`:

```
DBMS_SCHEDULER.DROP_JOB('update_log');
```

3.15.8 DROP_PROGRAM

The DROP_PROGRAM procedure

The signature of the DROP_PROGRAM procedure is:

```
DROP_PROGRAM (  
    <program_name> IN VARCHAR2,  
    <force> IN BOOLEAN DEFAULT FALSE)
```

Parameters

<program_name>

<program_name> specifies the name of the program that is being dropped.

<force>

<force> is a BOOLEAN value that instructs the server how to handle programs with dependent jobs.

Specify FALSE to instruct the server to return an error if the program is referenced by a job.

Specify TRUE to instruct the server to disable any jobs that reference the program before dropping the program.

The default value is FALSE.

Example

The following call to DROP_PROGRAM drops a job named update_emp:

```
DBMS_SCHEDULER.DROP_PROGRAM('update_emp');
```

3.15.9 DROP_PROGRAM_ARGUMENT

Use the `DROP_PROGRAM_ARGUMENT` procedure to drop a program argument. The `DROP_PROGRAM_ARGUMENT` procedure comes in two forms; the first form uses an argument position to specify which argument to drop:

```
drop_program_argument (  
  <program_name> IN VARCHAR2,  
  <argument_position> IN PLS_INTEGER)
```

The second form takes the argument name:

```
drop_program_argument (  
  <program_name> IN VARCHAR2,  
  <argument_name> IN VARCHAR2)
```

Parameters

<program_name>

<program_name> specifies the name of the program that is being modified.

<argument_position>

<argument_position> specifies the position of the argument that is being dropped.

<argument_name>

<argument_name> specifies the name of the argument that is being dropped.

Examples

The following call to `DROP_PROGRAM_ARGUMENT` drops the first argument in the `update_emp` program:

```
DBMS_SCHEDULER.DROP_PROGRAM_ARGUMENT ('update_emp', 1);
```

The following call to `DROP_PROGRAM_ARGUMENT` drops an argument named `emp_name`:

```
DBMS_SCHEDULER.DROP_PROGRAM_ARGUMENT ('update_emp', 'emp_name');
```

3.15.10 DROP_SCHEDULE

Use the DROP_SCHEDULE procedure to drop a schedule. The signature is:

```
DROP_SCHEDULE (  
  <schedule_name> IN VARCHAR2,  
  <force> IN BOOLEAN DEFAULT FALSE)
```

Parameters

<schedule_name>

<schedule_name> specifies the name of the schedule that is being dropped.

<force>

<force> specifies the behavior of the server if the specified schedule is referenced by any job:

- Specify FALSE to instruct the server to return an error if the specified schedule is referenced by a job. This is the default behavior.
- Specify TRUE to instruct the server to disable to any jobs that use the specified schedule before dropping the schedule. Any running jobs will be allowed to complete before the schedule is dropped.

Example

The following call to DROP_SCHEDULE drops a schedule named weeknights_at_5:

```
DBMS_SCHEDULER.DROP_SCHEDULE('weeknights_at_5', TRUE);
```

The server will disable any jobs that use the schedule before dropping the schedule.

3.15.11 ENABLE

Use the `ENABLE` procedure to enable a disabled program or job.

The signature of the `ENABLE` procedure is:

```
ENABLE (  
  <name> IN VARCHAR2,  
  <commit_semantics> IN VARCHAR2 DEFAULT 'STOP_ON_FIRST_ERROR')
```

Parameters

<name>

<name> specifies the name of the program or job that is being enabled.

<commit_semantics>

<commit_semantics> instructs the server how to handle an error encountered while enabling a program or job. By default, <commit_semantics> is set to `STOP_ON_FIRST_ERROR`, instructing the server to stop when it encounters an error.

The `TRANSACTIONAL` and `ABSORB_ERRORS` keywords are accepted for compatibility, and ignored.

Example

The following call to `DBMS_SCHEDULER.ENABLE` enables the `update_emp` program:

```
DBMS_SCHEDULER.ENABLE('update_emp');
```

3.15.12 EVALUATE_CALENDAR_STRING

Use the `EVALUATE_CALENDAR_STRING` procedure to evaluate the `<repeat_interval>` value specified when creating a schedule with the `CREATE_SCHEDULE` procedure. The `EVALUATE_CALENDAR_STRING` procedure will return the date and time that a specified schedule will execute without actually scheduling the job.

The signature of the `EVALUATE_CALENDAR_STRING` procedure is:

```
evaluate_calendar_string(
  <calendar_string> IN VARCHAR2,
  <start_date> IN TIMESTAMP WITH TIME ZONE,
  <return_date_after> IN TIMESTAMP WITH TIME ZONE,
  <next_run_date> OUT TIMESTAMP WITH TIME ZONE)
```

Parameters

`<calendar_string>`

`<calendar_string>` is the calendar string that describes a `<repeat_interval>` that is being evaluated.

`<start_date> IN TIMESTAMP WITH TIME ZONE`

`<start_date>` is the date and time after which the `<repeat_interval>` will become valid.

`<return_date_after>`

Use the `<return_date_after>` parameter to specify the date and time that `EVALUATE_CALENDAR_STRING` should use as a starting date when evaluating the `<repeat_interval>`.

For example, if you specify a `<return_date_after>` value of `01-APR-13 09.00.00.000000`, `EVALUATE_CALENDAR_STRING` will return the date and time of the first iteration of the schedule after April 1st, 2013.

`<next_run_date> OUT TIMESTAMP WITH TIME ZONE`

`<next_run_date>` is an OUT parameter that will contain the first occurrence of the schedule after the date specified by the `<return_date_after>` parameter.

Example

The following example evaluates a calendar string and returns the first date and time that the schedule will be executed after June 15, 2013:

```
DECLARE
  result      TIMESTAMP;
BEGIN

  DBMS_SCHEDULER.EVALUATE_CALENDAR_STRING
  (
    'FREQ=DAILY;BYDAY=MON,TUE,WED,THU,FRI;BYHOUR=17;',
    '15-JUN-2013', NULL, result
  );

  DBMS_OUTPUT.PUT_LINE('next_run_date: ' || result);
END;
/

next_run_date: 17-JUN-13 05.00.00.000000 PM
```

June 15, 2013 is a Saturday; the schedule will not execute until Monday, June 17, 2013 at 5:00 pm.

3.15.13 RUN_JOB

Use the RUN_JOB procedure to execute a job immediately. The signature of the RUN_JOB procedure is:

```
run_job(  
  <job_name> IN VARCHAR2,  
  <use_current_session> IN BOOLEAN DEFAULT TRUE
```

Parameters

<job_name>

<job_name> specifies the name of the job that will execute.

<use_current_session>

By default, the job will execute in the current session. If specified, <use_current_session> must be set to TRUE ; if <use_current_session> is set to FALSE, Advanced Server will return an error.

Example

The following call to RUN_JOB executes a job named update_log:

```
DBMS_SCHEDULER.RUN_JOB('update_log', TRUE);
```

Passing a value of TRUE as the second argument instructs the server to invoke the job in the current session.

3.15.14 SET_JOB_ARGUMENT_VALUE

Use the SET_JOB_ARGUMENT_VALUE procedure to specify a value for an argument. The SET_JOB_ARGUMENT_VALUE procedure comes in two forms; the first form specifies which argument should be modified by position:

```
set_job_argument_value(
  <job_name> IN VARCHAR2,
  <argument_position> IN PLS_INTEGER,
  <argument_value> IN VARCHAR2)
```

The second form uses an argument name to specify which argument to modify:

```
set_job_argument_value(
  <job_name> IN VARCHAR2,
  <argument_name> IN VARCHAR2,
  <argument_value> IN VARCHAR2)
```

Argument values set by the SET_JOB_ARGUMENT_VALUE procedure override any values set by default.

Parameters

<job_name>

<job_name> specifies the name of the job to which the modified argument belongs.

<argument_position>

Use <argument_position> to specify the argument position for which the value will be set.

<argument_name>

Use <argument_name> to specify the argument by name for which the value will be set.

<argument_value>

<argument_value> specifies the new value of the argument.

Examples

The following example assigns a value of 30 to the first argument in the update_emp job:

```
DBMS_SCHEDULER.SET_JOB_ARGUMENT_VALUE('update_emp', 1, '30');
```

The following example sets the emp_name argument to SMITH:

```
DBMS_SCHEDULER.SET_JOB_ARGUMENT_VALUE('update_emp', 'emp_name',
'SMITH');
```

3.16 DBMS_SESSION

Advanced Server provides support for the following DBMS_SESSION.SET_ROLE procedure:

Function/Procedure	Return Type	Description
SET_ROLE(<role_cmd>)	n/a	Executes a SET ROLE statement followed by the string value specified in <role_cmd>.

Advanced Server's implementation of DBMS_SESSION is a partial implementation when compared to Oracle's version. Only DBMS_SESSION.SET_ROLE is supported.

3.16.1 SET_ROLE

The SET_ROLE procedure sets the current session user to the role specified in <role_cmd>. After invoking the SET_ROLE procedure, the current session will use the permissions assigned to the specified role. The signature of the procedure is:

```
SET_ROLE (<role_cmd>)
```

The SET_ROLE procedure appends the value specified for <role_cmd> to the SET ROLE statement, and then invokes the statement.

Parameters

<role_cmd>

<role_cmd> specifies a role name in the form of a string value.

Example

The following call to the SET_ROLE procedure invokes the SET ROLE command to set the identity of the current session user to manager:

```
edb=# exec DBMS_SESSION.SET_ROLE('manager');
```

3.17 DBMS_SQL

The DBMS_SQL package provides an application interface compatible with Oracle databases to the EnterpriseDB dynamic SQL functionality. With DBMS_SQL you can construct queries and other commands at run time (rather than when you write the application). EnterpriseDB Advanced Server offers native support for dynamic SQL; DBMS_SQL provides a way to use dynamic SQL in a fashion compatible with Oracle databases without modifying your application.

DBMS_SQL assumes the privileges of the current user when executing dynamic SQL statements.

Function/Procedure	Function or Procedure	Return Type	Description
BIND_VARIABLE(c, name, value [, out_value_size])	Procedure	n/a	Bind a value to a variable.
BIND_VARIABLE_CHAR(c, name, value [, out_value_size])	Procedure	n/a	Bind a CHAR value to a variable.
BIND_VARIABLE_RAW(c, name, value [, out_value_size])	Procedure	n/a	Bind a RAW value to a variable.
CLOSE_CURSOR(c IN OUT)	Procedure	n/a	Close a cursor.
COLUMN_VALUE(c, position, value OUT [, column_error OUT [, actual_length OUT]])	Procedure	n/a	Return a column value into a variable.
COLUMN_VALUE_CHAR(c, position, value OUT [, column_error OUT [, actual_length OUT]])	Procedure	n/a	Return a CHAR column value into a variable.
COLUMN_VALUE_RAW(c, position, value OUT [, column_error OUT [, actual_length OUT]])	Procedure	n/a	Return a RAW column value into a variable.
DEFINE_COLUMN(c, position, column [, column_size])	Procedure	n/a	Define a column in the SELECT list.
DEFINE_COLUMN_CHAR(c, position, column, column_size)	Procedure	n/a	Define a CHAR column in the SELECT list.
DEFINE_COLUMN_RAW(c, position, column, column_size)	Procedure	n/a	Define a RAW column in the SELECT list.
DESCRIBE_COLUMNS	Procedure	n/a	Defines columns to hold a cursor result set.
EXECUTE(c)	Function	INTEGER	Execute a cursor.
EXECUTE_AND_FETCH(c [, exact])	Function	INTEGER	Execute a cursor and fetch a single row.
FETCH_ROWS(c)	Function	INTEGER	Fetch rows from the cursor.
IS_OPEN(c)	Function	BOOLEAN	Check if a cursor is open.
LAST_ROW_COUNT	Function	INTEGER	Return cumulative number of rows fetched.
OPEN_CURSOR	Function	INTEGER	Open a cursor.
PARSE(c, statement, language_flag)	Procedure	n/a	Parse a statement.

Advanced Server's implementation of DBMS_SQL is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The following table lists the public variable available in the DBMS_SQL package.

Public Variables	Data Type	Value	Description
native	INTEGER	1	Provided for compatibility with Oracle syntax. See DBMS_SQL.PARSE for more information.
V6	INTEGER	2	Provided for compatibility with Oracle syntax. See DBMS_SQL.PARSE for more information.
V7	INTEGER	3	Provided for compatibility with Oracle syntax. See DBMS_SQL.PARSE for more information.

3.17.1 BIND_VARIABLE

The `BIND_VARIABLE` procedure provides the capability to associate a value with an `IN` or `IN OUT` bind variable in a SQL command.

```

BIND_VARIABLE(c INTEGER, <name> VARCHAR2,
              <value> { BLOB | CLOB | DATE | FLOAT | INTEGER | NUMBER |
                        TIMESTAMP | VARCHAR2 }
              [, <out_value_size> INTEGER ])

```

Parameters

<c>

Cursor ID of the cursor for the SQL command with bind variables.

<name>

Name of the bind variable in the SQL command.

<value>

Value to be assigned.

<out_value_size>

If <name> is an `IN OUT` variable, defines the maximum length of the output value. If not specified, the length of <value> is assumed.

Examples

The following anonymous block uses bind variables to insert a row into the `emp` table.

```

DECLARE
    curid          INTEGER;
    v_sql          VARCHAR2(150) := 'INSERT INTO emp VALUES ' ||
                                     '(:p_empno, :p_ename, :p_job, :p_mgr, ' ||
                                     ':p_hiredate, :p_sal, :p_comm, :p_deptno)';
    v_empno        emp.empno%TYPE;
    v_ename        emp.ename%TYPE;
    v_job          emp.job%TYPE;
    v_mgr          emp.mgr%TYPE;
    v_hiredate     emp.hiredate%TYPE;
    v_sal          emp.sal%TYPE;
    v_comm         emp.comm%TYPE;
    v_deptno       emp.deptno%TYPE;
    v_status       INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.NATIVE);
    v_empno := 9001;
    v_ename := 'JONES';
    v_job   := 'SALESMAN';
    v_mgr   := 7369;
    v_hiredate := TO_DATE('13-DEC-07', 'DD-MON-YY');
    v_sal    := 8500.00;
    v_comm   := 1500.00;
    v_deptno := 40;
    DBMS_SQL.BIND_VARIABLE(curid, ':p_empno', v_empno);
    DBMS_SQL.BIND_VARIABLE(curid, ':p_ename', v_ename);
    DBMS_SQL.BIND_VARIABLE(curid, ':p_job', v_job);

```

```
DBMS_SQL.BIND_VARIABLE(curid,':p_mgr',v_mgr);
DBMS_SQL.BIND_VARIABLE(curid,':p_hiredate',v_hiredate);
DBMS_SQL.BIND_VARIABLE(curid,':p_sal',v_sal);
DBMS_SQL.BIND_VARIABLE(curid,':p_comm',v_comm);
DBMS_SQL.BIND_VARIABLE(curid,':p_deptno',v_deptno);
v_status := DBMS_SQL.EXECUTE(curid);
DBMS_OUTPUT.PUT_LINE('Number of rows processed: ' || v_status);
DBMS_SQL.CLOSE_CURSOR(curid);
END;
```

Number of rows processed: 1

3.17.2 BIND_VARIABLE_CHAR

The `BIND_VARIABLE_CHAR` procedure provides the capability to associate a `CHAR` value with an `IN` or `IN OUT` bind variable in a SQL command.

```
    BIND_VARIABLE_CHAR(<c> INTEGER, <name> VARCHAR2, <value> CHAR  
                      [, <out_value_size> INTEGER ])
```

Parameters

<c>

Cursor ID of the cursor for the SQL command with bind variables.

<name>

Name of the bind variable in the SQL command.

<value>

Value of type `CHAR` to be assigned.

<out_value_size>

If <name> is an `IN OUT` variable, defines the maximum length of the output value. If not specified, the length of <value> is assumed.

3.17.3 BIND VARIABLE RAW

The `BIND_VARIABLE_RAW` procedure provides the capability to associate a RAW value with an IN or IN OUT bind variable in a SQL command.

```
    BIND_VARIABLE_RAW(<c> INTEGER, <name> VARCHAR2, <value> RAW  
                    [, <out_value_size> INTEGER ])
```

Parameters

<c>

Cursor ID of the cursor for the SQL command with bind variables.

<name>

Name of the bind variable in the SQL command.

<value>

Value of type RAW to be assigned.

<out_value_size>

If <name> is an IN OUT variable, defines the maximum length of the output value. If not specified, the length of <value> is assumed.

3.17.4 CLOSE_CURSOR

The `CLOSE_CURSOR` procedure closes an open cursor. The resources allocated to the cursor are released and it can no longer be used.

```
CLOSE_CURSOR(<c> IN OUT INTEGER)
```

Parameters

<c>

Cursor ID of the cursor to be closed.

Examples

The following example closes a previously opened cursor:

```
DECLARE
    curid          INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    .
    .
    .
    DBMS_SQL.CLOSE_CURSOR(curid);
END;
```

3.17.5 COLUMN_VALUE

The COLUMN_VALUE procedure defines a variable to receive a value from a cursor.

```
COLUMN_VALUE(<c> INTEGER, <position> INTEGER, <value> OUT { BLOB |
              CLOB | DATE | FLOAT | INTEGER | NUMBER | TIMESTAMP | VARCHAR2 }
              [, <column_error> OUT NUMBER [, <actual_length> OUT INTEGER ]])
```

Parameters

<c>

Cursor id of the cursor returning data to the variable being defined.

<position>

Position within the cursor of the returned data. The first value in the cursor is position 1.

<value>

Variable receiving the data returned in the cursor by a prior fetch call.

<column_error>

Error number associated with the column, if any.

<actual_length>

Actual length of the data prior to any truncation.

Examples

The following example shows the portion of an anonymous block that receives the values from a cursor using the COLUMN_VALUE procedure.

```
DECLARE
    curid          INTEGER;
    v_empno        NUMBER(4);
    v_ename        VARCHAR2(10);
    v_hiredate     DATE;
    v_sal          NUMBER(7,2);
    v_comm         NUMBER(7,2);
    v_sql          VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, ' ||
                                   'comm FROM emp';
    v_status       INTEGER;
BEGIN
    .
    .
    .
    LOOP
        v_status := DBMS_SQL.FETCH_ROWS(curid);
        EXIT WHEN v_status = 0;
        DBMS_SQL.COLUMN_VALUE(curid,1,v_empno);
        DBMS_SQL.COLUMN_VALUE(curid,2,v_ename);
        DBMS_SQL.COLUMN_VALUE(curid,3,v_hiredate);
        DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
        DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
        DBMS_SQL.COLUMN_VALUE(curid,5,v_comm);
        DBMS_OUTPUT.PUT_LINE(v_empno || ' ' || RPAD(v_ename,10) || ' ' ||
                              TO_CHAR(v_hiredate,'yyyy-mm-dd') || ' ' ||
                              TO_CHAR(v_sal,'9,999.99') || ' ' ||
                              TO_CHAR(NVL(v_comm,0),'9,999.99'));
    END LOOP;
```

```
END LOOP;  
DBMS_SQL.CLOSE_CURSOR(curid);  
END;
```

3.17.6 COLUMN_VALUE_CHAR

The COLUMN_VALUE_CHAR procedure defines a variable to receive a CHAR value from a cursor.

```
COLUMN_VALUE_CHAR(<c> INTEGER, <position> INTEGER, <value> OUT CHAR  
                [, <column_error> OUT NUMBER [, <actual_length> OUT INTEGER ]])
```

Parameters

<c>

Cursor id of the cursor returning data to the variable being defined.

<position>

Position within the cursor of the returned data. The first value in the cursor is position 1.

<value>

Variable of data type CHAR receiving the data returned in the cursor by a prior fetch call.

<column_error>

Error number associated with the column, if any.

<actual_length>

Actual length of the data prior to any truncation.

3.17.7 COLUMN_VALUE_RAW

The COLUMN_VALUE_RAW procedure defines a variable to receive a RAW value from a cursor.

```
COLUMN_VALUE_RAW(<c> INTEGER, <position> INTEGER, <value> OUT RAW  
                [, <column_error> OUT NUMBER [, <actual_length> OUT INTEGER ]])
```

Parameters

<c>

Cursor id of the cursor returning data to the variable being defined.

<position>

Position within the cursor of the returned data. The first value in the cursor is position 1.

<value>

Variable of data type RAW receiving the data returned in the cursor by a prior fetch call.

<column_error>

Error number associated with the column, if any.

<actual_length>

Actual length of the data prior to any truncation.

3.17.8 DEFINE_COLUMN

The `DEFINE_COLUMN` procedure defines a column or expression in the `SELECT` list that is to be returned and retrieved in a cursor.

```
DEFINE_COLUMN(<c> INTEGER, <position> INTEGER, <column> { BLOB |
    CLOB | DATE | FLOAT | INTEGER | NUMBER | TIMESTAMP | VARCHAR2 }
    [, <column_size> INTEGER ])
```

Parameters

<c>

Cursor id of the cursor associated with the `SELECT` command.

<position>

Position of the column or expression in the `SELECT` list that is being defined.

<column>

A variable that is of the same data type as the column or expression in position <position> of the `SELECT` list.

<column_size>

The maximum length of the returned data. <column_size> must be specified only if <column> is `VARCHAR2`. Returned data exceeding <column_size> is truncated to <column_size> characters.

Examples

The following shows how the `empno`, `ename`, `hiredate`, `sal`, and `comm` columns of the `emp` table are defined with the `DEFINE_COLUMN` procedure.

```
DECLARE
    curid          INTEGER;
    v_empno        NUMBER(4);
    v_ename        VARCHAR2(10);
    v_hiredate     DATE;
    v_sal          NUMBER(7,2);
    v_comm         NUMBER(7,2);
    v_sql          VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, ' ||
                                'comm FROM emp';
    v_status       INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    DBMS_SQL.DEFINE_COLUMN(curid, 1, v_empno);
    DBMS_SQL.DEFINE_COLUMN(curid, 2, v_ename, 10);
    DBMS_SQL.DEFINE_COLUMN(curid, 3, v_hiredate);
    DBMS_SQL.DEFINE_COLUMN(curid, 4, v_sal);
    DBMS_SQL.DEFINE_COLUMN(curid, 5, v_comm);
    .
    .
    .
END;
```

The following shows an alternative to the prior example that produces the exact same results. Note that the lengths of the data types are irrelevant – the `empno`, `sal`, and `comm` columns will still return data equivalent to `NUMBER(4)` and `NUMBER(7,2)`, respectively, even though `v_num` is defined as `NUMBER(1)` (assuming the declarations in the `COLUMN_VALUE` procedure are of the appropriate maximum sizes). The `ename` column will return data up to ten characters in length as defined by the `<length>` parameter in the `DEFINE_COLUMN` call, not by the data type declaration, `VARCHAR2(1)` declared for `v_varchar`. The actual size of the returned data is dictated by the `COLUMN_VALUE` procedure.

```
DECLARE
    curid          INTEGER;
    v_num          NUMBER(1);
    v_varchar      VARCHAR2(1);
    v_date         DATE;
    v_sql          VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, ' ||
                                'comm FROM emp';
    v_status       INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    DBMS_SQL.DEFINE_COLUMN(curid, 1, v_num);
    DBMS_SQL.DEFINE_COLUMN(curid, 2, v_varchar, 10);
    DBMS_SQL.DEFINE_COLUMN(curid, 3, v_date);
    DBMS_SQL.DEFINE_COLUMN(curid, 4, v_num);
    DBMS_SQL.DEFINE_COLUMN(curid, 5, v_num);
    .
    .
    .
END;
```


3.17.9 DEFINE_COLUMN_CHAR

The `DEFINE_COLUMN_CHAR` procedure defines a CHAR column or expression in the `SELECT` list that is to be returned and retrieved in a cursor.

```
DEFINE_COLUMN_CHAR(<c> INTEGER, <position> INTEGER, <column>      CHAR,  
                  <column_size> INTEGER)
```

Parameters

<c>

Cursor id of the cursor associated with the `SELECT` command.

<position>

Position of the column or expression in the `SELECT` list that is being defined.

<column>

A CHAR variable.

<column_size>

The maximum length of the returned data. Returned data exceeding <column_size> is truncated to <column_size> characters.

3.17.10 DEFINE_COLUMN_RAW

The DEFINE_COLUMN_RAW procedure defines a RAW column or expression in the SELECT list that is to be returned and retrieved in a cursor.

```
DEFINE_COLUMN_RAW(<c> INTEGER, <position> INTEGER, <column> RAW,  
                 <column_size> INTEGER)
```

Parameters

<c>

Cursor id of the cursor associated with the SELECT command.

<position>

Position of the column or expression in the SELECT list that is being defined.

<column>

A RAW variable.

<column_size>

The maximum length of the returned data. Returned data exceeding <column_size> is truncated to <column_size> characters.

3.17.11 DESCRIBE COLUMNS

The DESCRIBE_COLUMNS procedure describes the columns returned by a cursor.

```
DESCRIBE_COLUMNS(c INTEGER, col_cnt OUT INTEGER, desc_t OUT  
DESC_TAB);
```

Parameters

<c>

The cursor ID of the cursor.

<col_cnt>

The number of columns in cursor result set.

<desc_tab>

The table that contains a description of each column returned by the cursor. The descriptions are of type DESC_REC, and contain the following values:

Column Name	Type
col_type	INTEGER
col_max_len	INTEGER
col_name	VARCHAR2(128)
col_name_len	INTEGER
col_schema_name	VARCHAR2(128)
col_schema_name_len	INTEGER
col_precision	INTEGER
col_scale	INTEGER
col_charsetid	INTEGER
col_charsetform	INTEGER
col_null_ok	BOOLEAN

3.17.12 EXECUTE

The EXECUTE function executes a parsed SQL command or SPL block.

```
<status> INTEGER EXECUTE(<c> INTEGER)
```

Parameters

<c>

Cursor ID of the parsed SQL command or SPL block to be executed.

<status>

Number of rows processed if the SQL command was DELETE, INSERT, or UPDATE. <status> is meaningless for all other commands.

Examples

The following anonymous block inserts a row into the dept table.

```
DECLARE
    curid          INTEGER;
    v_sql          VARCHAR2(50);
    v_status       INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    v_sql := 'INSERT INTO dept VALUES (50, ''HR'', ''LOS ANGELES'')';
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('Number of rows processed: ' || v_status);
    DBMS_SQL.CLOSE_CURSOR(curid);
END;
```

3.17.13 EXECUTE_AND_FETCH

Function EXECUTE_AND_FETCH executes a parsed SELECT command and fetches one row.

```
<status> INTEGER EXECUTE_AND_FETCH(<c> INTEGER
    [, <exact> BOOLEAN ])
```

Parameters

<c>

Cursor id of the cursor for the SELECT command to be executed.

<exact>

If set to TRUE, an exception is thrown if the number of rows in the result set is not exactly equal to 1. If set to FALSE, no exception is thrown. The default is FALSE. A NO_DATA_FOUND exception is thrown if <exact> is TRUE and there are no rows in the result set. A TOO_MANY_ROWS exception is thrown if <exact> is TRUE and there is more than one row in the result set.

<status>

Returns 1 if a row was successfully fetched, 0 if no rows to fetch. If an exception is thrown, no value is returned.

Examples

The following stored procedure uses the EXECUTE_AND_FETCH function to retrieve one employee using the employee's name. An exception will be thrown if the employee is not found, or there is more than one employee with the same name.

```
CREATE OR REPLACE PROCEDURE select_by_name(
    p_ename          emp.ename%TYPE
)
IS
    curid            INTEGER;
    v_empno          emp.empno%TYPE;
    v_hiredate       emp.hiredate%TYPE;
    v_sal            emp.sal%TYPE;
    v_comm           emp.comm%TYPE;
    v_dname          dept.dname%TYPE;
    v_disp_date      VARCHAR2(10);
    v_sql            VARCHAR2(120) := 'SELECT empno, hiredate, sal, ' ||
                                     'NVL(comm, 0), dname ' ||
                                     'FROM emp e, dept d ' ||
                                     'WHERE ename = :p_ename ' ||
                                     'AND e.deptno = d.deptno';

    v_status         INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    DBMS_SQL.BIND_VARIABLE(curid, ':p_ename', UPPER(p_ename));
    DBMS_SQL.DEFINE_COLUMN(curid, 1, v_empno);
    DBMS_SQL.DEFINE_COLUMN(curid, 2, v_hiredate);
    DBMS_SQL.DEFINE_COLUMN(curid, 3, v_sal);
    DBMS_SQL.DEFINE_COLUMN(curid, 4, v_comm);
    DBMS_SQL.DEFINE_COLUMN(curid, 5, v_dname, 14);
    v_status := DBMS_SQL.EXECUTE_AND_FETCH(curid, TRUE);
    DBMS_SQL.COLUMN_VALUE(curid, 1, v_empno);
    DBMS_SQL.COLUMN_VALUE(curid, 2, v_hiredate);
    DBMS_SQL.COLUMN_VALUE(curid, 3, v_sal);
```

```

DBMS_SQL.COLUMN_VALUE(curid,4,v_comm);
DBMS_SQL.COLUMN_VALUE(curid,5,v_dname);
v_disp_date := TO_CHAR(v_hiredate, 'MM/DD/YYYY');
DBMS_OUTPUT.PUT_LINE('Number      : ' || v_empno);
DBMS_OUTPUT.PUT_LINE('Name        : ' || UPPER(p_ename));
DBMS_OUTPUT.PUT_LINE('Hire Date  : ' || v_disp_date);
DBMS_OUTPUT.PUT_LINE('Salary     : ' || v_sal);
DBMS_OUTPUT.PUT_LINE('Commission: ' || v_comm);
DBMS_OUTPUT.PUT_LINE('Department: ' || v_dname);
DBMS_SQL.CLOSE_CURSOR(curid);
EXCEPTION
  WHEN NO_DATA_FOUND THEN
    DBMS_OUTPUT.PUT_LINE('Employee ' || p_ename || ' not found');
    DBMS_SQL.CLOSE_CURSOR(curid);
  WHEN TOO_MANY_ROWS THEN
    DBMS_OUTPUT.PUT_LINE('Too many employees named, ' ||
      p_ename || ', found');
    DBMS_SQL.CLOSE_CURSOR(curid);
  WHEN OTHERS THEN
    DBMS_OUTPUT.PUT_LINE('The following is SQLERRM:');
    DBMS_OUTPUT.PUT_LINE(SQLERRM);
    DBMS_OUTPUT.PUT_LINE('The following is SQLCODE:');
    DBMS_OUTPUT.PUT_LINE(SQLCODE);
    DBMS_SQL.CLOSE_CURSOR(curid);
END;

EXEC select_by_name('MARTIN')

Number      : 7654
Name        : MARTIN
Hire Date   : 09/28/1981
Salary      : 1250
Commission  : 1400
Department  : SALES

```

3.17.14 FETCH_ROWS

The `FETCH_ROWS` function retrieves a row from a cursor.

```
<status> INTEGER FETCH_ROWS (<c> INTEGER)
```

Parameters

<c>

Cursor ID of the cursor from which to fetch a row.

<status>

Returns 1 if a row was successfully fetched, 0 if no more rows to fetch.

Examples

The following examples fetches the rows from the `emp` table and displays the results.

```
DECLARE
    curid          INTEGER;
    v_empno        NUMBER(4);
    v_ename        VARCHAR2(10);
    v_hiredate     DATE;
    v_sal          NUMBER(7,2);
    v_comm         NUMBER(7,2);
    v_sql          VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, ' ||
                                'comm FROM emp';
    v_status       INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    DBMS_SQL.DEFINE_COLUMN(curid, 1, v_empno);
    DBMS_SQL.DEFINE_COLUMN(curid, 2, v_ename, 10);
    DBMS_SQL.DEFINE_COLUMN(curid, 3, v_hiredate);
    DBMS_SQL.DEFINE_COLUMN(curid, 4, v_sal);
    DBMS_SQL.DEFINE_COLUMN(curid, 5, v_comm);

    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('EMPNO   ENAME           HIREDATE       SAL           COMM');
    DBMS_OUTPUT.PUT_LINE('-----  -');
    LOOP
        v_status := DBMS_SQL.FETCH_ROWS(curid);
        EXIT WHEN v_status = 0;
        DBMS_SQL.COLUMN_VALUE(curid, 1, v_empno);
        DBMS_SQL.COLUMN_VALUE(curid, 2, v_ename);
        DBMS_SQL.COLUMN_VALUE(curid, 3, v_hiredate);
        DBMS_SQL.COLUMN_VALUE(curid, 4, v_sal);
        DBMS_SQL.COLUMN_VALUE(curid, 5, v_comm);
        DBMS_OUTPUT.PUT_LINE(v_empno || ' ' || RPAD(v_ename, 10) || ' ' ||
            TO_CHAR(v_hiredate, 'yyyy-mm-dd') || ' ' ||
            TO_CHAR(v_sal, '9,999.99') || ' ' ||
            TO_CHAR(NVL(v_comm, 0), '9,999.99'));
    END LOOP;
    DBMS_SQL.CLOSE_CURSOR(curid);
END;
```

EMPNO	ENAME	HIREDATE	SAL	COMM
7369	SMITH	1980-12-17	800.00	
7469	ALLEN	1981-01-21	1600.00	.15
7569	WARD	1981-02-21	1200.00	.15
7669	MARTIN	1981-03-24	1200.00	.15
7769	BLAKE	1981-05-01	2800.00	
7869	CLARK	1981-06-07	800.00	
7969	ADAMS	1981-07-13	750.00	
8069	JAMES	1981-09-17	950.00	
8169	TURNER	1981-09-28	1200.00	.15
8269	POWERS	1981-10-08	900.00	
8369	BEVER	1981-11-07	1100.00	.15
8469	COOPER	1981-12-09	850.00	
8569	PEROTT	1982-01-08	2900.00	
8669	FURNESS	1982-02-01	3000.00	
8769	STEINER	1982-02-22	2800.00	
8869	GLADSTONE	1982-03-15	2400.00	
8969	POPE	1982-04-07	2900.00	
9069	DESSA	1982-04-19	2400.00	
9169	CHANDLER	1982-05-12	2200.00	
9269	WATSON	1982-06-06	2100.00	
9369	FRANK	1982-06-23	1700.00	
9469	MACKEY	1982-07-17	1200.00	
9569	JEANS	1982-08-16	900.00	
9669	WATSON	1982-09-14	800.00	
9769	COOPER	1982-10-13	750.00	
9869	JOHNSON	1982-11-11	700.00	
9969	SMITH	1982-12-10	650.00	

7369	SMITH	1980-12-17	800.00	.00
7499	ALLEN	1981-02-20	1,600.00	300.00
7521	WARD	1981-02-22	1,250.00	500.00
7566	JONES	1981-04-02	2,975.00	.00
7654	MARTIN	1981-09-28	1,250.00	1,400.00
7698	BLAKE	1981-05-01	2,850.00	.00
7782	CLARK	1981-06-09	2,450.00	.00
7788	SCOTT	1987-04-19	3,000.00	.00
7839	KING	1981-11-17	5,000.00	.00
7844	TURNER	1981-09-08	1,500.00	.00
7876	ADAMS	1987-05-23	1,100.00	.00
7900	JAMES	1981-12-03	950.00	.00
7902	FORD	1981-12-03	3,000.00	.00
7934	MILLER	1982-01-23	1,300.00	.00

3.17.15 IS_OPEN

The IS_OPEN function provides the capability to test if the given cursor is open.

```
<status> BOOLEAN IS_OPEN(<c> INTEGER)
```

Parameters

<c>

Cursor ID of the cursor to be tested.

<status>

Set to TRUE if the cursor is open, set to FALSE if the cursor is not open.

3.17.16 LAST_ROW_COUNT

The LAST_ROW_COUNT function returns the number of rows that have been currently fetched.

<rowcnt> INTEGER LAST_ROW_COUNT

Parameters

<rowcnt>

Number of row fetched thus far.

Examples

The following example uses the LAST_ROW_COUNT function to display the total number of rows fetched in the query.

```

DECLARE
    curid          INTEGER;
    v_empno        NUMBER(4);
    v_ename        VARCHAR2(10);
    v_hiredate     DATE;
    v_sal          NUMBER(7,2);
    v_comm         NUMBER(7,2);
    v_sql          VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, ' ||
                                'comm FROM emp';
    v_status       INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQL.PARSE(curid,v_sql,DBMS_SQL.native);
    DBMS_SQL.DEFINE_COLUMN(curid,1,v_empno);
    DBMS_SQL.DEFINE_COLUMN(curid,2,v_ename,10);
    DBMS_SQL.DEFINE_COLUMN(curid,3,v_hiredate);
    DBMS_SQL.DEFINE_COLUMN(curid,4,v_sal);
    DBMS_SQL.DEFINE_COLUMN(curid,5,v_comm);

    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('EMPNO   ENAME       HIREDATE    SAL        COMM');
    DBMS_OUTPUT.PUT_LINE('-----  -');
    LOOP
        v_status := DBMS_SQL.FETCH_ROWS(curid);
        EXIT WHEN v_status = 0;
        DBMS_SQL.COLUMN_VALUE(curid,1,v_empno);
        DBMS_SQL.COLUMN_VALUE(curid,2,v_ename);
        DBMS_SQL.COLUMN_VALUE(curid,3,v_hiredate);
        DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
        DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
        DBMS_SQL.COLUMN_VALUE(curid,5,v_comm);
        DBMS_OUTPUT.PUT_LINE(v_empno || ' ' || RPAD(v_ename,10) || ' ' ||
            TO_CHAR(v_hiredate,'yyyy-mm-dd') || ' ' ||
            TO_CHAR(v_sal,'9,999.99') || ' ' ||
            TO_CHAR(NVL(v_comm,0),'9,999.99'));
    END LOOP;
    DBMS_OUTPUT.PUT_LINE('Number of rows: ' || DBMS_SQL.LAST_ROW_COUNT);
    DBMS_SQL.CLOSE_CURSOR(curid);
END;
```

EMPNO	ENAME	HIREDATE	SAL	COMM
7369	SMITH	1980-12-17	800.00	.00
7499	ALLEN	1981-02-20	1,600.00	300.00

7521	WARD	1981-02-22	1,250.00	500.00
7566	JONES	1981-04-02	2,975.00	.00
7654	MARTIN	1981-09-28	1,250.00	1,400.00
7698	BLAKE	1981-05-01	2,850.00	.00
7782	CLARK	1981-06-09	2,450.00	.00
7788	SCOTT	1987-04-19	3,000.00	.00
7839	KING	1981-11-17	5,000.00	.00
7844	TURNER	1981-09-08	1,500.00	.00
7876	ADAMS	1987-05-23	1,100.00	.00
7900	JAMES	1981-12-03	950.00	.00
7902	FORD	1981-12-03	3,000.00	.00
7934	MILLER	1982-01-23	1,300.00	.00

Number of rows: 14

3.17.17 OPEN_CURSOR

The `OPEN_CURSOR` function creates a new cursor. A cursor must be used to parse and execute any dynamic SQL statement. Once a cursor has been opened, it can be re-used with the same or different SQL statements. The cursor does not have to be closed and re-opened in order to be re-used.

`<c> INTEGER OPEN_CURSOR`

Parameters

`<c>`

Cursor ID number associated with the newly created cursor.

Examples

The following example creates a new cursor:

```
DECLARE
    curid          INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    .
    .
    .
END;
```

3.17.18 PARSE

The PARSE procedure parses a SQL command or SPL block. If the SQL command is a DDL command, it is immediately executed and does not require running the EXECUTE function.

```
PARSE(<c> INTEGER, <statement> VARCHAR2, <language_flag> INTEGER)
```

Parameters

<c>

Cursor ID of an open cursor.

<statement>

SQL command or SPL block to be parsed. A SQL command must not end with the semi-colon terminator, however an SPL block does require the semi-colon terminator.

<language_flag>

Language flag provided for compatibility with Oracle syntax. Use DBMS_SQL.V6, DBMS_SQL.V7 or DBMS_SQL.native. This flag is ignored, and all syntax is assumed to be in EnterpriseDB Advanced Server form.

Examples

The following anonymous block creates a table named, job. Note that DDL statements are executed immediately by the PARSE procedure and do not require a separate EXECUTE step.

```
DECLARE
    curid          INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQL.PARSE(curid, 'CREATE TABLE job (jobno NUMBER(3), ' ||
        'jname VARCHAR2(9))', DBMS_SQL.native);
    DBMS_SQL.CLOSE_CURSOR(curid);
END;
```

The following inserts two rows into the job table.

```
DECLARE
    curid          INTEGER;
    v_sql          VARCHAR2(50);
    v_status       INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    v_sql := 'INSERT INTO job VALUES (100, ''ANALYST'')';
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('Number of rows processed: ' || v_status);
    v_sql := 'INSERT INTO job VALUES (200, ''CLERK'')';
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('Number of rows processed: ' || v_status);
    DBMS_SQL.CLOSE_CURSOR(curid);
END;
```

Number of rows processed: 1
Number of rows processed: 1

The following anonymous block uses the DBMS_SQL package to execute a block containing two INSERT statements. Note that the end of the block contains a terminating semi-colon, while in the prior example, each individual INSERT statement does not have a terminating semi-colon.

```
DECLARE
    curid          INTEGER;
    v_sql          VARCHAR2(100);
    v_status       INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    v_sql := 'BEGIN ' ||
              'INSERT INTO job VALUES (300, ''MANAGER''); ' ||
              'INSERT INTO job VALUES (400, ''SALESMAN''); ' ||
              'END;';
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_SQL.CLOSE_CURSOR(curid);
END;
```

3.18 DBMS_UTILITY

The DBMS_UTILITY package provides support for the following various utility programs:

Function/Procedure	Function or Procedure	Return Type	Description
ANALYZE_DATABASE(method [, estimate_rows [, estimate_percent [, method_opt]])	Procedure	n/a	Analyze database tables.
ANALYZE_PART_OBJECT(schema, object_name [, object_type [, command_type [, command_opt [, sample_clause]]])	Procedure	n/a	Analyze a partitioned table.
ANALYZE_SCHEMA(schema, method [, estimate_rows [, estimate_percent [, method_opt]])	Procedure	n/a	Analyze schema tables.
CANONICALIZE(name, canon_name OUT, canon_len)	Procedure	n/a	Canonicalizes a string – e.g., strips off white space.
COMMA_TO_TABLE(list, tablen OUT, tab OUT)	Procedure	n/a	Convert a comma-delimited list of names to a table of names.
DB_VERSION(version OUT, compatibility OUT)	Procedure	n/a	Get the database version.
EXEC_DDL_STATEMENT(parse_string)	Procedure	n/a	Execute a DDL statement.
FORMAT_CALL_STACK	Function	TEXT	Formats the current call stack.
GET_CPU_TIME	Function	NUMBER	Get the current CPU time.
GET_DEPENDENCY(type, schema, name)	Procedure	n/a	Get objects that are dependent upon the given object..
GET_HASH_VALUE(name, base, hash_size)	Function	NUMBER	Compute a hash value.
GET_PARAMETER_VALUE(parnam, interval OUT, strval OUT)	Procedure	BINARY_INTEGER	Get database initialization parameter settings.
GET_TIME	Function	NUMBER	Get the current time.
NAME_TOKENIZE(name, a OUT, b OUT, c OUT, dblink OUT, nextpos OUT)	Procedure	n/a	Parse the given name into its component parts.
TABLE_TO_COMMA(tab, tablen OUT, list OUT)	Procedure	n/a	Convert a table of names to a comma-delimited list.

Advanced Server's implementation of DBMS_UTILITY is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The following table lists the public variables available in the DBMS_UTILITY package.

Public Variables	Data Type	Value	Description
inv_error_on_restrictions	PLS_INTEGER	1	Used by the INVALIDATE procedure.
lname_array	TABLE		For lists of long names.
uncl_array	TABLE		For lists of users and names.

3.18.1 LNAME_ARRAY

The LNAME_ARRAY is for storing lists of long names including fully-qualified names.

```
TYPE lname_array IS `TABLE OF VARCHAR2(4000) INDEX BY BINARY_INTEGER;
```

3.18.2 UNCL_ARRAY

The UNCL_ARRAY is for storing lists of users and names.

```
TYPE uncl_array IS TABLE OF VARCHAR2(227) INDEX BY BINARY_INTEGER;
```


3.18.3 ANALYZE_DATABASE, ANALYZE_SCHEMA and ANALYZE_PART_OBJECT

The `ANALYZE_DATABASE()`, `ANALYZE_SCHEMA()` and `ANALYZE_PART_OBJECT()` procedures provide the capability to gather statistics on tables in the database. When you execute the `ANALYZE` statement, Postgres samples the data in a table and records distribution statistics in the `pg_statistics` system table.

`ANALYZE_DATABASE`, `ANALYZE_SCHEMA`, and `ANALYZE_PART_OBJECT` differ primarily in the number of tables that are processed:

- `ANALYZE_DATABASE` analyzes all tables in all schemas within the current database.
- `ANALYZE_SCHEMA` analyzes all tables in a given schema (within the current database).
- `ANALYZE_PART_OBJECT` analyzes a single table.

The syntax for the `ANALYZE` commands are:

```
ANALYZE_DATABASE(<method> VARCHAR2 [, <estimate_rows> NUMBER
                [, <estimate_percent> NUMBER [, <method_opt> VARCHAR2 ]]])
ANALYZE_SCHEMA(<schema> VARCHAR2, <method> VARCHAR2
               [, <estimate_rows> NUMBER [, <estimate_percent> NUMBER
               [, <method_opt> VARCHAR2 ]]])
ANALYZE_PART_OBJECT(<schema> VARCHAR2, <object_name> VARCHAR2
                   [, <object_type> CHAR [, <command_type> CHAR
                   [, <command_opt> VARCHAR2 [, <sample_clause> ]]])
```

Parameters - `ANALYZE_DATABASE` and `ANALYZE_SCHEMA`

`<method>`

method determines whether the `ANALYZE` procedure populates the `pg_statistics` table or removes entries from the `pg_statistics` table. If you specify a method of `DELETE`, the `ANALYZE` procedure removes the relevant rows from `pg_statistics`. If you specify a method of `COMPUTE` or `ESTIMATE`, the `ANALYZE` procedure analyzes a table (or multiple tables) and records the distribution information in `pg_statistics`. There is no difference between `COMPUTE` and `ESTIMATE`; both methods execute the Postgres `ANALYZE` statement. All other parameters are validated and then ignored.

`<estimate_rows>`

Number of rows upon which to base estimated statistics. One of `<estimate_rows>` or `<estimate_percent>` must be specified if method is `ESTIMATE`.

This argument is ignored, but is included for compatibility.

`<estimate_percent>`

Percentage of rows upon which to base estimated statistics. One of `<estimate_rows>` or `<estimate_percent>` must be specified if method is `ESTIMATE`.

This argument is ignored, but is included for compatibility.

`<method_opt>`

Object types to be analyzed. Any combination of the following:

```
[ FOR TABLE ]
[ FOR ALL [ INDEXED ] COLUMNS ] [ SIZE n ]
[ FOR ALL INDEXES ]
```

This argument is ignored, but is included for compatibility.

Parameters - ANALYZE_PART_OBJECT

<schema>

Name of the schema whose objects are to be analyzed.

<object_name>

Name of the partitioned object to be analyzed.

<object_type>

Type of object to be analyzed. Valid values are: T – table, I – index.

This argument is ignored, but is included for compatibility.

<command_type>

Type of analyze functionality to perform. Valid values are: E - gather estimated statistics based upon on a specified number of rows or a percentage of rows in the <sample_clause> clause; C - compute exact statistics; or V – validate the structure and integrity of the partitions.

This argument is ignored, but is included for compatibility.

<command_opt>

For <command_type> C or E, can be any combination of:

[FOR TABLE]

[FOR ALL COLUMNS]

[FOR ALL LOCAL INDEXES]

For <command_type> V, can be CASCADE if <object_type> is T.

This argument is ignored, but is included for compatibility.

<sample_clause>

If <command_type> is E, contains the following clause to specify the number of rows or percentage or rows on which to base the estimate.

SAMPLE <n> { ROWS | PERCENT }

This argument is ignored, but is included for compatibility.

3.18.4 CANONICALIZE

The CANONICALIZE procedure performs the following operations on an input string:

- If the string is not double-quoted, verifies that it uses the characters of a legal identifier. If not, an exception is thrown. If the string is double-quoted, all characters are allowed.
- If the string is not double-quoted and does not contain periods, uppercases all alphabetic characters and eliminates leading and trailing spaces.
- If the string is double-quoted and does not contain periods, strips off the double quotes.
- If the string contains periods and no portion of the string is double-quoted, uppercases each portion of the string and encloses each portion in double quotes.
- If the string contains periods and portions of the string are double-quoted, returns the double-quoted portions unchanged including the double quotes and returns the non-double-quoted portions uppercased and enclosed in double quotes.

```
CANONICALIZE(<name> VARCHAR2, <canon_name> OUT VARCHAR2,
<canon_len> BINARY_INTEGER)
```

Parameters

<name>

String to be canonicalized.

<canon_name>

The canonicalized string.

<canon_len>

Number of bytes in <name> to canonicalize starting from the first character.

Examples

The following procedure applies the CANONICALIZE procedure on its input parameter and displays the results.

```
CREATE OR REPLACE PROCEDURE canonicalize (
    p_name      VARCHAR2,
    p_length    BINARY_INTEGER DEFAULT 30
)
IS
    v_canon     VARCHAR2(100);
BEGIN
    DBMS_UTILITY.CANONICALIZE(p_name,v_canon,p_length);
    DBMS_OUTPUT.PUT_LINE('Canonicalized name ==>' || v_canon || '<==');
    DBMS_OUTPUT.PUT_LINE('Length: ' || LENGTH(v_canon));
EXCEPTION
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
        DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;

EXEC canonicalize('Identifier')
Canonicalized name ==>IDENTIFIER<==
Length: 10

EXEC canonicalize('"Identifier"')
Canonicalized name ==>Identifier<==
Length: 10
```

```
EXEC canonicalize('_'+142%')
Canonicalized name ==>_+142%<==
Length: 6

EXEC canonicalize('abc.def.ghi')
Canonicalized name ==>"ABC"."DEF"."GHI"<==
Length: 17

EXEC canonicalize('"abc.def.ghi"')
Canonicalized name ==>abc.def.ghi<==
Length: 11

EXEC canonicalize('"abc".def."ghi"')
Canonicalized name ==>"abc"."DEF"."ghi"<==
Length: 17

EXEC canonicalize('"abc.def".ghi')
Canonicalized name ==>"abc.def"."GHI"<==
Length: 15
```

3.18.5 COMMA_TO_TABLE

The `COMMA_TO_TABLE` procedure converts a comma-delimited list of names into a table of names. Each entry in the list becomes a table entry. The names must be formatted as valid identifiers.

```
COMMA_TO_TABLE(<list> VARCHAR2, <tablen> OUT BINARY_INTEGER,
               <tab> OUT { LNAME_ARRAY | UNCL_ARRAY })
```

Parameters

<list>

Comma-delimited list of names.

<tablen>

Number of entries in <tab>.

<tab>

Table containing the individual names in <list>.

LNAME_ARRAY

A DBMS_UTILITY LNAME_ARRAY (as described in the [LNAME_ARRAY](#) section).

<UNCL_ARRAY>

A DBMS_UTILITY UNCL_ARRAY (as described in the [UNCL_ARRAY](#) section).

Examples

The following procedure uses the `COMMA_TO_TABLE` procedure to convert a list of names to a table. The table entries are then displayed.

```
CREATE OR REPLACE PROCEDURE comma_to_table (
    p_list      VARCHAR2
)
IS
    r_lname     DBMS_UTILITY.LNAME_ARRAY;
    v_length    BINARY_INTEGER;
BEGIN
    DBMS_UTILITY.COMMA_TO_TABLE(p_list,v_length,r_lname);
    FOR i IN 1..v_length LOOP
        DBMS_OUTPUT.PUT_LINE(r_lname(i));
    END LOOP;
END;

EXEC comma_to_table('edb.dept, edb.emp, edb.jobhist')

edb.dept
edb.emp
edb.jobhist
```

3.18.6 DB_VERSION

The DB_VERSION procedure returns the version number of the database.

```
DB_VERSION(<version> OUT VARCHAR2, <compatibility> OUT VARCHAR2)
```

Parameters

<version>

Database version number.

<compatibility>

Compatibility setting of the database. (To be implementation-defined as to its meaning.)

Examples

The following anonymous block displays the database version information.

```
DECLARE
    v_version      VARCHAR2(150);
    v_compat       VARCHAR2(150);
BEGIN
    DBMS_UTILITY.DB_VERSION(v_version,v_compat);
    DBMS_OUTPUT.PUT_LINE('Version: ' || v_version);
    DBMS_OUTPUT.PUT_LINE('Compatibility: ' || v_compat);
END;
```

Version: EnterpriseDB 10.0.0 on i686-pc-linux-gnu, compiled by GCC gcc (GCC) 4.1.2_↵
↵20080704 (Red Hat 4.1.2-48), 32-bit
Compatibility: EnterpriseDB 10.0.0 on i686-pc-linux-gnu, compiled by GCC gcc (GCC) 4.
↵1.22080704 (Red Hat 4.1.2-48), 32-bit

3.18.7 EXEC_DDL_STATEMENT

The EXEC_DDL_STATEMENT provides the capability to execute a DDL command.

```
EXEC_DDL_STATEMENT(<parse_string> VARCHAR2)
```

Parameters

<parse_string>

The DDL command to be executed.

Examples

The following anonymous block creates the job table.

```
BEGIN
  DBMS_UTILITY.EXEC_DDL_STATEMENT(
    'CREATE TABLE job (' ||
      'jobno NUMBER(3), ' ||
      'jname VARCHAR2(9)) '
  );
END;
```

If the <parse_string> does not include a valid DDL statement, Advanced Server returns the following error:

```
edb=# exec dbms_utility.exec_ddl_statement('select rownum from dual');
ERROR:  EDB-20001: 'parse_string' must be a valid DDL statement
```

In this case, Advanced Server's behavior differs from Oracle's; Oracle accepts the invalid <parse_string> without complaint.

3.18.8 FORMAT_CALL_STACK

The `FORMAT_CALL_STACK` function returns the formatted contents of the current call stack.

```
DBMS_UTILITY.FORMAT_CALL_STACK  
return VARCHAR2
```

This function can be used in a stored procedure, function or package to return the current call stack in a readable format. This function is useful for debugging purposes.

3.18.9 GET_CPU_TIME

The GET_CPU_TIME function returns the CPU time in hundredths of a second from some arbitrary point in time.

<cpu_time> NUMBER GET_CPU_TIME

Parameters

<cpu_time>

Number of hundredths of a second of CPU time.

Examples

The following SELECT command retrieves the current CPU time, which is 603 hundredths of a second or .0603 seconds.

```
SELECT DBMS_UTILITY.GET_CPU_TIME FROM DUAL;
```

```
get_cpu_time
```

```
-----
```

```
603
```

3.18.10 GET_DEPENDENCY

The GET_DEPENDENCY procedure provides the capability to list the objects that are dependent upon the specified object. GET_DEPENDENCY does not show dependencies for functions or procedures.

```
GET_DEPENDENCY(<type> VARCHAR2, <schema> VARCHAR2,
               <name> VARCHAR2)
```

Parameters

<type>

The object type of <name>. Valid values are INDEX, PACKAGE, PACKAGE BODY, SEQUENCE, TABLE, TRIGGER, TYPE and VIEW.

<schema>

Name of the schema in which <name> exists.

<name>

Name of the object for which dependencies are to be obtained.

Examples

The following anonymous block finds dependencies on the EMP table.

```
BEGIN
    DBMS_UTILITY.GET_DEPENDENCY('TABLE','public','EMP');
END;

DEPENDENCIES ON public.EMP
-----
*TABLE public.EMP()
*  CONSTRAINT c public.emp()
*  CONSTRAINT f public.emp()
*  CONSTRAINT p public.emp()
*  TYPE public.emp()
*  CONSTRAINT c public.emp()
*  CONSTRAINT f public.jobhist()
*  VIEW .empname_view()
```

3.18.11 GET_HASH_VALUE

The GET_HASH_VALUE function provides the capability to compute a hash value for a given string.

```
<hash> NUMBER GET_HASH_VALUE (<name> VARCHAR2, <base> NUMBER,
                                <hash_size> NUMBER)
```

Parameters

<name>

The string for which a hash value is to be computed.

<base>

Starting value at which hash values are to be generated.

<hash_size>

The number of hash values for the desired hash table.

<hash>

The generated hash value.

Examples

The following anonymous block creates a table of hash values using the `ename` column of the `emp` table and then displays the key along with the hash value. The hash values start at 100 with a maximum of 1024 distinct values.

```
DECLARE
    v_hash          NUMBER;
    TYPE hash_tab IS TABLE OF NUMBER INDEX BY VARCHAR2(10);
    r_hash          HASH_TAB;
    CURSOR emp_cur IS SELECT ename FROM emp;
BEGIN
    FOR r_emp IN emp_cur LOOP
        r_hash(r_emp.ename.. code-block:: text) :=
            DBMS_UTILITY.GET_HASH_VALUE(r_emp.ename,100,1024);
    END LOOP;
    FOR r_emp IN emp_cur LOOP
        DBMS_OUTPUT.PUT_LINE(RPAD(r_emp.ename,10) || ' ' ||
                               r_hash(r_emp.ename));
    END LOOP;
END;
```

SMITH	377
ALLEN	740
WARD	718.. code-block:: text
JONES	131
MARTIN	176
BLAKE	568
CLARK	621
SCOTT	1097
KING	235
TURNER	850
ADAMS	156
JAMES	942
FORD	775
MILLER	148

3.18.12 GET_PARAMETER_VALUE

The GET_PARAMETER_VALUE procedure provides the capability to retrieve database initialization parameter settings.

```
<status> BINARY_INTEGER GET_PARAMETER_VALUE (<parnam> VARCHAR2,
      <intval> OUT INTEGER, <strval> OUT VARCHAR2)
```

Parameters

<parnam>

Name of the parameter whose value is to be returned. The parameters are listed in the pg_settings system view.

<intval>

Value of an integer parameter or the length of <strval>.

<strval>

Value of a string parameter.

<status>

Returns 0 if the parameter value is INTEGER or BOOLEAN. Returns 1 if the parameter value is a string.

Examples

The following anonymous block shows the values of two initialization parameters.

```
DECLARE
    v_intval      INTEGER;
    v_strval      VARCHAR2(80);
BEGIN
    DBMS_UTILITY.GET_PARAMETER_VALUE('max_fsm_pages', v_intval, v_strval);
    DBMS_OUTPUT.PUT_LINE('max_fsm_pages' || ': ' || v_intval);
    DBMS_UTILITY.GET_PARAMETER_VALUE('client_encoding', v_intval, v_strval);
    DBMS_OUTPUT.PUT_LINE('client_encoding' || ': ' || v_strval);
END;
```

```
max_fsm_pages: 72625
client_encoding: SQL_ASCII
```

3.18.13 GET_TIME

The GET_TIME function provides the capability to return the current time in hundredths of a second.

<time> NUMBER GET_TIME

Parameters

<time>

Number of hundredths of a second from the time in which the program is started.

Examples

The following example shows calls to the GET_TIME function.

```
SELECT DBMS_UTILITY.GET_TIME FROM DUAL;

get_time
-----
1555860

SELECT DBMS_UTILITY.GET_TIME FROM DUAL;

get_time
-----
1556037
```

3.18.14 NAME_TOKENIZE

The NAME_TOKENIZE procedure parses a name into its component parts. Names without double quotes are upper-cased. The double quotes are stripped from names with double quotes.

```
NAME_TOKENIZE(<name> VARCHAR2, <a> OUT VARCHAR2,
             <b> OUT VARCHAR2, <c> OUT VARCHAR2, <dblink> OUT VARCHAR2,
             <nextpos> OUT BINARY_INTEGER)
```

Parameters

<name>

String containing a name in the following format:

```
<a> [.<b> [.<c>]][@<dblink> ]
```

<a>

Returns the leftmost component.

Returns the second component, if any.

<c>

Returns the third component, if any.

<dblink>

Returns the database link name.

<nextpos>

Position of the last character parsed in name.

Examples

The following stored procedure is used to display the returned parameter values of the NAME_TOKENIZE procedure for various names.

```
CREATE OR REPLACE PROCEDURE name_tokenize (
    p_name          VARCHAR2
)
IS
    v_a             VARCHAR2(30);
    v_b             VARCHAR2(30);
    v_c             VARCHAR2(30);
    v_dblink        VARCHAR2(30);
    v_nextpos       BINARY_INTEGER;
BEGIN
    DBMS_UTILITY.NAME_TOKENIZE(p_name,v_a,v_b,v_c,v_dblink,v_nextpos);
    DBMS_OUTPUT.PUT_LINE('name      : ' || p_name);
    DBMS_OUTPUT.PUT_LINE('a       : ' || v_a);
    DBMS_OUTPUT.PUT_LINE('b       : ' || v_b);
    DBMS_OUTPUT.PUT_LINE('c       : ' || v_c);
    DBMS_OUTPUT.PUT_LINE('dblink  : ' || v_dblink);
    DBMS_OUTPUT.PUT_LINE('nextpos : ' || v_nextpos);
END;
```

Tokenize the name, emp:

```
BEGIN
name_tokenize('emp');
END;
name
: emp
a
: EMP
b
:
c
:
dblink :
nextpos: 3
```

Tokenize the name, edb.list_emp:

```
BEGIN
name_tokenize('edb.list_emp');
END;
name
:
a
:
b
:
c
:
dblink :
nextpos:
edb.list_emp
EDB
LIST_EMP
12
```

Tokenize the name, "edb"."Emp_Admin".update_emp_sal :

```
BEGIN
name_tokenize('"edb"."Emp_Admin".update_emp_sal');
END;
name
:
a
:
b
:
c
:
dblink :
nextpos:
"edb"."Emp_Admin".update_emp_sal
edb
Emp_Admin
UPDATE_EMP_SAL
32
```

Tokenize the name `edb.emp@edb_dblink`:

```
BEGIN
Copyright © 2007 - 2019 EnterpriseDB Corporation. All rights reserved.
255Database Compatibility for Oracle® Developers
Built-in Package Guide
name_tokenize('edb.emp@edb_dblink');
END;
name
:
a
:
b
:
c
:
dblink :
nextpos:
edb.emp@edb_dblink
EDB
EMP
EDB_DBLINK
18
```


3.18.15 TABLE_TO_COMMA

The TABLE_TO_COMMA procedure converts table of names into a comma-delimited list of names. Each table entry becomes a list entry. The names must be formatted as valid identifiers.

```
TABLE_TO_COMMA(<tab> { LNAME_ARRAY | UNCL_ARRAY },
               <tablen> OUT BINARY_INTEGER, <list> OUT VARCHAR2)
```

Parameters

<tab>

Table containing names.

LNAME_ARRAY

A DBMS_UTILITY LNAME_ARRAY (as described in the [LNAME_ARRAY](#) section).

UNCL_ARRAY

A DBMS_UTILITY UNCL_ARRAY (as described the [UNCL_ARRAY](#) section).

<tablen>

Number of entries in <list>.

<list>

Comma-delimited list of names from <tab>.

Examples

The following example first uses the COMMA_TO_TABLE procedure to convert a comma-delimited list to a table. The TABLE_TO_COMMA procedure then converts the table back to a comma-delimited list that is displayed.

```
CREATE OR REPLACE PROCEDURE table_to_comma (
    p_list      VARCHAR2
)
IS
    r_lname     DBMS_UTILITY.LNAME_ARRAY;
    v_length    BINARY_INTEGER;
    v_listlen   BINARY_INTEGER;
    v_list      VARCHAR2(80);
BEGIN
    DBMS_UTILITY.COMMA_TO_TABLE(p_list,v_length,r_lname);
    DBMS_OUTPUT.PUT_LINE('Table Entries');
    DBMS_OUTPUT.PUT_LINE('-----');
    FOR i IN 1..v_length LOOP
        DBMS_OUTPUT.PUT_LINE(r_lname(i));
    END LOOP;
    DBMS_OUTPUT.PUT_LINE('-----');
    DBMS_UTILITY.TABLE_TO_COMMA(r_lname,v_listlen,v_list);
    DBMS_OUTPUT.PUT_LINE('Comma-Delimited List: ' || v_list);
END;

EXEC table_to_comma('edb.dept, edb.emp, edb.jobhist')

Table Entries
-----
edb.dept
edb.emp
edb.jobhist
```

```
-----  
Comma-Delimited List: edb.dept, edb.emp, edb.jobhist
```

3.19 UTL_ENCODE

The UTL_ENCODE package provides a way to encode and decode data. Advanced Serve supports the following functions and procedures:

Function/Procedure	Return Type	Description
BASE64_DECODE(r)	RAW	Use the BASE64_DECODE function to translate a Base64 encoded string to the original RAW value.
BASE64_ENCODE(r)	RAW	Use the BASE64_ENCODE function to translate a RAW string to an encoded Base64 value.
BASE64_ENCODE(loid)	TEXT	Use the BASE64_ENCODE function to translate a TEXT string to an encoded Base64 value.
MIMEHEADER_DECODE(buf)	VARCHAR2	Use the MIMEHEADER_DECODE function to translate an encoded MIME-HEADER formatted string to its original value.
MIMEHEADER_ENCODE(buf, encode_charset, encoding)	VARCHAR2	Use the MIMEHEADER_ENCODE function to convert and encode a string in MIMEHEADER format.
QUOTED_PRINTABLE_DECODE(r)	RAW	Use the QUOTED_PRINTABLE_DECODE function to translate an encoded string to a RAW value.
QUOTED_PRINTABLE_ENCODE(r)	RAW	Use the QUOTED_PRINTABLE_ENCODE function to translate an input string to a quoted-printable formatted RAW value.
TEXT_DECODE(buf, encode_charset, encoding)	VARCHAR2	Use the TEXT_DECODE function to decode a string encoded by TEXT_ENCODE.
TEXT_ENCODE(buf, encode_charset, encoding)	VARCHAR2	Use the TEXT_ENCODE function to translate a string to a user-specified character set, and then encode the string.
UUDECODE(r)	RAW	Use the UUDECODE function to translate a uuencode encoded string to a RAW value.
UUENCODE(r, type, filename, permission)	RAW	Use the UUENCODE function to translate a RAW string to an encoded uuencode value.

3.19.1 BASE64_DECODE

Use the BASE64_DECODE function to translate a Base64 encoded string to the original value originally encoded by BASE64_ENCODE. The signature is:

```
BASE64_DECODE (<r> IN RAW)
```

This function returns a RAW value.

Parameters

<r>

<r> is the string that contains the Base64 encoded data that will be translated to RAW form.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display BYTEA or RAW values onscreen in readable form. For more information, please refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/12/static/datatype-binary.html>

The following example first encodes (using BASE64_ENCODE), and then decodes (using BASE64_DECODE) a string that contains the text abc:

```
edb=# SELECT UTL_ENCODE.BASE64_ENCODE (CAST ('abc' AS RAW));
base64_encode
-----
YWJj
(1 row)

edb=# SELECT UTL_ENCODE.BASE64_DECODE (CAST ('YWJj' AS RAW));
base64_decode
-----
abc
(1 row)
```

3.19.2 BASE64_ENCODE

Use the `BASE64_ENCODE` function to translate and encode a string in Base64 format (as described in RFC 4648). This function can be useful when composing MIME email that you intend to send using the `UTL_SMTP` package. The `BASE64_ENCODE` function has two signatures:

```
BASE64_ENCODE(<r> IN RAW)
```

and

```
BASE64_ENCODE(<loid> IN OID)
```

This function returns a RAW value or an OID.

Parameters

<r>

<r> specifies the RAW string that will be translated to Base64.

<loid>

<loid> specifies the object ID of a large object that will be translated to Base64.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display BYTEA or RAW values onscreen in readable form. For more information, please refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/12/static/datatype-binary.html>

The following example first encodes (using `BASE64_ENCODE`), and then decodes (using `BASE64_DECODE`) a string that contains the text `abc`:

```
edb=# SELECT UTL_ENCODE.BASE64_ENCODE(CAST ('abc' AS RAW));
base64_encode
-----
YWJj
(1 row)

edb=# SELECT UTL_ENCODE.BASE64_DECODE(CAST ('YWJj' AS RAW));
base64_decode
-----
abc
(1 row)
```

3.19.3 MIMEHEADER_DECODE

Use the MIMEHEADER_DECODE function to decode values that are encoded by the MIMEHEADER_ENCODE function. The signature is:

```
MIMEHEADER_DECODE (<buf> IN VARCHAR2)
```

This function returns a VARCHAR2 value.

Parameters

<buf>

<buf> contains the value (encoded by MIMEHEADER_ENCODE) that will be decoded.

Examples

The following examples use the MIMEHEADER_ENCODE and MIMEHEADER_DECODE functions to first encode, and then decode a string:

```
edb=# SELECT UTL_ENCODE.MIMEHEADER_ENCODE('What is the date?') FROM DUAL;
      mimeheader_encode
-----
=?UTF8?Q?What is the date??=
(1 row)

edb=# SELECT UTL_ENCODE.MIMEHEADER_DECODE('=?UTF8?Q?What is the date??=') FROM DUAL;
      mimeheader_decode
-----
What is the date?
(1 row)
```

3.19.4 MIMEHEADER_ENCODE

Use the MIMEHEADER_ENCODE function to convert a string into mime header format, and then encode the string. The signature is:

```
MIMEHEADER_ENCODE(<buf> IN VARCHAR2, <encode_charset> IN VARCHAR2
DEFAULT NULL, <encoding> IN INTEGER DEFAULT NULL)
```

This function returns a VARCHAR2 value.

Parameters

<buf>

<buf> contains the string that will be formatted and encoded. The string is a VARCHAR2 value.

<encode_charset>

<encode_charset> specifies the character set to which the string will be converted before being formatted and encoded. The default value is NULL.

<encoding>

<encoding> specifies the encoding type used when encoding the string. You can specify:

- Q to enable quoted-printable encoding. If you do not specify a value, MIMEHEADER_ENCODE will use quoted-printable encoding.
- B to enable base-64 encoding.

Examples

The following examples use the MIMEHEADER_ENCODE and MIMEHEADER_DECODE functions to first encode, and then decode a string:

```
edb=# SELECT UTL_ENCODE.MIMEHEADER_ENCODE('What is the date?') FROM DUAL;
      mimeheader_encode
-----
=?UTF8?Q?What is the date??=
(1 row)

edb=# SELECT UTL_ENCODE.MIMEHEADER_DECODE('=?UTF8?Q?What is the date??=') FROM DUAL;
      mimeheader_decode
-----
What is the date?
(1 row)
```

3.19.5 QUOTED_PRINTABLE_DECODE

Use the QUOTED_PRINTABLE_DECODE function to translate an encoded quoted-printable string into a decoded RAW string.

The signature is:

```
QUOTED_PRINTABLE_DECODE(<r> IN RAW)
```

This function returns a RAW value.

Parameters

<r>

<r> contains the encoded string that will be decoded. The string is a RAW value, encoded by QUOTED_PRINTABLE_ENCODE.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display BYTEA or RAW values onscreen in readable form. For more information, please refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/12/static/datatype-binary.html>

The following example first encodes and then decodes a string:

```
edb=# SELECT UTL_ENCODE.QUOTED_PRINTABLE_ENCODE('E=mc2') FROM DUAL;  quoted_printable_
↪encode
-----
E=3Dmc2
(1 row)

edb=# SELECT UTL_ENCODE.QUOTED_PRINTABLE_DECODE('E=3Dmc2') FROM DUAL;
quoted_printable_decode
-----
E=mc2
(1 row)
```


3.19.6 QUOTED_PRINTABLE_ENCODE

Use the `QUOTED_PRINTABLE_ENCODE` function to translate and encode a string in quoted-printable format. The signature is:

```
QUOTED_PRINTABLE_ENCODE (<r> IN RAW)
```

This function returns a RAW value.

Parameters

<r>

<r> contains the string (a RAW value) that will be encoded in a quoted-printable format.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display BYTEA or RAW values onscreen in readable form. For more information, please refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/12/static/datatype-binary.html>

The following example first encodes and then decodes a string:

```
edb=# SELECT UTL_ENCODE.QUOTED_PRINTABLE_ENCODE('E=mc2') FROM DUAL; quoted_printable_
↪encode
-----
E=3Dmc2
(1 row)

edb=# SELECT UTL_ENCODE.QUOTED_PRINTABLE_DECODE('E=3Dmc2') FROM DUAL;
quoted_printable_decode
-----
E=mc2
(1 row)
```

3.19.7 TEXT_DECODE

Use the TEXT_DECODE function to translate and decode an encoded string to the VARCHAR2 value that was originally encoded by the TEXT_ENCODE function. The signature is:

```
TEXT_DECODE(<buf> IN VARCHAR2, <encode_charset> IN VARCHAR2 DEFAULT
NULL, <encoding> IN PLS_INTEGER DEFAULT NULL)
```

This function returns a VARCHAR2 value.

Parameters

<buf>

<buf> contains the encoded string that will be translated to the original value encoded by TEXT_ENCODE.

<encode_charset>

<encode_charset> specifies the character set to which the string will be translated before encoding. The default value is NULL.

<encoding>

<encoding> specifies the encoding type used by TEXT_DECODE. Specify:

- UTL_ENCODE.BASE64 to specify base-64 encoding.
- UTL_ENCODE.QUOTED_PRINTABLE to specify quoted printable encoding. This is the default.

Examples

The following example uses the TEXT_ENCODE and TEXT_DECODE functions to first encode, and then decode a string:

```
edb=# SELECT UTL_ENCODE.TEXT_ENCODE('What is the date?', 'BIG5',
UTL_ENCODE.BASE64) FROM DUAL;
text_encode
-----
V2hhdCBpcyB0aGUgZGF0ZT8=
(1 row)
edb=# SELECT UTL_ENCODE.TEXT_DECODE('V2hhdCBpcyB0aGUgZGF0ZT8=', 'BIG5',
UTL_ENCODE.BASE64) FROM DUAL;
text_decode
-----
What is the date?
(1 row)
```

3.19.8 TEXT_ENCODE

Use the TEXT_ENCODE function to translate a string to a user-specified character set, and then encode the string. The signature is:

```
TEXT_ENCODE(<buf> IN VARCHAR2, <encode_charset> IN VARCHAR2 DEFAULT
NULL, <encoding> IN PLS_INTEGER DEFAULT NULL)
```

This function returns a VARCHAR2 value.

Parameters

<buf>

<buf> contains the encoded string that will be translated to the specified character set and encoded by TEXT_ENCODE.

<encode_charset>

<encode_charset> specifies the character set to which the value will be translated before encoding. The default value is NULL.

<encoding>

<encoding> specifies the encoding type used by TEXT_ENCODE. Specify:

- UTL_ENCODE.BASE64 to specify base-64 encoding.
- UTL_ENCODE.QUOTED_PRINTABLE to specify quoted printable encoding. This is the default.

Examples

The following example uses the TEXT_ENCODE and TEXT_DECODE functions to first encode, and then decode a string:

```
edb=# SELECT UTL_ENCODE.TEXT_ENCODE('What is the date?', 'BIG5',
UTL_ENCODE.BASE64) FROM DUAL;
text_encode
-----
V2hhdCBpcyB0aGUgZGF0ZT8=
(1 row)
edb=# SELECT UTL_ENCODE.TEXT_DECODE('V2hhdCBpcyB0aGUgZGF0ZT8=', 'BIG5',
UTL_ENCODE.BASE64) FROM DUAL;
text_decode
-----
What is the date?
(1 row)
```

3.19.9 UUDECODE

Use the UUDECODE function to translate and decode a uuencode encoded string to the RAW value that was originally encoded by the UUENCODE function. The signature is:

```
UUDECODE (<r> IN RAW)
```

This function returns a RAW value.

If you are using the Advanced Server UUDECODE function to decode uuencoded data that was created by the Oracle implementation of the UTL_ENCODE.UUENCODE function, then you must first set the Advanced Server configuration parameter `utl_encode.uudecode_redwood` to TRUE before invoking the Advanced Server UUDECODE function on the Oracle-created data. (For example, this situation may occur if you migrated Oracle tables containing uuencoded data to an Advanced Server database.)

The uuencoded data created by the Oracle version of the UUENCODE function results in a format that differs from the uuencoded data created by the Advanced Server UUENCODE function. As a result, attempting to use the Advanced Server UUDECODE function on the Oracle uuencoded data results in an error unless the configuration parameter `utl_encode.uudecode_redwood` is set to TRUE.

However, if you are using the Advanced Server UUDECODE function on uuencoded data created by the Advanced Server UUENCODE function, then `utl_encode.uudecode_redwood` must be set to FALSE, which is the default setting.

Parameters

<r>

<r> contains the uuencoded string that will be translated to RAW.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display BYTEA or RAW values onscreen in readable form. For more information, please refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/12/static/datatype-binary.html>

The following example uses UUENCODE and UUDECODE to first encode and then decode a string:

```
edb=# SET bytea_output = escape;
SET
edb=# SELECT UTL_ENCODE.UUENCODE('What is the date?') FROM DUAL;
uuencode
-----
begin 0 uuencode.txt\01215VAA="!I<R!T:&4@9&%T93\\`012`012end\012
(1 row)
edb=# SELECT UTL_ENCODE.UUDECODE
edb-# ('begin 0 uuencode.txt\01215VAA="!I<R!T:&4@9&%T93\\`012`012end\012')
edb-# FROM DUAL;
uudecode
-----
What is the date?
(1 row)
```

3.19.10 UUENCODE

Use the UUENCODE function to translate RAW data into a uuencode formatted encoded string. The signature is:

```
UUENCODE(<r> IN RAW, <type> IN INTEGER DEFAULT 1, <filename> IN
VARCHAR2 DEFAULT NULL, <permission> IN VARCHAR2 DEFAULT NULL)
```

This function returns a RAW value.

Parameters

<r>

<r> contains the RAW string that will be translated to uuencode format.

<type>

<type> is an INTEGER value or constant that specifies the type of uuencoded string that will be returned; the default value is 1. The possible values are:

Value	Constant
1	complete
2	header_piece
3	middle_piece
4	end_piece

<filename>

<filename> is a VARCHAR2 value that specifies the file name that you want to embed in the encoded form; if you do not specify a file name, UUENCODE will include a filename of uuencode.txt in the encoded form.

<permission>

<permission> is a VARCHAR2 that specifies the permission mode; the default value is NULL.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display BYTEA or RAW values onscreen in readable form. For more information, please refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/12/static/datatype-binary.html>

The following example uses UUENCODE and UUDECODE to first encode and then decode a string:

```
edb=# SET bytea_output = escape;
SET
edb=# SELECT UTL_ENCODE.UUENCODE('What is the date?') FROM DUAL;
                                uencode
-----
begin 0 uuencode.txt\01215VAA="!I<R!T:&4@9&%T93\\\012`\012end\012
(1 row)

edb=# SELECT UTL_ENCODE.UUDECODE
edb-# ('begin 0 uuencode.txt\01215VAA="!I<R!T:&4@9&%T93\\\012`\012end\012')
edb-# FROM DUAL;
                                uudecode
-----
What is the date?
(1 row)
```

3.20 UTL_FILE

The UTL_FILE package provides the capability to read from, and write to files on the operating system's file system. Non-superusers must be granted EXECUTE privilege on the UTL_FILE package by a superuser before using any of the functions or procedures in the package. For example the following command grants the privilege to user mary:

```
GRANT EXECUTE ON PACKAGE SYS.UTL_FILE TO mary;
```

Also, the operating system username, enterprisedb, must have the appropriate read and/or write permissions on the directories and files to be accessed using the UTL_FILE functions and procedures. If the required file permissions are not in place, an exception is thrown in the UTL_FILE function or procedure.

A handle to the file to be written to, or read from is used to reference the file. The *file handle* is defined by a public variable in the UTL_FILE package named, UTL_FILE.FILE_TYPE. A variable of type FILE_TYPE must be declared to receive the file handle returned by calling the FOPEN function. The file handle is then used for all subsequent operations on the file.

References to directories on the file system are done using the directory name or alias that is assigned to the directory using the CREATE DIRECTORY command. The procedures and functions available in the UTL_FILE package are listed in the following table:

Function/Procedure	Return Type	Description
FCLOSE(file IN OUT)	n/a	Closes the specified file identified by file.
FCLOSE_ALL	n/a	Closes all open files.
FCOPY(location, filename, dest_dir, dest_file [, start_line [, end_line]])	n/a	Copies filename in the directory identified by location to file, dest_file, in directory, dest_dir, starting from line, start_line, to line, end_line.
FFLUSH(file)	n/a	Forces data in the buffer to be written to disk in the file identified by file.
FOPEN(location, filename, open_mode [, max_linesize])	FILE_TYPE	Opens file, filename, in the directory identified by location.
FREMOVE(location, filename)	n/a	Removes the specified file from the file system.
FRENAME(location, filename, dest_dir, dest_file [, overwrite])	n/a	Renames the specified file.
GET_LINE(file, buffer OUT)	n/a	Reads a line of text into variable, buffer, from the file identified by file.
IS_OPEN(file)	BOOLEAN	Determines whether or not the given file is open.
NEW_LINE(file [, lines])	n/a	Writes an end-of-line character sequence into the file.
PUT(file, buffer)	n/a	Writes buffer to the given file. PUT does not write an end-of-line character sequence.
PUT_LINE(file, buffer)	n/a	Writes buffer to the given file. An end-of-line character sequence is added by the PUT_LINE procedure.
PUTF(file, format [, arg1] [, ...])	n/a	Writes a formatted string to the given file. Up to five substitution parameters, arg1,...arg5 may be specified for replacement in format.

Advanced Server's implementation of UTL_FILE is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

UTL_FILE Exception Codes

If a call to a UTL_FILE procedure or function raises an exception, you can use the condition name to catch the exception. The UTL_FILE package reports the following exception codes compatible with Oracle databases:

Exception Code	Condition name
-29283	invalid_operation
-29285	write_error
-29284	read_error
-29282	invalid_filehandle
-29287	invalid_maxlinesize
-29281	invalid_mode
-29280	invalid_path

3.20.1 Setting File Permissions with utl_file.umask

When a UTL_FILE function or procedure creates a file, there are default file permissions as shown by the following.

```
-rw----- 1 enterprisedb enterprisedb 21 Jul 24 16:08 utlfile
```

Note that all permissions are denied on users belonging to the `enterprisedb` group as well as all other users. Only the `enterprisedb` user has read and write permissions on the created file.

If you wish to have a different set of file permissions on files created by the UTL_FILE functions and procedures, you can accomplish this by setting the `utl_file.umask` configuration parameter.

The `utl_file.umask` parameter sets the *file mode creation mask* or simply, the *mask*, in a manner similar to the Linux `umask` command. This is for usage only within the Advanced Server UTL_FILE package.

Note: The `utl_file.umask` parameter is not supported on Windows systems.

The value specified for `utl_file.umask` is a 3 or 4-character octal string that would be valid for the Linux `umask` command. The setting determines the permissions on files created by the UTL_FILE functions and procedures. (Refer to any information source regarding Linux or Unix systems for information on file permissions and the usage of the `umask` command.)

The following is an example of setting the file permissions with `utl_file.umask`.

First, set up the directory in the file system to be used by the UTL_FILE package. Be sure the operating system account, `enterprisedb` or `postgres`, whichever is applicable, can read and write in the directory.

```
mkdir /tmp/utldir
chmod 777 /tmp/utldir
```

The `CREATE DIRECTORY` command is issued in `psql` to create the directory database object using the file system directory created in the preceding step.

```
CREATE DIRECTORY utldir AS '/tmp/utldir';
```

Set the `utl_file.umask` configuration parameter. The following setting allows the file owner any permission. Group users and other users are permitted any permission except for the execute permission.

```
SET utl_file.umask TO '0011';
```


In the same session during which the `utl_file.umask` parameter is set to the desired value, run the `UTL_FILE` functions and procedures.

```
DECLARE
    v_utlfile      UTL_FILE.FILE_TYPE;
    v_directory    VARCHAR2(50) := 'utldir';
    v_filename     VARCHAR2(20) := 'utlfile';
BEGIN
    v_utlfile := UTL_FILE.FOPEN(v_directory, v_filename, 'w');
    UTL_FILE.PUT_LINE(v_utlfile, 'Simple one-line file');
    DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
    UTL_FILE.FCLOSE(v_utlfile);
END;
```

The permission settings on the resulting file show that group users and other users have read and write permissions on the file as well as the file owner.

```
$ pwd
/tmp/utldir
$ ls -l
total 4
-rw-rw-rw- 1 enterprisedb enterprisedb 21 Jul 24 16:04 utlfile
```

This parameter can also be set on a per role basis with the `ALTER ROLE` command, on a per database basis with the `ALTER DATABASE` command, or for the entire database server instance by setting it in the `postgresql.conf` file.

3.20.2 FCLOSE

The `FCLOSE` procedure closes an open file.

```
FCLOSE(<file> IN OUT FILE_TYPE)
```

Parameters

<file>

Variable of type `FILE_TYPE` containing a file handle of the file to be closed.

3.20.3 FCLOSE_ALL

The `FCLOSE_ALL` procedure closes all open files. The procedure executes successfully even if there are no open files to close.

```
FCLOSE_ALL
```

3.20.4 FCOPY

The `FCOPY` procedure copies text from one file to another.

```
FCOPY(<location> VARCHAR2, <filename> VARCHAR2,
      <dest_dir> VARCHAR2, <dest_file> VARCHAR2
      [, <start_line> PLS_INTEGER [, <end_line> PLS_INTEGER ]])
```

Parameters

<location>

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory containing the file to be copied.

<filename>

Name of the source file to be copied.

<dest_dir>

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory to which the file is to be copied.

<dest_file>

Name of the destination file.

<start_line>

Line number in the source file from which copying will begin. The default is 1.

<end_line>

Line number of the last line in the source file to be copied. If omitted or null, copying will go to the last line of the file.

Examples

The following makes a copy of a file, `C:\TEMP\EMPDIR\empfile.csv`, containing a comma-delimited list of employees from the `emp` table. The copy, `empcopy.csv`, is then listed.

```
CREATE DIRECTORY empdir AS 'C:/TEMP/EMPDIR';

DECLARE
    v_empfile          UTL_FILE.FILE_TYPE;
    v_src_dir           VARCHAR2(50) := 'empdir';
    v_src_file          VARCHAR2(20) := 'empfile.csv';
    v_dest_dir          VARCHAR2(50) := 'empdir';
    v_dest_file         VARCHAR2(20) := 'empcopy.csv';
    v_emprec            VARCHAR2(120);
    v_count             INTEGER := 0;
BEGIN
    UTL_FILE.FCOPY(v_src_dir,v_src_file,v_dest_dir,v_dest_file);
    v_empfile := UTL_FILE.FOPEN(v_dest_dir,v_dest_file,'r');
    DBMS_OUTPUT.PUT_LINE('The following is the destination file, '' ||
        v_dest_file || ''');
    LOOP
        UTL_FILE.GET_LINE(v_empfile,v_emprec);
        DBMS_OUTPUT.PUT_LINE(v_emprec);
        v_count := v_count + 1;
    END LOOP;
    EXCEPTION
        WHEN NO_DATA_FOUND THEN
            UTL_FILE.FCLOSE(v_empfile);
            DBMS_OUTPUT.PUT_LINE(v_count || ' records retrieved');
        WHEN OTHERS THEN
            DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
            DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;
```

```

The following is the destination file, 'empcopy.csv'
7369,SMITH,CLERK,7902,17-DEC-80,800,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81,1600,300,30
7521,WARD,SALESMAN,7698,22-FEB-81,1250,500,30
7566,JONES,MANAGER,7839,02-APR-81,2975,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81,1250,1400,30
7698,BLAKE,MANAGER,7839,01-MAY-81,2850,,30
7782,CLARK,MANAGER,7839,09-JUN-81,2450,,10
7788,SCOTT,ANALYST,7566,19-APR-87,3000,,20
7839,KING,PRESIDENT,,17-NOV-81,5000,,10
7844,TURNER,SALESMAN,7698,08-SEP-81,1500,0,30
7876,ADAMS,CLERK,7788,23-MAY-87,1100,,20
7900,JAMES,CLERK,7698,03-DEC-81,950,,30
7902,FORD,ANALYST,7566,03-DEC-81,3000,,20
7934,MILLER,CLERK,7782,23-JAN-82,1300,,10
14 records retrieved

```

3.20.5 FFLUSH

The FFLUSH procedure flushes unwritten data from the write buffer to the file.

```
FFLUSH(<file> FILE_TYPE)
```

Parameters

<file>

Variable of type FILE_TYPE containing a file handle.

Examples

Each line is flushed after the NEW_LINE procedure is called.

```

DECLARE
    v_empfile          UTL_FILE.FILE_TYPE;
    v_directory         VARCHAR2(50) := 'empdir';
    v_filename          VARCHAR2(20) := 'empfile.csv';
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');
    FOR i IN emp_cur LOOP
        UTL_FILE.PUT(v_empfile,i.empno);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.ename);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.job);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.mgr);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.hiredate);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.sal);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.comm);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.deptno);
        UTL_FILE.NEW_LINE(v_empfile);
        UTL_FILE.FFLUSH(v_empfile);
    END LOOP;
END;

```

```

END LOOP;
DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
UTL_FILE.FCLOSE(v_empfile);
END;

```

3.20.6 FOPEN

The FOPEN function opens a file for I/O.

```

<filetype> FILE_TYPE FOPEN(<location> VARCHAR2,
<filename> VARCHAR2, <open_mode> VARCHAR2
[, <max_linesize> BINARY_INTEGER ])

```

Parameters

<location>

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory containing the file to be opened.

<filename>

Name of the file to be opened.

<open_mode>

Mode in which the file will be opened. Modes are: `a` - append to file; `r` - read from file; `w` - write to file.

<max_linesize>

Maximum size of a line in characters. In read mode, an exception is thrown if an attempt is made to read a line exceeding `<max_linesize>`. In write and append modes, an exception is thrown if an attempt is made to write a line exceeding `<max_linesize>`. The end-of-line character(s) are not included in determining if the maximum line size is exceeded. This behavior is not compatible with Oracle databases; Oracle does count the end-of-line character(s).

<filetype>

Variable of type `FILE_TYPE` containing the file handle of the opened file.

3.20.7 FREMOVE

The FREMOVE procedure removes a file from the system.

```

FREMOVE(<location> VARCHAR2, <filename> VARCHAR2)

```

An exception is thrown if the file to be removed does not exist.

Parameters

<location>

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory containing the file to be removed.

<filename>

Name of the file to be removed.

Examples

The following removes file `empfile.csv`.

```

DECLARE
    v_directory    VARCHAR2(50) := 'empdir';
    v_filename     VARCHAR2(20) := 'empfile.csv';
BEGIN
    UTL_FILE.FREMOVE(v_directory,v_filename);
    DBMS_OUTPUT.PUT_LINE('Removed file: ' || v_filename);
    EXCEPTION
        WHEN OTHERS THEN
            DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
            DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;

Removed file: empfile.csv

```

3.20.8 FRENAME

The `FRENAME` procedure renames a given file. This effectively moves a file from one location to another.

```

FRENAME(<location> VARCHAR2, <filename> VARCHAR2,

<dest_dir> VARCHAR2, <dest_file> VARCHAR2,

[ <overwrite> BOOLEAN ])

```

Parameters

<location>

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory containing the file to be renamed.

<filename>

Name of the source file to be renamed.

<dest_dir>

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory to which the renamed file is to exist.

<dest_file>

New name of the original file.

<overwrite>

Replaces any existing file named `<dest_file>` in `<dest_dir>` if set to `TRUE`, otherwise an exception is thrown if set to `FALSE`. This is the default.

Examples

The following renames a file, C:\TEMP\EMPDIR\empfile.csv, containing a comma-delimited list of employees from the emp table. The renamed file, C:\TEMP\NEWDIR\newemp.csv, is then listed.

```
CREATE DIRECTORY "newdir" AS 'C:/TEMP/NEWDIR';

DECLARE
    v_empfile      UTL_FILE.FILE_TYPE;
    v_src_dir      VARCHAR2(50) := 'empdir';
    v_src_file     VARCHAR2(20) := 'empfile.csv';
    v_dest_dir     VARCHAR2(50) := 'newdir';
    v_dest_file    VARCHAR2(50) := 'newemp.csv';
    v_replace      BOOLEAN := FALSE;
    v_emprec       VARCHAR2(120);
    v_count        INTEGER := 0;
BEGIN
    UTL_FILE.FRENAME(v_src_dir,v_src_file,v_dest_dir,
        v_dest_file,v_replace);
    v_empfile := UTL_FILE.FOPEN(v_dest_dir,v_dest_file,'r');
    DBMS_OUTPUT.PUT_LINE('The following is the renamed file, '' ||
        v_dest_file || ''');
    LOOP
        UTL_FILE.GET_LINE(v_empfile,v_emprec);
        DBMS_OUTPUT.PUT_LINE(v_emprec);
        v_count := v_count + 1;
    END LOOP;
    EXCEPTION
        WHEN NO_DATA_FOUND THEN
            UTL_FILE.FCLOSE(v_empfile);
            DBMS_OUTPUT.PUT_LINE(v_count || ' records retrieved');
        WHEN OTHERS THEN
            DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
            DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;
```

```
The following is the renamed file, 'newemp.csv'
7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
14 records retrieved
```

3.20.9 GET_LINE

The GET_LINE procedure reads a line of text from a given file up to, but not including the end-of-line terminator. A NO_DATA_FOUND exception is thrown when there are no more lines to read.

```
GET_LINE(<file> FILE_TYPE, <buffer> OUT VARCHAR2)
```

Parameters

<file>

Variable of type FILE_TYPE containing the file handle of the opened file.

<buffer>

Variable to receive a line from the file.

Examples

The following anonymous block reads through and displays the records in file empfile.csv.

```
DECLARE
    v_empfile      UTL_FILE.FILE_TYPE;
    v_directory    VARCHAR2(50) := 'empdir';
    v_filename     VARCHAR2(20) := 'empfile.csv';
    v_emprec       VARCHAR2(120);
    v_count        INTEGER := 0;
BEGIN
    v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'r');
    LOOP
        UTL_FILE.GET_LINE(v_empfile,v_emprec);
        DBMS_OUTPUT.PUT_LINE(v_emprec);
        v_count := v_count + 1;
    END LOOP;
    EXCEPTION
        WHEN NO_DATA_FOUND THEN
            UTL_FILE.FCLOSE(v_empfile);
            DBMS_OUTPUT.PUT_LINE('End of file ' || v_filename || ' - ' ||
                v_count || ' records retrieved');
        WHEN OTHERS THEN
            DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
            DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;
```

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
End of file empfile.csv - 14 records retrieved

3.20.10 IS_OPEN

The IS_OPEN function determines whether or not the given file is open.

```
<status> BOOLEAN IS_OPEN(<file> FILE_TYPE)
```

Parameters

<file>

Variable of type FILE_TYPE containing the file handle of the file to be tested.

<status>

TRUE if the given file is open, FALSE otherwise.

3.20.11 NEW_LINE

The NEW_LINE procedure writes an end-of-line character sequence in the file.

```
NEW_LINE(<file> FILE_TYPE [, <lines> INTEGER ])
```

Parameters

<file>

Variable of type FILE_TYPE containing the file handle of the file to which end-of-line character sequences are to be written.

<lines>

Number of end-of-line character sequences to be written. The default is one.

Examples

A file containing a double-spaced list of employee records is written.

```
DECLARE
    v_empfile      UTL_FILE.FILE_TYPE;
    v_directory    VARCHAR2(50) := 'empdir';
    v_filename     VARCHAR2(20) := 'empfile.csv';
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');
    FOR i IN emp_cur LOOP
        UTL_FILE.PUT(v_empfile,i.empno);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.ename);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.job);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.mgr);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.hiredate);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.sal);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.comm);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.deptno);
        UTL_FILE.NEW_LINE(v_empfile,2);
    END LOOP;
```



```

    DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
    UTL_FILE.FCLOSE(v_empfile);
END;

Created file: empfile.csv

```

This file is then displayed:

```

C:\TEMP\EMPDIR>TYPE empfile.csv

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10

```

3.20.12 PUT

The PUT procedure writes a string to the given file. No end-of-line character sequence is written at the end of the string. Use the NEW_LINE procedure to add an end-of-line character sequence.

```

PUT(<file> FILE_TYPE, <buffer> { DATE | NUMBER | TIMESTAMP |
    VARCHAR2 })

```

Parameters

<file>

Variable of type FILE_TYPE containing the file handle of the file to which the given string is to be written.

<buffer>

Text to be written to the specified file.

Examples

The following example uses the PUT procedure to create a comma-delimited file of employees from the emp table.

```
DECLARE
    v_empfile          UTL_FILE.FILE_TYPE;
    v_directory        VARCHAR2(50) := 'empdir';
    v_filename         VARCHAR2(20) := 'empfile.csv';
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');
    FOR i IN emp_cur LOOP
        UTL_FILE.PUT(v_empfile,i.empno);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.ename);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.job);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.mgr);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.hiredate);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.sal);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.comm);
        UTL_FILE.PUT(v_empfile,',');
        UTL_FILE.PUT(v_empfile,i.deptno);
        UTL_FILE.NEW_LINE(v_empfile);
    END LOOP;
    DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
    UTL_FILE.FCLOSE(v_empfile);
END;
```

Created file: empfile.csv

The following is the contents of empfile.csv created above:

```
C:\TEMP\EMPDIR>TYPE empfile.csv

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
```

3.20.13 PUT_LINE

The PUT_LINE procedure writes a single line to the given file including an end-of-line character sequence.

```
PUT_LINE(<file> FILE_TYPE,
        <buffer> {DATE|NUMBER|TIMESTAMP|VARCHAR2})
```

Parameters

<file>

Variable of type FILE_TYPE containing the file handle of the file to which the given line is to be written.

<buffer>

Text to be written to the specified file.

Examples

The following example uses the PUT_LINE procedure to create a comma-delimited file of employees from the emp table.

```
DECLARE
    v_empfile      UTL_FILE.FILE_TYPE;
    v_directory    VARCHAR2(50) := 'empdir';
    v_filename     VARCHAR2(20) := 'empfile.csv';
    v_emptrec      VARCHAR2(120);
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');
    FOR i IN emp_cur LOOP
        v_emptrec := i.empno || ',' || i.ename || ',' || i.job || ',' ||
            NVL(LTRIM(TO_CHAR(i.mgr,'9999')),')') || ',' || i.hiredate ||
            ',' || i.sal || ',' ||
            NVL(LTRIM(TO_CHAR(i.comm,'9990.99')),')') || ',' || i.deptno;
        UTL_FILE.PUT_LINE(v_empfile,v_emptrec);
    END LOOP;
    DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
    UTL_FILE.FCLOSE(v_empfile);
END;
```

The following is the contents of empfile.csv created above:

```
C:\TEMP\EMPDIR>TYPE empfile.csv

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
```

3.20.14 PUTF

The PUTF procedure writes a formatted string to the given file.

```
PUTF(<file> FILE_TYPE, <format> VARCHAR2 [, <arg1> VARCHAR2]
    [, ...])
```

Parameters

<file>

Variable of type FILE_TYPE containing the file handle of the file to which the formatted line is to be written.

<format>

String to format the text written to the file. The special character sequence, %s, is substituted by the value of arg. The special character sequence, \n, indicates a new line. Note, however, in Advanced Server, a new line character must be specified with two consecutive backslashes instead of one - \\n. This characteristic is not compatible with Oracle databases.

<arg1>

Up to five arguments, <arg1>,...<arg5>, to be substituted in the format string for each occurrence of %s. The first arg is substituted for the first occurrence of %s, the second arg is substituted for the second occurrence of %s, etc.

Examples

The following anonymous block produces formatted output containing data from the emp table. Note the use of the E literal syntax and double backslashes for the new line character sequence in the format string which are not compatible with Oracle databases.

```
DECLARE
    v_empfile      UTL_FILE.FILE_TYPE;
    v_directory    VARCHAR2(50) := 'empdir';
    v_filename     VARCHAR2(20) := 'empfile.csv';
    v_format       VARCHAR2(200);
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    v_format := E'%s %s, %s\\nSalary: %s Commission: %s\\n\\n';
    v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');
    FOR i IN emp_cur LOOP
        UTL_FILE.PUTF(v_empfile,v_format,i.empno,i.ename,i.job,i.sal,
            NVL(i.comm,0));
    END LOOP;
    DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
    UTL_FILE.FCLOSE(v_empfile);
EXCEPTION
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
        DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;
```

Created file: empfile.csv

The following is the contents of empfile.csv created above:

```
C:\TEMP\EMPDIR>TYPE empfile.csv
7369 SMITH, CLERK
Salary: $800.00 Commission: $0
```

```
7499 ALLEN, SALESMAN
Salary: $1600.00 Commission: $300.00
7521 WARD, SALESMAN
Salary: $1250.00 Commission: $500.00
7566 JONES, MANAGER
Salary: $2975.00 Commission: $0
7654 MARTIN, SALESMAN
Salary: $1250.00 Commission: $1400.00
7698 BLAKE, MANAGER
Salary: $2850.00 Commission: $0
7782 CLARK, MANAGER
Salary: $2450.00 Commission: $0
7788 SCOTT, ANALYST
Salary: $3000.00 Commission: $0
7839 KING, PRESIDENT
Salary: $5000.00 Commission: $0
7844 TURNER, SALESMAN
Salary: $1500.00 Commission: $0.00
7876 ADAMS, CLERK
Salary: $1100.00 Commission: $0
7900 JAMES, CLERK
Salary: $950.00 Commission: $0
7902 FORD, ANALYST
Salary: $3000.00 Commission: $0
7934 MILLER, CLERK
Salary: $1300.00 Commission: $0
```

3.21 UTL_HTTP

The UTL_HTTP package provides a way to use the HTTP or HTTPS protocol to retrieve information found at an URL. Advanced Server supports the following functions and procedures:

Function/Procedure	Return Type	Description
BEGIN_REQUEST(url, method, http_version)	UTL_HTTP.REQ	Initiates a new HTTP request.
END_REQUEST(r IN OUT)	n/a	Ends an HTTP request before allowing it to complete.
END_RESPONSE(r IN OUT)	n/a	Ends the HTTP response.
GET_BODY_CHARSET	VARCHAR2	Returns the default character set of the body of future HTTP requests.
GET_BODY_CHARSET(charset OUT)	n/a	Returns the default character set of the body of future HTTP requests.
GET_FOLLOW_REDIRECT(max_redirects OUT)	n/a	Current setting for the maximum number of redirections allowed.
GET_HEADER(r IN OUT, n, name OUT, value OUT)	n/a	Returns the nth header of the HTTP response.
GET_HEADER_BY_NAME(r IN OUT, name, value OUT, n)	n/a	Returns the HTTP response header for the specified name.
GET_HEADER_COUNT(r IN OUT)	INTEGER	Returns the number of HTTP response headers.
GET_RESPONSE(r IN OUT)	UTL_HTTP.RESP	Returns the HTTP response.
GET_RESPONSE_ERROR_CHECK(enable OUT)	n/a	Returns whether or not response error check is set.
GET_TRANSFER_TIMEOUT(timeout OUT)	n/a	Returns the transfer timeout setting for HTTP requests.
(r IN OUT, data OUT, remove_crlf)	n/a	Returns the HTTP response body in text form until the end of line.
READ_RAW(r IN OUT, data OUT, len)	n/a	Returns the HTTP response body in binary form for a specified number of bytes.
READ_TEXT(r IN OUT, data OUT, len)	n/a	Returns the HTTP response body in text form for a specified number of characters.
REQUEST(url)	VARCHAR2	Returns the content of a web page.
REQUEST_PIECES(url, max_pieces)	UTL_HTTP.HTML_PIECES	Returns a table of 2000-byte segments retrieved from an URL.
SET_BODY_CHARSET(charset)	n/a	Sets the default character set of the body of future HTTP requests.
SET_FOLLOW_REDIRECT(max_redirects)	n/a	Sets the maximum number of times to follow the redirect instruction.
SET_FOLLOW_REDIRECT(r IN OUT, max_redirects)	n/a	Sets the maximum number of times to follow the redirect instruction for an individual request.
SET_HEADER(r IN OUT, name, value)	n/a	Sets the HTTP request header.
SET_RESPONSE_ERROR_CHECK(enable)	n/a	Determines whether or not HTTP 4xx and 5xx status codes are to be treated as errors.
SET_TRANSFER_TIMEOUT(timeout)	n/a	Sets the default, transfer timeout value for HTTP requests.
SET_TRANSFER_TIMEOUT(r IN OUT, timeout)	n/a	Sets the transfer timeout value for an individual HTTP request.

Continued on next page

Table 3.9 – continued from previous page

WRITE_LINE(r IN OUT, data)	n/a	Writes CRLF terminated data to the HTTP request body in TEXT form.
WRITE_RAW(r IN OUT, data)	n/a	Writes data to the HTTP request body in BINARY form.
WRITE_TEXT(r IN OUT, data)	n/a	Writes data to the HTTP request body in TEXT form.

Advanced Server's implementation of UTL_HTTP is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Note: In Advanced Server, an HTTP 4xx or HTTP 5xx response produces a database error; in Oracle, this is configurable but FALSE by default.

In Advanced Server, the UTL_HTTP text interfaces expect the downloaded data to be in the database encoding. All currently-available interfaces are text interfaces. In Oracle, the encoding is detected from HTTP headers; in the absence of the header, the default is configurable and defaults to ISO-8859-1.

Advanced Server ignores all cookies it receives.

The UTL_HTTP exceptions that can be raised in Oracle are not recognized by Advanced Server. In addition, the error codes returned by Advanced Server are not the same as those returned by Oracle.

There are various public constants available with UTL_HTTP. These are listed in the following tables.

The following table contains UTL_HTTP public constants defining HTTP versions and port assignments.

HTTP VERSIONS	
HTTP_VERSION_1_0	CONSTANT VARCHAR2(64) := 'HTTP/1.0';
HTTP_VERSION_1_1	CONSTANT VARCHAR2(64) := 'HTTP/1.1';
STANDARD PORT ASSIGNMENTS	
DEFAULT_HTTP_PORT	CONSTANT INTEGER := 80;
DEFAULT_HTTPS_PORT	CONSTANT INTEGER := 443;

The following table contains UTL_HTTP public status code constants.

1XX INFORMATIONAL	
HTTP_CONTINUE	CONSTANT INTEGER := 100;
HTTP_SWITCHING_PROTOCOLS	CONSTANT INTEGER := 101;
HTTP_PROCESSING	CONSTANT INTEGER := 102;
2XX SUCCESS	
HTTP_OK	CONSTANT INTEGER := 200;
HTTP_CREATED	CONSTANT INTEGER := 201;
HTTP_ACCEPTED	CONSTANT INTEGER := 202;
HTTP_NON_AUTHORITATIVE_INFO	CONSTANT INTEGER := 203;
HTTP_NO_CONTENT	CONSTANT INTEGER := 204;
HTTP_RESET_CONTENT	CONSTANT INTEGER := 205;
HTTP_PARTIAL_CONTENT	CONSTANT INTEGER := 206;
HTTP_MULTI_STATUS	CONSTANT INTEGER := 207;
HTTP_ALREADY_REPORTED	CONSTANT INTEGER := 208;
HTTP_IM_USED	CONSTANT INTEGER := 226;
3XX REDIRECTION	
HTTP_MULTIPLE_CHOICES	CONSTANT INTEGER := 300;
HTTP_MOVED_PERMANENTLY	CONSTANT INTEGER := 301;
HTTP_FOUND	CONSTANT INTEGER := 302;
HTTP_SEE_OTHER	CONSTANT INTEGER := 303;
HTTP_NOT_MODIFIED	CONSTANT INTEGER := 304;
HTTP_USE_PROXY	CONSTANT INTEGER := 305;
HTTP_SWITCH_PROXY	CONSTANT INTEGER := 306;
HTTP_TEMPORARY_REDIRECT	CONSTANT INTEGER := 307;
HTTP_PERMANENT_REDIRECT	CONSTANT INTEGER := 308;

4XX CLIENT ERROR	
HTTP_BAD_REQUEST	CONSTANT INTEGER := 400;
HTTP_UNAUTHORIZED	CONSTANT INTEGER := 401;
HTTP_PAYMENT_REQUIRED	CONSTANT INTEGER := 402;
HTTP_FORBIDDEN	CONSTANT INTEGER := 403;
HTTP_NOT_FOUND	CONSTANT INTEGER := 404;
HTTP_METHOD_NOT_ALLOWED	CONSTANT INTEGER := 405;
HTTP_NOT_ACCEPTABLE	CONSTANT INTEGER := 406;
HTTP_PROXY_AUTH_REQUIRED	CONSTANT INTEGER := 407;
HTTP_REQUEST_TIME_OUT	CONSTANT INTEGER := 408;
HTTP_CONFLICT	CONSTANT INTEGER := 409;
HTTP_GONE	CONSTANT INTEGER := 410;
HTTP_LENGTH_REQUIRED	CONSTANT INTEGER := 411;
HTTP_PRECONDITION_FAILED	CONSTANT INTEGER := 412;
HTTP_REQUEST_ENTITY_TOO_LARGE	CONSTANT INTEGER := 413;
HTTP_REQUEST_URI_TOO_LARGE	CONSTANT INTEGER := 414;
HTTP_UNSUPPORTED_MEDIA_TYPE	CONSTANT INTEGER := 415;
HTTP_REQ_RANGE_NOT_SATISFIABLE	CONSTANT INTEGER := 416;
HTTP_EXPECTATION_FAILED	CONSTANT INTEGER := 417;
HTTP_I_AM_A_TEAPOT	CONSTANT INTEGER := 418;
HTTP_AUTHENTICATION_TIME_OUT	CONSTANT INTEGER := 419;
HTTP_ENHANCE_YOUR_CALM	CONSTANT INTEGER := 420;
HTTP_UNPROCESSABLE_ENTITY	CONSTANT INTEGER := 422;
HTTP_LOCKED	CONSTANT INTEGER := 423;
HTTP_FAILED_DEPENDENCY	CONSTANT INTEGER := 424;
HTTP_UNORDERED_COLLECTION	CONSTANT INTEGER := 425;
HTTP_UPGRADE_REQUIRED	CONSTANT INTEGER := 426;
HTTP_PRECONDITION_REQUIRED	CONSTANT INTEGER := 428;
HTTP_TOO_MANY_REQUESTS	CONSTANT INTEGER := 429;
HTTP_REQUEST_HEADER_FIELDS_TOO_LARGE	CONSTANT INTEGER := 431;
HTTP_NO_RESPONSE	CONSTANT INTEGER := 444;
HTTP_RETRY_WITH	CONSTANT INTEGER := 449;
HTTP_BLOCKED_BY_WINDOWS_PARENTAL_CONTROLS	CONSTANT INTEGER := 450;
HTTP_REDIRECT	CONSTANT INTEGER := 451;
HTTP_REQUEST_HEADER_TOO_LARGE	CONSTANT INTEGER := 494;
HTTP_CERT_ERROR	CONSTANT INTEGER := 495;
HTTP_NO_CERT	CONSTANT INTEGER := 496;
HTTP_HTTP_TO_HTTPS	CONSTANT INTEGER := 497;
HTTP_CLIENT_CLOSED_REQUEST	CONSTANT INTEGER := 499;

5XX SERVER ERROR	
HTTP_INTERNAL_SERVER_ERROR	CONSTANT INTEGER := 500;
HTTP_NOT_IMPLEMENTED	CONSTANT INTEGER := 501;
HTTP_BAD_GATEWAY	CONSTANT INTEGER := 502;
HTTP_SERVICE_UNAVAILABLE	CONSTANT INTEGER := 503;
HTTP_GATEWAY_TIME_OUT	CONSTANT INTEGER := 504;
HTTP_VERSION_NOT_SUPPORTED	CONSTANT INTEGER := 505;
HTTP_VARIANT_ALSO_NEGOTIATES	CONSTANT INTEGER := 506;
HTTP_INSUFFICIENT_STORAGE	CONSTANT INTEGER := 507;
HTTP_LOOP_DETECTED	CONSTANT INTEGER := 508;
HTTP_BANDWIDTH_LIMIT_EXCEEDED	CONSTANT INTEGER := 509;
HTTP_NOT_EXTENDED	CONSTANT INTEGER := 510;
HTTP_NETWORK_AUTHENTICATION_REQUIRED	CONSTANT INTEGER := 511;
HTTP_NETWORK_READ_TIME_OUT_ERROR	CONSTANT INTEGER := 598;
HTTP_NETWORK_CONNECT_TIME_OUT_ERROR	CONSTANT INTEGER := 599;

3.21.1 HTML_PIECES

The UTL_HTTP package declares a type named HTML_PIECES, which is a table of type VARCHAR2 (2000) indexed by BINARY_INTEGER. A value of this type is returned by the REQUEST_PIECES function.

```
TYPE html_pieces IS TABLE OF VARCHAR2(2000) INDEX BY BINARY_INTEGER;
```

3.21.2 REQ

The REQ record type holds information about each HTTP request.

```
TYPE req IS RECORD (
    url            VARCHAR2(32767),      -- URL to be accessed
    method         VARCHAR2(64),        -- HTTP method
    http_version   VARCHAR2(64),        -- HTTP version
    private_hndl   INTEGER               -- Holds handle for this request
);
```

3.21.3 RESP

The RESP record type holds information about the response from each HTTP request.

```
TYPE resp IS RECORD (
    status_code    INTEGER,              -- HTTP status code
    reason_phrase   VARCHAR2(256),       -- HTTP response reason phrase
    http_version   VARCHAR2(64),        -- HTTP version
    private_hndl   INTEGER               -- Holds handle for this response
);
```

3.21.4 BEGIN_REQUEST

The BEGIN_REQUEST function initiates a new HTTP request. A network connection is established to the web server with the specified URL. The signature is:

```
BEGIN_REQUEST(<url> IN VARCHAR2, <method> IN VARCHAR2 DEFAULT
'GET ', <http_version> IN VARCHAR2 DEFAULT NULL) RETURN
UTL_HTTP.REQ
```

The `BEGIN_REQUEST` function returns a record of type `UTL_HTTP.REQ`.

Parameters

`<url>`

`<url>` is the Uniform Resource Locator from which `UTL_HTTP` will return content.

`<method>`

`<method>` is the HTTP method to be used. The default is `GET`.

`<http_version>`

`<http_version>` is the HTTP protocol version sending the request. The specified values should be either `HTTP/1.0` or `HTTP/1.1`. The default is null in which case the latest HTTP protocol version supported by the `UTL_HTTP` package is used which is `1.1`.

3.21.5 END_REQUEST

The `END_REQUEST` procedure terminates an HTTP request. Use the `END_REQUEST` procedure to terminate an HTTP request without completing it and waiting for the response. The normal process is to begin the request, get the response, then close the response. The signature is:

```
END_REQUEST(<r> IN OUT UTL_HTTP.REQ)
```

Parameters

`<r>`

`<r>` is the HTTP request record.

3.21.6 END_RESPONSE

The `END_RESPONSE` procedure terminates the HTTP response. The `END_RESPONSE` procedure completes the HTTP request and response. This is the normal method to end the request and response process. The signature is:

```
END_RESPONSE(<r> IN OUT UTL_HTTP.RESP)
```

Parameters

`<r>`

`<r>` is the HTTP response record.

3.21.7 GET_BODY_CHARSET

The `GET_BODY_CHARSET` program is available in the form of both a procedure and a function. A call to `GET_BODY_CHARSET` returns the default character set of the body of future HTTP requests.

The procedure signature is:

```
GET_BODY_CHARSET(<charset> OUT VARCHAR2)
```

The function signature is:

```
GET_BODY_CHARSET() RETURN VARCHAR2
```

This function returns a VARCHAR2 value.

Parameters

<charset>

<charset> is the character set of the body.

Examples

The following is an example of the GET_BODY_CHARSET function.

```
edb=# SELECT UTL_HTTP.GET_BODY_CHARSET() FROM DUAL;
get_body_charset
-----
ISO-8859-1
(1 row)
```

3.21.8 GET_FOLLOW_REDIRECT

The GET_FOLLOW_REDIRECT procedure returns the current setting for the maximum number of redirections allowed. The signature is:

```
GET_FOLLOW_REDIRECT(<max_redirects> OUT INTEGER)
```

Parameters

<max_redirects>

<max_redirects> is maximum number of redirections allowed.

3.21.9 GET_HEADER

The GET_HEADER procedure returns the *n*th header of the HTTP response. The signature is:

```
GET_HEADER(<r> IN OUT UTL_HTTP.RESP, <n> INTEGER, <name> OUT
VARCHAR2, <value> OUT VARCHAR2)
```

Parameters

<r>

<r> is the HTTP response record.

<n>

<n> is the *n*th header of the HTTP response record to retrieve.

<name>

<name> is the name of the response header.

<value>

<value> is the value of the response header.

Examples

The following example retrieves the header count, then the headers.

```
DECLARE
    v_req          UTL_HTTP.REQ;
    v_resp         UTL_HTTP.RESP;
    v_name         VARCHAR2(30);
    v_value        VARCHAR2(200);
    v_header_cnt   INTEGER;
BEGIN
    -- Initiate request and get response
    v_req := UTL_HTTP.BEGIN_REQUEST('www.enterprisedb.com');
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);

    -- Get header count
    v_header_cnt := UTL_HTTP.GET_HEADER_COUNT(v_resp);
    DBMS_OUTPUT.PUT_LINE('Header Count: ' || v_header_cnt);

    -- Get all headers
    FOR i IN 1 .. v_header_cnt LOOP
        UTL_HTTP.GET_HEADER(v_resp, i, v_name, v_value);
        DBMS_OUTPUT.PUT_LINE(v_name || ': ' || v_value);
    END LOOP;

    -- Terminate request
    UTL_HTTP.END_RESPONSE(v_resp);
END;
```

The following is the output from the example.

```
Header Count: 23
Age: 570
Cache-Control: must-revalidate
Content-Type: text/html; charset=utf-8
Date: Wed, 30 Apr 2015 14:57:52 GMT
ETag: "aab02f2bd2d696eed817ca89ef411dda"
Expires: Sun, 19 Nov 1978 05:00:00 GMT
Last-Modified: Wed, 30 Apr 2015 14:15:49 GMT
RTSS: 1-1307-3
Server: Apache/2.2.3 (Red Hat)
Set-Cookie: SESS2771d0952de2a1a84d322a262e0c173c=jnlulj1etmdi5gg4lh8hakvs01;␣
↪expires=Fri, 23-May-2015 18:21:43 GMT; path=/; domain=.enterprisedb.com
Vary: Accept-Encoding
Via: 1.1 varnish
X-EDB-Backend: ec
X-EDB-Cache: HIT
X-EDB-Cache-Address: 10.31.162.212
X-EDB-Cache-Server: ip-10-31-162-212
X-EDB-Cache-TTL: 600.000
X-EDB-Cacheable: MAYBE: The user has a cookie of some sort. Maybe it's double choc-
↪chip!
X-EDB-Do-GZIP: false
X-Powered-By: PHP/5.2.17
X-Varnish: 484508634 484506789
transfer-encoding: chunked
Connection: keep-alive
```

3.21.10 GET_HEADER_BY_NAME

The GET_HEADER_BY_NAME procedure returns the header of the HTTP response according to the specified name. The signature is:

```
GET_HEADER_BY_NAME(<r> IN OUT UTL_HTTP.RESP, <name> VARCHAR2,
<value> OUT VARCHAR2, <n> INTEGER DEFAULT 1)
```

Parameters

<r>

<r> is the HTTP response record.

<name>

<name> is the name of the response header to retrieve.

<value>

<value> is the value of the response header.

<n>

<n> is the *n*th header of the HTTP response record to retrieve according to the values specified by <name>. The default is 1.

Examples

The following example retrieves the header for Content-Type.

```
DECLARE
    v_req      UTL_HTTP.REQ;
    v_resp     UTL_HTTP.RESP;
    v_name     VARCHAR2(30) := 'Content-Type';
    v_value    VARCHAR2(200);
BEGIN
    v_req := UTL_HTTP.BEGIN_REQUEST('www.enterprisedb.com');
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);
    UTL_HTTP.GET_HEADER_BY_NAME(v_resp, v_name, v_value);
    DBMS_OUTPUT.PUT_LINE(v_name || ': ' || v_value);
    UTL_HTTP.END_RESPONSE(v_resp);
END;
```

Content-Type: text/html; charset=utf-8

3.21.11 GET_HEADER_COUNT

The GET_HEADER_COUNT function returns the number of HTTP response headers. The signature is:

```
GET_HEADER_COUNT(<r> IN OUT UTL_HTTP.RESP) RETURN INTEGER
```

This function returns an INTEGER value.

Parameters

<r>

<r> is the HTTP response record.

3.21.12 GET_RESPONSE

The GET_RESPONSE function sends the network request and returns any HTTP response. The signature is:

```
GET_RESPONSE(<r> IN OUT UTL_HTTP.REQ) RETURN UTL_HTTP.RESP
```

This function returns a UTL_HTTP.RESP record.

Parameters

<r>

<r> is the HTTP request record.

3.21.13 GET_RESPONSE_ERROR_CHECK

The GET_RESPONSE_ERROR_CHECK procedure returns whether or not response error check is set. The signature is:

```
GET_RESPONSE_ERROR_CHECK(<enable> OUT BOOLEAN)
```

Parameters

<enable>

<enable> returns TRUE if response error check is set, otherwise it returns FALSE.

3.21.14 GET_TRANSFER_TIMEOUT

The GET_TRANSFER_TIMEOUT procedure returns the current, default transfer timeout setting for HTTP requests. The signature is:

```
GET_TRANSFER_TIMEOUT(<timeout> OUT INTEGER)
```

Parameters

<timeout>

<timeout> is the transfer timeout setting in seconds.

3.21.15 READ_LINE

The `READ_LINE` procedure returns the data from the HTTP response body in text form until the end of line is reached. A CR character, a LF character, a CR LF sequence, or the end of the response body constitutes the end of line. The signature is:

```
READ_LINE(<r> IN OUT UTL_HTTP.RESP, <data> OUT VARCHAR2, <remove_crlf>
          BOOLEAN DEFAULT FALSE)
```

Parameters

<r>

<r> is the HTTP response record.

<data>

<data> is the response body in text form.

<remove_crlf>

Set <remove_crlf> to TRUE to remove new line characters, otherwise set to FALSE. The default is FALSE.

Examples

The following example retrieves and displays the body of the specified website.

```
DECLARE
    v_req      UTL_HTTP.REQ;
    v_resp     UTL_HTTP.RESP;
    v_value    VARCHAR2(1024);
BEGIN
    v_req := UTL_HTTP.BEGIN_REQUEST('http://www.enterprisedb.com');
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);
    LOOP
        UTL_HTTP.READ_LINE(v_resp, v_value, TRUE);
        DBMS_OUTPUT.PUT_LINE(v_value);
    END LOOP;
    EXCEPTION
        WHEN OTHERS THEN
            UTL_HTTP.END_RESPONSE(v_resp);
END;
```

The following is the output.

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
  "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html xmlns="http://www.w3.org/1999/xhtml" xml:lang="en" lang="en" dir="ltr">

  <!-- _____ HEAD _____ -->

  <head>
<meta http-equiv="Content-Type" content="text/html; charset=utf-8" />

    <title>EnterpriseDB | The Postgres Database Company</title>

    <meta http-equiv="Content-Type" content="text/html; charset=utf-8" />
<meta name="keywords" content="postgres, postgresql, postgresql installer, mysql,
migration, open source database, training, replication" />
<meta name="description" content="The leader in open source database products,
services, support, training and expertise based on PostgreSQL. Free downloads,
documentation, and tutorials." />
```



```
<meta name="abstract" content="The Enterprise PostgreSQL Company" />
<link rel="EditURI" type="application/rsd+xml" title="RSD" href="http://www.
↪enterprisedb.com/blogapi/rsd" />
<link rel="alternate" type="application/rss+xml" title="EnterpriseDB RSS" href="http:/
↪www.enterprisedb.com/rss.xml" />
<link rel="shortcut icon" href="/sites/all/themes/edb_pixelcrayons/favicon.ico" type=
↪"image/x-icon" />
    <link type="text/css" rel="stylesheet" media="all" href="/sites/default/files/css/
↪css_db11adabae0aed6b79a2c3c52def4754.css" />
<!--[if IE 6]>
<link type="text/css" rel="stylesheet" media="all" href="/sites/all/themes/oho_basic/
↪css/ie6.css?g" />
<![endif]-->
<!--[if IE 7]>
<link type="text/css" rel="stylesheet" media="all" href="/sites/all/themes/oho_basic/
↪css/ie7.css?g" />
<![endif]-->
    <script type="text/javascript" src="/sites/default/files/js/js_
↪74d97b1176812e2fd6e43d62503a5204.js"></script>
<script type="text/javascript">
<!--//--><![CDATA[//><!--
```

3.21.16 READ_RAW

The `READ_RAW` procedure returns the data from the HTTP response body in binary form. The number of bytes returned is specified by the `<len>` parameter. The signature is:

```
READ_RAW(<r> IN OUT UTL_HTTP.RESP, <data> OUT RAW, <len> INTEGER)
```

Parameters

`<r>`

`<r>` is the HTTP response record.

`<data>`

`<data>` is the response body in binary form.

`<len>`

Set `<len>` to the number of bytes of data to be returned.

Examples

The following example retrieves and displays the first 150 bytes in binary form.

```
DECLARE
    v_req      UTL_HTTP.REQ;
    v_resp     UTL_HTTP.RESP;
    v_data     RAW;
BEGIN
    v_req := UTL_HTTP.BEGIN_REQUEST('http://www.enterprisedb.com');
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);
    UTL_HTTP.READ_RAW(v_resp, v_data, 150);
    DBMS_OUTPUT.PUT_LINE(v_data);
    UTL_HTTP.END_RESPONSE(v_resp);
END;
```

The following is the output from the example.

```
\x3c21444f43545950452068746d6c205055424c494320222d2f2f5733432f2f4454442058485
44d4c20312e30205374726963742f2f454e220d0a202022687474703a2f2f7777772e77332e6f
72672f54522f7868746d6c312f4454442f7868746d6c312d7374726963742e647464223e0d0a3
c68746d6c20786d6c6e733d22687474703a2f2f7777772e77332e6f72672f313939392f
```

3.21.17 READ_TEXT

The `READ_TEXT` procedure returns the data from the HTTP response body in text form. The maximum number of characters returned is specified by the `<len>` parameter. The signature is:

```
READ_TEXT(<r> IN OUT UTL_HTTP.RESP, <data> OUT VARCHAR2, <len>
INTEGER)
```

Parameters

`<r>`

`<r>` is the HTTP response record.

`<data>`

`<data>` is the response body in text form.

`<len>`

Set <len> to the maximum number of characters to be returned.

Examples

The following example retrieves the first 150 characters.

```
DECLARE
    v_req          UTL_HTTP.REQ;
    v_resp         UTL_HTTP.RESP;
    v_data         VARCHAR2(150);
BEGIN
    v_req := UTL_HTTP.BEGIN_REQUEST('http://www.enterprisedb.com');
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);
    UTL_HTTP.READ_TEXT(v_resp, v_data, 150);
    DBMS_OUTPUT.PUT_LINE(v_data);
    UTL_HTTP.END_RESPONSE(v_resp);
END;
```

The following is the output.

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
    "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html xmlns="http://www.w3.org/1999/
```

3.21.18 REQUEST

The REQUEST function returns the first 2000 bytes retrieved from a user-specified URL. The signature is:

```
REQUEST(<url> IN VARCHAR2) RETURN VARCHAR2
```

If the data found at the given URL is longer than 2000 bytes, the remainder will be discarded. If the data found at the given URL is shorter than 2000 bytes, the result will be shorter than 2000 bytes.

Parameters

<url>

<url> is the Uniform Resource Locator from which UTL_HTTP will return content.

Example

The following command returns the first 2000 bytes retrieved from the EnterpriseDB website:

```
SELECT UTL_HTTP.REQUEST('http://www.enterprisedb.com/') FROM DUAL;
```

3.21.19 REQUEST_PIECES

The REQUEST_PIECES function returns a table of 2000-byte segments retrieved from an URL. The signature is:

```
REQUEST_PIECES(<url> IN VARCHAR2, <max_pieces> NUMBER IN
DEFAULT 32767) RETURN UTL_HTTP.HTML_PIECES
```

Parameters

<url>

<url> is the Uniform Resource Locator from which UTL_HTTP will return content.

<max_pieces>

<max_pieces> specifies the maximum number of 2000-byte segments that the REQUEST_PIECES function will return. If <max_pieces> specifies more units than are available at the specified <url>, the final unit will contain fewer bytes.

Example

The following example returns the first four 2000 byte segments retrieved from the EnterpriseDB website:

```
DECLARE
    result UTL_HTTP.HTML_PIECES;
BEGIN
    result := UTL_HTTP.REQUEST_PIECES('http://www.enterprisedb.com/', 4);
END
```

3.21.20 SET_BODY_CHARSET

The SET_BODY_CHARSET procedure sets the default character set of the body of future HTTP requests. The signature is:

```
SET_BODY_CHARSET(<charset> VARCHAR2 DEFAULT NULL)
```

Parameters

<charset>

<charset> is the character set of the body of future requests. The default is null in which case the database character set is assumed.

3.21.21 SET_FOLLOW_REDIRECT

The SET_FOLLOW_REDIRECT procedure sets the maximum number of times the HTTP redirect instruction is to be followed in the response to this request or future requests. This procedures has two signatures:

```
SET_FOLLOW_REDIRECT(<max_redirects> IN INTEGER DEFAULT 3)
```

and

```
SET_FOLLOW_REDIRECT(<r> IN OUT UTL_HTTP.REQ, <max_redirects> IN
INTEGER DEFAULT 3)
```

Use the second form to change the maximum number of redirections for an individual request that a request inherits from the session default settings.

Parameters

<r>

<r> is the HTTP request record.

<max_redirects>

<max_redirects> is maximum number of redirections allowed. Set to 0 to disable redirections. The default is 3.

3.21.22 SET_HEADER

The SET_HEADER procedure sets the HTTP request header. The signature is:

```
SET_HEADER(<r> IN OUT UTL_HTTP.REQ, <name> IN VARCHAR2, <value>
IN VARCHAR2 DEFAULT NULL)
```

Parameters

<r>

<r> is the HTTP request record.

<name>

<name> is the name of the request header.

<value>

<value> is the value of the request header. The default is null.

3.21.23 SET_RESPONSE_ERROR_CHECK

The SET_RESPONSE_ERROR_CHECK procedure determines whether or not HTTP 4xx and 5xx status codes returned by the GET_RESPONSE function should be interpreted as errors. The signature is:

```
SET_RESPONSE_ERROR_CHECK(<enable> IN BOOLEAN DEFAULT FALSE)
```

Parameters

<enable>

Set <enable> to TRUE if HTTP 4xx and 5xx status codes are to be treated as errors, otherwise set to FALSE. The default is FALSE.

3.21.24 SET_TRANSFER_TIMEOUT

The SET_TRANSFER_TIMEOUT procedure sets the default, transfer timeout setting for waiting for a response from an HTTP request. This procedure has two signatures:

```
SET_TRANSFER_TIMEOUT(<timeout> IN INTEGER DEFAULT 60)
```

and

```
SET_TRANSFER_TIMEOUT(<r> IN OUT UTL_HTTP.REQ, <timeout> IN
INTEGER
DEFAULT 60)
```

Use the second form to change the transfer timeout setting for an individual request that a request inherits from the session default settings.

Parameters

<r>

<r> is the HTTP request record.

<timeout>

<timeout> is the transfer timeout setting in seconds for HTTP requests. The default is 60 seconds.

3.21.25 WRITE_LINE

The WRITE_LINE procedure writes data to the HTTP request body in text form; the text is terminated with a CRLF character pair. The signature is:

```
WRITE_LINE(<r> IN OUT UTL_HTTP.REQ, <data> IN VARCHAR2)
```

Parameters

<r>

<r> is the HTTP request record.

<data>

<data> is the request body in TEXT form.

Example

The following example writes data (Account balance \$500.00) in text form to the request body to be sent using the HTTP POST method. The data is sent to a hypothetical web application (post.php) that accepts and processes data.

```

DECLARE
    v_req          UTL_HTTP.REQ;
    v_resp         UTL_HTTP.RESP;
BEGIN
    v_req := UTL_HTTP.BEGIN_REQUEST('http://www.example.com/post.php',
    'POST');
    UTL_HTTP.SET_HEADER(v_req, 'Content-Length', '23');
    UTL_HTTP.WRITE_LINE(v_req, 'Account balance $500.00');
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);
    DBMS_OUTPUT.PUT_LINE('Status Code: ' || v_resp.status_code);
    DBMS_OUTPUT.PUT_LINE('Reason Phrase: ' || v_resp.reason_phrase);
    UTL_HTTP.END_RESPONSE(v_resp);
END;
```

Assuming the web application successfully processed the POST method, the following output would be displayed:

```

Status Code: 200
Reason Phrase: OK
```

3.21.26 WRITE_RAW

The WRITE_RAW procedure writes data to the HTTP request body in binary form. The signature is:

```
WRITE_RAW(<r> IN OUT UTL_HTTP.REQ, <data> IN RAW)
```

Parameters

<r>

<r> is the HTTP request record.

<data>

<data> is the request body in binary form.

Example

The following example writes data in binary form to the request body to be sent using the HTTP POST method to a hypothetical web application that accepts and processes such data.

```
DECLARE
    v_req          UTL_HTTP.REQ;
    v_resp         UTL_HTTP.RESP;
BEGIN
    v_req := UTL_HTTP.BEGIN_REQUEST('http://www.example.com/post.php',
        'POST');
    UTL_HTTP.SET_HEADER(v_req, 'Content-Length', '23');
    UTL_HTTP.WRITE_RAW(v_req, HEXTORAW
('54657374696e6720504f5354206d6574686f6420696e20485454502072657175657374'));
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);
    DBMS_OUTPUT.PUT_LINE('Status Code: ' || v_resp.status_code);
    DBMS_OUTPUT.PUT_LINE('Reason Phrase: ' || v_resp.reason_phrase);
    UTL_HTTP.END_RESPONSE(v_resp);
END;
```

The text string shown in the HEXTORAW function is the hexadecimal translation of the text Testing POST method in HTTP request.

Assuming the web application successfully processed the POST method, the following output would be displayed:

```
Status Code: 200
Reason Phrase: OK
```

3.21.27 WRITE_TEXT

The `WRITE_TEXT` procedure writes data to the HTTP request body in text form. The signature is:

```
WRITE_TEXT(<r> IN OUT UTL_HTTP.REQ, <data> IN VARCHAR2)
```

Parameters

<r>

<r> is the HTTP request record.

<data>

<data> is the request body in text form.

Example

The following example writes data (Account balance \$500.00) in text form to the request body to be sent using the HTTP POST method. The data is sent to a hypothetical web application (`post.php`) that accepts and processes data.

```
DECLARE
    v_req      UTL_HTTP.REQ;
    v_resp     UTL_HTTP.RESP;
BEGIN
    v_req := UTL_HTTP.BEGIN_REQUEST('http://www.example.com/post.php',
        'POST');
    UTL_HTTP.SET_HEADER(v_req, 'Content-Length', '23');
    UTL_HTTP.WRITE_TEXT(v_req, 'Account balance $500.00');
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);
    DBMS_OUTPUT.PUT_LINE('Status Code: ' || v_resp.status_code);
    DBMS_OUTPUT.PUT_LINE('Reason Phrase: ' || v_resp.reason_phrase);
    UTL_HTTP.END_RESPONSE(v_resp);
END;
```

Assuming the web application successfully processed the POST method, the following output would be displayed:

```
Status Code: 200
Reason Phrase: OK
```


3.22 UTL_MAIL

The UTL_MAIL package provides the capability to manage e-mail. Advanced Server supports the following procedures:

Function/Procedure	Return Type	Description
SEND(<sender>, <recipients>, <cc>, <bcc>, <subject>, <message> [, <mime_type> [, <priority>]])	n/a	Packages and sends an e-mail to an SMTP server.
SEND_ATTACH_RAW(<sender>, <recipients>, <cc>, <bcc>, <subject>, <message>, <mime_type>, <priority>, <attachment> [, <att_inline> [, <att_mime_type> [, <att_filename>]]])	n/a	Same as the SEND procedure, but with BYTEA or large object attachments.
SEND_ATTACH_VARCHAR2(<sender>, <recipients>, <cc>, <bcc>, <subject>, <message>, <mime_type>, <priority>, <attachment> [, <att_inline> [, <att_mime_type> [, <att_filename>]]])	n/a	Same as the SEND procedure, but with VARCHAR2 attachments.

Note: An administrator must grant execute privileges to each user or group before they can use this package.

3.22.1 SEND

The SEND procedure provides the capability to send an e-mail to an SMTP server.

```
SEND(<sender> VARCHAR2, <recipients> VARCHAR2, <cc> VARCHAR2, <bcc> VARCHAR2,
<subject> VARCHAR2, <message> VARCHAR2 [, <mime_type> VARCHAR2 [, <priority>
PLS_INTEGER ]])
```

Parameters

<sender>

E-mail address of the sender.

<recipients>

Comma-separated e-mail addresses of the recipients.

<cc>

Comma-separated e-mail addresses of copy recipients.

<bcc>

Comma-separated e-mail addresses of blind copy recipients.

<subject>

Subject line of the e-mail.

<message>

Body of the e-mail.

<mime_type>

Mime type of the message. The default is text/plain; charset=us-ascii.

<priority>

Priority of the e-mail The default is 3.

Examples

The following anonymous block sends a simple e-mail message.

```
DECLARE
    v_sender      VARCHAR2(30);
    v_recipients  VARCHAR2(60);
    v_subj        VARCHAR2(20);
    v_msg         VARCHAR2(200);
BEGIN
    v_sender := 'jsmith@enterprisedb.com';
    v_recipients := 'ajones@enterprisedb.com,rrogers@enterprisedb.com';
    v_subj := 'Holiday Party';
    v_msg := 'This year''s party is scheduled for Friday, Dec. 21 at ' ||
            '6:00 PM. Please RSVP by Dec. 15th.';
    UTL_MAIL.SEND(v_sender,v_recipients,NULL,NULL,v_subj,v_msg);
END;
```

3.22.2 SEND_ATTACH_RAW

The SEND_ATTACH_RAW procedure provides the capability to send an e-mail to an SMTP server with an attachment containing either BYTEA data or a large object (identified by the large object's OID). The call to SEND_ATTACH_RAW can be written in two ways:

```
SEND_ATTACH_RAW(<sender> VARCHAR2, <recipients> VARCHAR2, <cc> VARCHAR2,
<bcc> VARCHAR2, <subject> VARCHAR2, <message> VARCHAR2, <mime_type> VARCHAR2,
<priority> PLS_INTEGER, <attachment> BYTEA[, <att_inline> BOOLEAN [,
<att_mime_type> VARCHAR2[, <att_filename> VARCHAR2 ]]])
```

or

```
SEND_ATTACH_RAW(<sender> VARCHAR2, <recipients> VARCHAR2, <cc> VARCHAR2,
<bcc> VARCHAR2, <subject> VARCHAR2, <message> VARCHAR2, <mime_type>
VARCHAR2, <priority> PLS_INTEGER, <attachment> OID [, <att_inline> BOOLEAN
[, <att_mime_type> VARCHAR2 [, <att_filename> VARCHAR2 ]]])
```

Parameters

<sender>

E-mail address of the sender.

<recipients>

Comma-separated e-mail addresses of the recipients.

<cc>

Comma-separated e-mail addresses of copy recipients.

<bcc>

Comma-separated e-mail addresses of blind copy recipients.

<subject>

Subject line of the e-mail.

<message>

Body of the e-mail.

<mime_type>

Mime type of the message. The default is text/plain; charset=us-ascii.

<priority>

Priority of the e-mail. The default is 3.

<attachment>

The attachment.

<att_inline>

If set to TRUE, then the attachment is viewable inline, FALSE otherwise. The default is TRUE.

<att_mime_type>

Mime type of the attachment. The default is application/octet.

<att_filename>

The file name containing the attachment. The default is null.

3.22.3 SEND_ATTACH_VARCHAR2

The SEND_ATTACH_VARCHAR2 procedure provides the capability to send an e-mail to an SMTP server with a text attachment.

```
SEND_ATTACH_VARCHAR2(<sender> VARCHAR2, <recipients> VARCHAR2, <cc> VARCHAR2,
<bcc> VARCHAR2, <subject> VARCHAR2, <message> VARCHAR2, <mime_type> VARCHAR2,
<priority> PLS_INTEGER, <attachment> VARCHAR2 [, <att_inline> BOOLEAN [,
<att_mime_type> VARCHAR2 [, <att_filename> VARCHAR2 ]]))
```

Parameters

<sender>

E-mail address of the sender.

<recipients>

Comma-separated e-mail addresses of the recipients.

<cc>

Comma-separated e-mail addresses of copy recipients.

<bcc>

Comma-separated e-mail addresses of blind copy recipients.

<subject>

Subject line of the e-mail.

<message>

Body of the e-mail.

<mime_type>

Mime type of the message. The default is text/plain; charset=us-ascii.

<priority>

Priority of the e-mail The default is 3 .

<attachment>

The VARCHAR2 attachment.

<att_inline>

If set to TRUE, then the attachment is viewable inline, FALSE otherwise. The default is TRUE.

<att_mime_type>

Mime type of the attachment. The default is text/plain; charset=us-ascii.

<att_filename>

The file name containing the attachment. The default is null.

3.23 UTL_RAW

The UTL_RAW package allows you to manipulate or retrieve the length of raw data types.

Note: An administrator must grant execute privileges to each user or group before they can use this package.

Function/Procedure	Function or Procedure	Return Type	Description
CAST_TO_RAW(c IN VARCHAR2)	Function	RAW	Converts a VARCHAR2 string to a RAW value.
CAST_TO_VARCHAR2(r IN RAW)	Function	VARCHAR2	Converts a RAW value to a VARCHAR2 string.
CONCAT(r1 IN RAW, r2 IN RAW, r3 IN RAW,...)	Function	RAW	Concatenate multiple RAW values into a single RAW value.
CONVERT(r IN RAW, to_charset IN VARCHAR2, from_charset IN VARCHAR2)	Function	RAW	Converts encoded data from one encoding to another, and returns the result as a RAW value.
LENGTH(r IN RAW)	Function	NUMBER	Returns the length of a RAW value.
SUBSTR(r IN RAW, pos IN INTEGER, len IN INTEGER)	Function	RAW	Returns a portion of a RAW value.

Advanced Server's implementation of UTL_RAW is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

3.23.1 CAST_TO_RAW

The CAST_TO_RAW function converts a VARCHAR2 string to a RAW value. The signature is:

```
CAST_TO_RAW (<c> VARCHAR2)
```

The function returns a RAW value if you pass a non-NULL value; if you pass a NULL value, the function will return NULL.

Parameters

<c>

The VARCHAR2 value that will be converted to RAW.

Example

The following example uses the CAST_TO_RAW function to convert a VARCHAR2 string to a RAW value:

```
DECLARE
  v VARCHAR2;
  r RAW;
BEGIN
  v := 'Accounts';
  dbms_output.put_line(v);
  r := UTL_RAW.CAST_TO_RAW(v);
  dbms_output.put_line(r);
END;
```

The result set includes the content of the original string and the converted RAW value:

```
Accounts
\x4163636f756e7473
```

3.23.2 CAST_TO_VARCHAR2

The CAST_TO_VARCHAR2 function converts RAW data to VARCHAR2 data. The signature is:

```
CAST_TO_VARCHAR2 (<r> RAW)
```

The function returns a VARCHAR2 value if you pass a non-NULL value; if you pass a NULL value, the function will return NULL.

Parameters

<r>

The RAW value that will be converted to a VARCHAR2 value.

Example

The following example uses the CAST_TO_VARCHAR2 function to convert a RAW value to a VARCHAR2 string:

```
DECLARE
  r RAW;
  v VARCHAR2;
BEGIN
  r := '\x4163636f756e7473'
  dbms_output.put_line(v);
  v := UTL_RAW.CAST_TO_VARCHAR2(r);
  dbms_output.put_line(r);
END;
```

The result set includes the content of the original string and the converted RAW value:

```
\x4163636f756e7473
Accounts
```

3.23.3 CONCAT

The CONCAT function concatenates multiple RAW values into a single RAW value. The signature is:

```
CONCAT(<r1> RAW, <r2> RAW, <r3> RAW, ...)
```

The function returns a RAW value. Unlike the Oracle implementation, the Advanced Server implementation is a variadic function, and does not place a restriction on the number of values that can be concatenated.

Parameters

<r1, r2, r3, ...>

The RAW values that CONCAT will concatenate.

Example

The following example uses the CONCAT function to concatenate multiple RAW values into a single RAW value:

```
| SELECT UTL_RAW.CAST_TO_VARCHAR2(UTL_RAW.CONCAT('\x61', '\x62',  
| '\x63')) FROM DUAL;  
| concat  
| -----  
| abc  
| (1 row)
```

The result (the concatenated values) is then converted to VARCHAR2 format by the CAST_TO_VARCHAR2 function.

3.23.4 CONVERT

The CONVERT function converts a string from one encoding to another encoding and returns the result as a RAW value. The signature is:

```
CONVERT(<r> RAW, <to_charset> VARCHAR2, <from_charset> VARCHAR2)
```

The function returns a RAW value.

Parameters

<r>

The RAW value that will be converted.

<to_charset>

The name of the encoding to which <r> will be converted.

<from_charset>

The name of the encoding from which <r> will be converted.

Example

The following example uses the UTL_RAW.CAST_TO_RAW function to convert a VARCHAR2 string (Accounts) to a raw value, and then convert the value from UTF8 to LATIN7, and then from LATIN7 to UTF8:

```
DECLARE
  r RAW;
  v VARCHAR2;
BEGIN
  v:= 'Accounts';
  dbms_output.put_line(v);
  r:= UTL_RAW.CAST_TO_RAW(v);
  dbms_output.put_line(r);
  r:= UTL_RAW.CONVERT(r, 'UTF8', 'LATIN7');
  dbms_output.put_line(r);
  r:= UTL_RAW.CONVERT(r, 'LATIN7', 'UTF8');
  dbms_output.put_line(r);
```

The example returns the VARCHAR2 value, the RAW value, and the converted values:

```
Accounts
\x4163636f756e7473
\x4163636f756e7473
\x4163636f756e7473
```

3.23.5 LENGTH

The LENGTH function returns the length of a RAW value. The signature is:

```
LENGTH (<r> RAW)
```

The function returns a RAW value.

Parameters

<r>

The RAW value that LENGTH will evaluate.

Example

The following example uses the LENGTH function to return the length of a RAW value:

```
SELECT UTL_RAW.LENGTH(UTL_RAW.CAST_TO_RAW('Accounts')) FROM DUAL;
length
-----
8
(1 row)
```

The following example uses the LENGTH function to return the length of a RAW value that includes multi-byte characters:

```
SELECT UTL_RAW.LENGTH(UTL_RAW.CAST_TO_RAW('hello'));
length
-----
5
(1 row)
```

3.23.6 SUBSTR

The SUBSTR function returns a substring of a RAW value. The signature is:

```
SUBSTR (<r> RAW, <pos> INTEGER, <len> INTEGER)
```

This function returns a RAW value.

Parameters

<r>

The RAW value from which the substring will be returned.

<pos>

The position within the RAW value of the first byte of the returned substring.

- If <pos> is 0 or 1, the substring begins at the first byte of the RAW value.
- If <pos> is greater than one, the substring begins at the first byte specified by <pos>. For example, if <pos> is 3, the substring begins at the third byte of the value.
- If <pos> is negative, the substring begins at <pos> bytes from the end of the source value. For example, if <pos> is -3, the substring begins at the third byte from the end of the value.

<len>

The maximum number of bytes that will be returned.

Example

The following example uses the SUBSTR function to select a substring that begins 3 bytes from the start of a RAW value:

```
SELECT UTL_RAW.SUBSTR(UTL_RAW.CAST_TO_RAW('Accounts'), 3, 5) FROM DUAL;
substr
-----
count
(1 row)
```

The following example uses the SUBSTR function to select a substring that starts 5 bytes from the end of a RAW value:

```
SELECT UTL_RAW.SUBSTR(UTL_RAW.CAST_TO_RAW('Accounts'), -5 , 3) FROM
DUAL;
substr
-----
oun
(1 row)
```

3.24 UTL_SMTP

The UTL_SMTP package provides the capability to send e-mails over the Simple Mail Transfer Protocol (SMTP).

Note: An administrator must grant execute privileges to each user or group before they can use this package.

Function/Procedure	Function or Procedure	Return Type	Description
CLOSE_DATA(c IN OUT)	Procedure	n/a	Ends an e-mail message.
COMMAND(c IN OUT, cmd [, arg])	Both	REPLY	Execute an SMTP command.
COMMAND_REPLIES(c IN OUT, cmd [, arg])	Function	REPLIES	Execute an SMTP command where multiple reply lines are expected.
DATA(c IN OUT, body VARCHAR2)	Procedure	n/a	Specify the body of an e-mail message.
EHLO(c IN OUT, domain)	Procedure	n/a	Perform initial handshaking with an SMTP server and return extended information.
HELO(c IN OUT, domain)	Procedure	n/a	Perform initial handshaking with an SMTP server
HELP(c IN OUT [, command])	Function	REPLIES	Send the HELP command.
MAIL(c IN OUT, sender [, parameters])	Procedure	n/a	Start a mail transaction.
NOOP(c IN OUT)	Both	REPLY	Send the null command.
OPEN_CONNECTION(host [, port [, tx_timeout]])	Function	CONNECTION	Open a connection.
OPEN_DATA(c IN OUT)	Both	REPLY	Send the DATA command.
QUIT(c IN OUT)	Procedure	n/a	Terminate the SMTP session and disconnect.
RCPT(c IN OUT, recipient [, parameters])	Procedure	n/a	Specify the recipient of an e-mail message.
RSET(c IN OUT)	Procedure	n/a	Terminate the current mail transaction.
VRFY(c IN OUT, recipient)	Function	REPLY	Validate an e-mail address.
WRITE_DATA(c IN OUT, data)	Procedure	n/a	Write a portion of the e-mail message.

Advanced Server's implementation of UTL_SMTP is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The following table lists the public variables available in the UTL_SMTP package.

Public Variables	Data Type	Value	Description
connection	RECORD		Description of an SMTP connection.
reply	RECORD		SMTP reply line.

3.24.1 CONNECTION

The CONNECTION record type provides a description of an SMTP connection.

```
TYPE connection IS RECORD (
    host          VARCHAR2(255),
    port          PLS_INTEGER,
    tx_timeout    PLS_INTEGER
);
```

3.24.2 REPLY/REPLIES

The REPLY record type provides a description of an SMTP reply line. REPLIES is a table of multiple SMTP reply lines.

```
TYPE reply IS RECORD (
    code          INTEGER,
    text          VARCHAR2(508)
);
TYPE replies IS TABLE OF reply INDEX BY BINARY_INTEGER;
```

3.24.3 CLOSE_DATA

The CLOSE_DATA procedure terminates an e-mail message by sending the following sequence:

```
<CR><LF>.<CR><LF>
```

This is a single period at the beginning of a line.

```
CLOSE_DATA(<c> IN OUT CONNECTION)
```

Parameters

```
<c>
```

The SMTP connection to be closed.

3.24.4 COMMAND

The COMMAND procedure provides the capability to execute an SMTP command. If you are expecting multiple reply lines, use COMMAND_REPLIES.

```
<reply> REPLY COMMAND(<c> IN OUT CONNECTION, <cmd> VARCHAR2
[, <arg> VARCHAR2 ])
COMMAND(<c> IN OUT CONNECTION, <cmd> VARCHAR2 [, <arg> VARCHAR2 ])
```

Parameters

```
<c>
```

The SMTP connection to which the command is to be sent.

```
<cmd>
```

The SMTP command to be processed.

<arg>

An argument to the SMTP command. The default is null.

<reply>

SMTP reply to the command. If SMTP returns multiple replies, only the last one is returned in <reply>.

See [Reply/Replies](#) for a description of `REPLY` and `REPLIES`.

3.24.5 COMMAND_REPLIES

The `COMMAND_REPLIES` function processes an SMTP command that returns multiple reply lines. Use `COMMAND` if only a single reply line is expected.

```
<replies> REPLIES COMMAND (<c> IN OUT CONNECTION, <cmd> VARCHAR2  
[, <arg> VARCHAR2 ])
```

Parameters

<c>

The SMTP connection to which the command is to be sent.

<cmd>

The SMTP command to be processed.

<arg>

An argument to the SMTP command. The default is null.

<replies>

SMTP reply lines to the command. See [Reply/Replies](#) for a description of `REPLY` and `REPLIES`.

3.24.6 DATA

The `DATA` procedure provides the capability to specify the body of the e-mail message. The message is terminated with a <CR><LF> .<CR><LF> sequence.

```
DATA (<c> IN OUT CONNECTION, <body> VARCHAR2)
```

Parameters

<c>

The SMTP connection to which the command is to be sent.

<body>

Body of the e-mail message to be sent.

3.24.7 EHLO

The EHLO procedure performs initial handshaking with the SMTP server after establishing the connection. The EHLO procedure allows the client to identify itself to the SMTP server according to RFC 821. RFC 1869 specifies the format of the information returned in the server's reply. The *HELO* procedure performs the equivalent functionality, but returns less information about the server.

```
EHLO(<c> IN OUT CONNECTION, <domain> VARCHAR2)
```

Parameters

<c>

The connection to the SMTP server over which to perform handshaking.

<domain>

Domain name of the sending host.

3.24.8 HELO

The HELO procedure performs initial handshaking with the SMTP server after establishing the connection. The HELO procedure allows the client to identify itself to the SMTP server according to RFC 821. The *EHLO* procedure performs the equivalent functionality, but returns more information about the server.

```
HELO(<c> IN OUT, <domain*> VARCHAR2)
```

Parameters

<c>

The connection to the SMTP server over which to perform handshaking.

<domain>

Domain name of the sending host.

3.24.9 HELP

The HELP function provides the capability to send the HELP command to the SMTP server.

```
<replies> REPLIES HELP(<c> IN OUT CONNECTION [, <command> VARCHAR2])
```

Parameters

<c>

The SMTP connection to which the command is to be sent.

<command>

Command on which help is requested.

<replies>

SMTP reply lines to the command. See [Reply/Replies](#) for a description of REPLY and REPLIES.

3.24.10 MAIL

The MAIL procedure initiates a mail transaction.

```
MAIL(<c> IN OUT CONNECTION, <sender> VARCHAR2
[, <parameters> VARCHAR2 ])
```

Parameters

<c>

Connection to SMTP server on which to start a mail transaction.

<sender>

The sender's e-mail address.

<parameters>

Mail command parameters in the format, key=value as defined in RFC 1869.

3.24.11 NOOP

The NOOP function/procedure sends the null command to the SMTP server. The NOOP has no effect upon the server except to obtain a successful response.

```
<reply> REPLY NOOP(<c> IN OUT CONNECTION)
NOOP(<c> IN OUT CONNECTION)
```

Parameters

<c>

The SMTP connection on which to send the command.

<reply>

SMTP reply to the command. If SMTP returns multiple replies, only the last one is returned in <reply>. See [Reply/Replies](#) for a description of REPLY and REPLIES.

3.24.12 OPEN_CONNECTION

The OPEN_CONNECTION functions open a connection to an SMTP server.

```
<c> CONNECTION OPEN_CONNECTION(<host> VARCHAR2 [, <port>
PLS_INTEGER [, <tx_timeout> PLS_INTEGER DEFAULT NULL])
```

Parameters

<host>

Name of the SMTP server.

<port>

Port number on which the SMTP server is listening. The default is 25.

<tx_timeout>

Time out value in seconds. Do not wait is indicated by specifying 0. Wait indefinitely is indicated by setting timeout to null. The default is null.

<c>

Connection handle returned by the SMTP server.

3.24.13 OPEN_DATA

The OPEN_DATA procedure sends the DATA command to the SMTP server.

```
OPEN_DATA (<c> IN OUT CONNECTION)
```

Parameters

<c>

SMTP connection on which to send the command.

3.24.14 QUIT

The QUIT procedure closes the session with an SMTP server.

```
QUIT (<c> IN OUT CONNECTION)
```

Parameters

<c>

SMTP connection to be terminated.

3.24.15 RCPT

The RCPT procedure provides the e-mail address of the recipient. To schedule multiple recipients, invoke RCPT multiple times.

```
RCPT (<c> IN OUT CONNECTION, <recipient> VARCHAR2  
[, <parameters> VARCHAR2 ])
```

Parameters

<c>

Connection to SMTP server on which to add a recipient.

<recipient>

The recipient's e-mail address.

<parameters>

Mail command parameters in the format, key=value as defined in RFC 1869.

3.24.16 RSET

The RSET procedure provides the capability to terminate the current mail transaction.

```
RSET (<c> IN OUT CONNECTION)
```

Parameters

<c>

SMTP connection on which to cancel the mail transaction.

3.24.17 VRFY

The VRFY function provides the capability to validate and verify the recipient's e-mail address. If valid, the recipient's full name and fully qualified mailbox is returned.

```
<reply> REPLY VRFY(<c> IN OUT CONNECTION, <recipient> VARCHAR2)
```

Parameters

<c>

The SMTP connection on which to verify the e-mail address.

<recipient>

The recipient's e-mail address to be verified.

<reply>

SMTP reply to the command. If SMTP returns multiple replies, only the last one is returned in *reply*. See [Reply/Replies](#) for a description of REPLY and REPLIES.

3.24.18 WRITE_DATA

The WRITE_DATA procedure provides the capability to add VARCHAR2 data to an e-mail message. The WRITE_DATA procedure may be repetitively called to add data.

```
WRITE_DATA(<c> IN OUT CONNECTION, <data> VARCHAR2)
```

Parameters

<c>

The SMTP connection on which to add data.

<data>

Data to be added to the e-mail message. The data must conform to the RFC 822 specification.

3.24.19 Comprehensive Example

The following procedure constructs and sends a text e-mail message using the UTL_SMTP package.

```
CREATE OR REPLACE PROCEDURE send_mail (
    p_sender      VARCHAR2,
    p_recipient   VARCHAR2,
    p_subj        VARCHAR2,
    p_msg          VARCHAR2,
    p_mailhost     VARCHAR2
)
IS
    v_conn        UTL_SMTP.CONNECTION;
    v_crlf        CONSTANT VARCHAR2(2) := CHR(13) || CHR(10);
    v_port        CONSTANT PLS_INTEGER := 25;
BEGIN
    v_conn := UTL_SMTP.OPEN_CONNECTION(p_mailhost,v_port);
    UTL_SMTP.HELO(v_conn,p_mailhost);
    UTL_SMTP.MAIL(v_conn,p_sender);
    UTL_SMTP.RCPT(v_conn,p_recipient);
    UTL_SMTP.DATA(v_conn, SUBSTR(
        'Date: ' || TO_CHAR(SYSDATE,
```

```

        'Dy, DD Mon YYYY HH24:MI:SS') || v_crlf
        || 'From: ' || p_sender || v_crlf
        || 'To: ' || p_recipient || v_crlf
        || 'Subject: ' || p_subj || v_crlf
        || p_msg
        , 1, 32767));
    UTL_SMTP.QUIT(v_conn);
END;

EXEC send_mail('asmith@enterprisedb.com','pjones@enterprisedb.com','Holiday Party',
→'Are you planning to attend?','smtp.enterprisedb.com');
```

The following example uses the OPEN_DATA, WRITE_DATA, and CLOSE_DATA procedures instead of the DATA procedure.

```

CREATE OR REPLACE PROCEDURE send_mail_2 (
    p_sender      VARCHAR2,
    p_recipient   VARCHAR2,
    p_subj        VARCHAR2,
    p_msg         VARCHAR2,
    p_mailhost    VARCHAR2
)
IS
    v_conn        UTL_SMTP.CONNECTION;
    v_crlf        CONSTANT VARCHAR2(2) := CHR(13) || CHR(10);
    v_port        CONSTANT PLS_INTEGER := 25;
BEGIN
    v_conn := UTL_SMTP.OPEN_CONNECTION(p_mailhost,v_port);
    UTL_SMTP.HELO(v_conn,p_mailhost);
    UTL_SMTP.MAIL(v_conn,p_sender);
    UTL_SMTP.RCPT(v_conn,p_recipient);
    UTL_SMTP.OPEN_DATA(v_conn);
    UTL_SMTP.WRITE_DATA(v_conn,'From: ' || p_sender || v_crlf);
    UTL_SMTP.WRITE_DATA(v_conn,'To: ' || p_recipient || v_crlf);
    UTL_SMTP.WRITE_DATA(v_conn,'Subject: ' || p_subj || v_crlf);
    UTL_SMTP.WRITE_DATA(v_conn,v_crlf || p_msg);
    UTL_SMTP.CLOSE_DATA(v_conn);
    UTL_SMTP.QUIT(v_conn);
END;

EXEC send_mail_2('asmith@enterprisedb.com','pjones@enterprisedb.com','Holiday Party',
→'Are you planning to attend?','smtp.enterprisedb.com');
```

3.25 UTL_URL

The UTL_URL package provides a way to escape illegal and reserved characters within an URL.

Function/Procedure	Return Type	Description
ESCAPE(url, escape_reserved_chars, url_charset)	VAR-CHAR2	Use the ESCAPE function to escape any illegal and reserved characters in a URL.
UNESCAPE(url, url_charset)	VAR-CHAR2	The UNESCAPE function to convert an URL to it's original form.

The UTL_URL package will return the BAD_URL exception if the call to a function includes an incorrectly-formed URL.

3.25.1 ESCAPE

Use the ESCAPE function to escape illegal and reserved characters within an URL. The signature is:

```
ESCAPE(<url> VARCHAR2, <escape_reserved_chars> BOOLEAN,
      <url_charset> VARCHAR2)
```

Reserved characters are replaced with a percent sign, followed by the two-digit hex code of the ascii value for the escaped character.

Parameters

<url>

<url> specifies the Uniform Resource Locator that UTL_URL will escape.

<escape_reserved_chars>

<escape_reserved_chars> is a BOOLEAN value that instructs the ESCAPE function to escape reserved characters as well as illegal characters:

- If <escaped_reserved_chars> is FALSE, ESCAPE will escape only the illegal characters in the specified URL.
- If <escape_reserved_chars> is TRUE, ESCAPE will escape both the illegal characters and the reserved characters in the specified URL.

By default, <escape_reserved_chars> is FALSE.

Within an URL, legal characters are:

Uppercase A through Z	Lowercase a through z	0 through 9
asterisk (*)	exclamation point (!)	hyphen (-)
left parenthesis ((period (.)	right parenthesis ())
single-quote (')	tilde (~)	underscore (_)

Some characters are legal in some parts of an URL, while illegal in others; to review comprehensive rules about illegal characters, please refer to RFC 2396. Some *examples* of characters that are considered illegal in any part of an URL are:

Illegal Character	Escape Sequence
a blank space ()	%20
curly braces ({ or })	%7b and %7d
hash mark (#)	%23

The ESCAPE function considers the following characters to be reserved, and will escape them if <escape_reserved_chars> is set to TRUE:

Reserved Character	Escape Sequence
ampersand (&)	%5C
at sign (@)	%25
colon (:)	%3a
comma (,)	%2c
dollar sign (\$)	%24
equal sign (=)	%3d
plus sign (+)	%2b
question mark (?)	%3f
semi-colon (;)	%3b
slash (/)	%2f

<url_charset>

<url_charset> specifies a character set to which a given character will be converted before it is escaped. If <url_charset> is NULL, the character will not be converted. The default value of <url_charset> is ISO-8859-1.

Examples

The following anonymous block uses the ESCAPE function to escape the blank spaces in the URL:

```
DECLARE
result varchar2(400);
BEGIN
result := UTL_URL.ESCAPE('http://www.example.com/Using the ESCAPE function.html');
DBMS_OUTPUT.PUT_LINE(result);
END;
```

The resulting (escaped) URL is:

```
http://www.example.com/Using%20the%20ESCAPE%20function.html
```

If you include a value of TRUE for the <escape_reserved_chars> parameter when invoking the function:

```
DECLARE
result varchar2(400);
BEGIN
result := UTL_URL.ESCAPE('http://www.example.com/Using the ESCAPE function.html',_
↳TRUE);
DBMS_OUTPUT.PUT_LINE(result);
END;
```

The ESCAPE function escapes the reserved characters as well as the illegal characters in the URL:

```
http%3A%2F%2Fwww.example.com%2FUsing%20the%20ESCAPE%20function.html
```

3.25.2 UNESCAPE

The UNESCAPE function removes escape characters added to an URL by the ESCAPE function, converting the URL to it's original form.

The signature is:

```
UNESCAPE(<url> VARCHAR2, <url_charset> VARCHAR2)
```

Parameters

<url>

<url> specifies the Uniform Resource Locator that UTL_URL will unescape.

<url_charset>

After unescaping a character, the character is assumed to be in <url_charset> encoding, and will be converted from that encoding to database encoding before being returned. If <url_charset> is NULL, the character will not be converted. The default value of <url_charset> is ISO-8859-1.

Examples

The following anonymous block uses the ESCAPE function to escape the blank spaces in the URL:

```
DECLARE
result varchar2(400);
BEGIN
result := UTL_URL.UNESCAPE('http://www.example.com/Using%20the%20UNESCAPE%20function.
↪html');
DBMS_OUTPUT.PUT_LINE(result);
END;
```

The resulting (unescaped) URL is:

```
http://www.example.com/Using the UNESCAPE function.htm
```

Acknowledgements

The PostgreSQL 8.3, 8.4, 9.0, 9.1, 9.2, 9.3, 9.4, 9.5, 9.6, 10, 11, and 12 Documentation provided the baseline for the portions of this guide that are common to PostgreSQL, and is hereby acknowledged:

Portions of this EnterpriseDB Software and Documentation may utilize the following copyrighted material, the use of which is hereby acknowledged.

PostgreSQL Documentation, Database Management System

PostgreSQL is Copyright © 1996-2020 by the PostgreSQL Global Development Group and is distributed under the terms of the license of the University of California below.

Postgres95 is Copyright © 1994-5 by the Regents of the University of California.

Permission to use, copy, modify, and distribute this software and its documentation for any purpose, without fee, and without a written agreement is hereby granted, provided that the above copyright notice and this paragraph and the following two paragraphs appear in all copies.

IN NO EVENT SHALL THE UNIVERSITY OF CALIFORNIA BE LIABLE TO ANY PARTY FOR DIRECT, INDIRECT, SPECIAL, INCIDENTAL, OR CONSEQUENTIAL DAMAGES, INCLUDING LOST PROFITS, ARISING OUT OF THE USE OF THIS SOFTWARE AND ITS DOCUMENTATION, EVEN IF THE UNIVERSITY OF CALIFORNIA HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

THE UNIVERSITY OF CALIFORNIA SPECIFICALLY DISCLAIMS ANY WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE SOFTWARE PROVIDED HEREUNDER IS ON AN “AS-IS” BASIS, AND THE UNIVERSITY OF CALIFORNIA HAS NO OBLIGATIONS TO PROVIDE MAINTENANCE, SUPPORT, UPDATES, ENHANCEMENTS, OR MODIFICATIONS.

CHAPTER 5

Conclusion

Database Compatibility for Oracle® Developers Built-in Packages Guide

Copyright © 2007 - 2020 EnterpriseDB Corporation. All rights reserved.

EnterpriseDB® Corporation 34 Crosby Drive, Suite 201, Bedford, MA 01730, USA

T +1 781 357 3390 F +1 978 467 1307 E info@enterprisedb.com www.enterprisedb.com

- EnterpriseDB and Postgres Enterprise Manager are registered trademarks of EnterpriseDB Corporation. EDB and EDB Postgres are trademarks of EnterpriseDB Corporation. Oracle is a registered trademark of Oracle, Inc. Other trademarks may be trademarks of their respective owners.
- EDB designs, establishes coding best practices, reviews, and verifies input validation for the logon UI for EDB Postgres Enterprise Manager where present. EDB follows the same approach for additional input components, however the nature of the product may require that it accepts freeform SQL, WMI or other strings to be entered and submitted by trusted users for which limited validation is possible. In such cases it is not possible to prevent users from entering incorrect or otherwise dangerous inputs.
- EDB reserves the right to add features to products that accept freeform SQL, WMI or other potentially dangerous inputs from authenticated, trusted users in the future, but will ensure all such features are designed and tested to ensure they provide the minimum possible risk, and where possible, require superuser or equivalent privileges.
- EDB does not warrant that we can or will anticipate all potential threats and therefore our process cannot fully guarantee that all potential vulnerabilities have been addressed or considered.

A

ALTER_POLICY, 171
 ALTER_QUEUE, 45
 ALTER_QUEUE_TABLE, 46
 ANALYZE_DATABASE, ANALYZE_SCHEMA and
 ANALYZE_PART_OBJECT, 242
 APPEND, 78

B

BASE64_DECODE, 261
 BASE64_ENCODE, 262
 BEGIN_REQUEST, 291
 BIND_VARIABLE_RAW, 218
 BIND_VARIABLE, 215
 BIND_VARIABLE_CHAR, 217
 BROKEN, 68
 Built-In Packages, 21

C

CANONICALIZE, 244
 CAST_TO_RAW, 311
 CAST_TO_VARCHAR2, 312
 CHANGE, 69
 CHARARR, 100
 CLOSE_CURSOR, 219
 CLOSE_DATA, 318
 COLUMN_VALUE, 220
 COLUMN_VALUE_CHAR, 222
 COLUMN_VALUE_RAW, 223
 COMMA_TO_TABLE, 246
 COMMAND, 318
 COMMAND_REPLIES, 319
 COMPARE, 79
 CONCAT, 313
 Conclusion, 329
 CONNECTION, 318
 CONVERT, 314
 CONVERTTLOB, 80
 CONVERTTOCLOB, 81

COPY, 82
 CREATE_JOB, 196
 CREATE_PIPE, 112
 CREATE_PROGRAM, 198
 CREATE_QUEUE, 47
 CREATE_QUEUE_TABLE, 49
 CREATE_SCHEDULE, 200
 Creating Packages, 12
 Creating the Package Body, 13
 Creating the Package Specification, 12

D

DATA, 319
 DB_VERSION, 247
 DBMS_ALERT, 22
 DBMS_ALERT_Comprehensive_example, 27
 DBMS_ALERT_Register, 22
 DBMS_ALERT_Remove, 23
 DBMS_AQ, 29
 DBMS_AQ_Register, 39
 DBMS_AQADM, 43
 DBMS_CRYPT, 56
 DBMS_JOB, 66
 DBMS_JOB_Remove, 72
 DBMS_LOB, 76
 DBMS_LOB_SUBSTR, 88
 DBMS_LOCK, 92
 DBMS_MVIEW, 93
 DBMS_OUTPUT, 100
 DBMS_OUTPUT_DISABLE, 101
 DBMS_OUTPUT_ENABLE, 102
 DBMS_OUTPUT_PUT, 108
 DBMS_OUTPUT_PUT_LINE, 109
 DBMS_PIPE, 111
 DBMS_PIPE_Comprehensive_example, 125
 DBMS_PROFILER, 128
 DBMS_PROFILER - Reference, 140
 DBMS_RANDOM, 151
 DBMS_RANDOM_VALUE_FIRST_FORM, 153
 DBMS_RANDOM_VALUE_SECOND_FORM, 154

DBMS_REDACT, 155
 DBMS_REDACT_ADD_POLICY, 168
 DBMS_REDACT_DROP_POLICY, 176
 DBMS_REDACT_ENABLE_POLICY, 175
 DBMS_RLS, 180
 DBMS_RLS_DROP_POLICY, 190
 DBMS_RLS_ENABLE_POLICY, 191
 DBMS_RLSADD_POLICY:, 183
 DBMS_SCHEDULER, 192
 DBMS_SCHEDULER_DISABLE, 203
 DBMS_SCHEDULER_ENABLE, 208
 DBMS_SESSION, 212
 DBMS_SQL, 213
 DBMS_UTILITY, 240
 DECRYPT, 57
 DEFINE_COLUMN, 224
 DEFINE_COLUMN_CHAR, 226
 DEFINE_COLUMN_RAW, 227
 DEFINE_PROGRAM_ARGUMENT, 201
 DEQUEUE, 35
 DESCRIBE COLUMNS, 228
 DISABLE_POLICY, 174
 DROP_JOB, 204
 DROP_PROGRAM, 205
 DROP_PROGRAM_ARGUMENT, 206
 DROP_QUEUE, 51
 DROP_QUEUE_TABLE, 52
 DROP_SCHEDULE, 207
 Dropping a Package, 20

E

EHLO, 320
 ENCRYPT, 59
 END_REQUEST, 292
 END_RESPONSE, 292
 ENQUEUE, 32
 ERASE, 83
 ESCAPE, 325
 EVALUATE_CALENDAR_STRING, 209
 EXEC_DDL_STATEMENT, 248
 EXECUTE, 229
 EXECUTE_AND_FETCH, 230

F

FCLOSE, 274
 FCLOSE_ALL, 274
 FCOPY, 274
 FETCH_ROWS, 232
 FFLUSH, 276
 FLUSH_DATA, 129
 FOPEN, 277
 FORMAT_CALL_STACK, 249
 FREMOVE, 277
 FRENAME, 278

G

GET_BODY_CHARSET, 292
 GET_CPU_TIME, 250
 GET_DEPENDENCY, 251
 GET_FOLLOW_REDIRECT, 293
 GET_HASH_VALUE, 252
 GET_HEADER, 293
 GET_HEADER_BY_NAME, 295
 GET_HEADER_COUNT, 295
 GET_LINE, 103, 280
 GET_LINES, 105
 GET_MV_DEPENDENCIES, 94
 GET_PARAMETER_VALUE, 253
 GET_RESPONSE, 296
 GET_RESPONSE_ERROR_CHECK, 296
 GET_STORAGE_LIMIT, 84
 GET_TIME, 254
 GET_TRANSFER_TIMEOUT, 296
 GET_VERSION, 129
 GETLENGTH, 85

H

HASH, 61
 HELO, 320
 HELP, 320
 HTML_PIECES, 291

I

INITIALIZE, 151
 INSTR, 86
 INTERNAL_VERSION_CHECK, 129
 INTERVAL, 70
 Introduction, 1
 IS_OPEN, 234, 281

L

LAST_ROW_COUNT, 235
 LENGTH, 315
 LNAME_ARRAY, 241

M

MAC, 62
 MAIL, 321
 MIMEHEADER_DECODE, 263
 MIMEHEADER_ENCODE, 264

N

NAME_TOKENIZE, 255
 NEW_LINE, 107, 281
 NEXT_DATE, 71
 NEXT_ITEM_TYPE, 113
 NOOP, 321
 NORMAL, 151

O

OPEN_CONNECTION, 321
 OPEN_CURSOR, 237
 OPEN_DATA, 322

P

PACK_MESSAGE, 115
 Package Body Syntax, 7
 Package Components, 4
 Package Specification Syntax, 4
 Packages, 3
 PARSE, 238
 PAUSE_PROFILER, 129
 PLSQL_PROFILER_DATA, 141
 PLSQL_PROFILER_RAWDATA, 141
 PLSQL_PROFILER_RUNS, 140
 PLSQL_PROFILER_UNITS, 140
 PURGE, 116
 PURGE_QUEUE_TABLE, 53
 PUTF, 285

Q

QUIT, 322
 QUOTED_PRINTABLE_DECODE, 265
 QUOTED_PRINTABLE_ENCODE, 266

R

RANDOM, 152
 RANDOMBYTES, 63
 RANDOMINTEGER, 64
 RANDOMNUMBER, 65
 RCPT, 322
 READ, 87
 READ_LINE, 297
 READ_RAW, 299
 READ_TEXT, 299
 RECEIVE_MESSAGE, 118
 Referencing a Package, 15
 REFRESH, 95
 REFRESH_ALL_MVIEWS, 97
 REFRESH_DEPENDENT, 98
 REMOVE_PIPE, 119
 REMOVEALL, 23
 REPLY/REPLIES, 318
 REQ, 291
 REQUEST, 300
 REQUEST_PIECES, 301
 RESET_BUFFER, 121
 RESP, 291
 RESUME_PROFILER, 130
 RSET, 322
 RUN, 73
 RUN_JOB, 210

S

SEED, 152
 SEND, 307
 SEND_ATTACH_RAW, 308
 SEND_ATTACH_VARCHAR2, 310
 SEND_MESSAGE, 122
 SERVEROUTPUT, 110
 SET_BODY_CHARSET, 301
 SET_FOLLOW_REDIRECT, 301
 SET_HEADER, 302
 SET_JOB_ARGUMENT_VALUE, 211
 SET_RESPONSE_ERROR_CHECK, 302
 SET_ROLE, 212
 SET_TRANSFER_TIMEOUT, 302
 Setting File Permissions with utl_file.umask, 273
 SIGNAL, 23
 SLEEP, 92
 START_PROFILER, 130
 START_QUEUE, 54
 STOP_PROFILER, 130
 STOP_QUEUE, 55
 STRING, 153
 SUBMIT, 74

T

TABLE_TO_COMMA, 258
 TERMINATE, 153
 TEXT_DECODE, 267
 TEXT_ENCODE, 268
 TRIM, 89

U

UNCL_ARRAY, 241
 UNESCAPE, 327
 UNIQUE_SESSION_NAME, 123
 UNPACK_MESSAGE, 124
 UNREGISTER, 41
 UPDATE_FULL_REDACTION_VALUES, 177
 Using Calendar Syntax to Specify a Repeating Interval, 194
 Using DBMS_PROFILER, 131
 Using DBMS_REDACT Constants and Function Parameters, 156
 Using Packages With User Defined Types, 16
 UTL_ENCODE, 260
 UTL_FILE, 272
 UTL_FILE_PUT, 282
 UTL_FILE_PUT_LINE, 284
 UTL_HTTP, 287
 UTL_MAIL, 306
 UTL_RAW, 311
 UTL_RAW_SUBSTR, 316
 UTL_SMTP, 317
 UTL_SMTP_Comprehensive_example, 323

UTL_URL, 325
UUDECODE, 269
UUENCODE, 270

V

VERFY, 323

W

WAITANY, 24
WAITONE, 25
WHAT, 75
What's New, 2
WRITE, 90
WRITE_DATA, 323
WRITE_LINE, 303
WRITE_RAW, 304
WRITE_TEXT, 305
WRITEAPPEND, 91